

RBE550

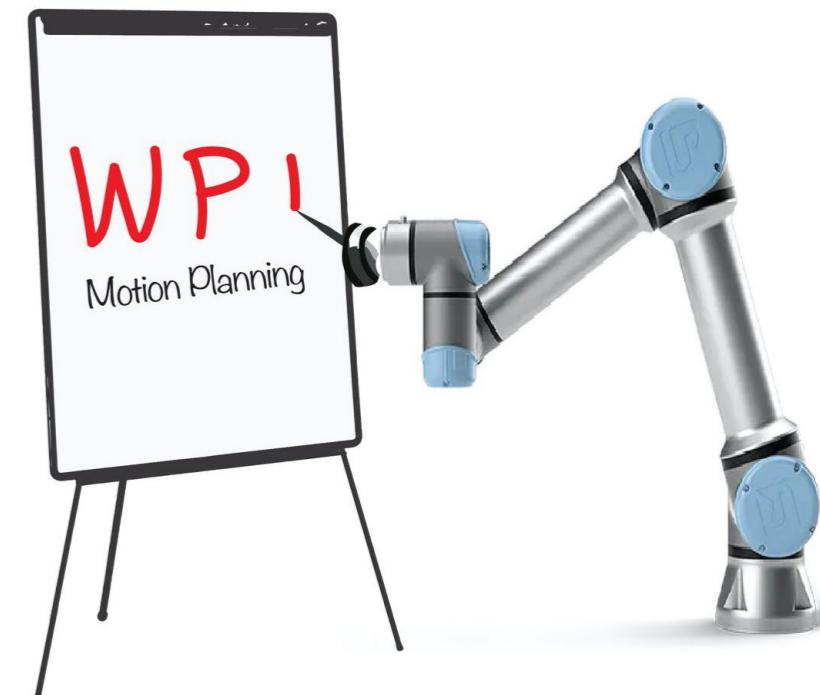
Motion Planning

Learning and Motion Planning II

Constantinos Chamzas

www.cchamzas.com

www.elpislab.org



Last Lecture Overview

- Why Combine Learning and Motion Planning?
 - To improve planning efficiency in challenging problems
- Learning For Motion Planning Archetypes
 - Retrieve and Repair, Biased Sampling
- Discussed 4 Papers
 - Lightning(2013) Thunder (2015), Rep Sampling (2017), Rep (Roadmap) (2019)

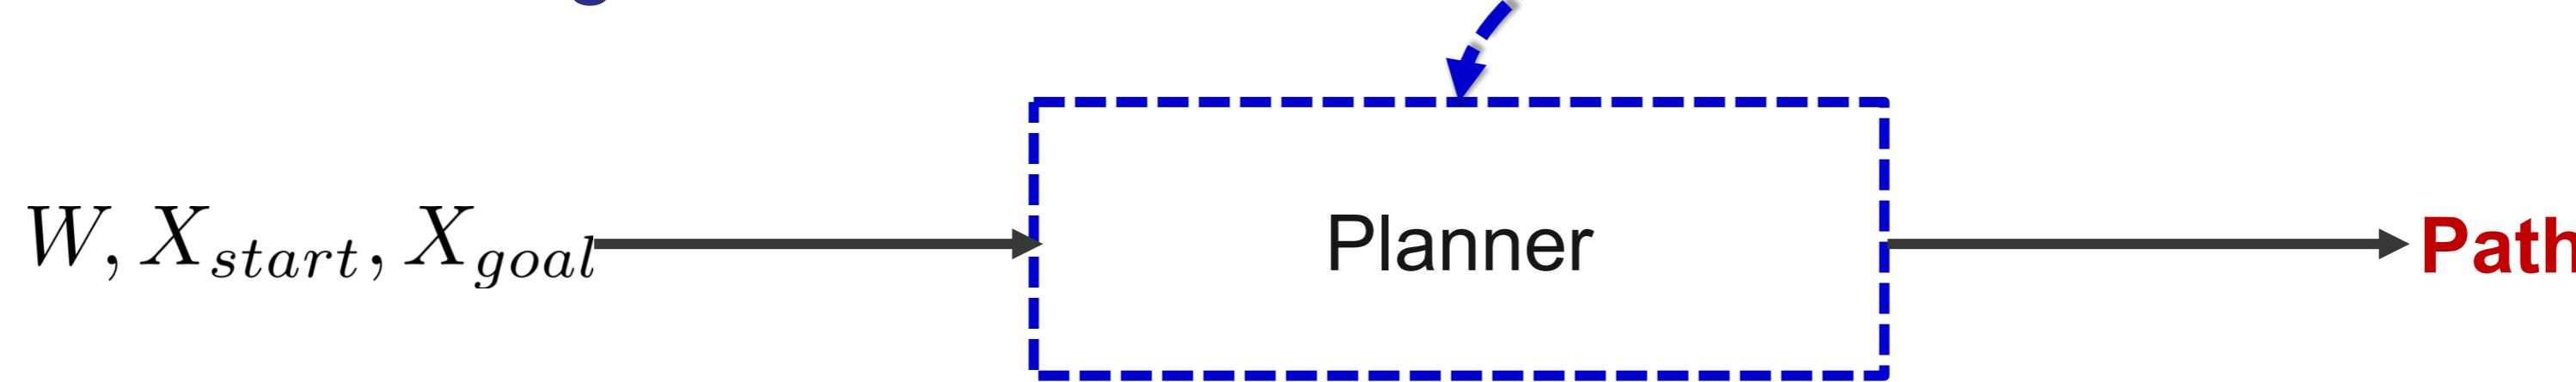
Today: We will focus only on Biased Sampling Methods

Machine Learning and Motion Planning

Training/Learning Phase



Testing/Inference Phase



Biasing the Sampling Distribution

Algorithm 1: General Sampling Based Planner

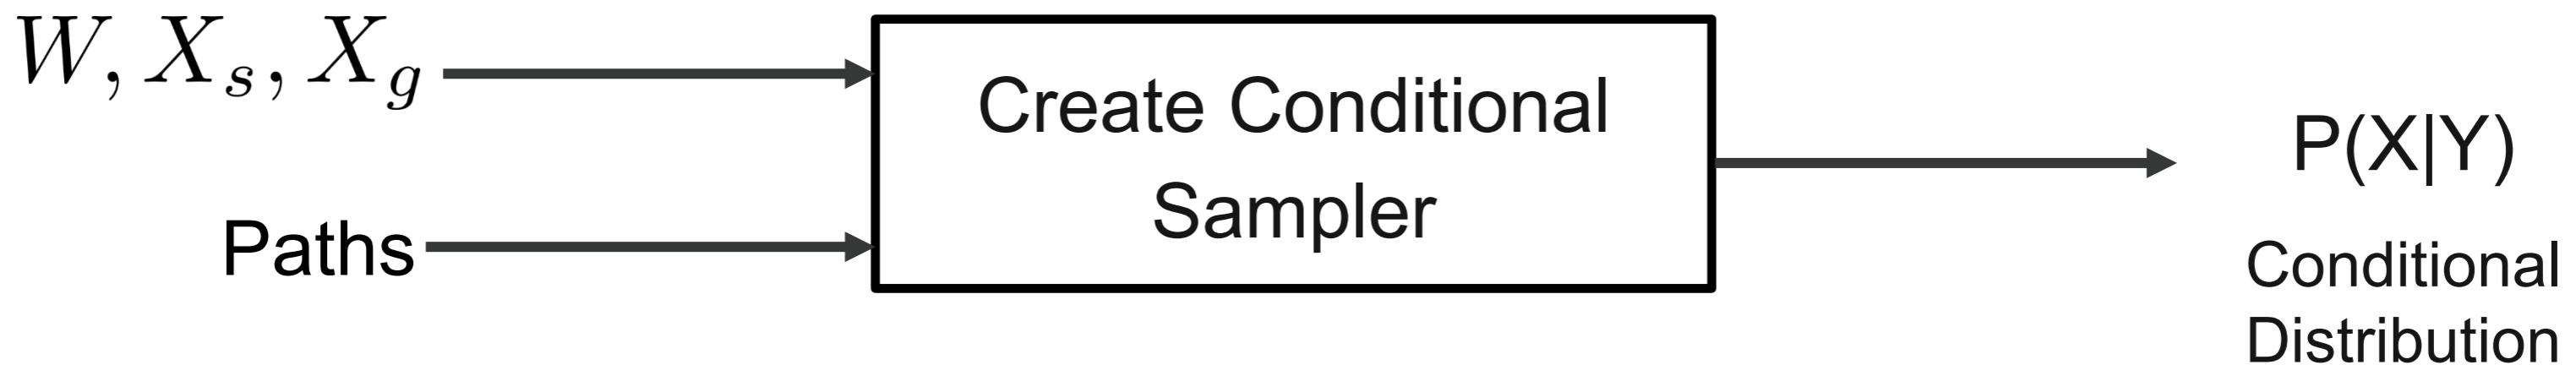
Input : Number of iterations N
Output : Graph structure G

```
1 while  $i \leq N$  or  $solutionFound()$  do
2   |   x  $\sim$  Uniform()
3   |   update( $G$ ,  $x$ )
4 end
5 return  $G$ 
```

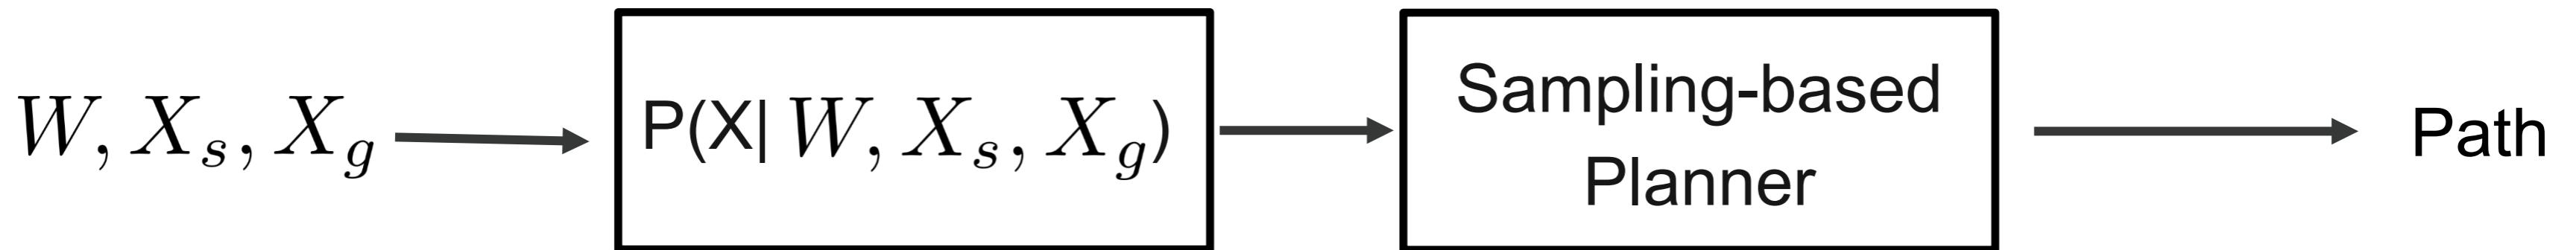
- Theoretical Insights supporting biased sampling
- Any sampling-based planner can be used

Biased Samplers Archetype

Training/Learning Phase



Testing/Inference



Learning Archetypes (Methodologies)

Learning Archetypes

1. Retrieve and Repair

2. Biased Samplers

Corresponding Papers

**Lightning[1], Thunder [2],
ERT[3],
Sim-Obstacles[4], Traj-Pred[5]**

**Rep-Sampling[6], Rep-Roadmaps[7], AWS[8],
SPARK2D[9], CVAE[10],
FLAME[11], FIRE[12]**

Today's Agenda

- Important Considerations
 - Locality of narrow passages
 - Discontinuity
- Biased Samplers (Cont.)
 - 2D Workspaces
 - 3D Workspaces
 - Workplace, Start and Goal

Learning Archetypes (Methodologies)

Learning Archetypes

1. Retrieve and Repair

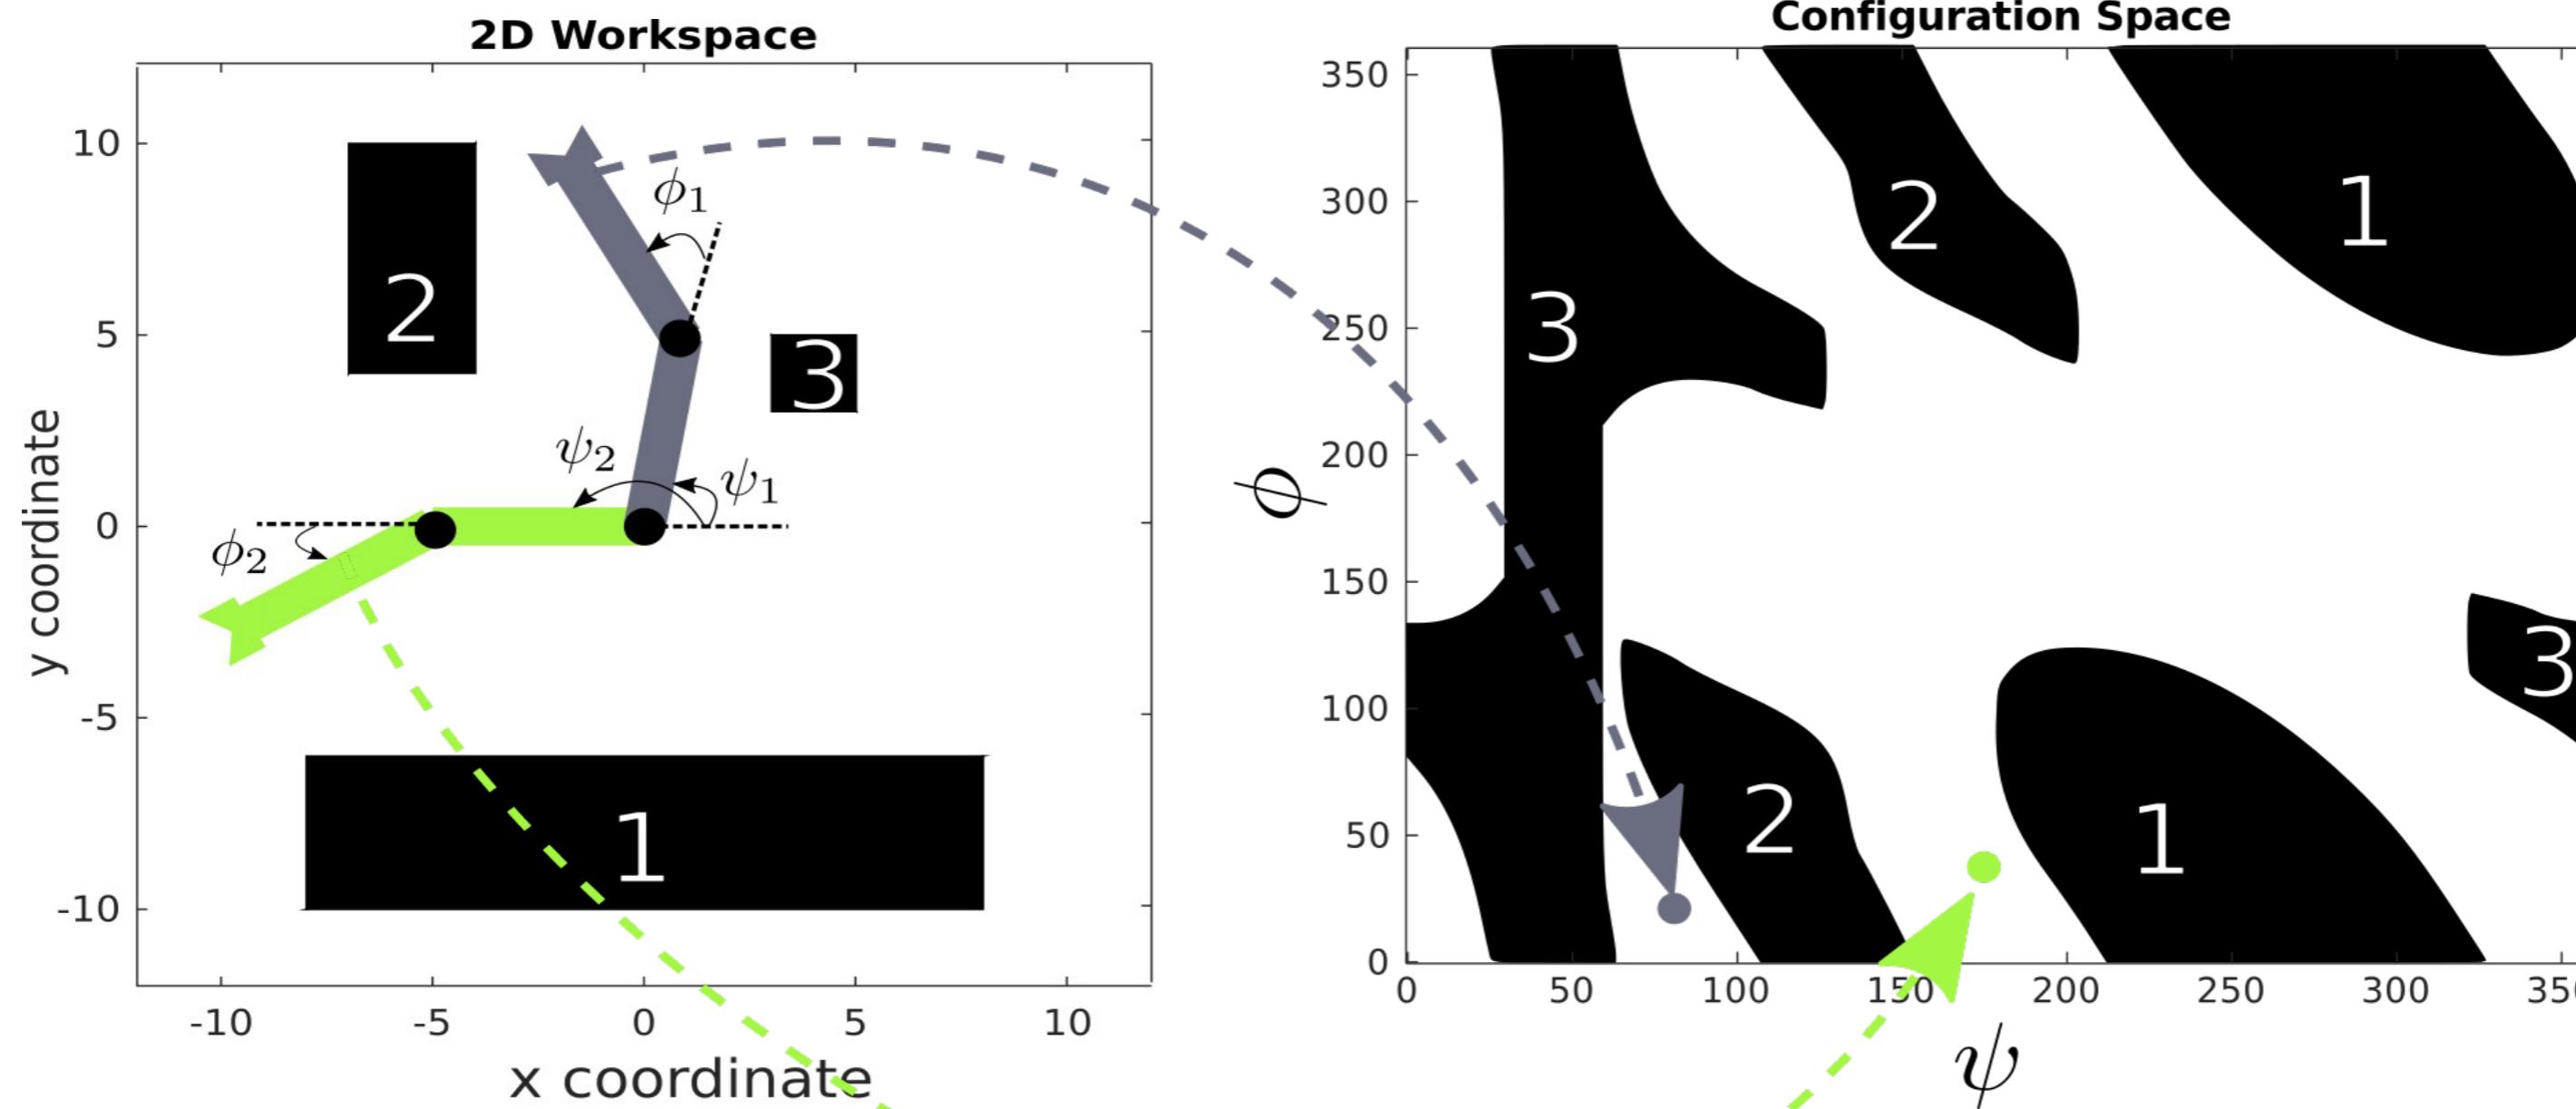
2. Biased Samplers

Corresponding Papers

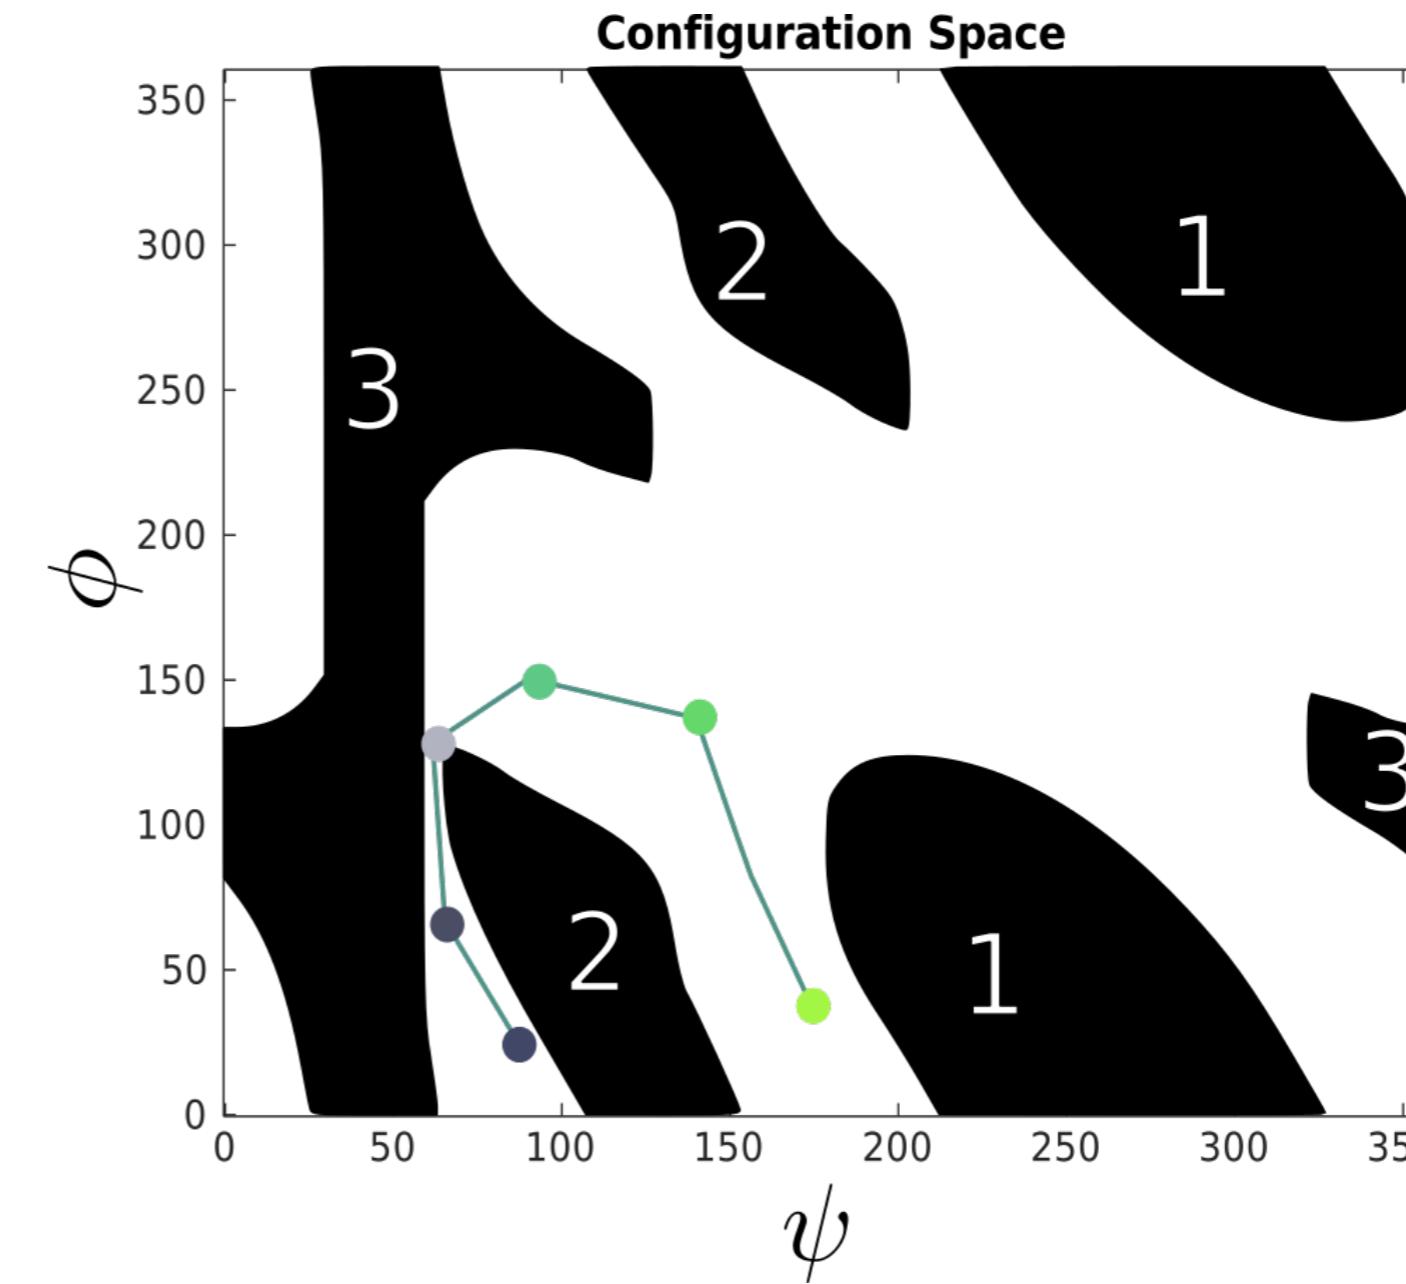
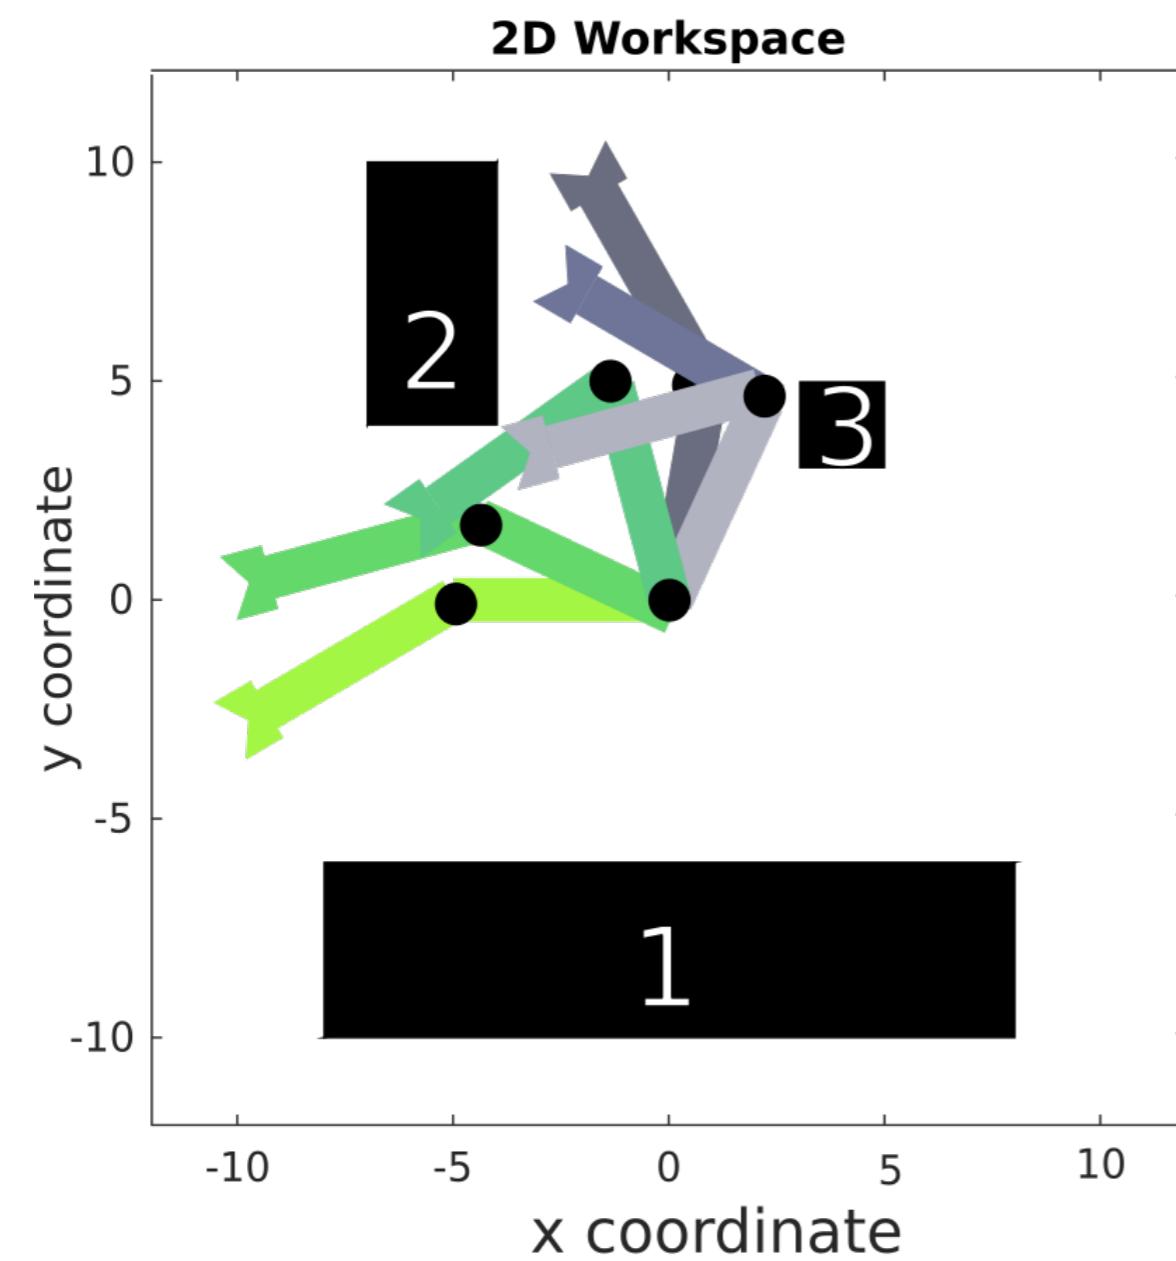
Lightning[1], Thunder [2],
ERT[3], Sim-Obstacles[4], Traj-
Pred[5]

Rep-Sampling[6], Rep-
Roadmaps[7], AWS[8],
SPARK2D[9], **CVAE[10]**,
FLAME[11], **FIRE[12]**

A 2D Manipulator Problem



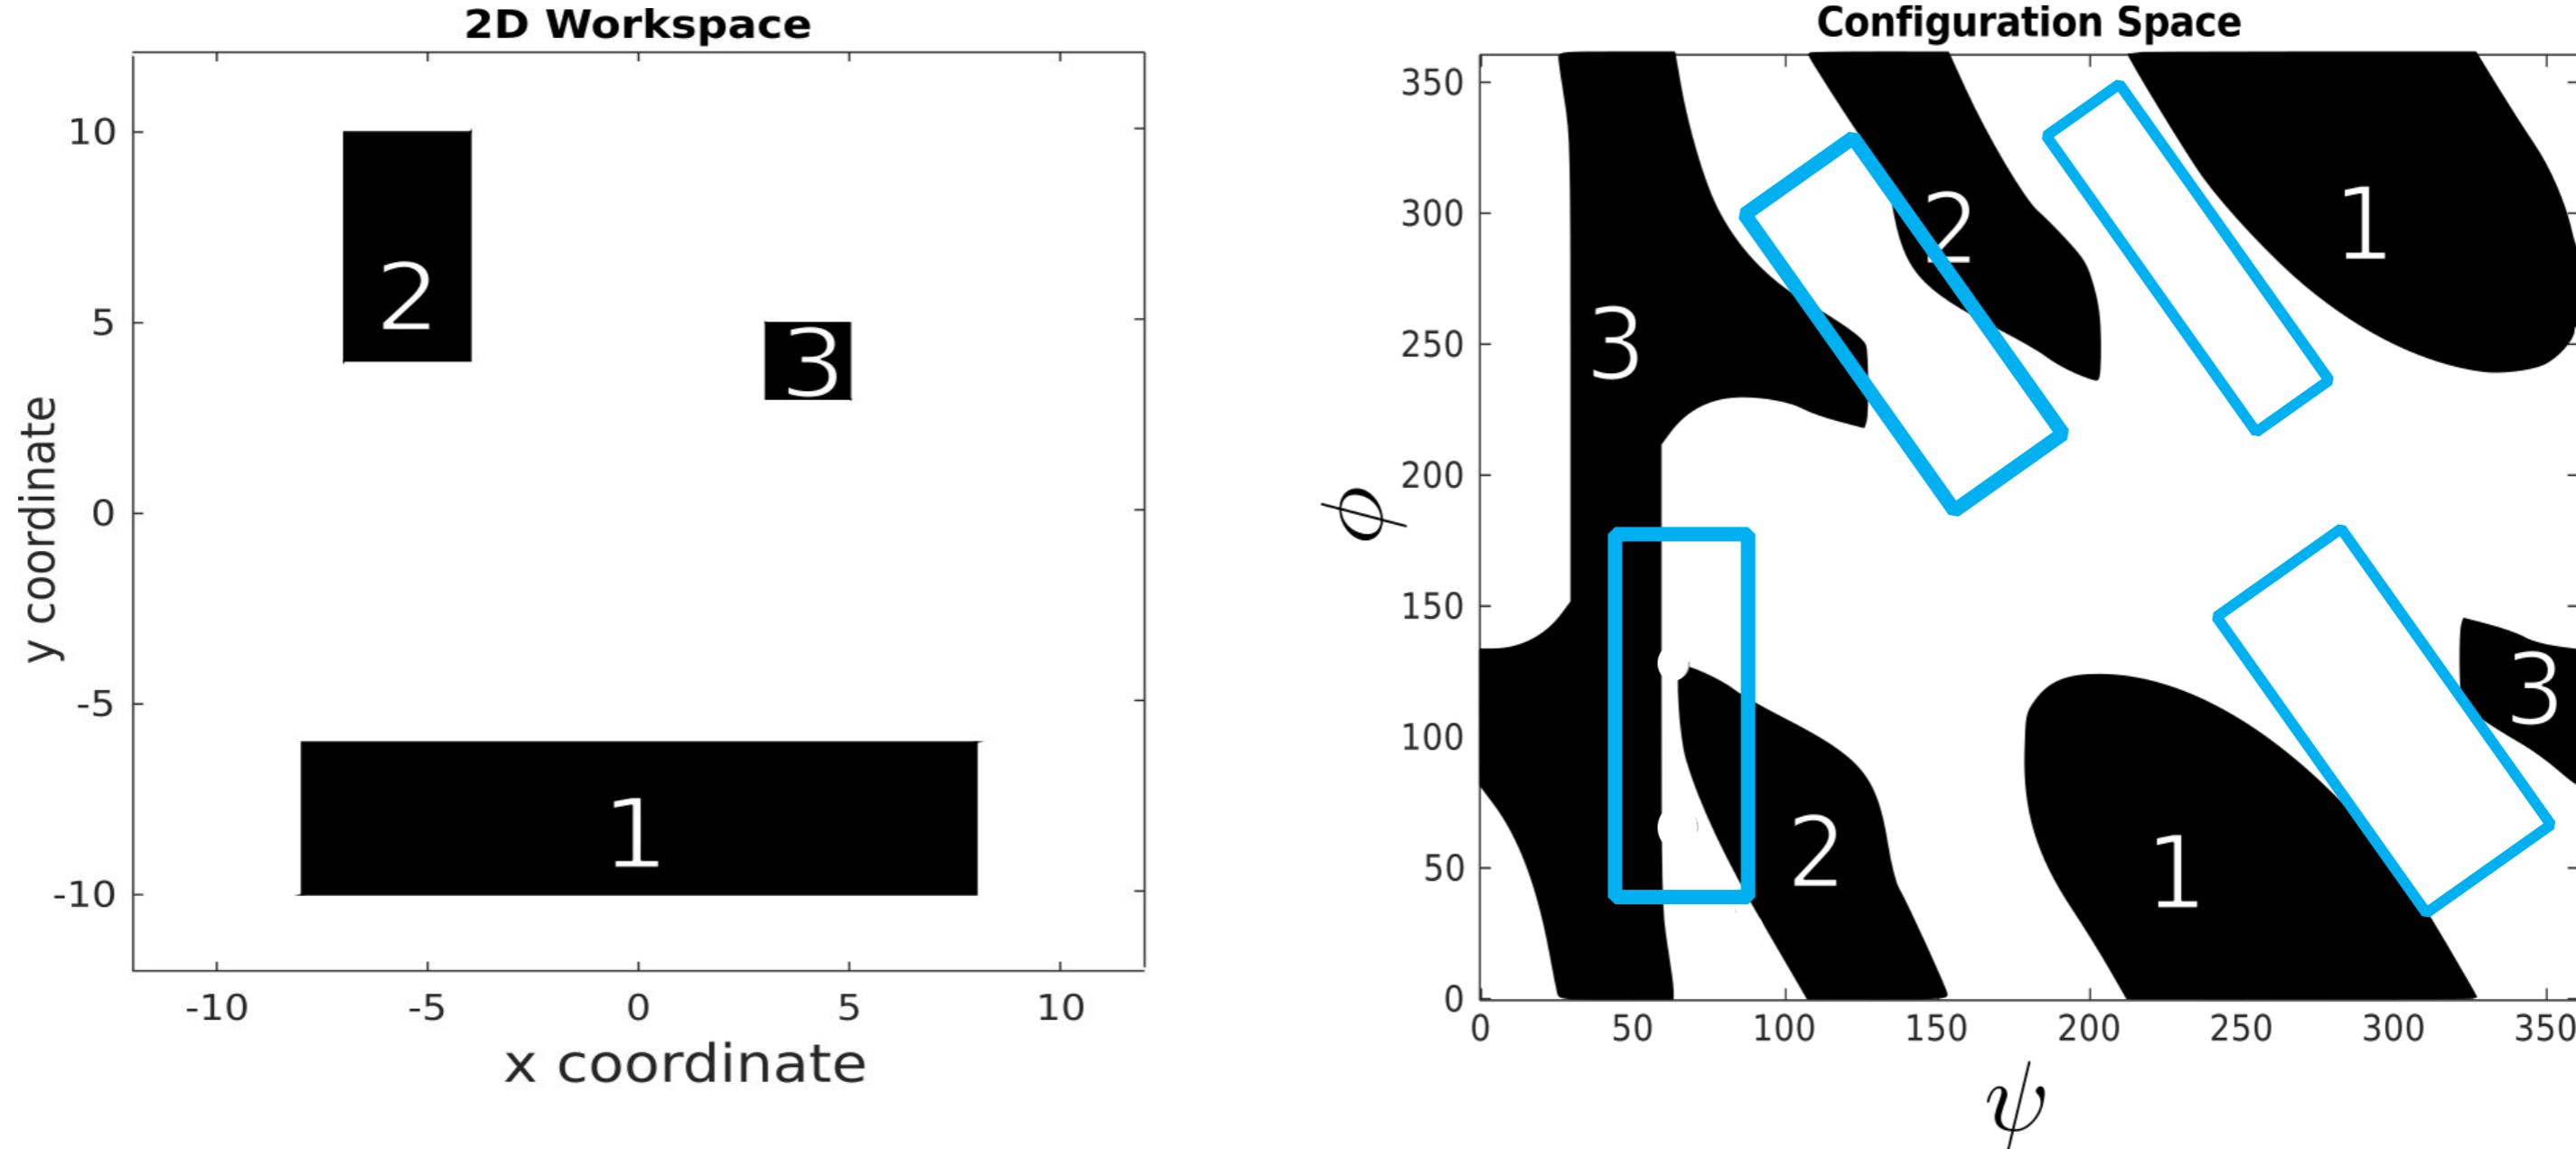
A 2D Manipulator Path/Solution



Conditioned Biased Sampling Distributions

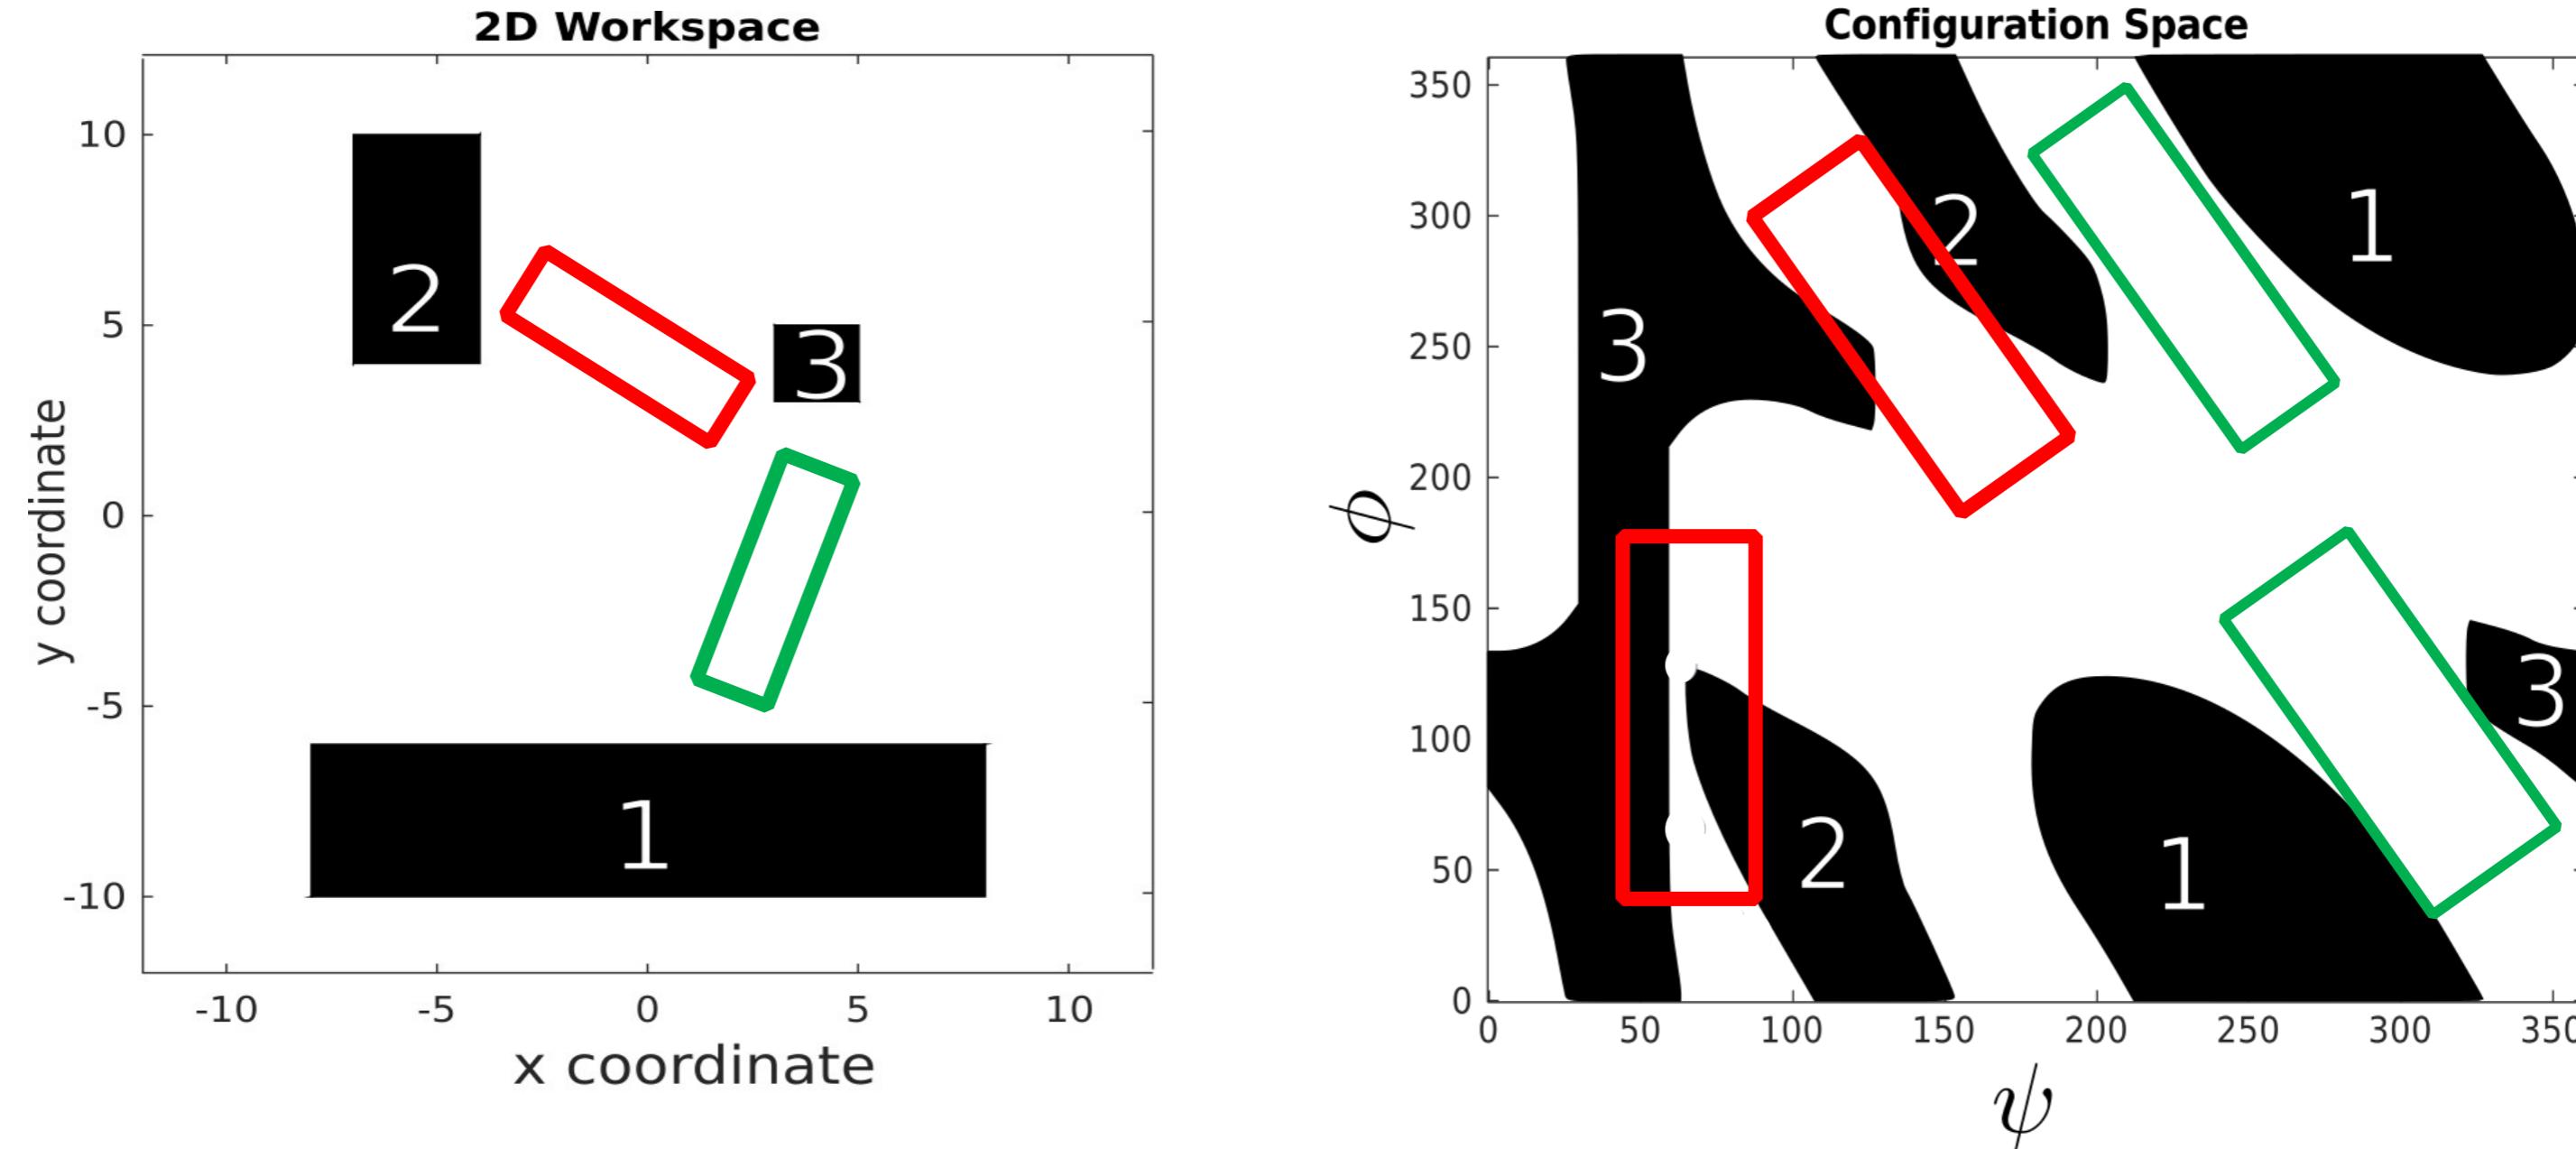
- $P(x)$: Learning Only Invariants
- $P(x| \text{Start, Goal})$: Sampling conditioned on start and goal
- **$P(x| W)$: Sampling biased from workspace features**
- $P(x| \text{Start, Goal, W})$: Sampling leveraging both workspace, start and goal information

Using only workspace features



Since we don't have start/goal information we will sample on all narrow passages

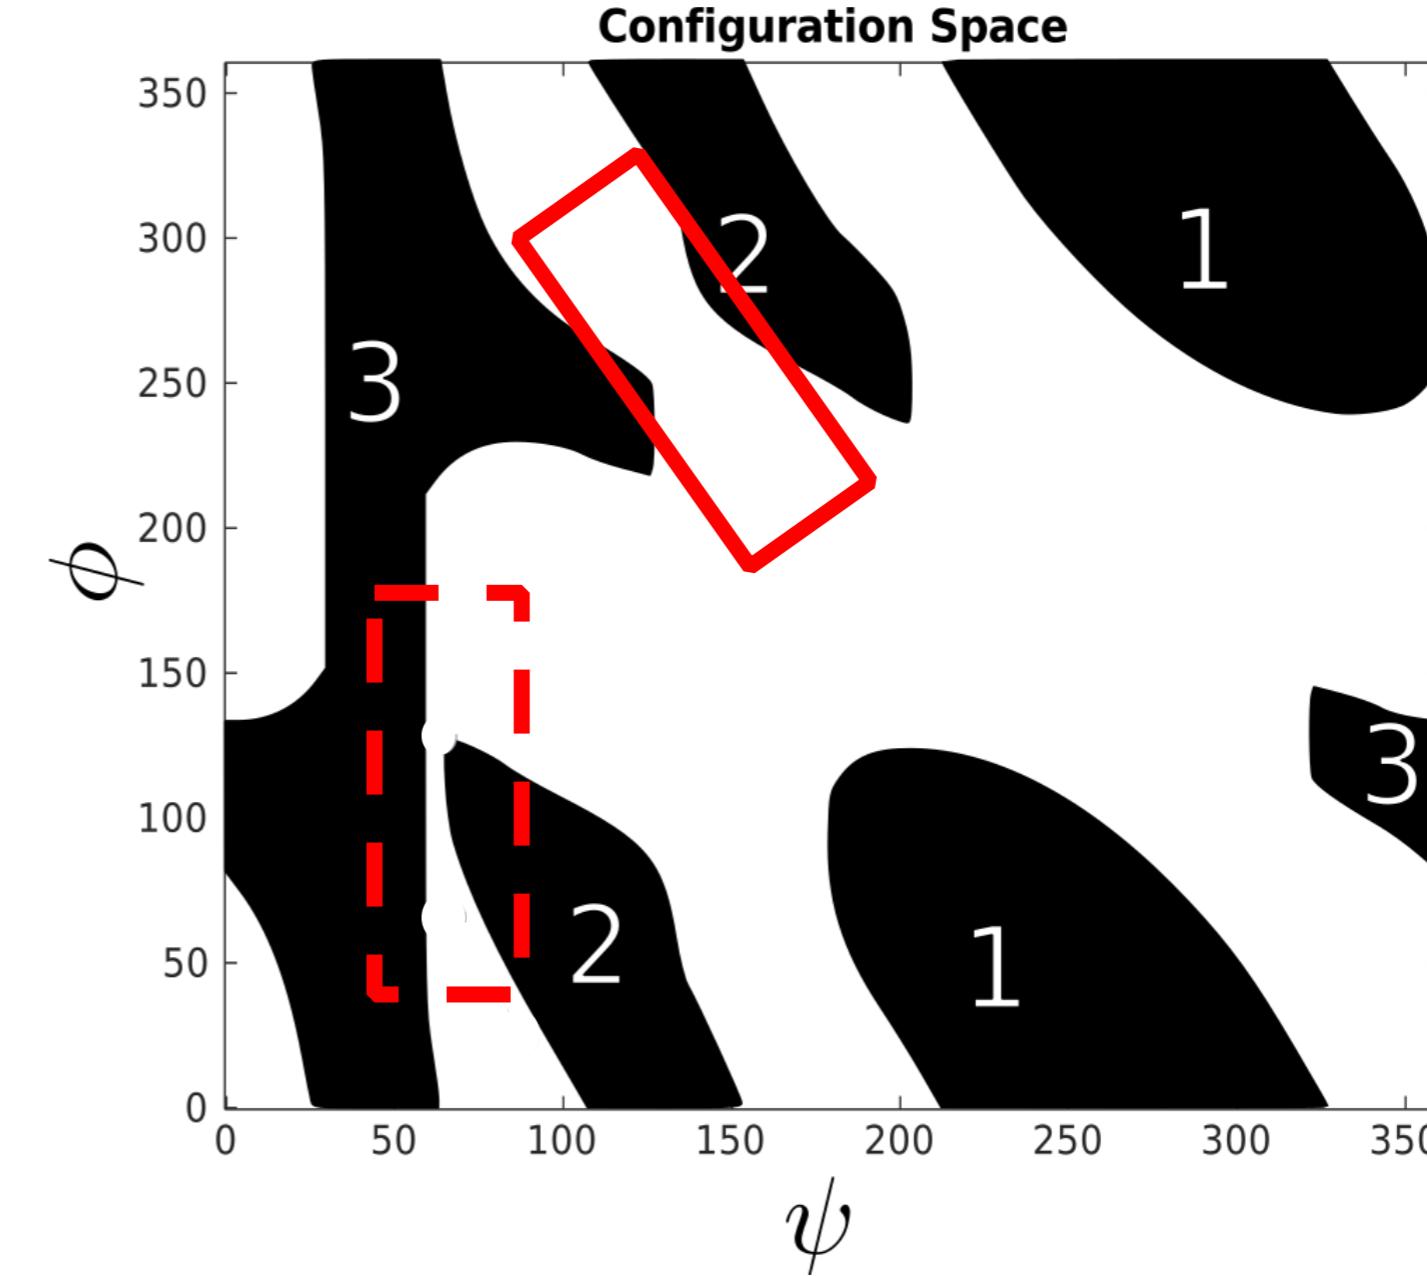
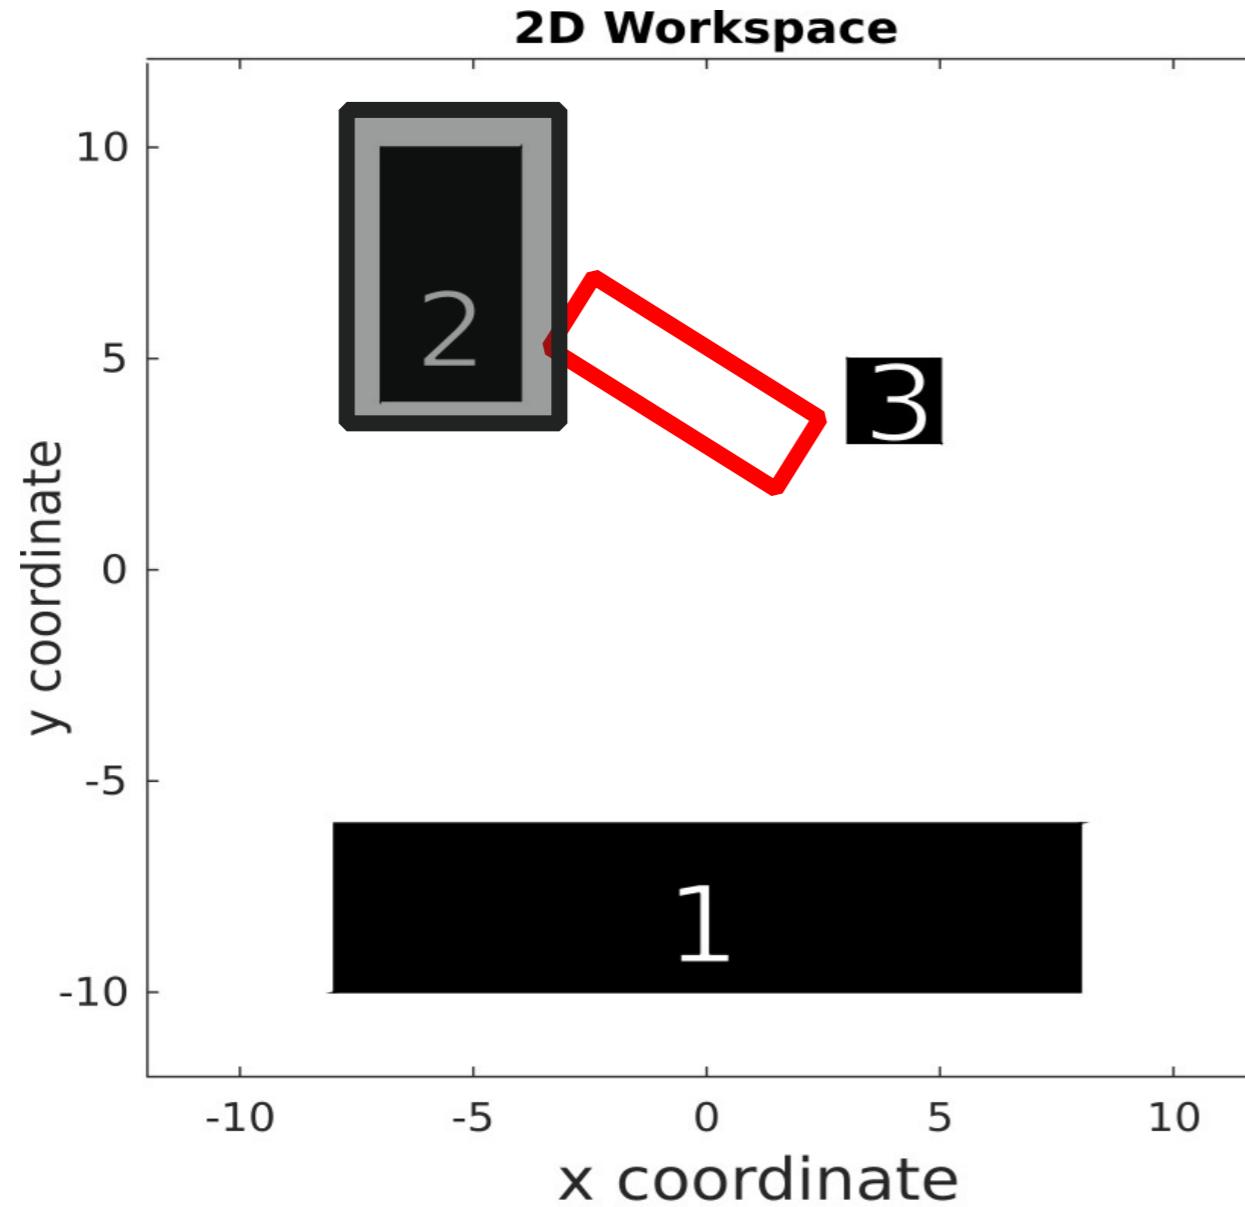
Consideration 1: Locality of narrows passages



Local workspace areas create different narrow passages !

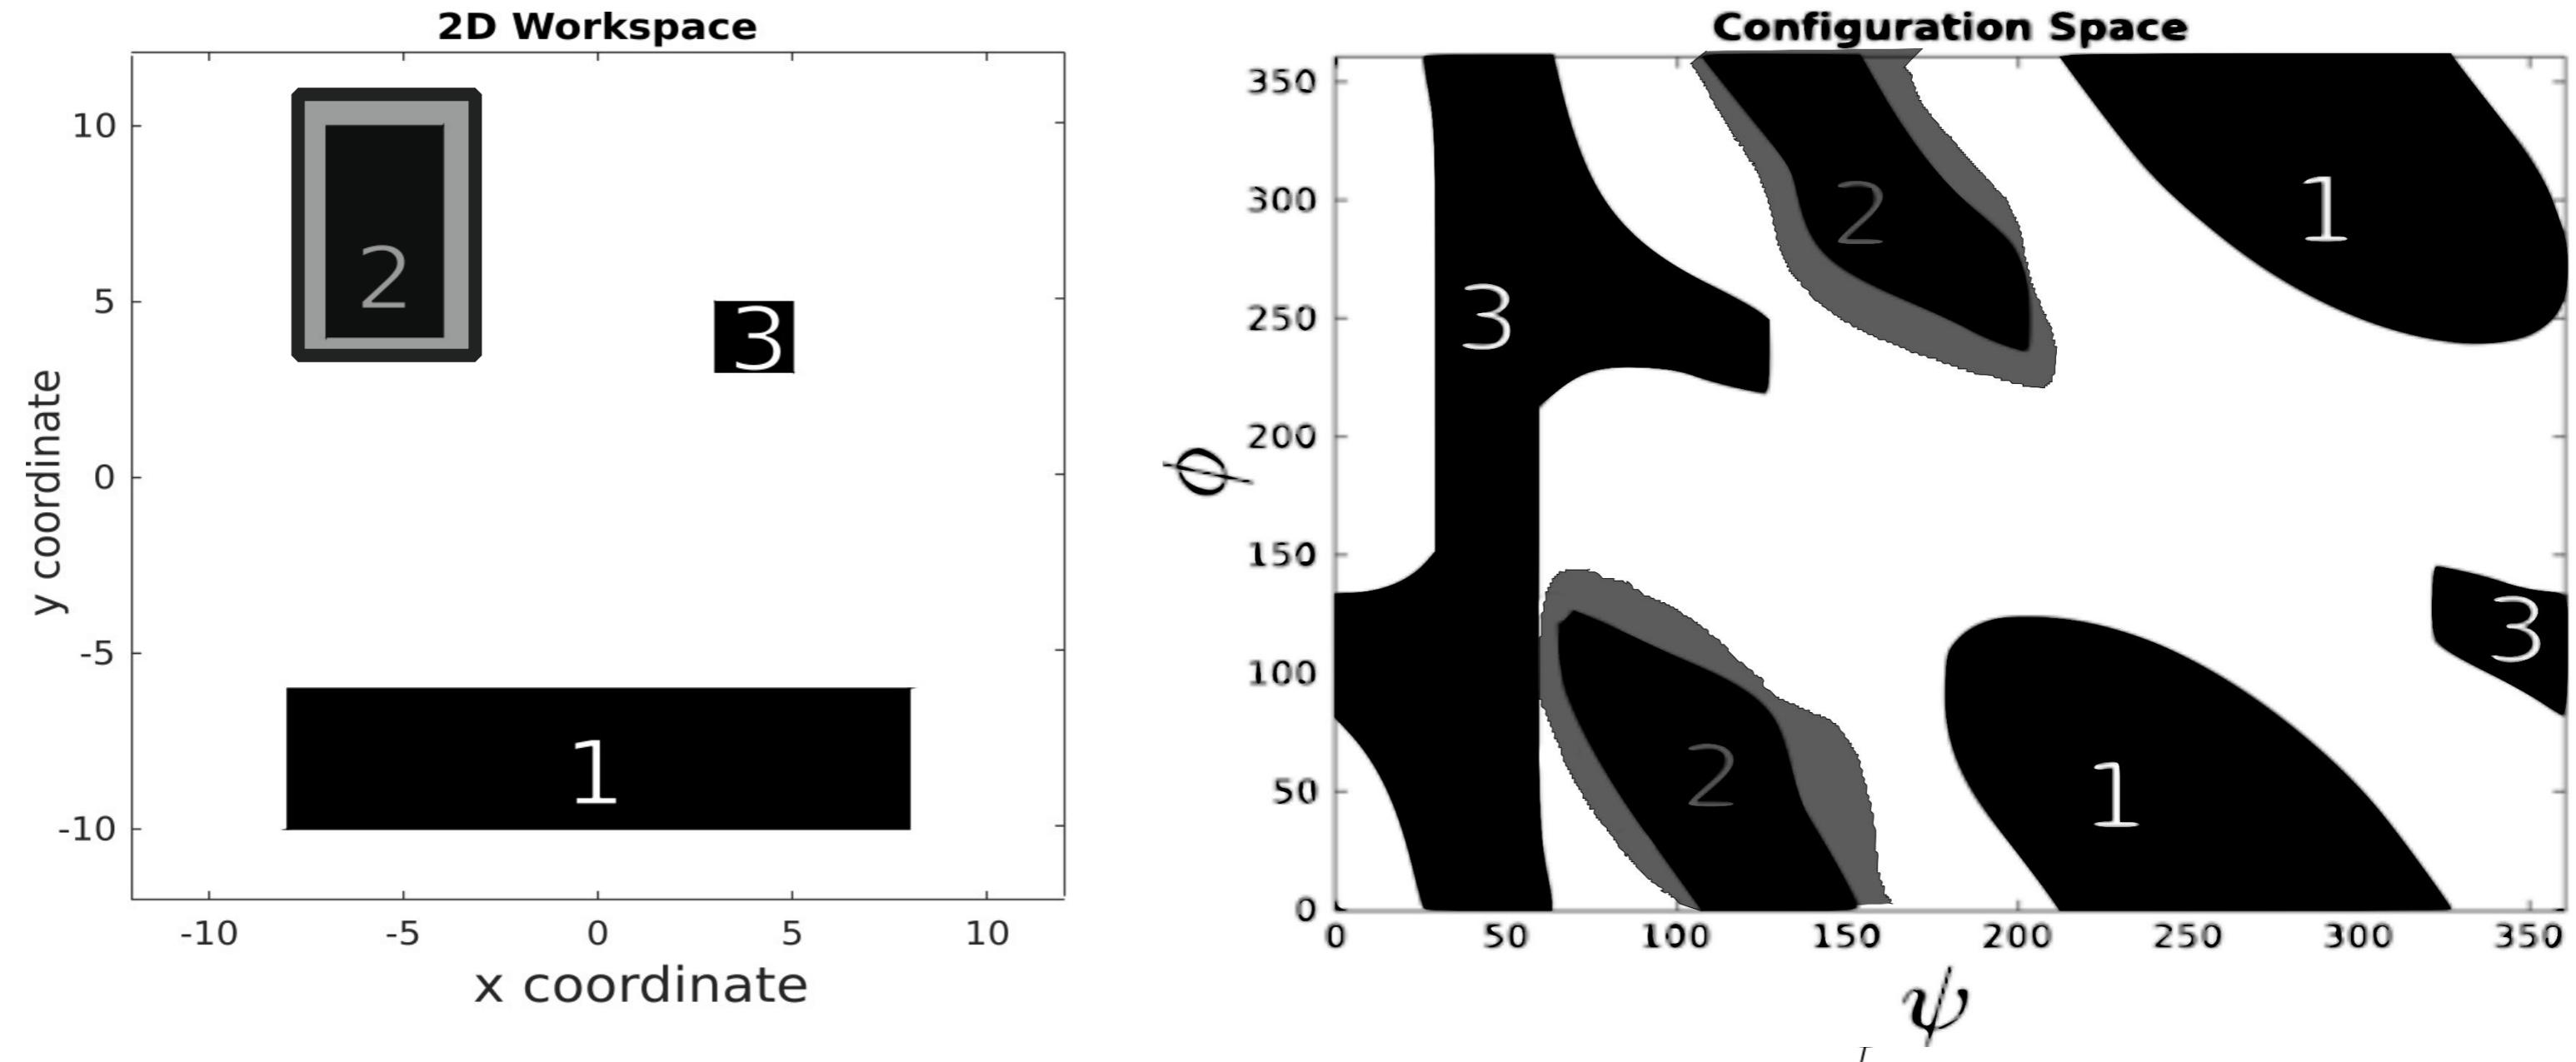
Consideration 2: Discontinuity

Lets make obstacle 2 a bit larger



What will happen to the narrow passages in the c-space?

Consideration 2: Discontinuity

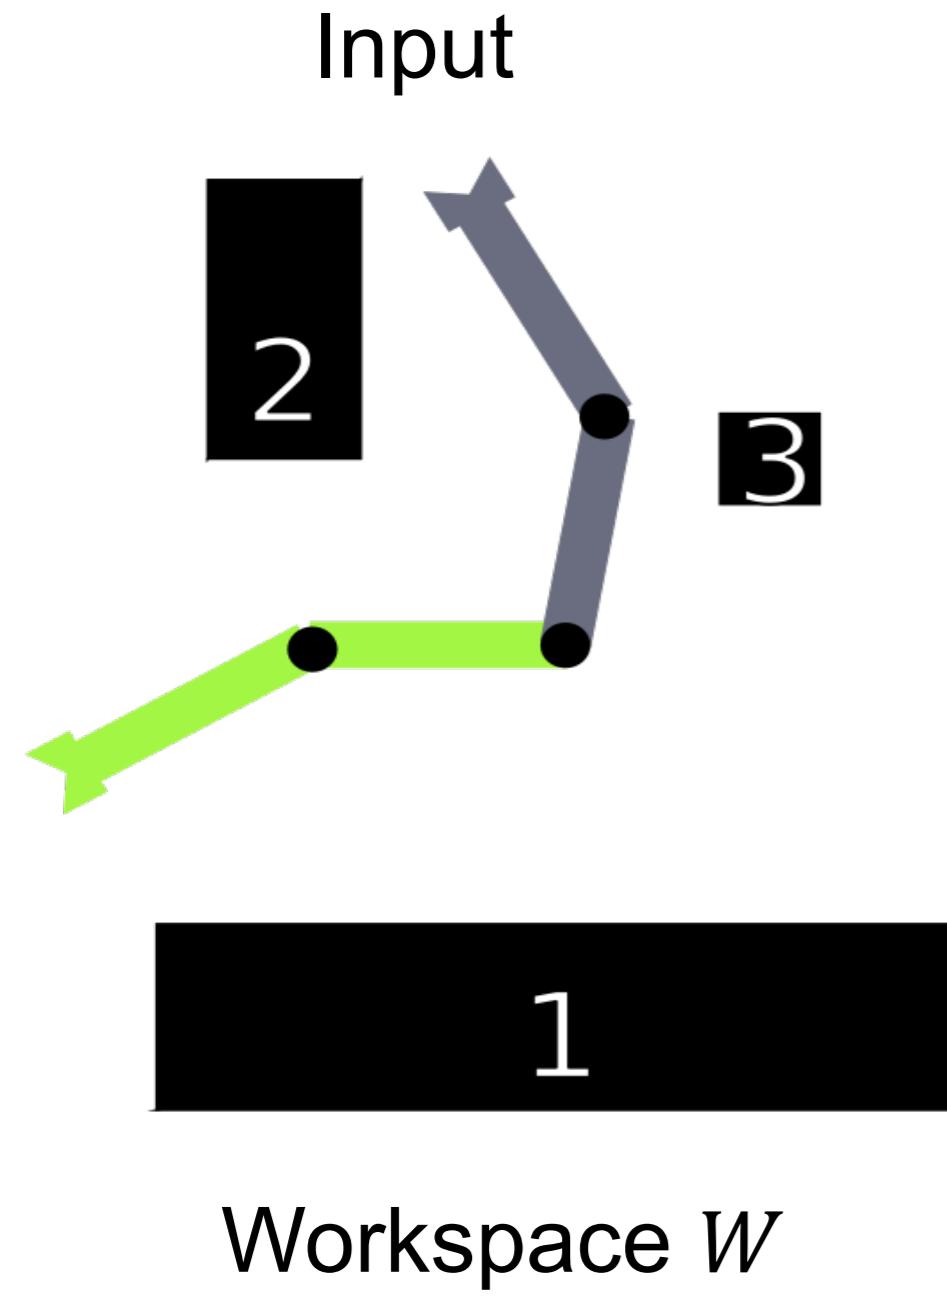


Small changes in the workspace can affect C-Space solutions a lot

Today's Agenda

- Important Considerations
 - Locality of narrow passages
 - Discontinuity
- Biased Samplers (Cont.)
 - **2D Workspaces**
 - 3D Workspaces
 - Workplace, Start and Goal

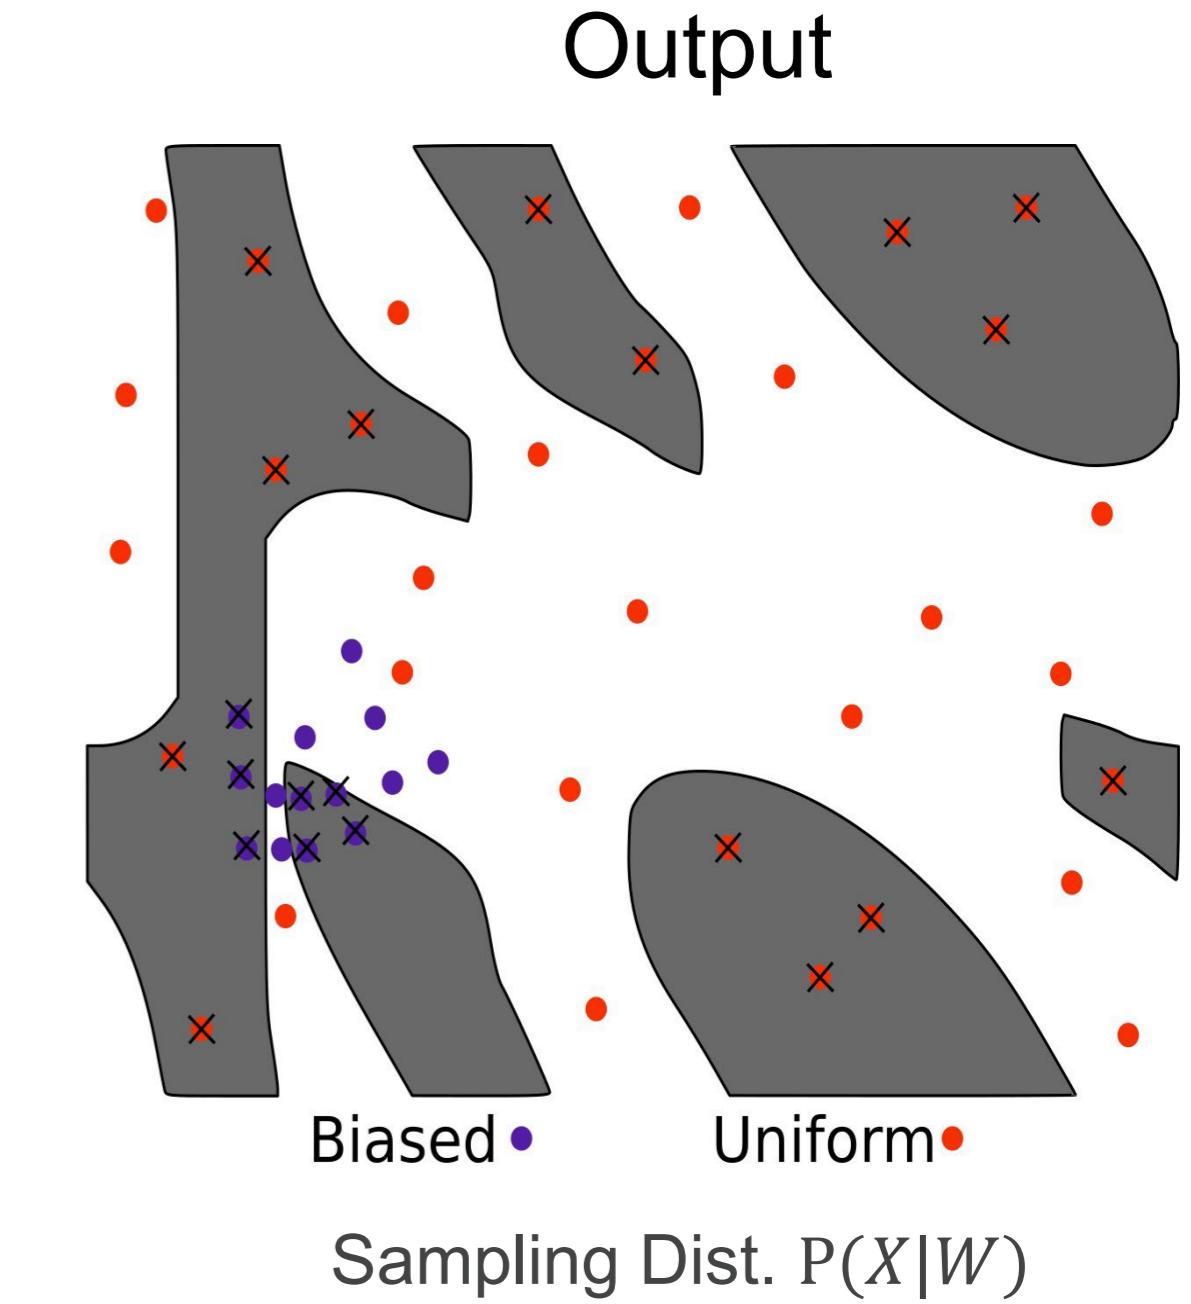
Learning a biased sampling distribution (Workspace)



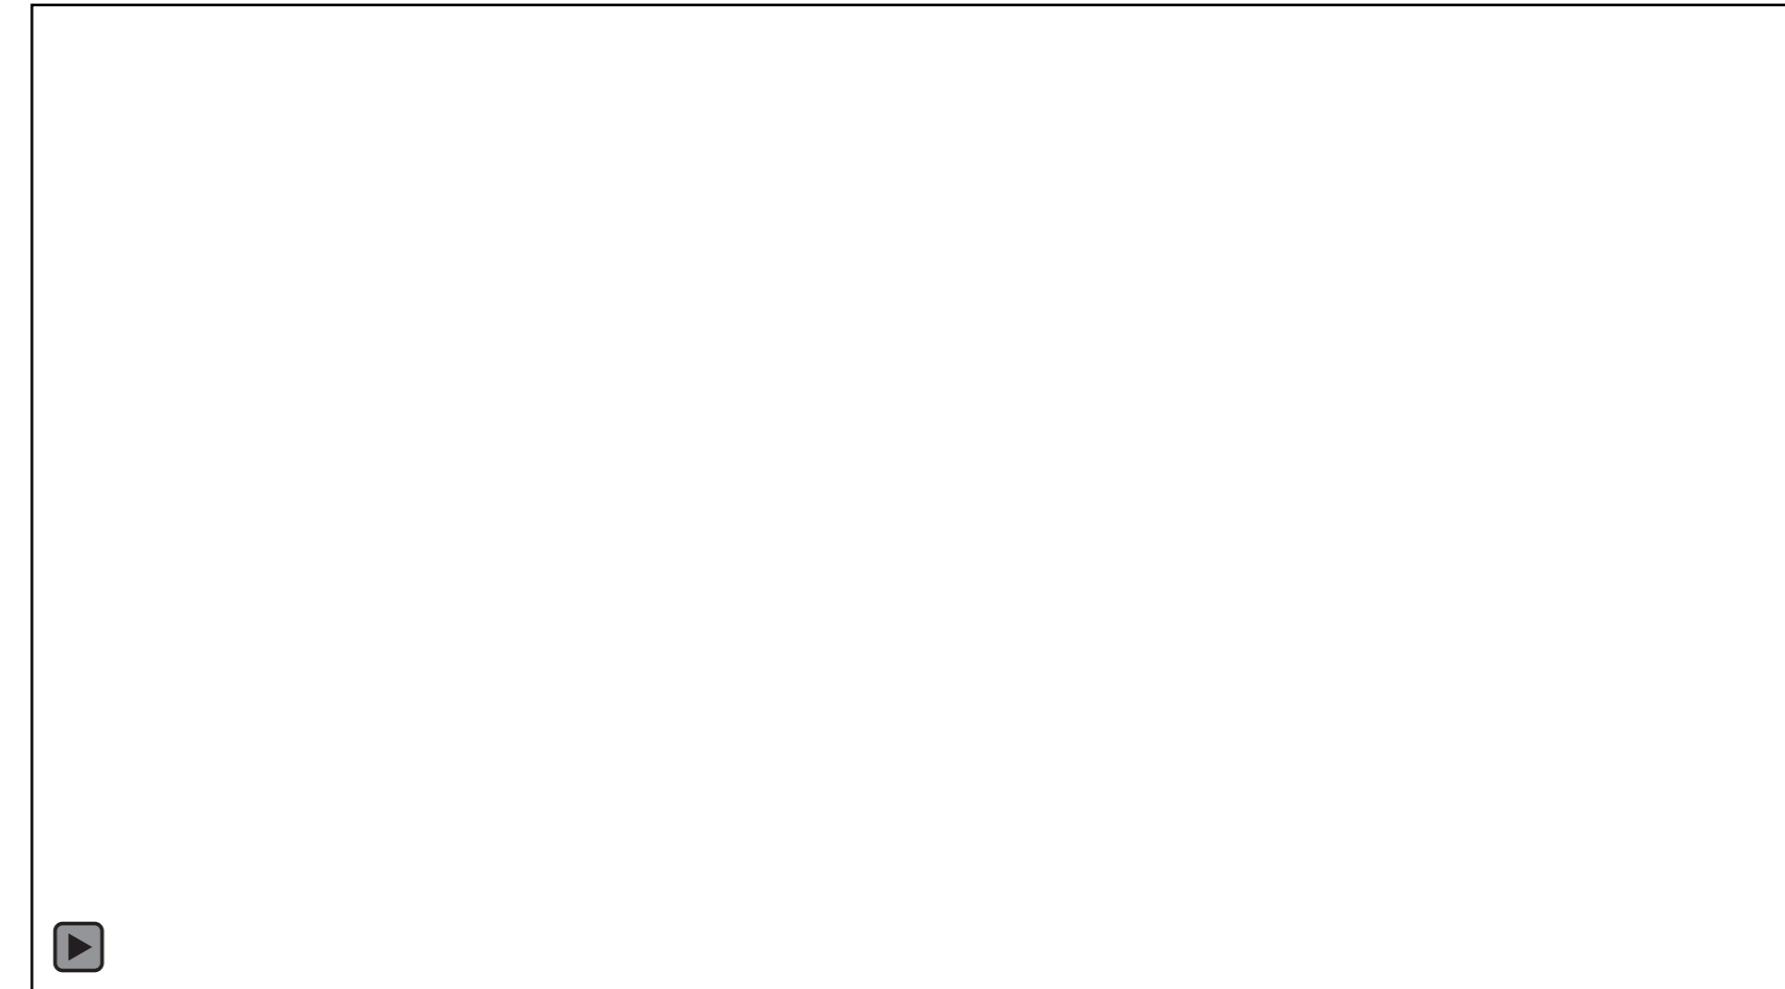
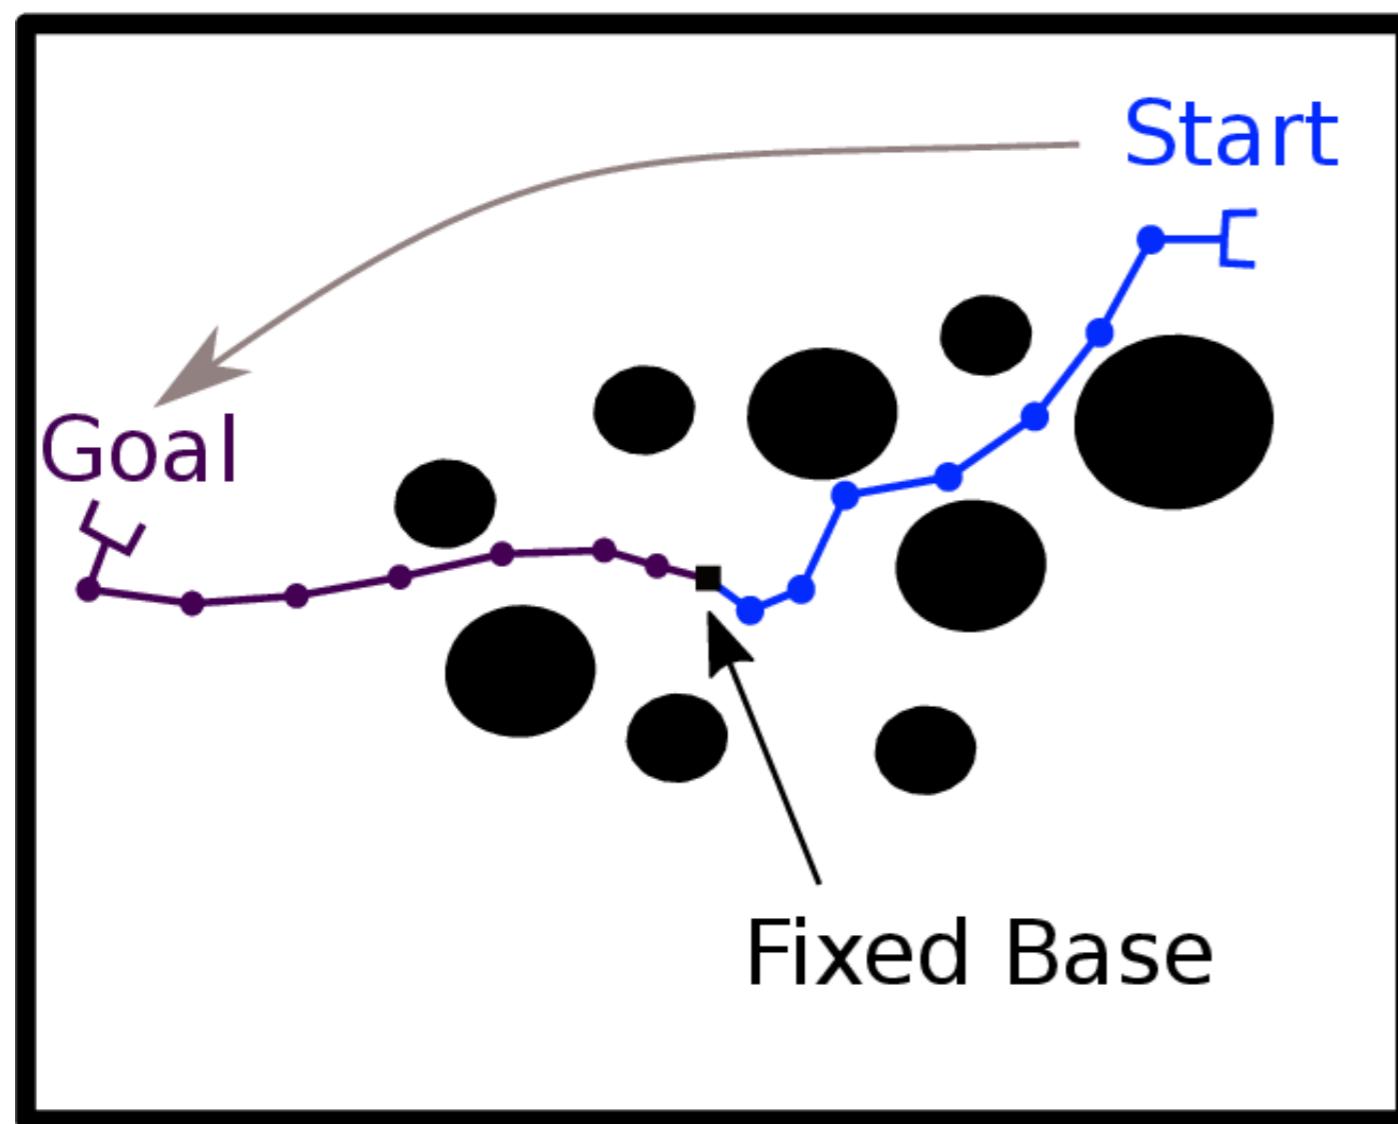
Learning Model

The diagram shows a learning model consisting of three interlocking black gears. Below the gears, the text $Map: [W] \rightarrow X$ is displayed, indicating the mapping function from the workspace to the output space X .

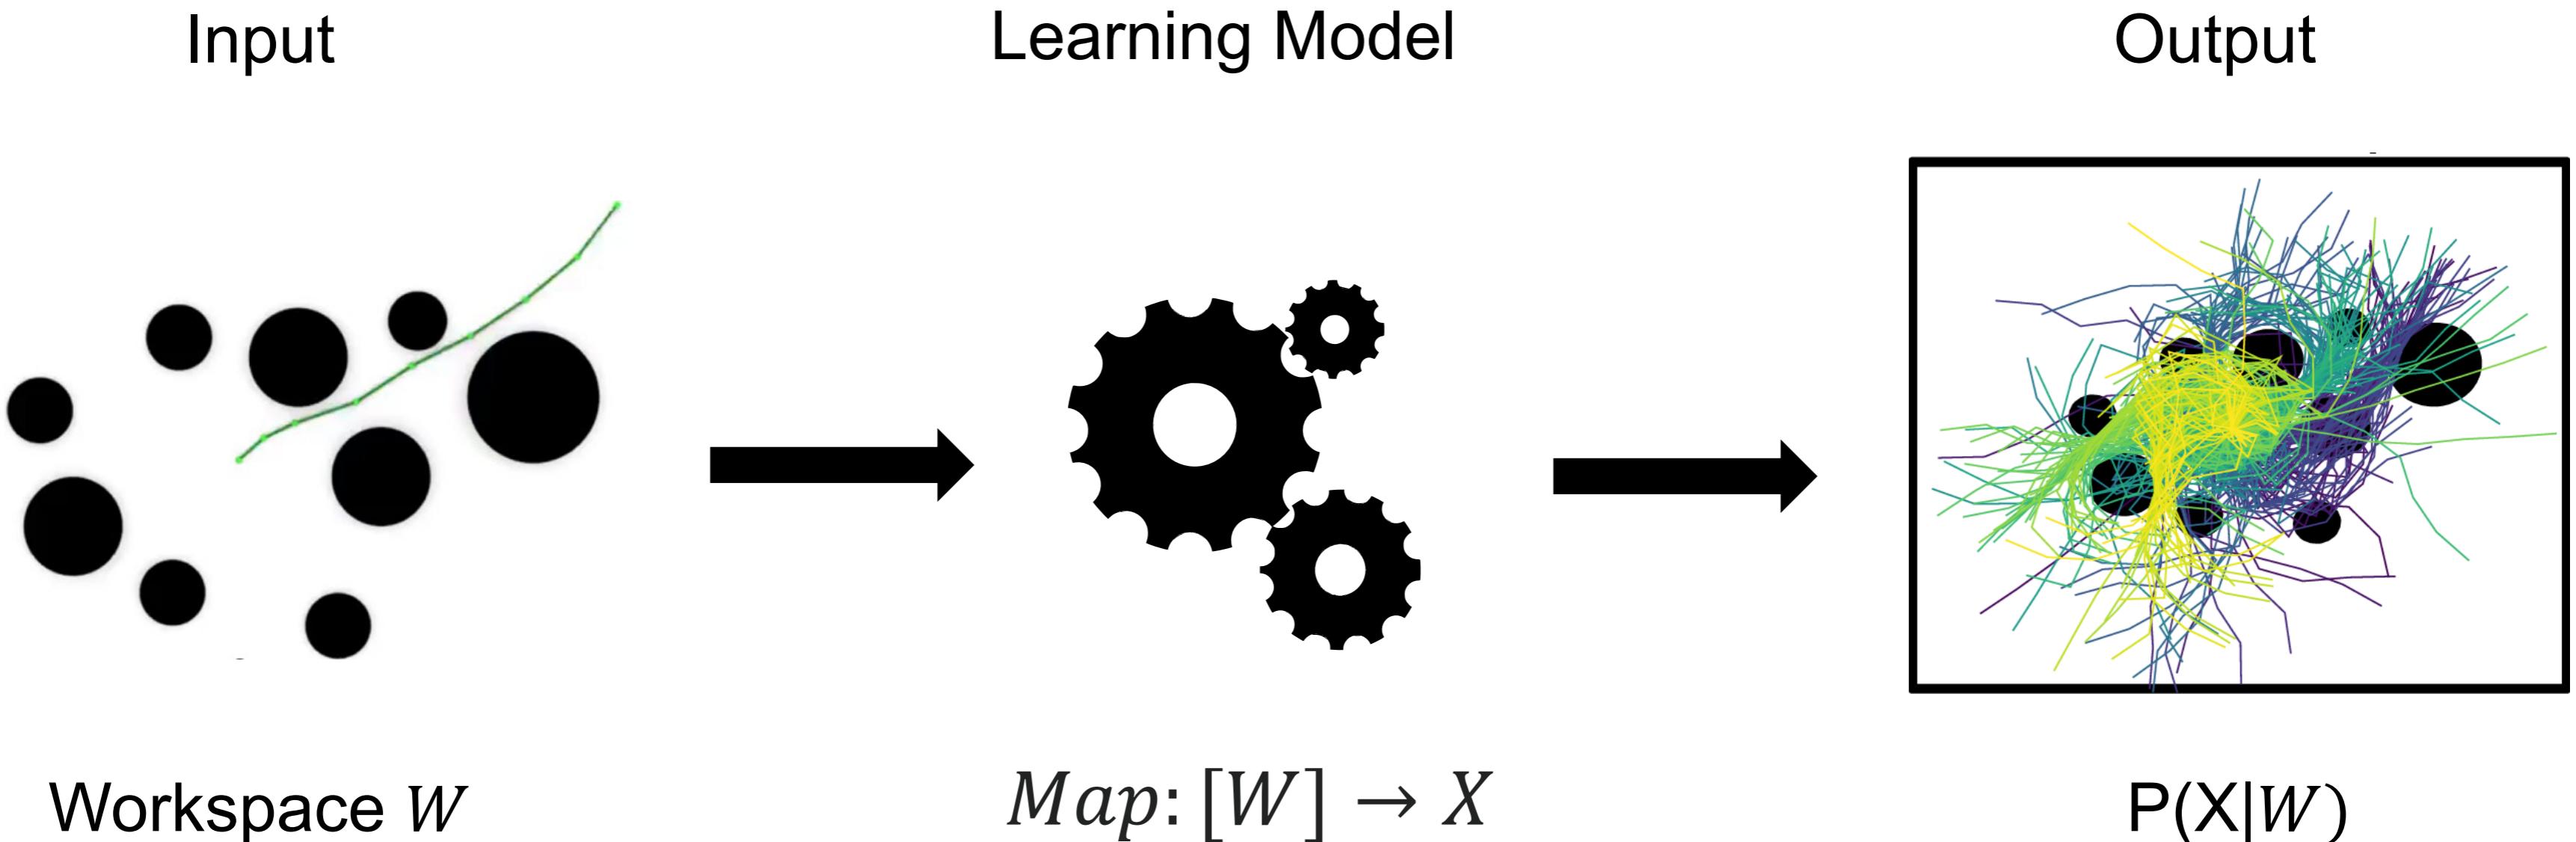
$Map: [W] \rightarrow X$



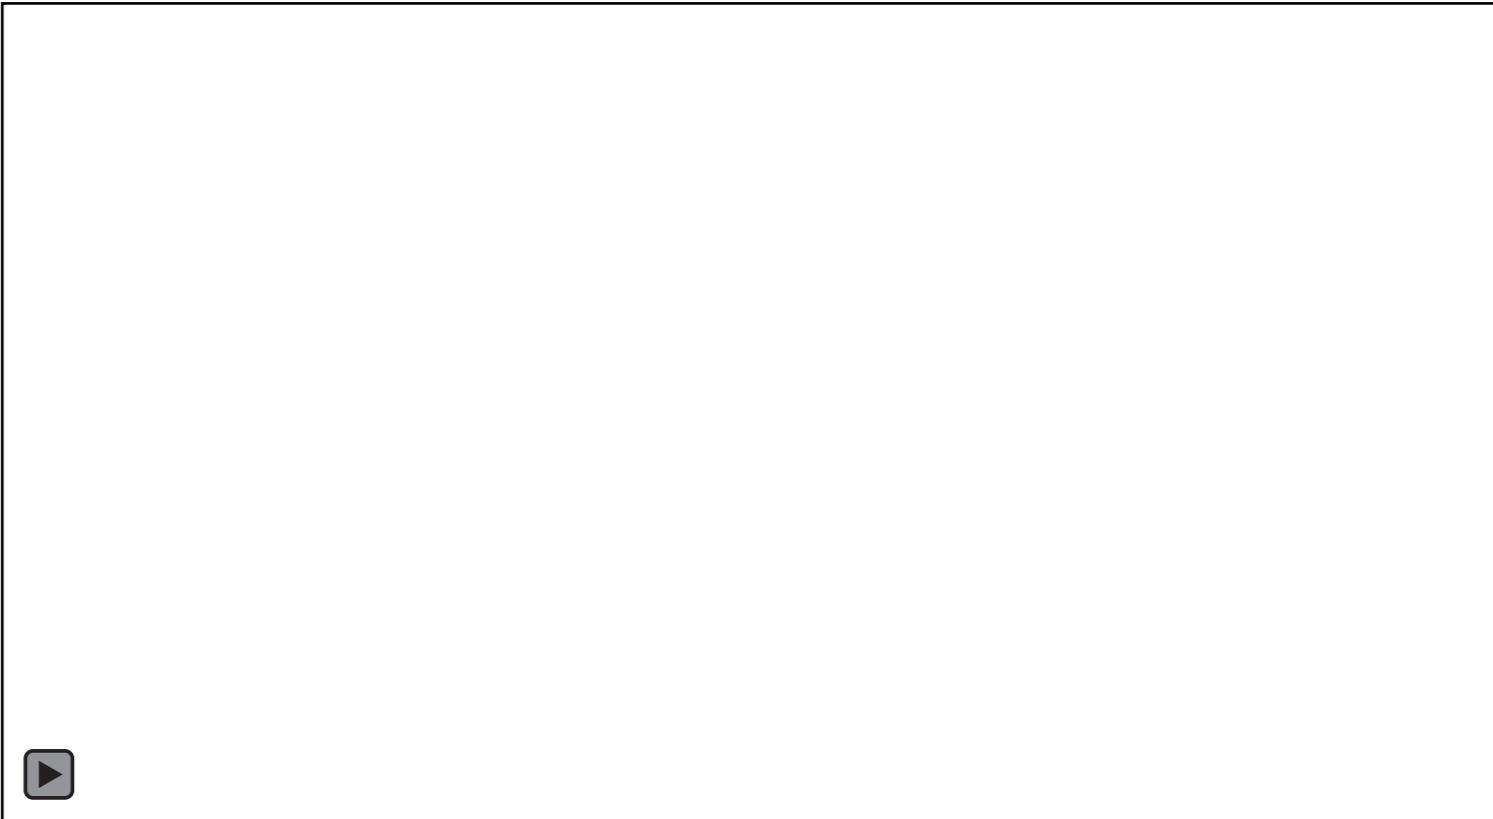
A 2D workspace challenging problem



Problem Formulation



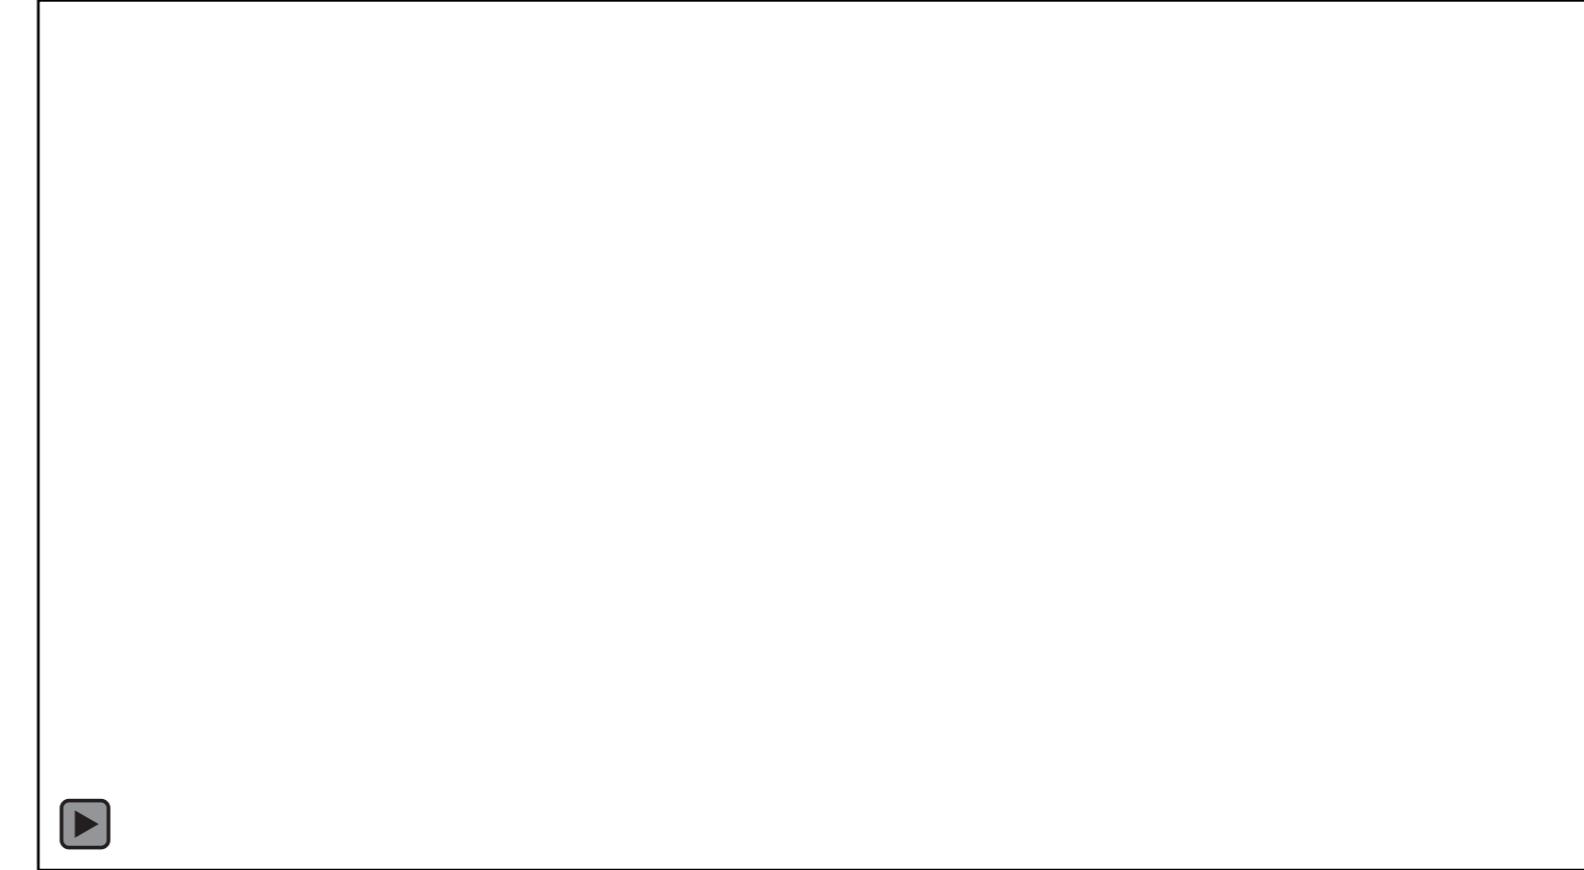
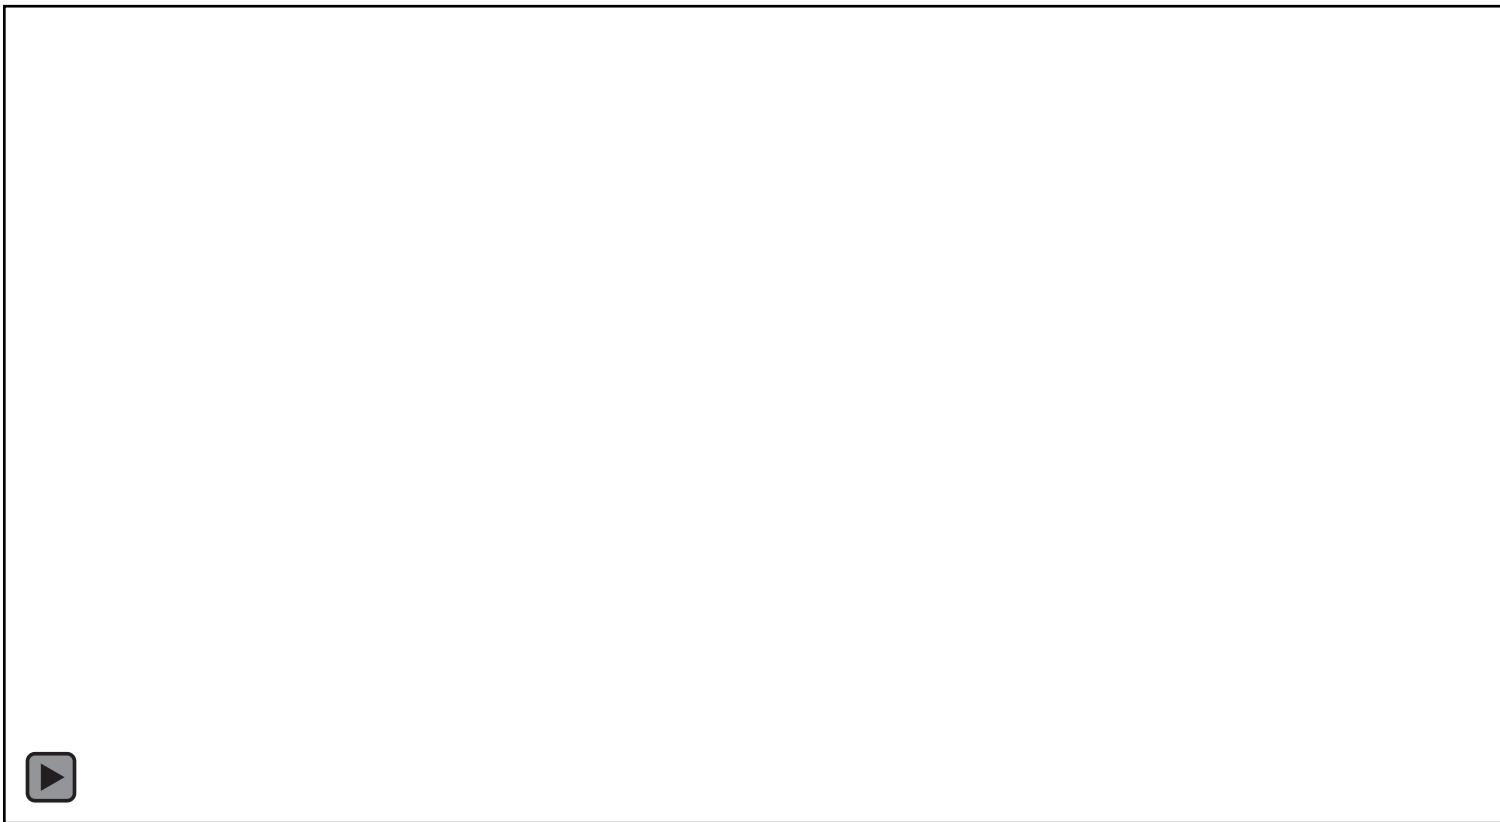
Dealing with Locality (Consideration 1)



A potential solution:

We will only consider pairs of obstacles (local workspaces)

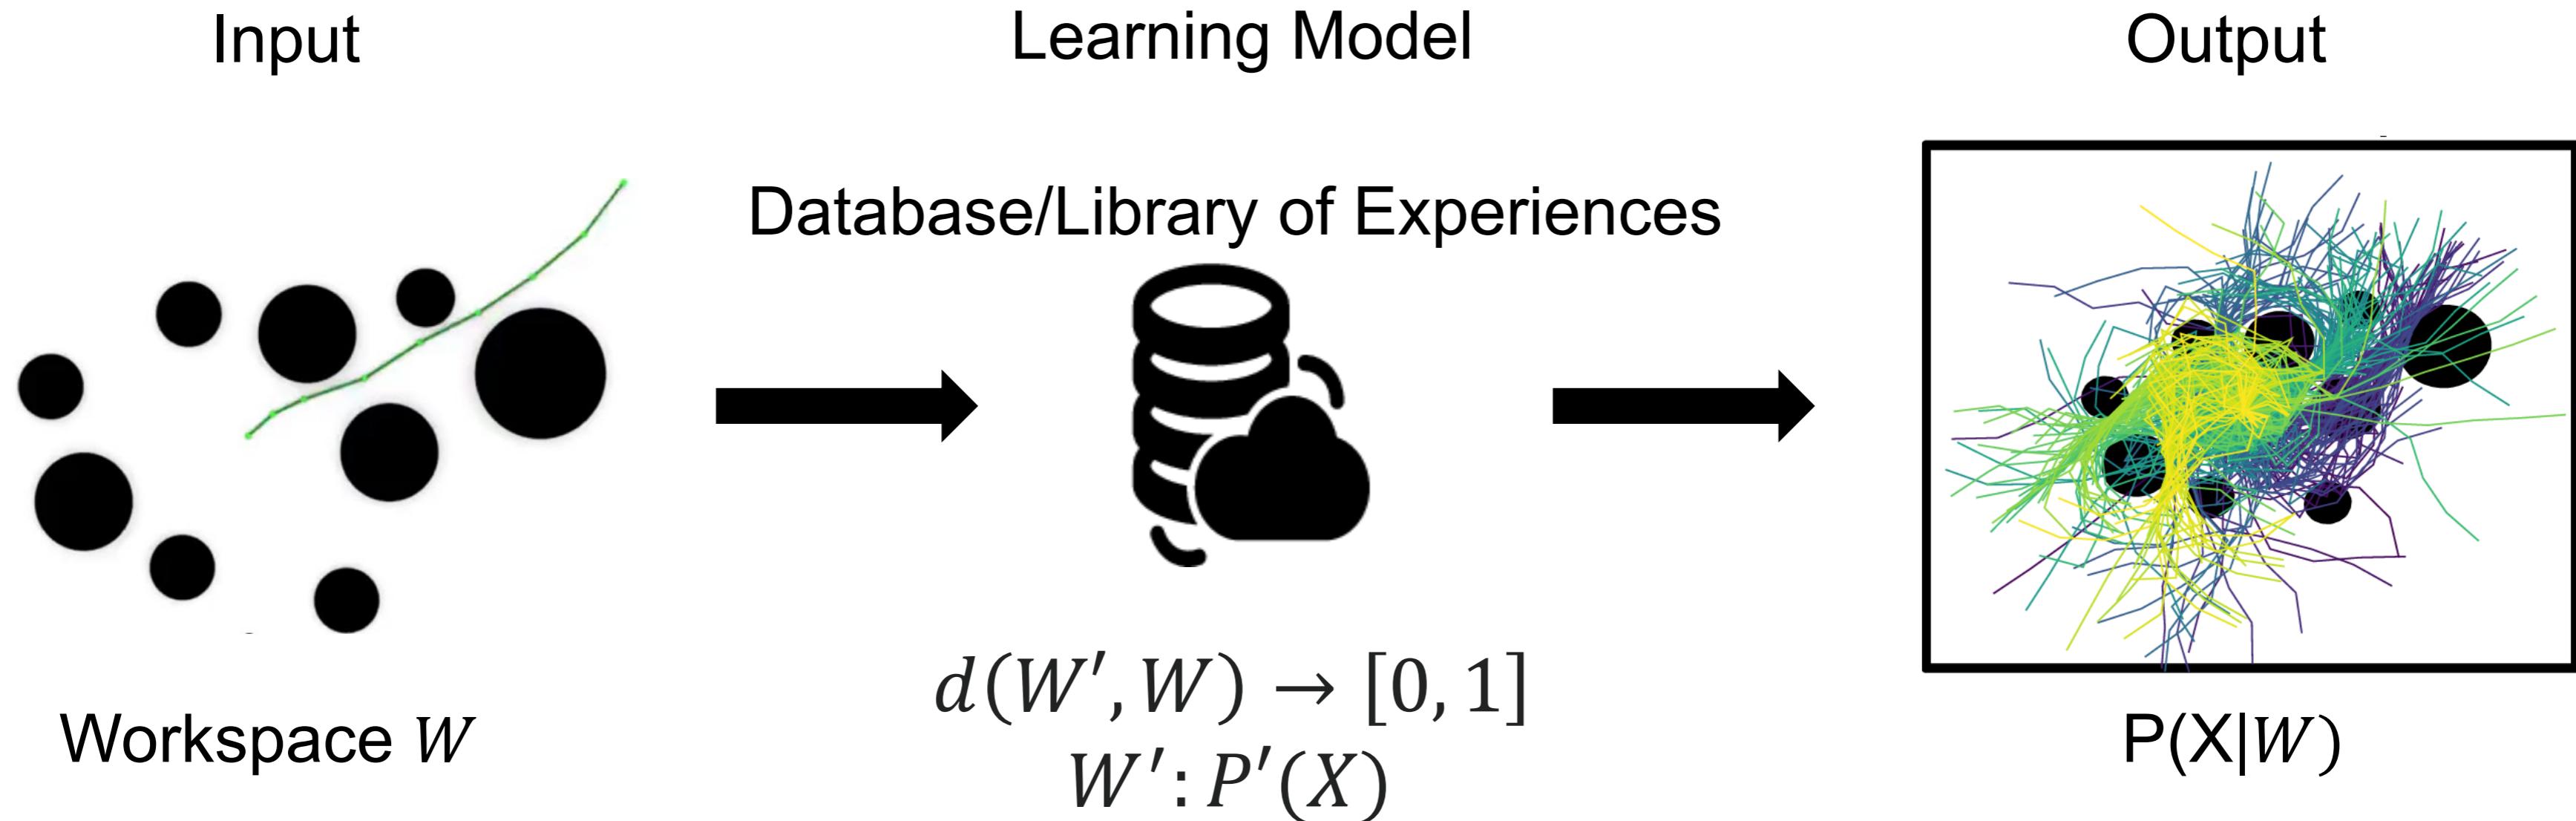
Dealing with Discontinuity (Consideration 2)



Problem:

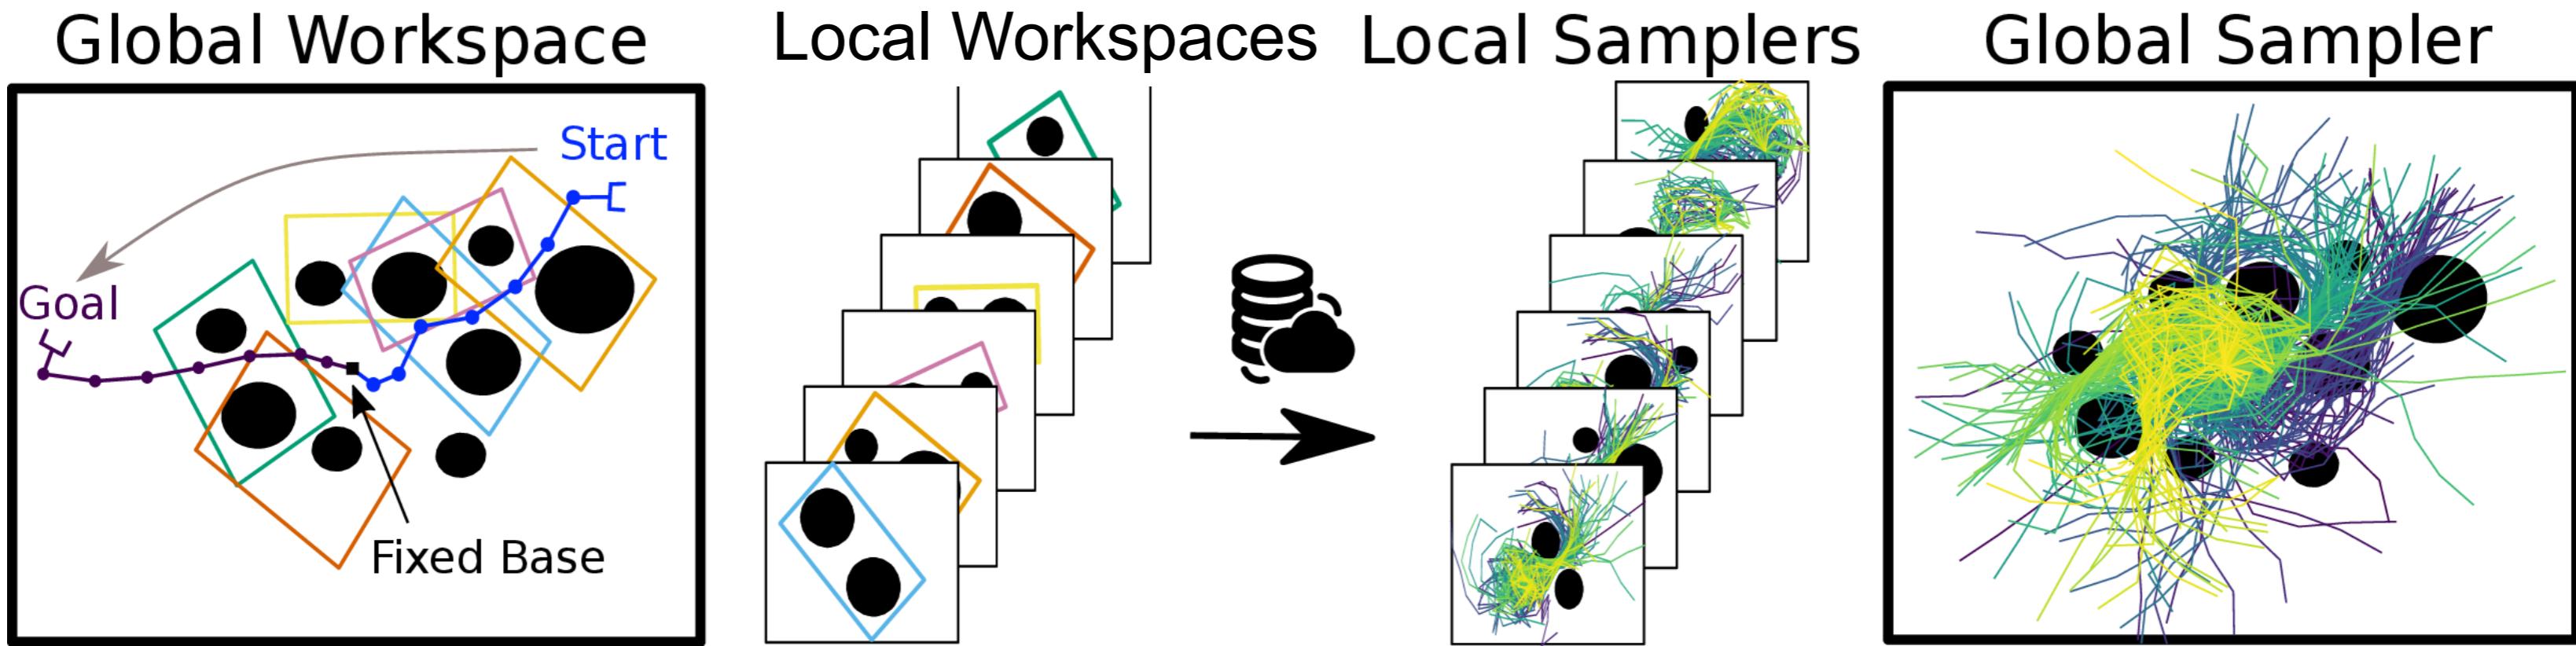
Small changes in the workspace can affect C-Space solutions a lot

Dealing with Discontinuity (Consideration 2)



A potential solution: retrieval methods are discontinuous/multimodal by nature

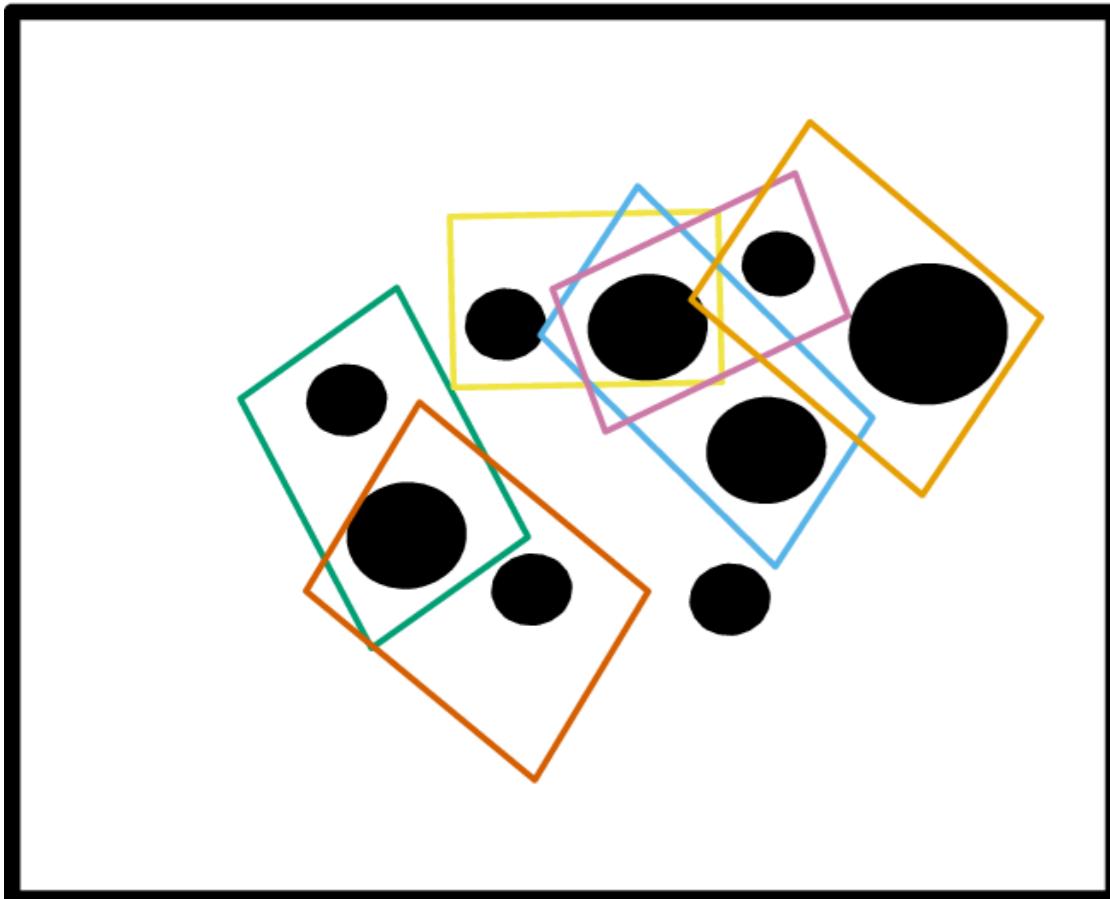
SPARK2D Framework



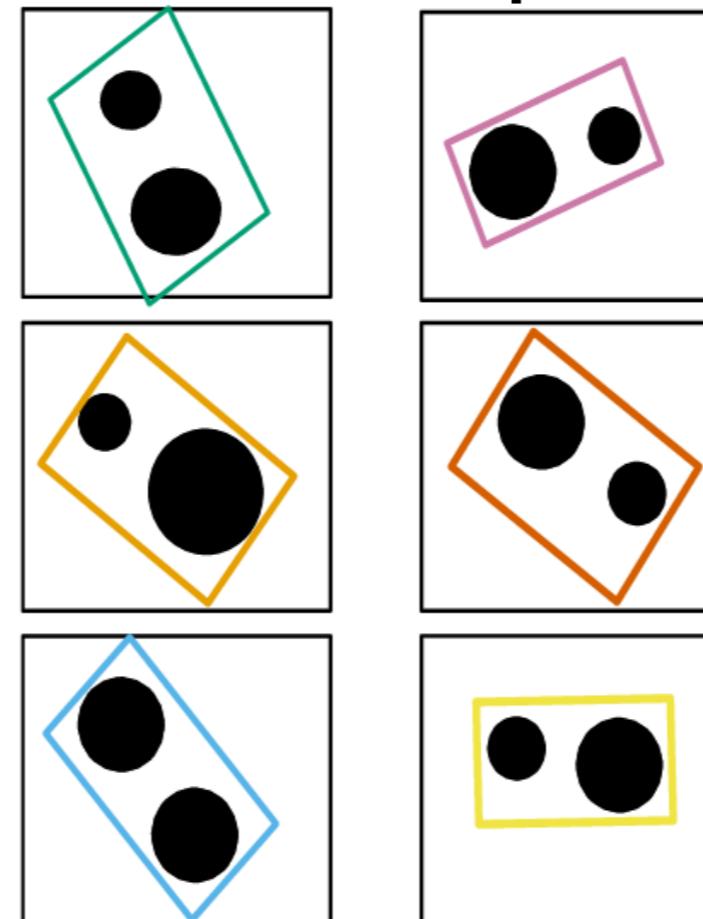
1. Decompose to local primitives
2. Creating a local sampler for every local workspaces [OFFLINE]
3. Retrieve local samplers given the local workspaces [ONLINE]
4. Synthesize a global sampler from the local samplers

1) Decomposition to local primitives

Global Workspace



Local Workspaces



We represent each local primitive with 6 numbers

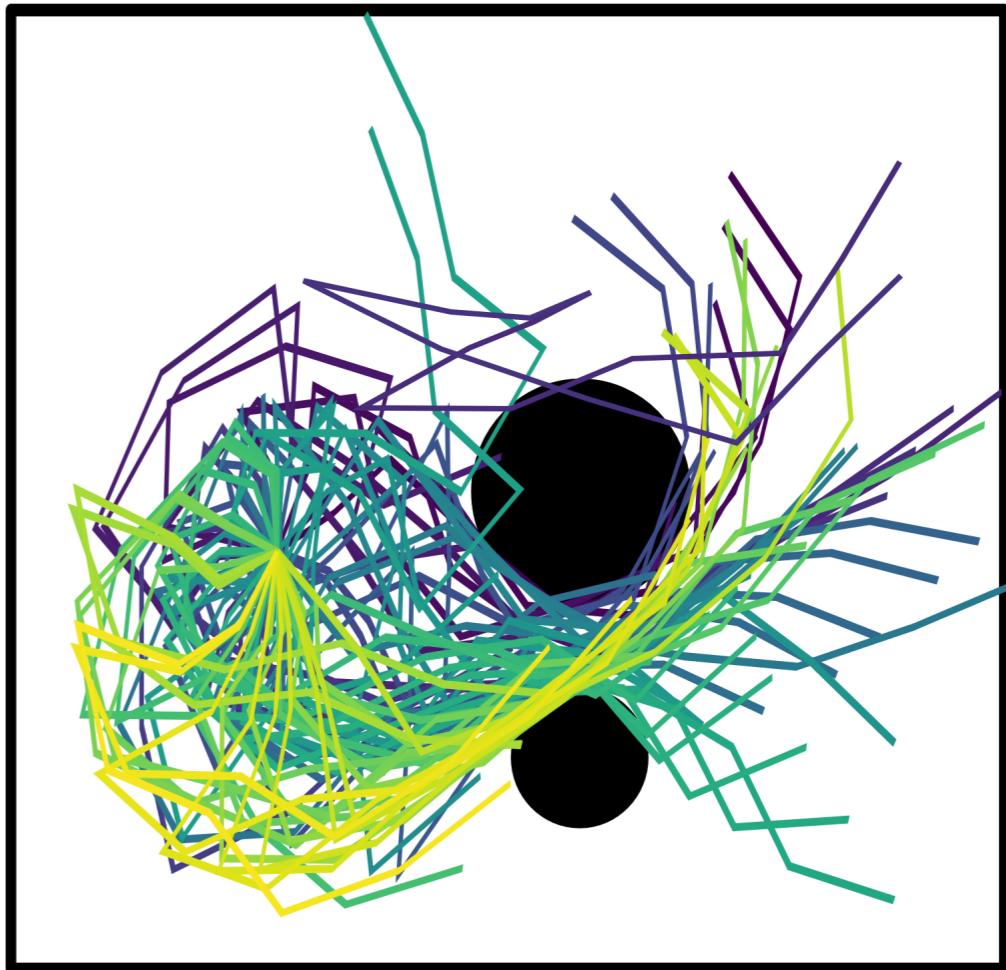
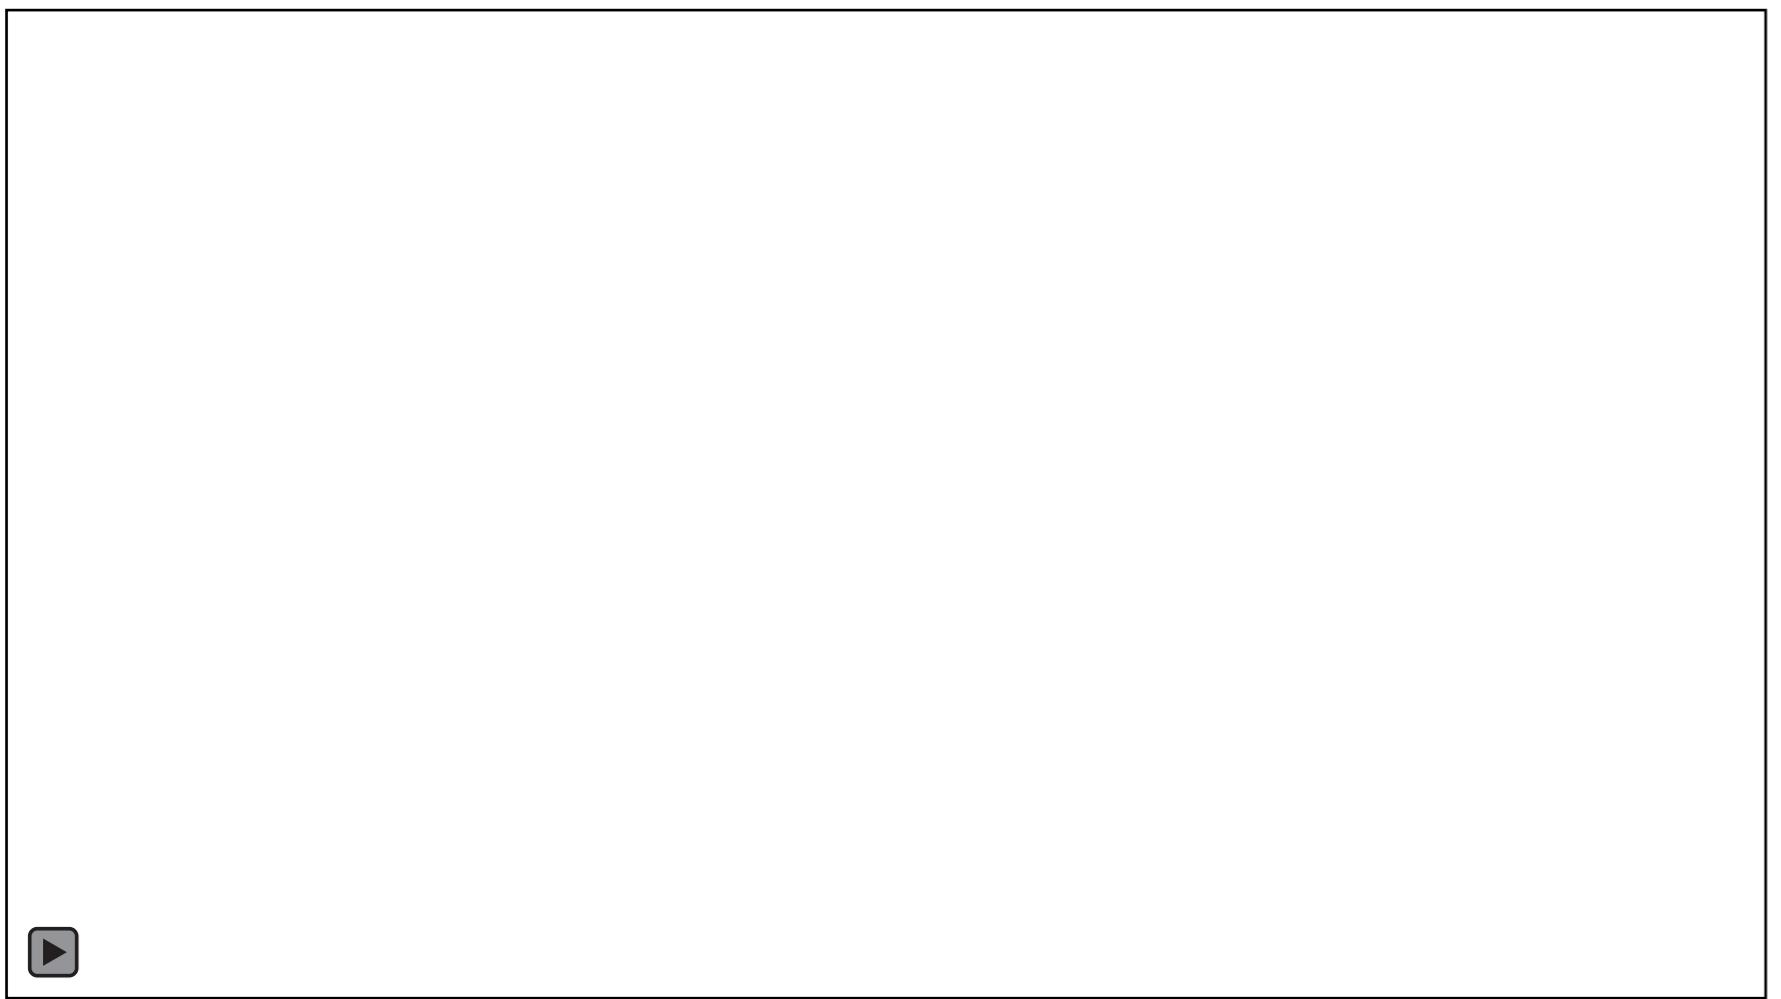
$$lw = \{x_a, y_a, r_a, x_b, y_b, r_b\}$$

Local primitives space is
6 dimensional

2) Creating Local Samplers

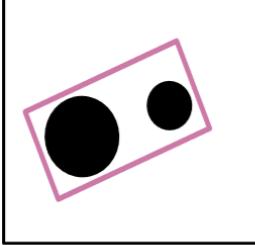
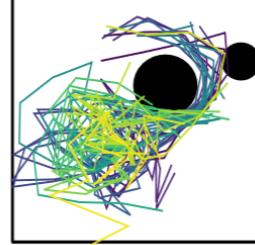
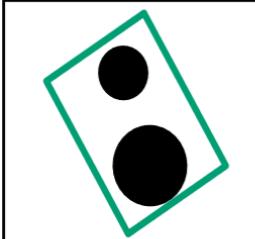
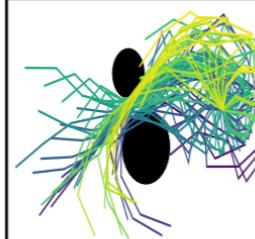
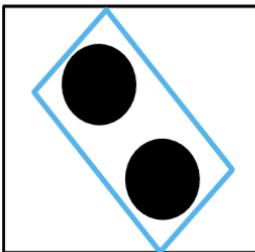
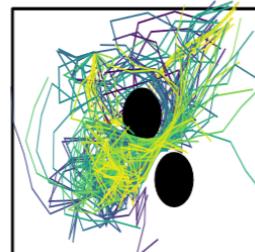
Create a representative set of samples
for each local workspace

Fit a Gaussian Mixture Model using the
samples of the solution



$$p(x|lw) = GMM(\mu, \Sigma) = \frac{1}{M} \sum_{j=1}^M N(q_j, \Sigma_j)$$

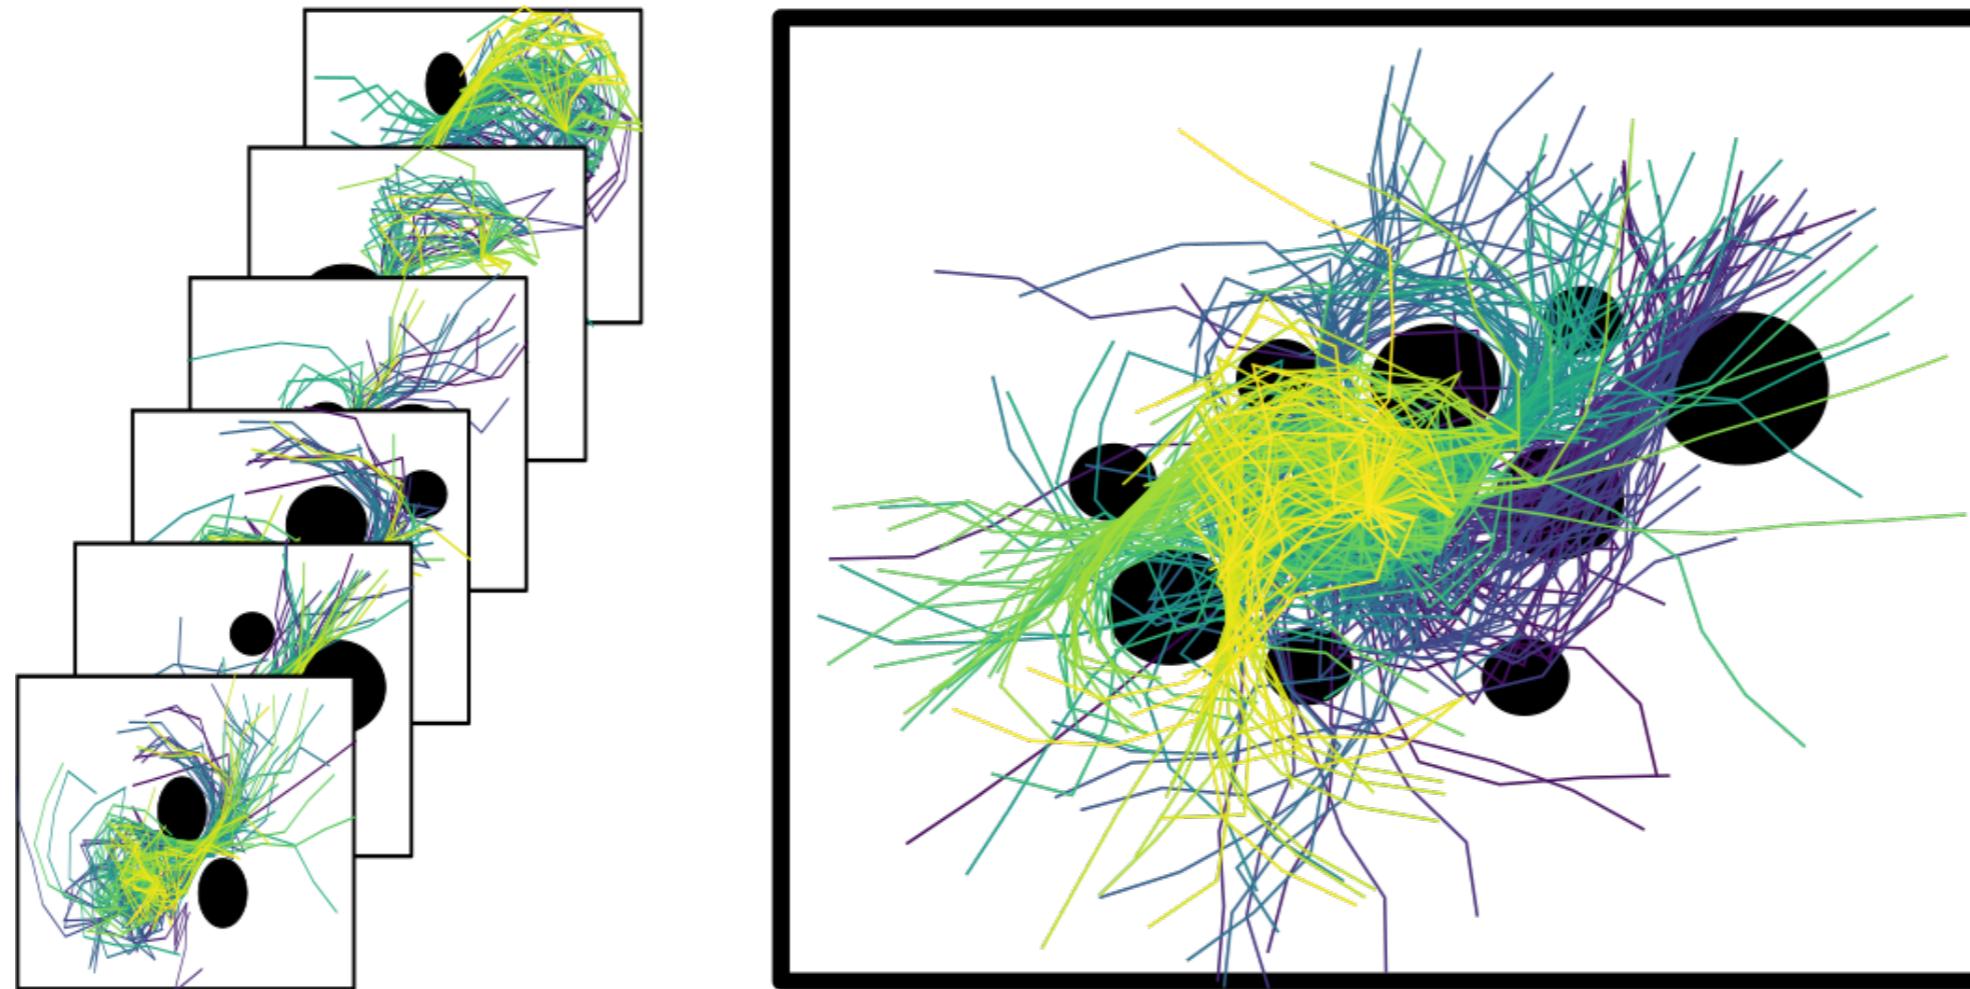
3) Retrieve a Local sampler given local primitive

		
		Keys Values
lw_1		 $p_1(x lw_1)$
lw_2		 $p_2(x lw_2)$
	⋮	⋮
lw_n		 $p_n(x lw_n)$

We define a simple similarity function

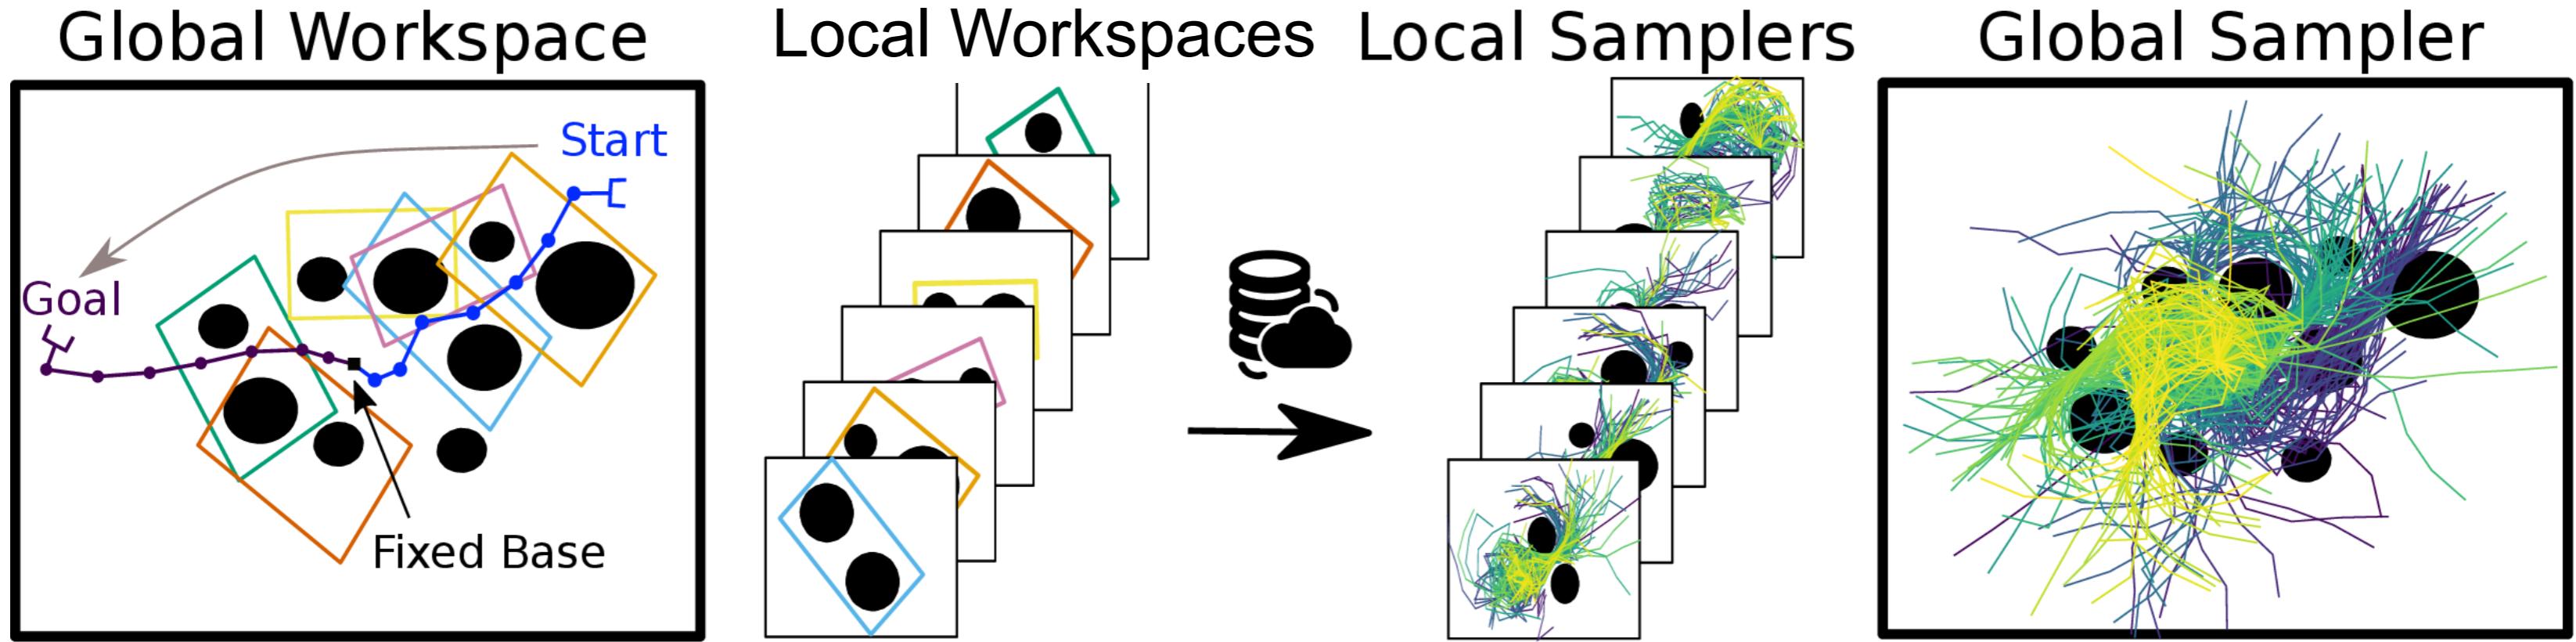
$$k(lw_i, lw_j) = (lw_i - lw_j)^T (lw_i - lw_j)$$

4) Synthesize the local samplers to a global sampler



We simply create a bigger GMM

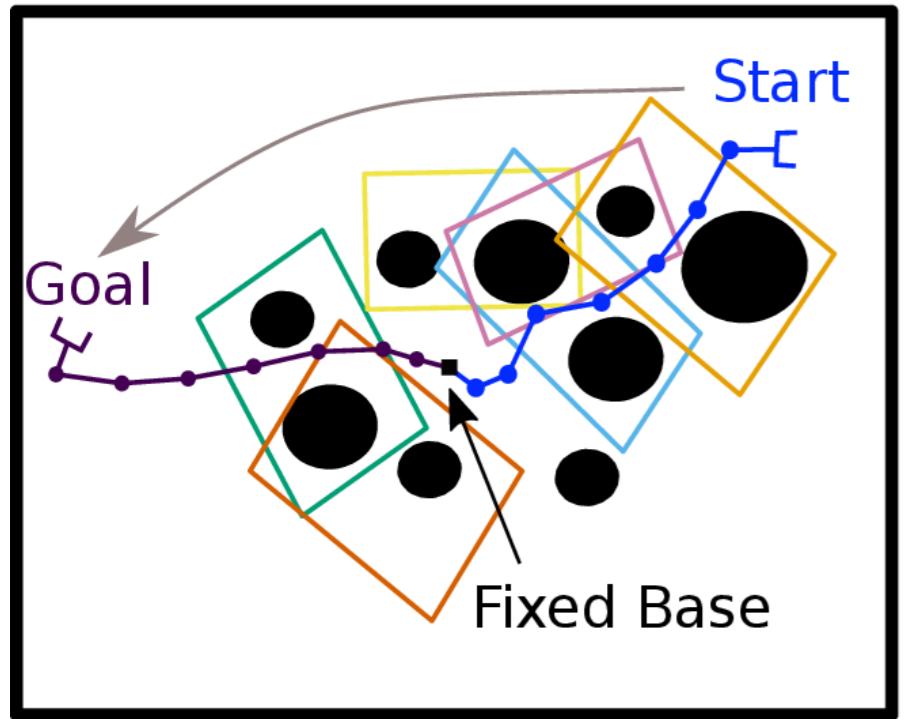
SPARK2D Framework



1. Decompose to local primitives
2. Creating a local sampler for every local workspaces [OFFLINE]
3. Retrieve local samplers given the local workspaces [ONLINE]
4. Synthesize a global sampler from the local samplers

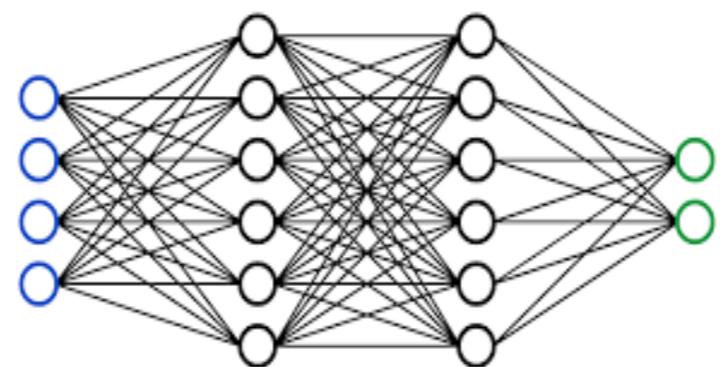
A neural network baseline

Global Workspace

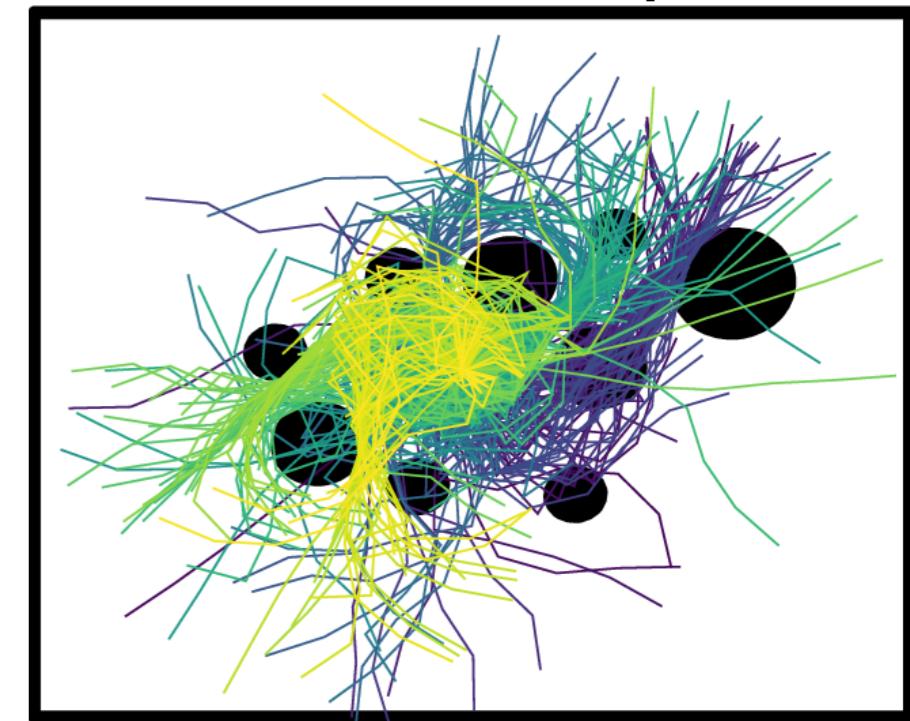


Learning Model

Neural Network (CVAE)



Output



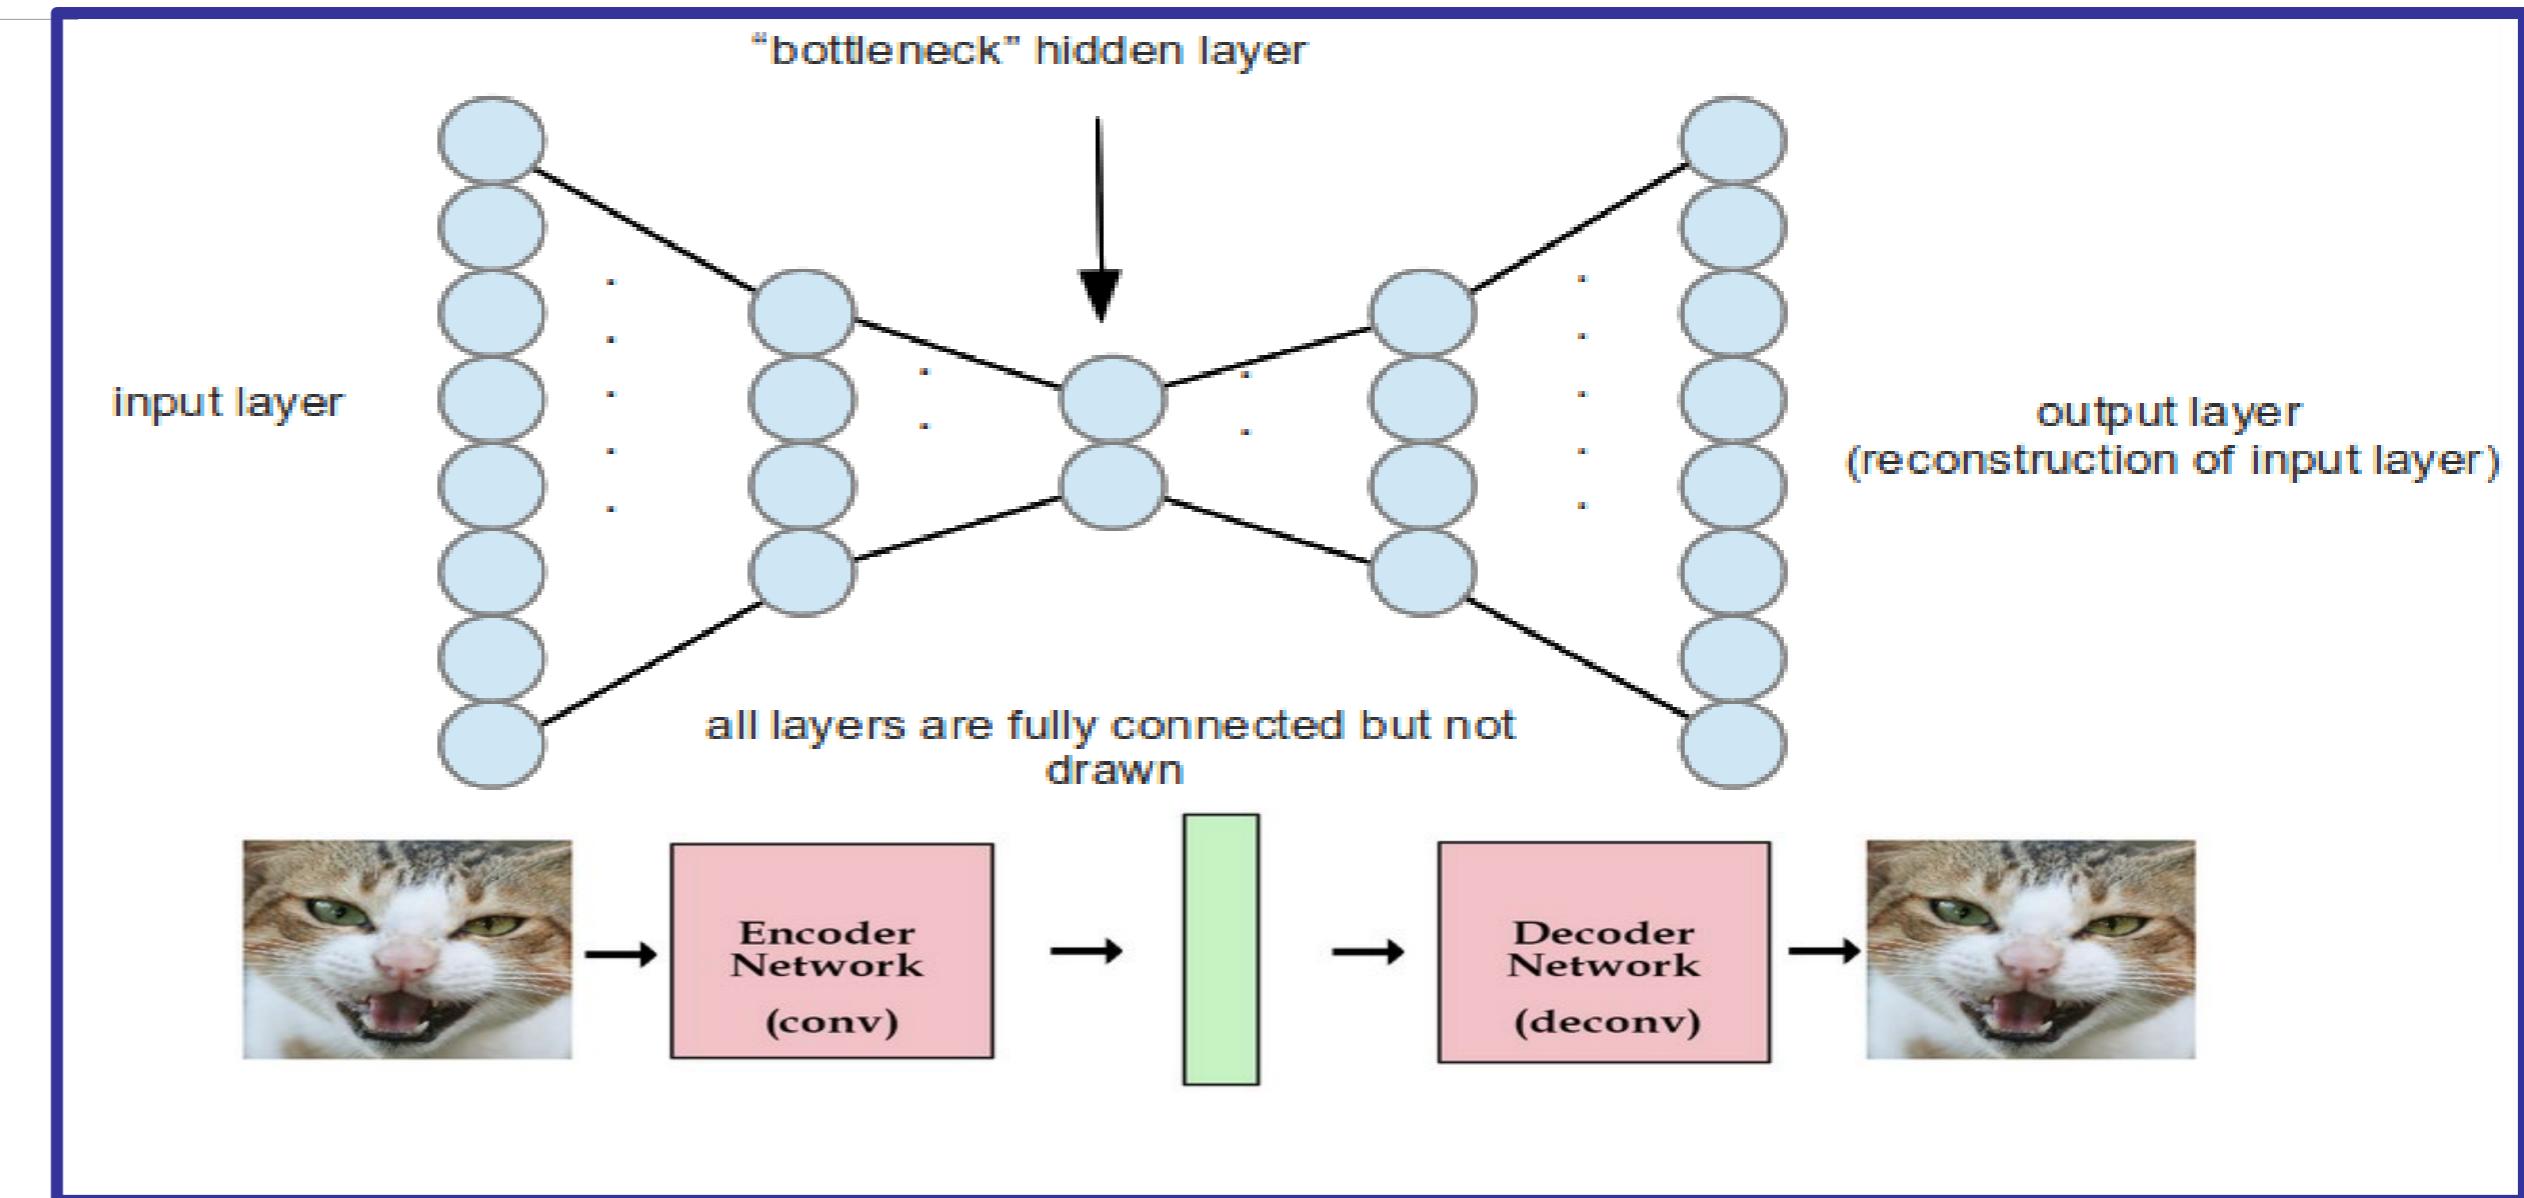
Workspace W

$P(X|W)$

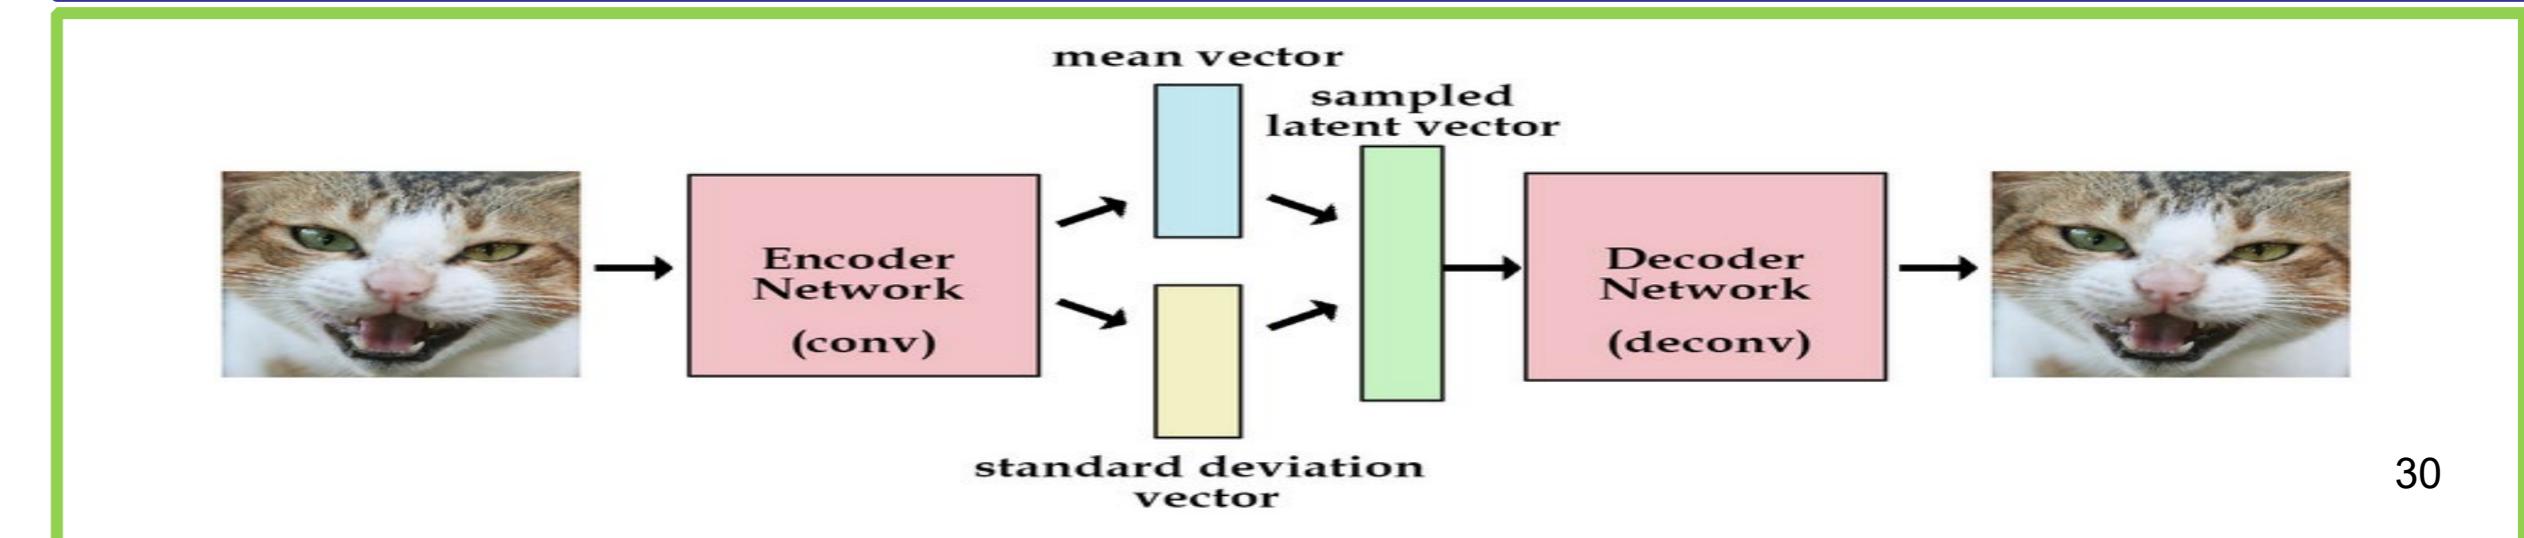
Variational Autoencoders

How to learn a sampling distribution with a neural network?

Autoencoder →



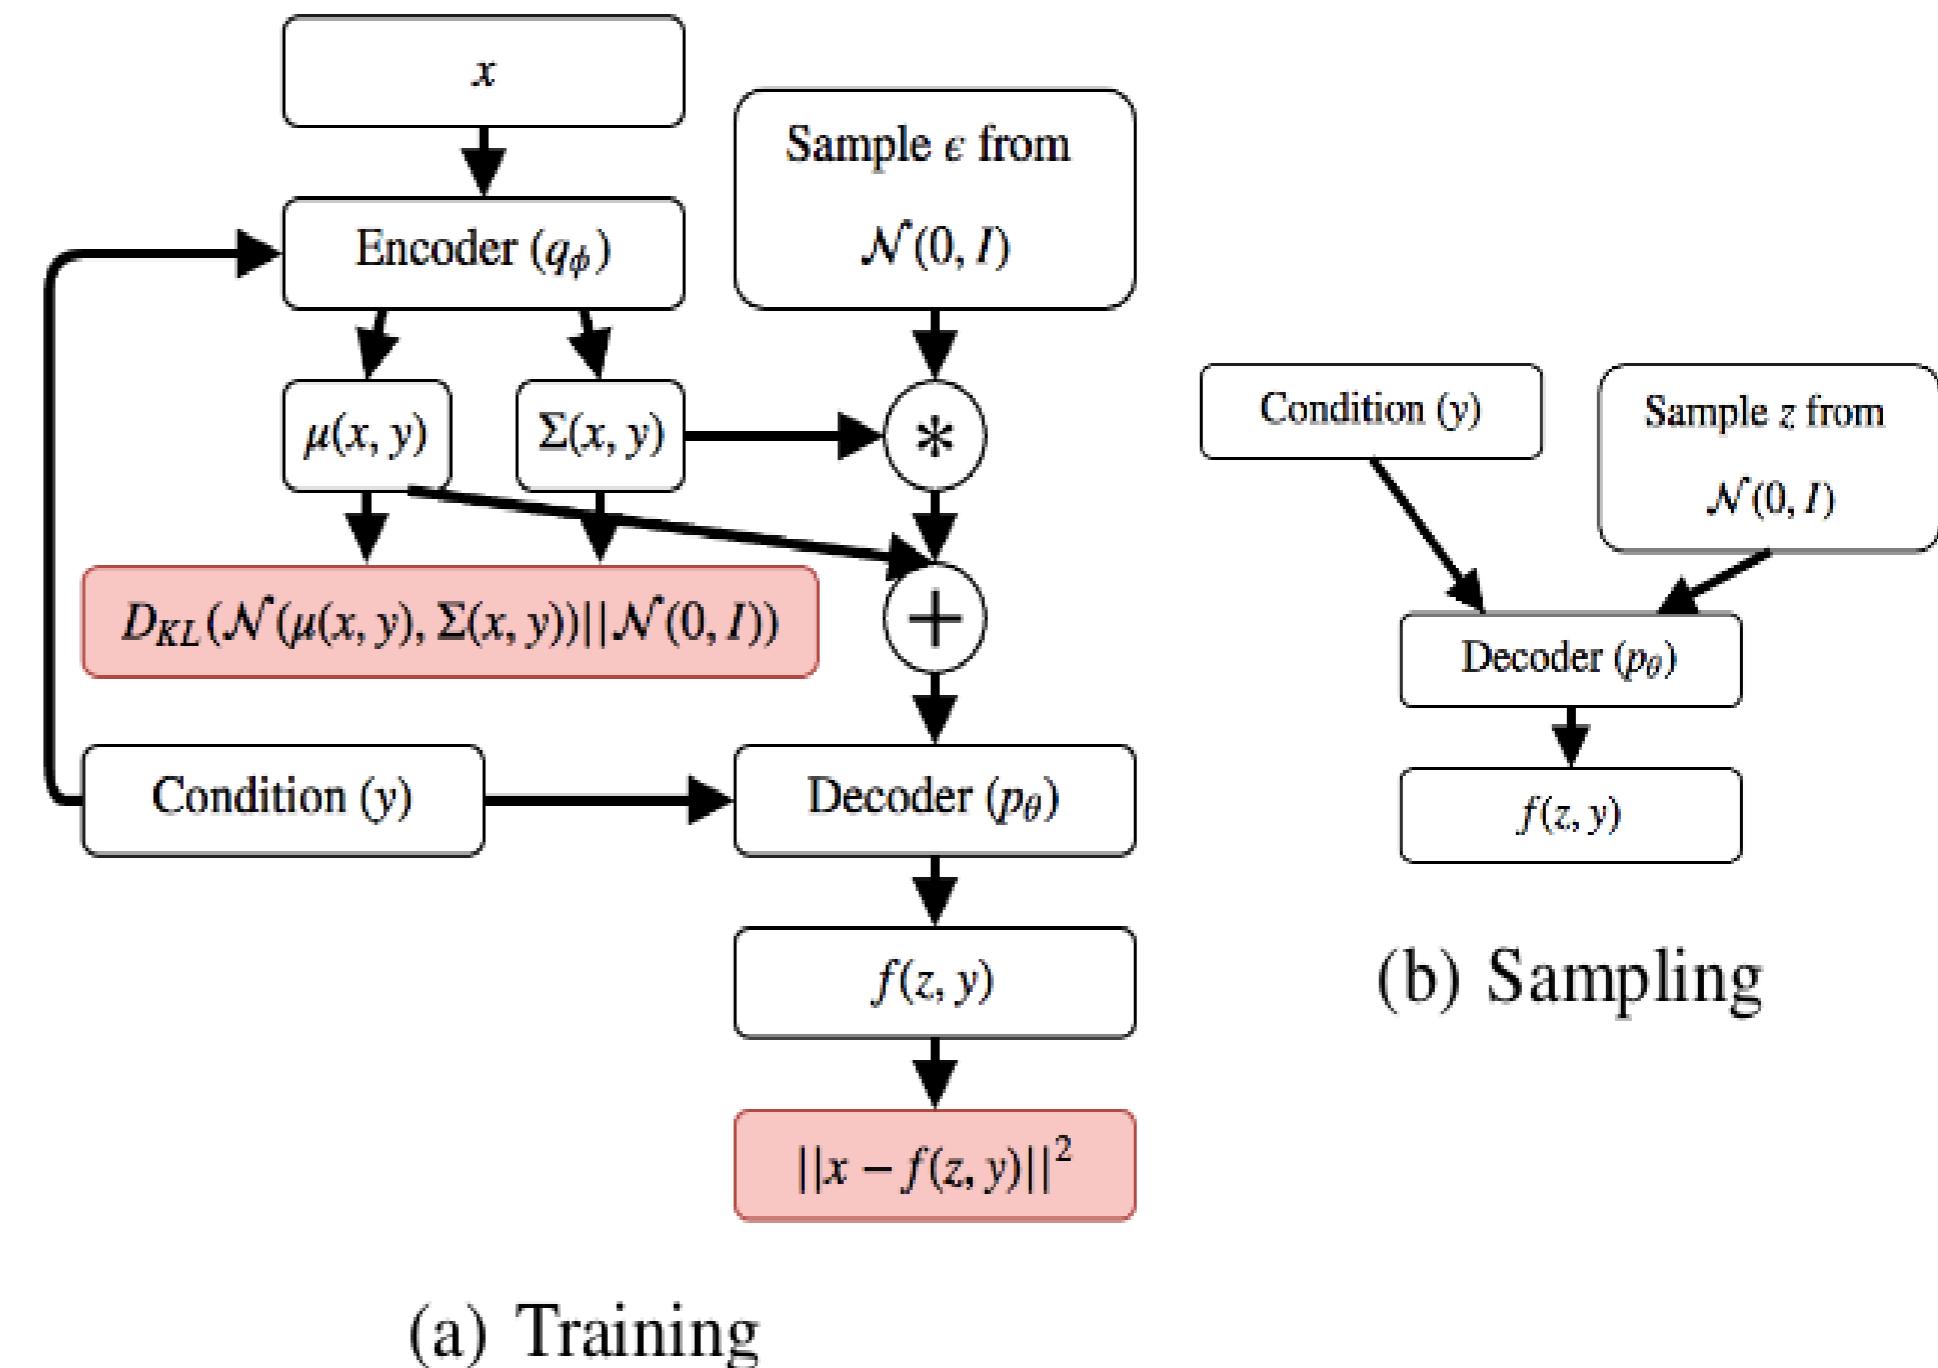
Variational
Autoencoder →



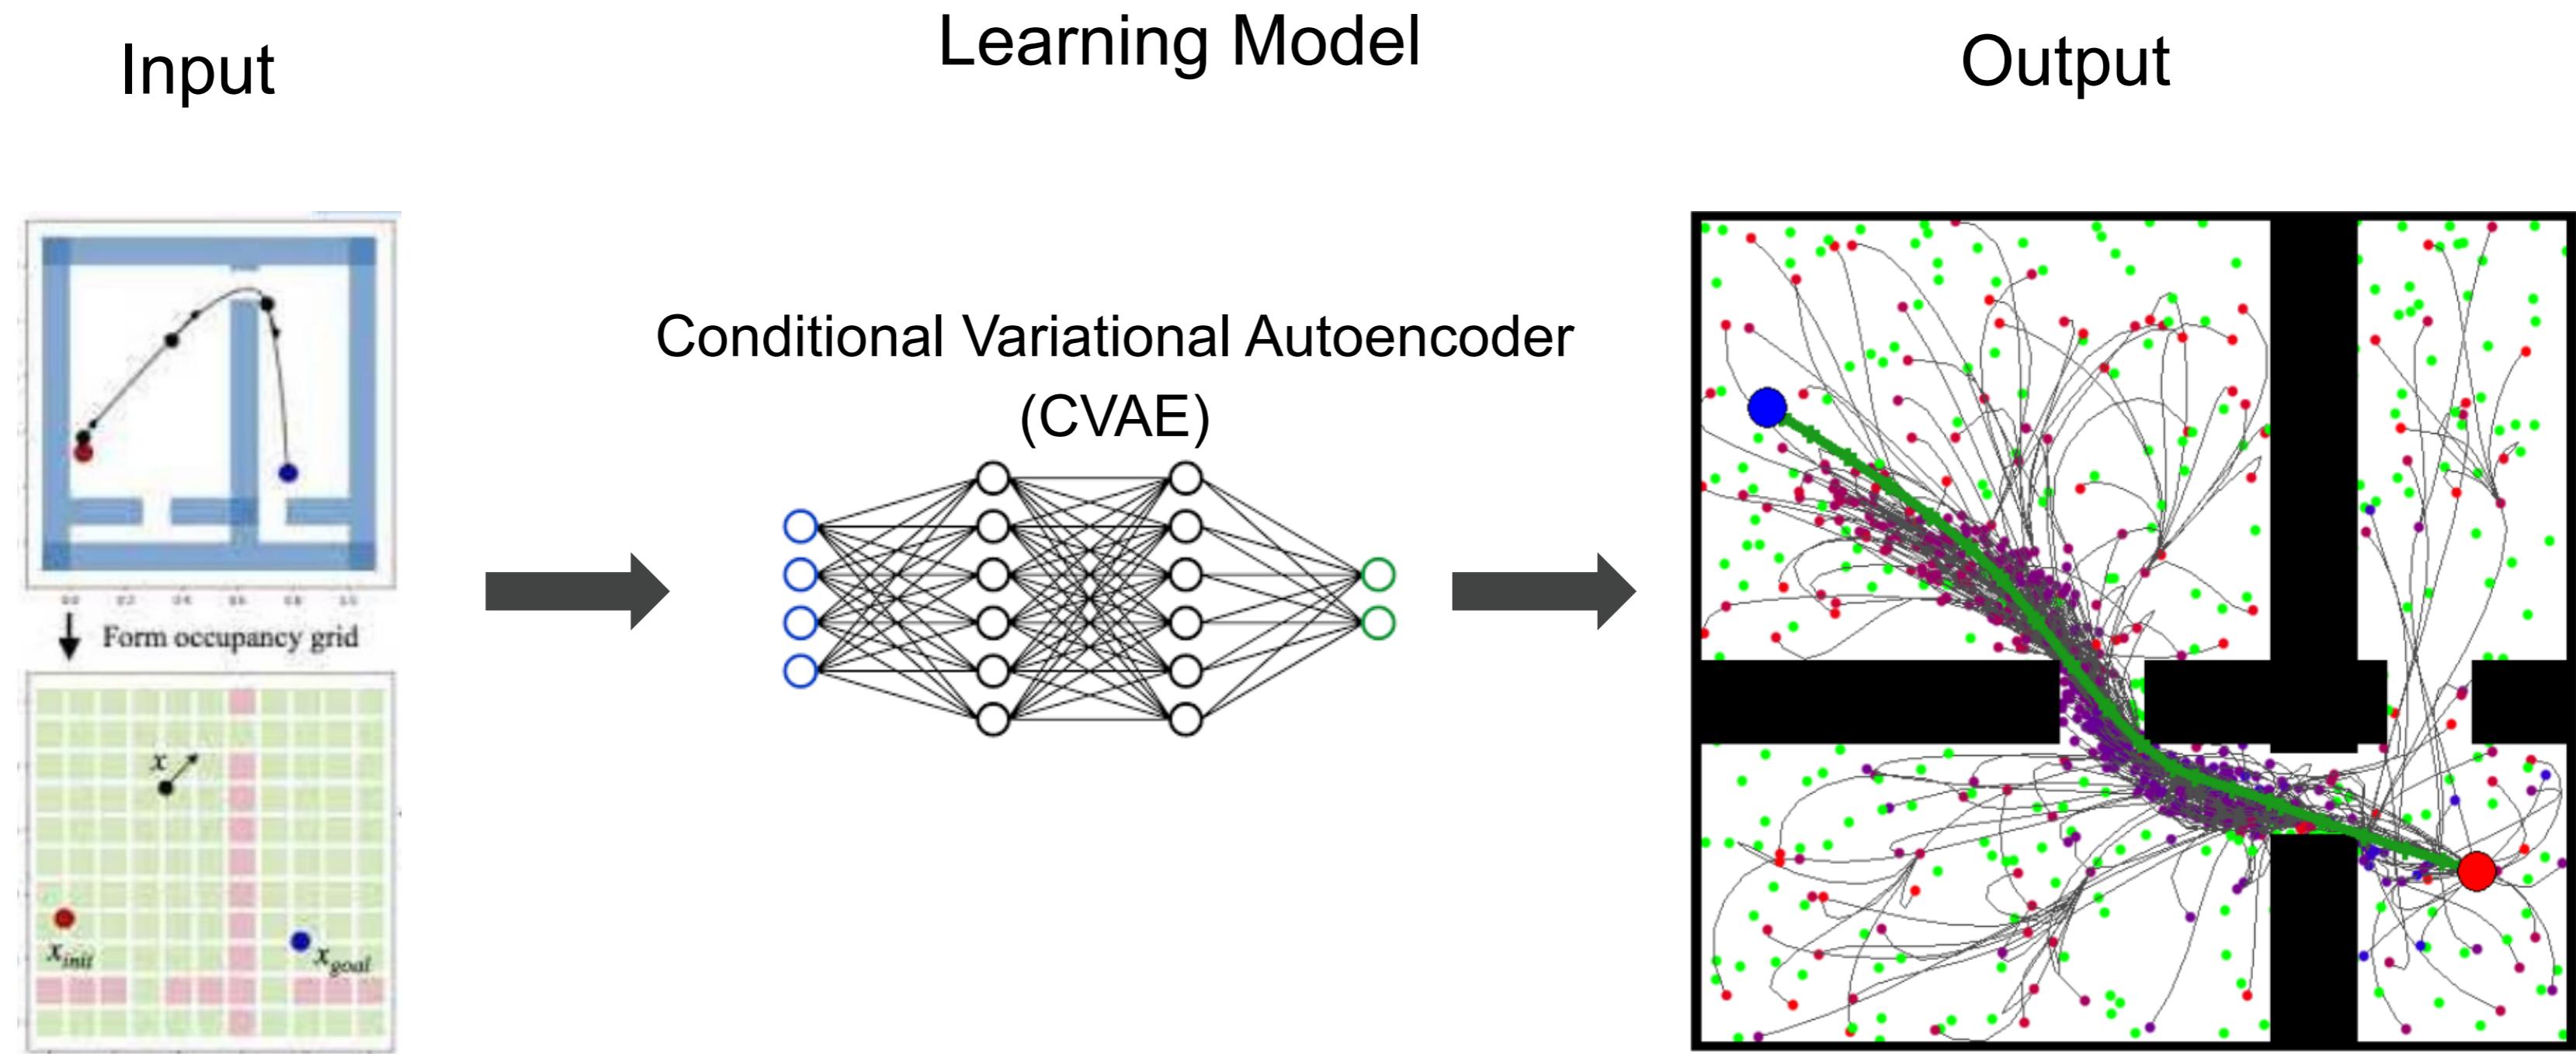
Conditional Variational Autoencoder (CVAE) [14]

Key Idea :

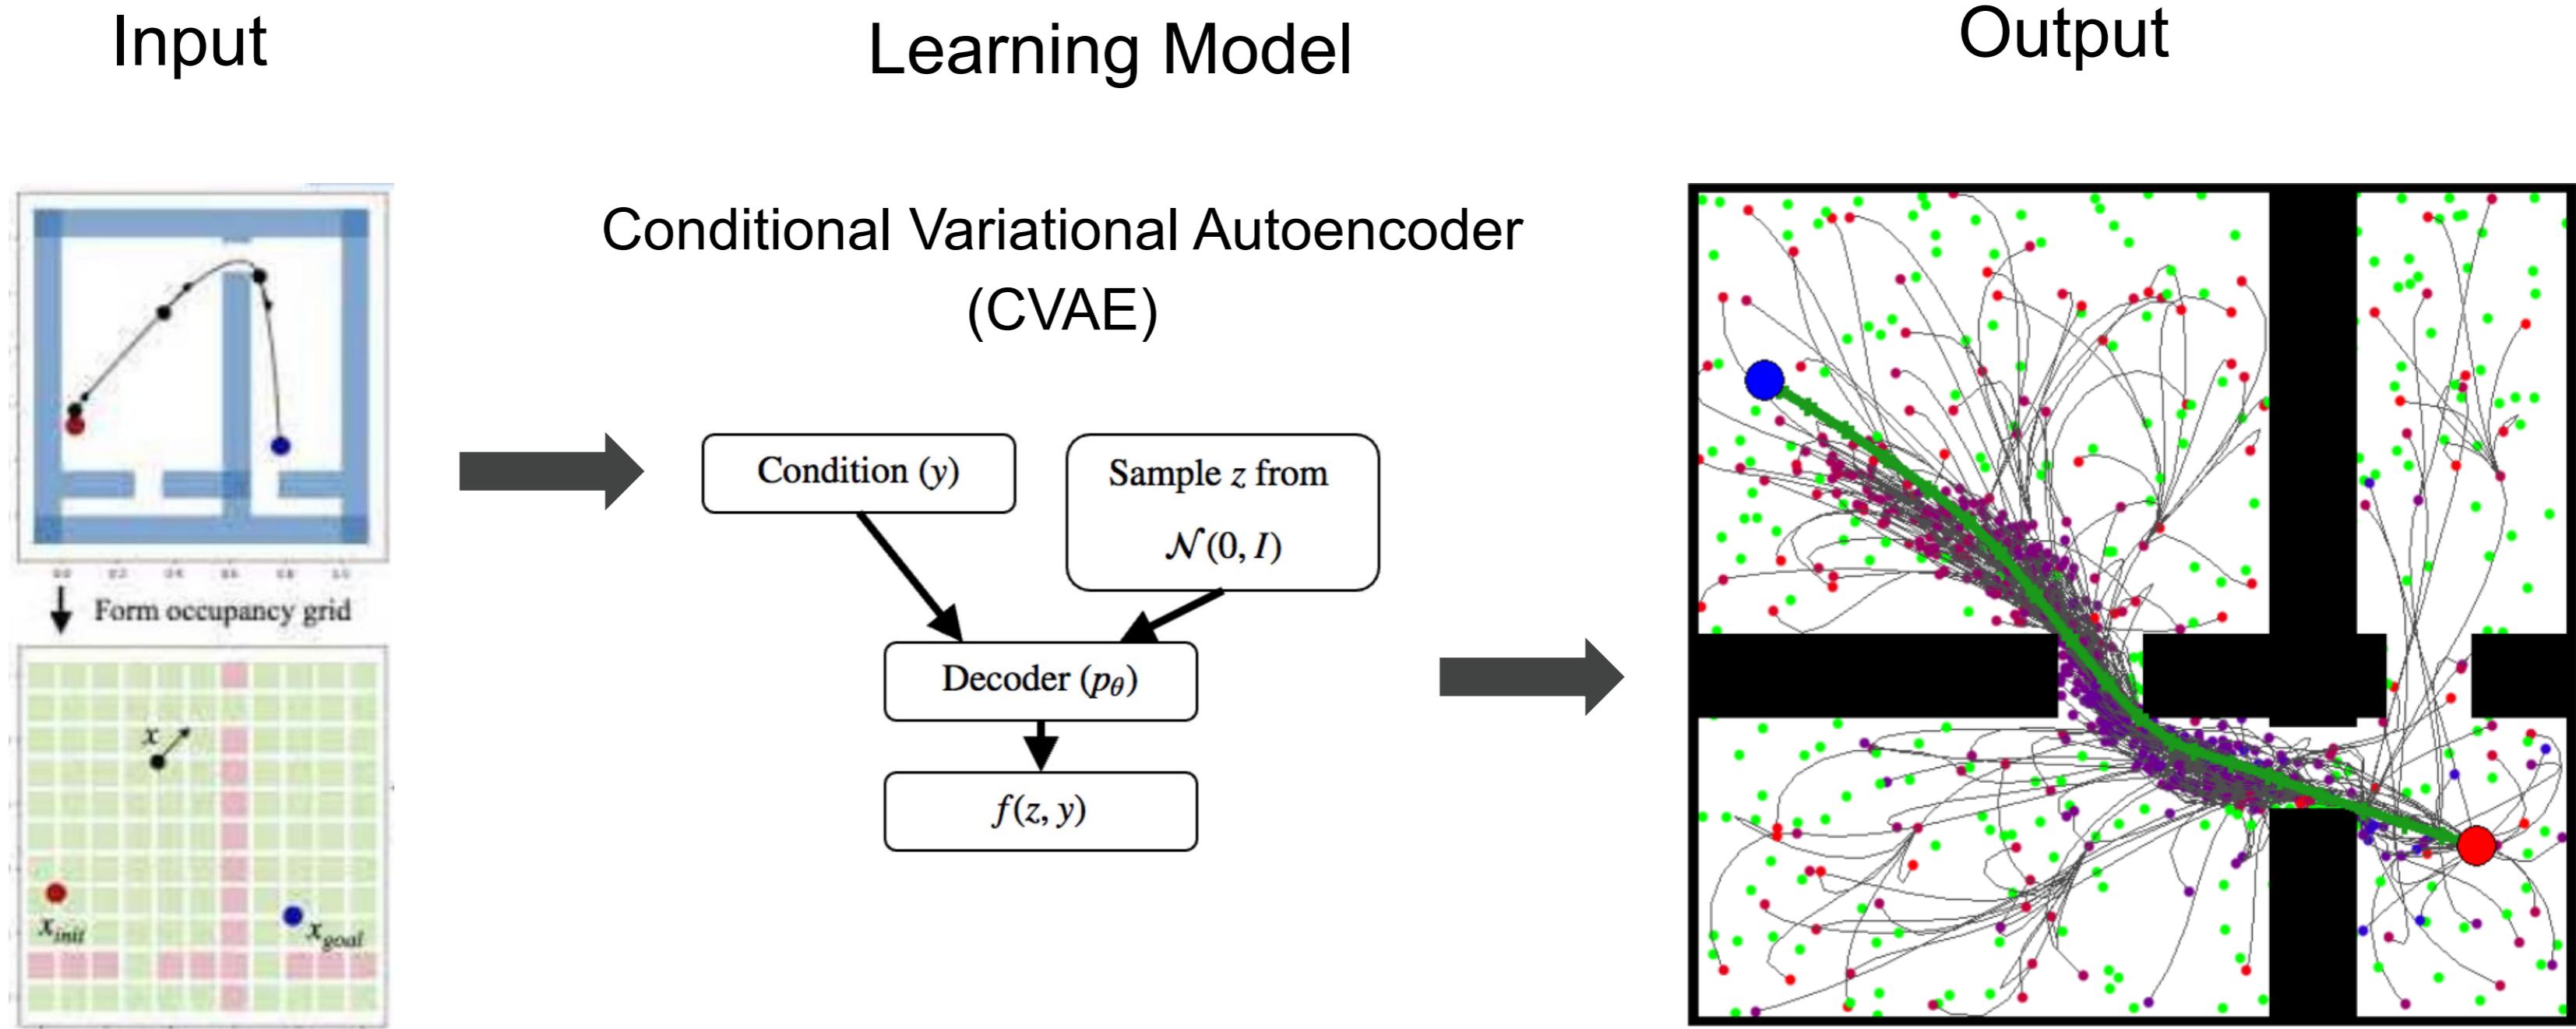
Learn a sampling distribution from demonstrations and use it to bias sampling of any sampling-based planner



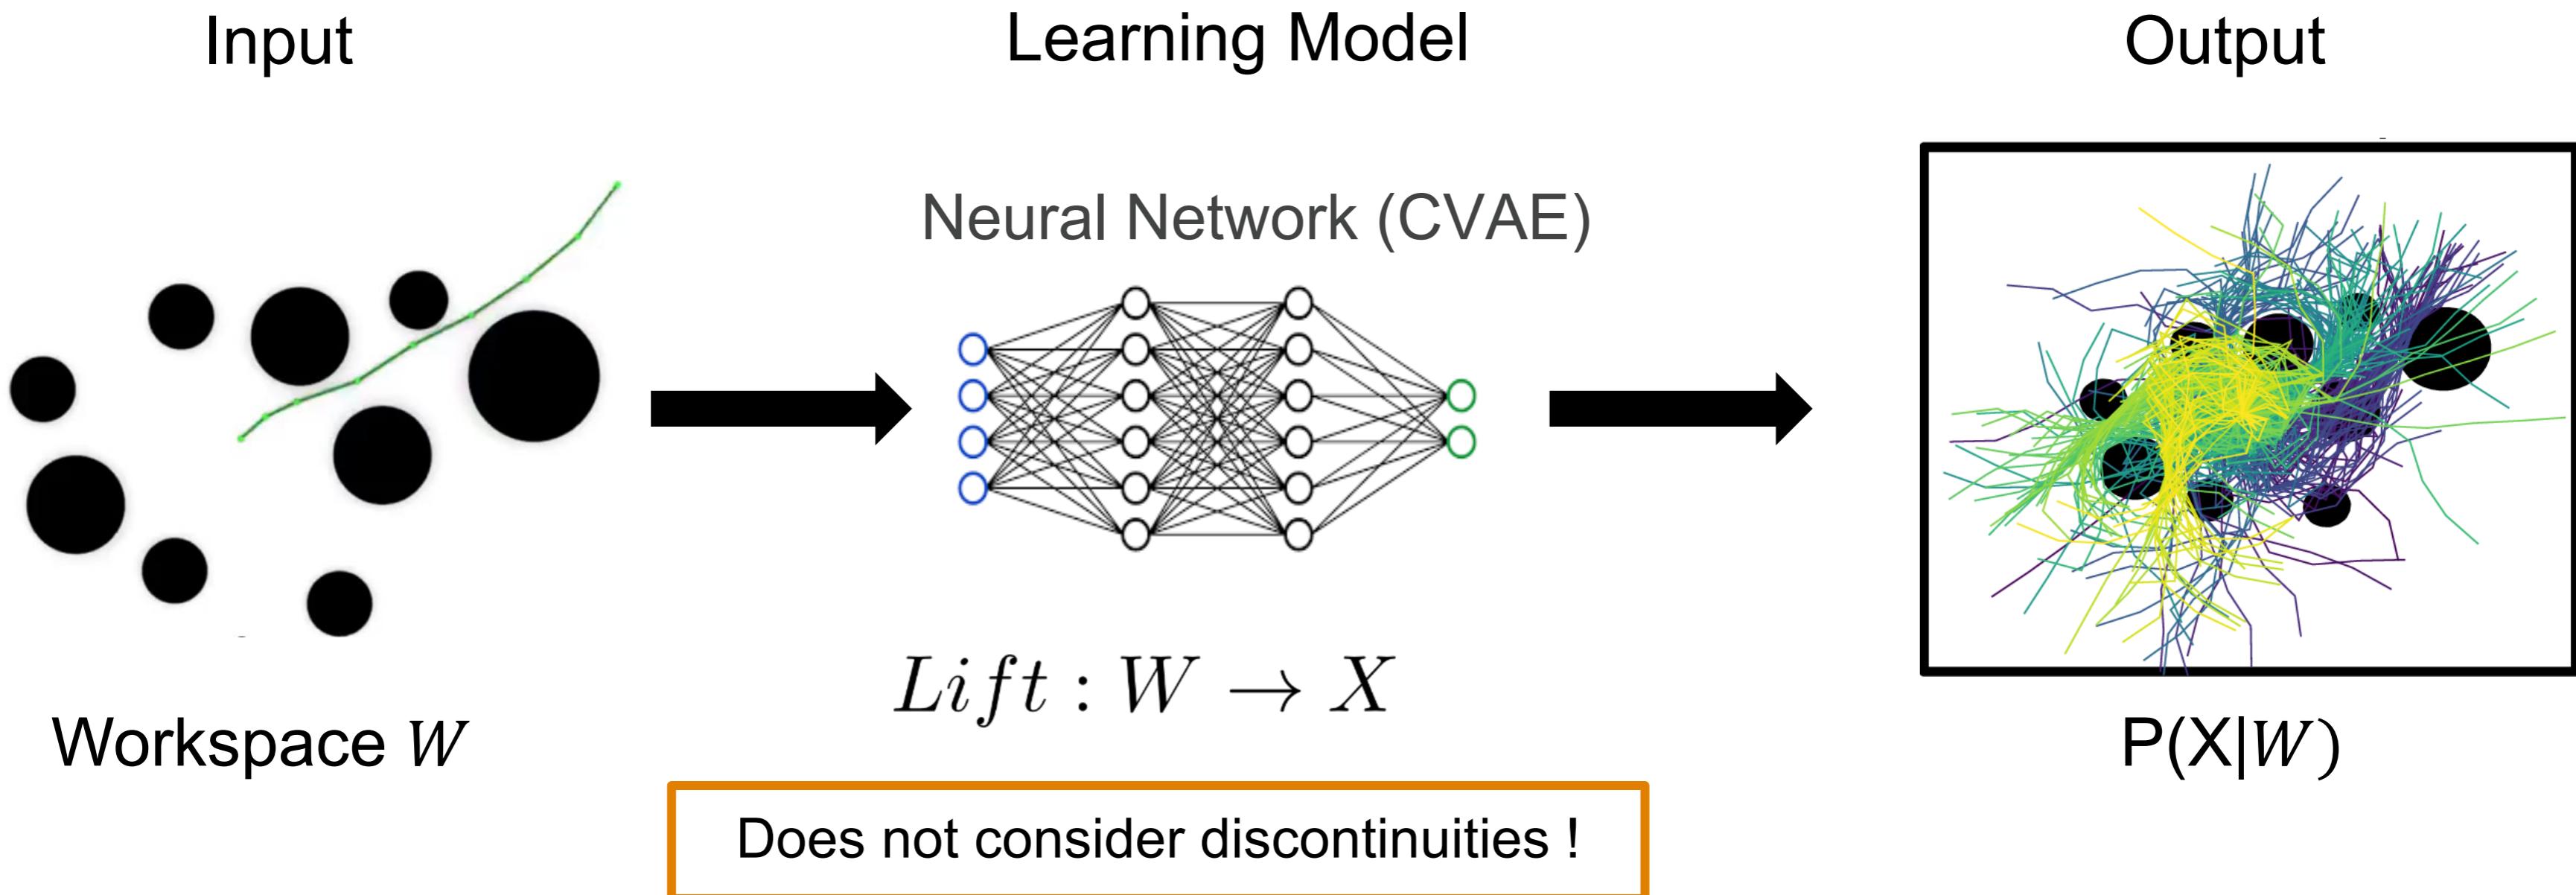
A neural network baseline



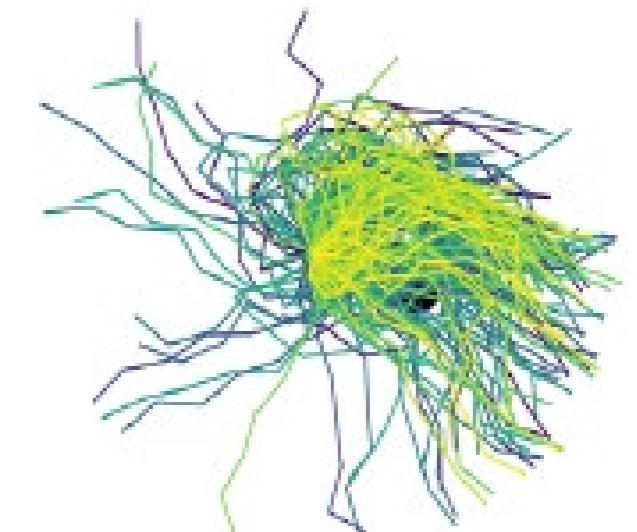
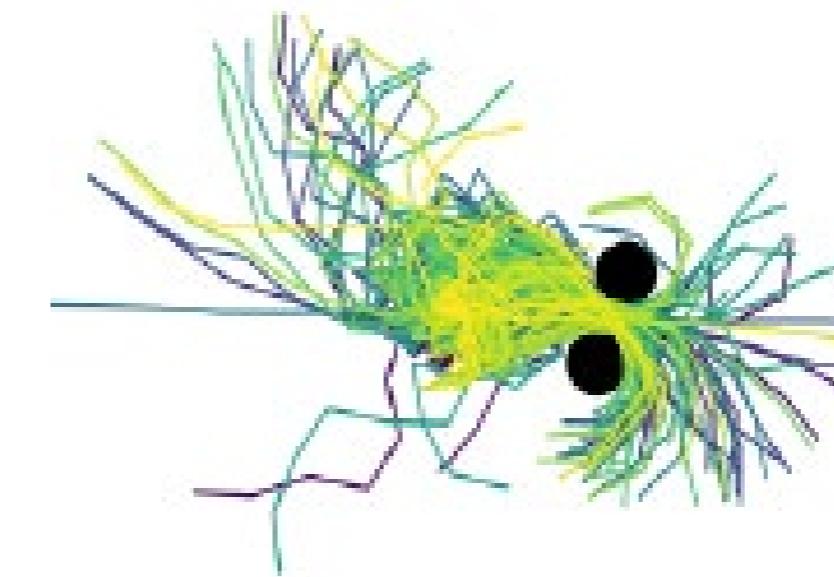
A neural network baseline (CVAE)



A neural network baseline (CVAE)



Qualitative comparison SPARK and CVAE



Which ones is SPARK and which one is CVAE

Which model corresponds to which method

Model A - SPARK, Model B - CVAE

Model A - CVAE, Model B - SPARK

Model A - CVAE, Model B - CVAE

Model A -SPARK, Model B - SPARK

Which model corresponds to which method

Model A - SPARK, Model B - CVAE

0%

Model A - CVAE, Model B - SPARK

0%

Model A - CVAE, Model B - CVAE

0%

Model A -SPARK, Model B - SPARK

0%

Which model corresponds to which method

Model A - SPARK, Model B - CVAE

0%

Model A - CVAE, Model B - SPARK

0%

Model A - CVAE, Model B - CVAE

0%

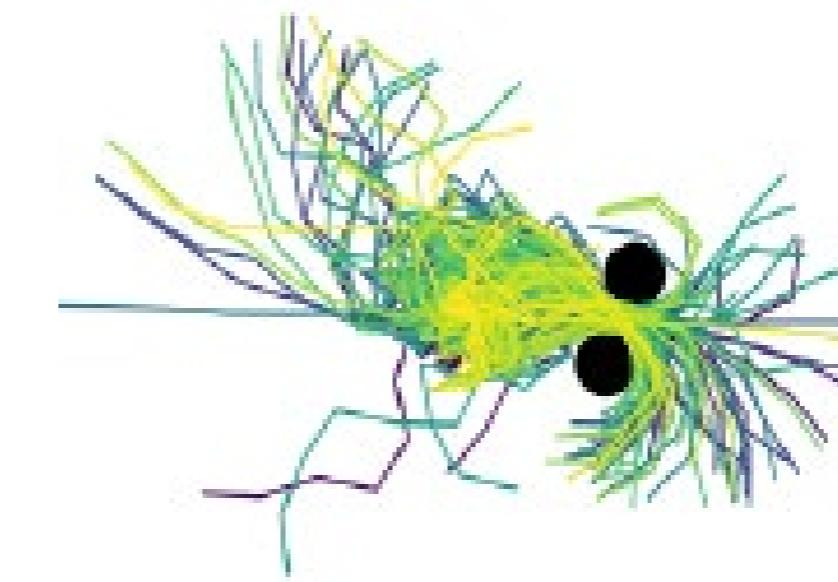
Model A -SPARK, Model B - SPARK

0%

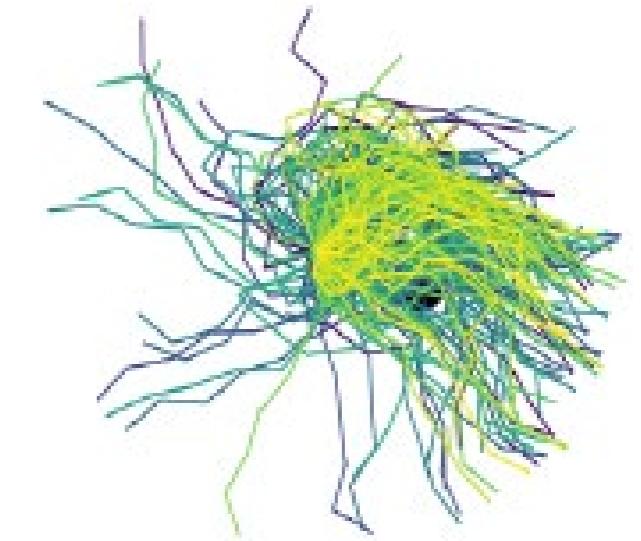
Qualitative comparison SPARK and CVAE



Problem



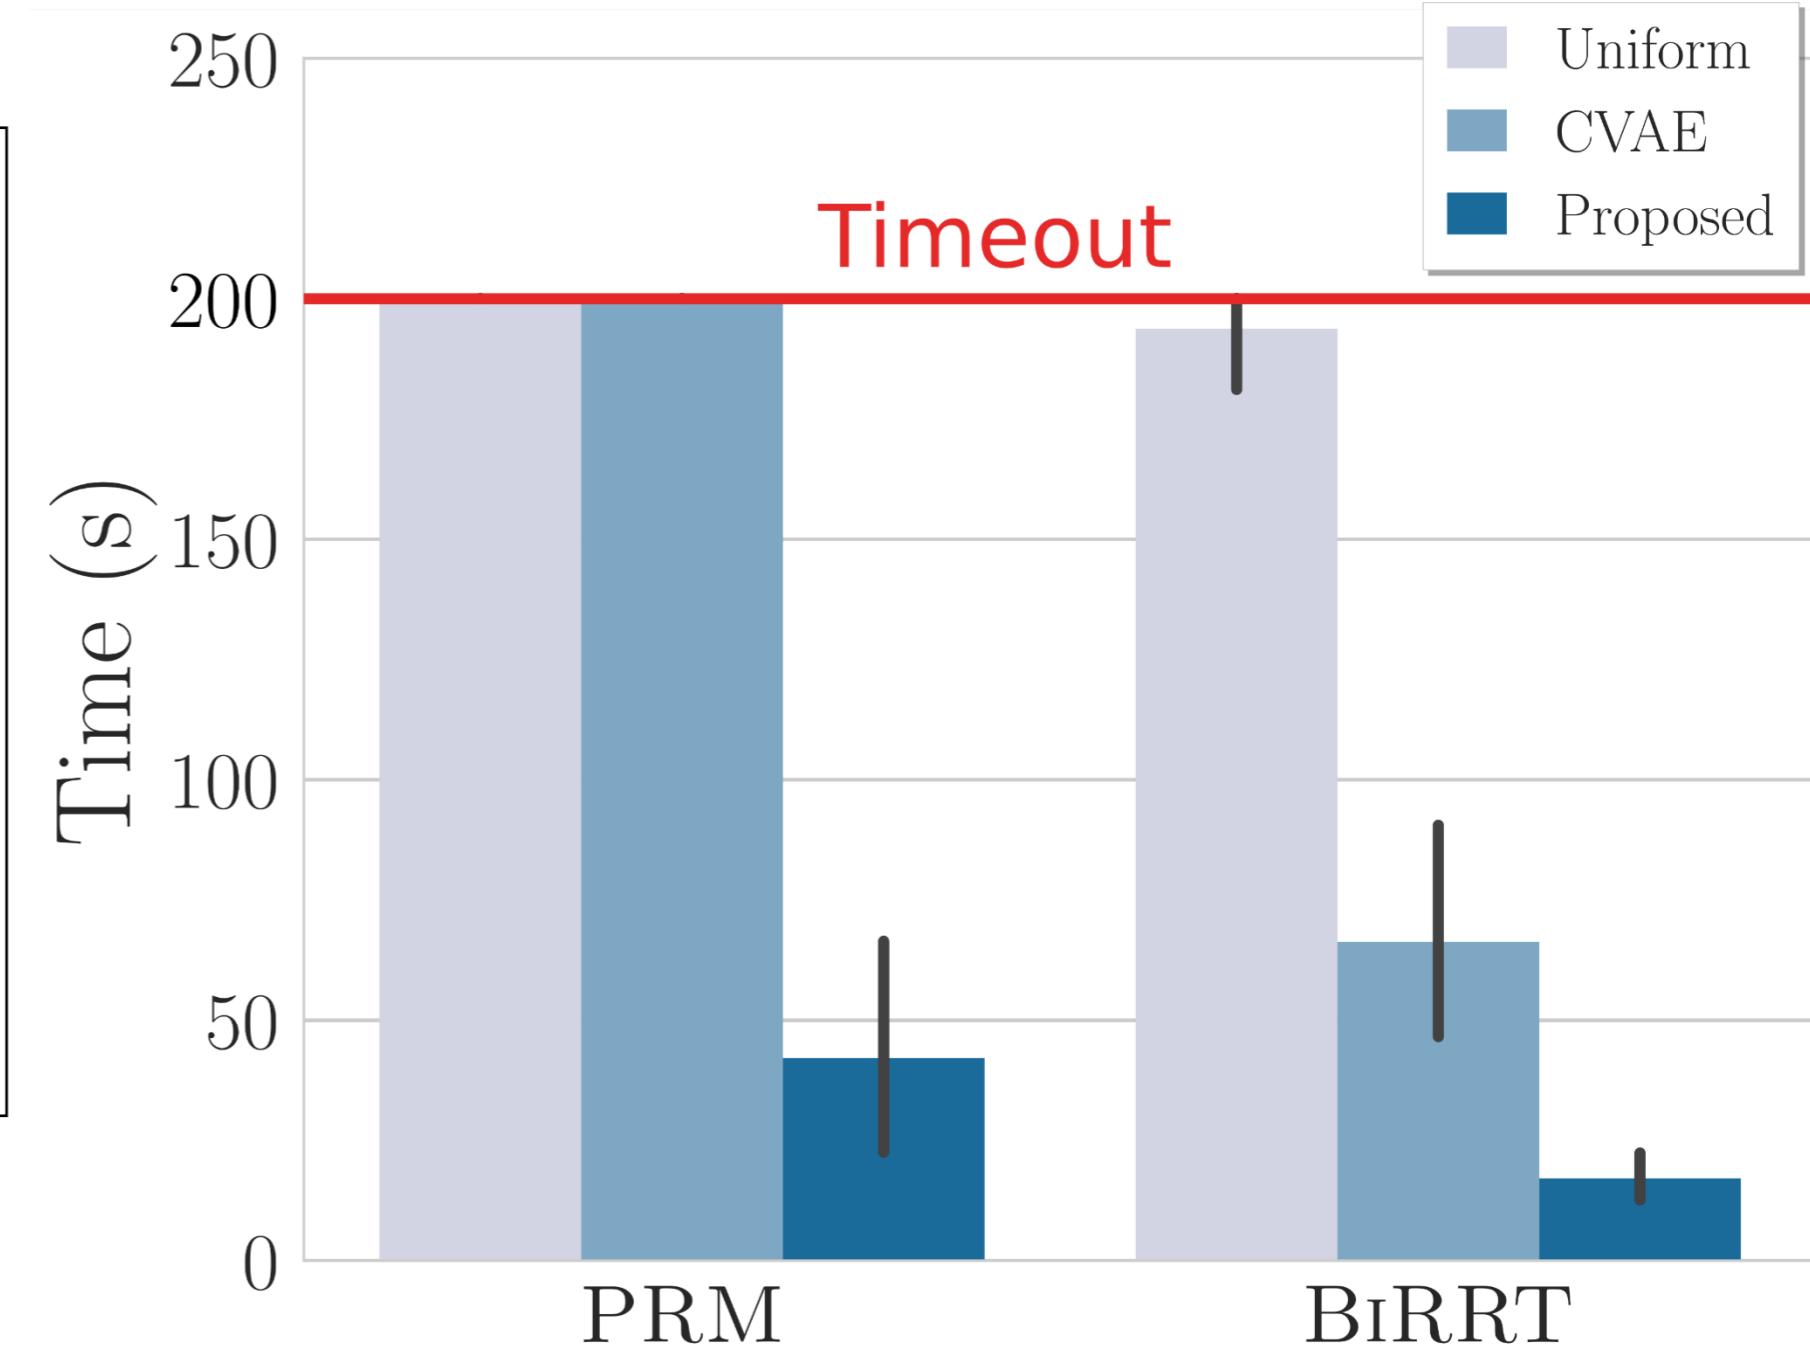
SPARK samples



CVAE samples

Discontinuity seems to affect CVAE samples

Quantitative comparison SPARK and CVAE



SPARD2D Overview

Advantages:

- Solved intractable problems
- Decomposition enables generalization with few samples

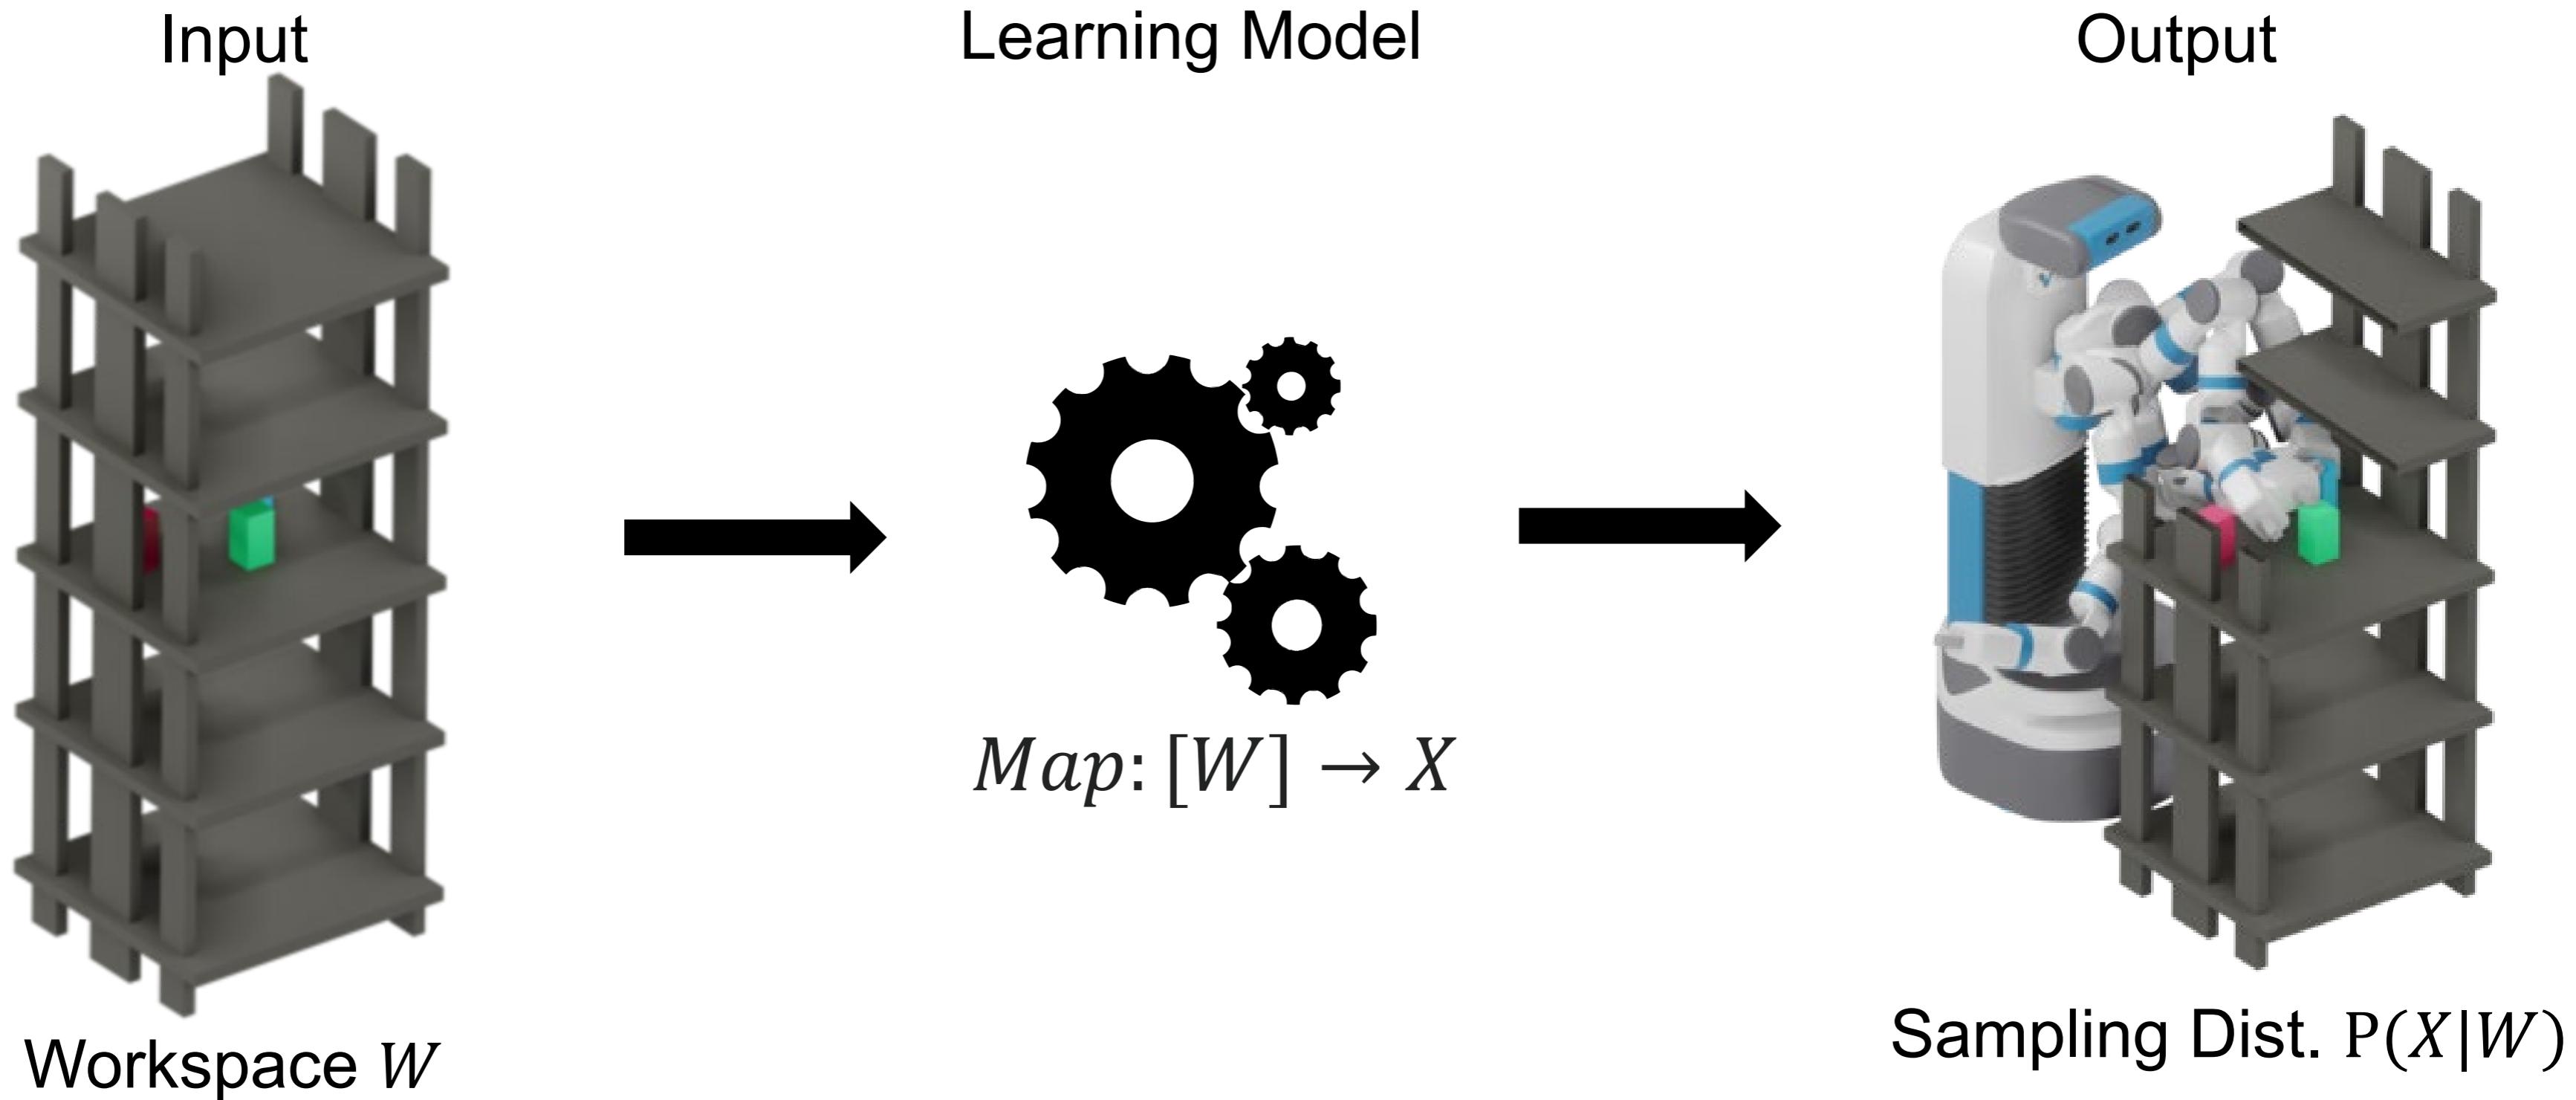
Disadvantages:

- Tested in a geometrically simple case
- Obstacles known beforehand
- Retrieval was trivial with exhaustive database

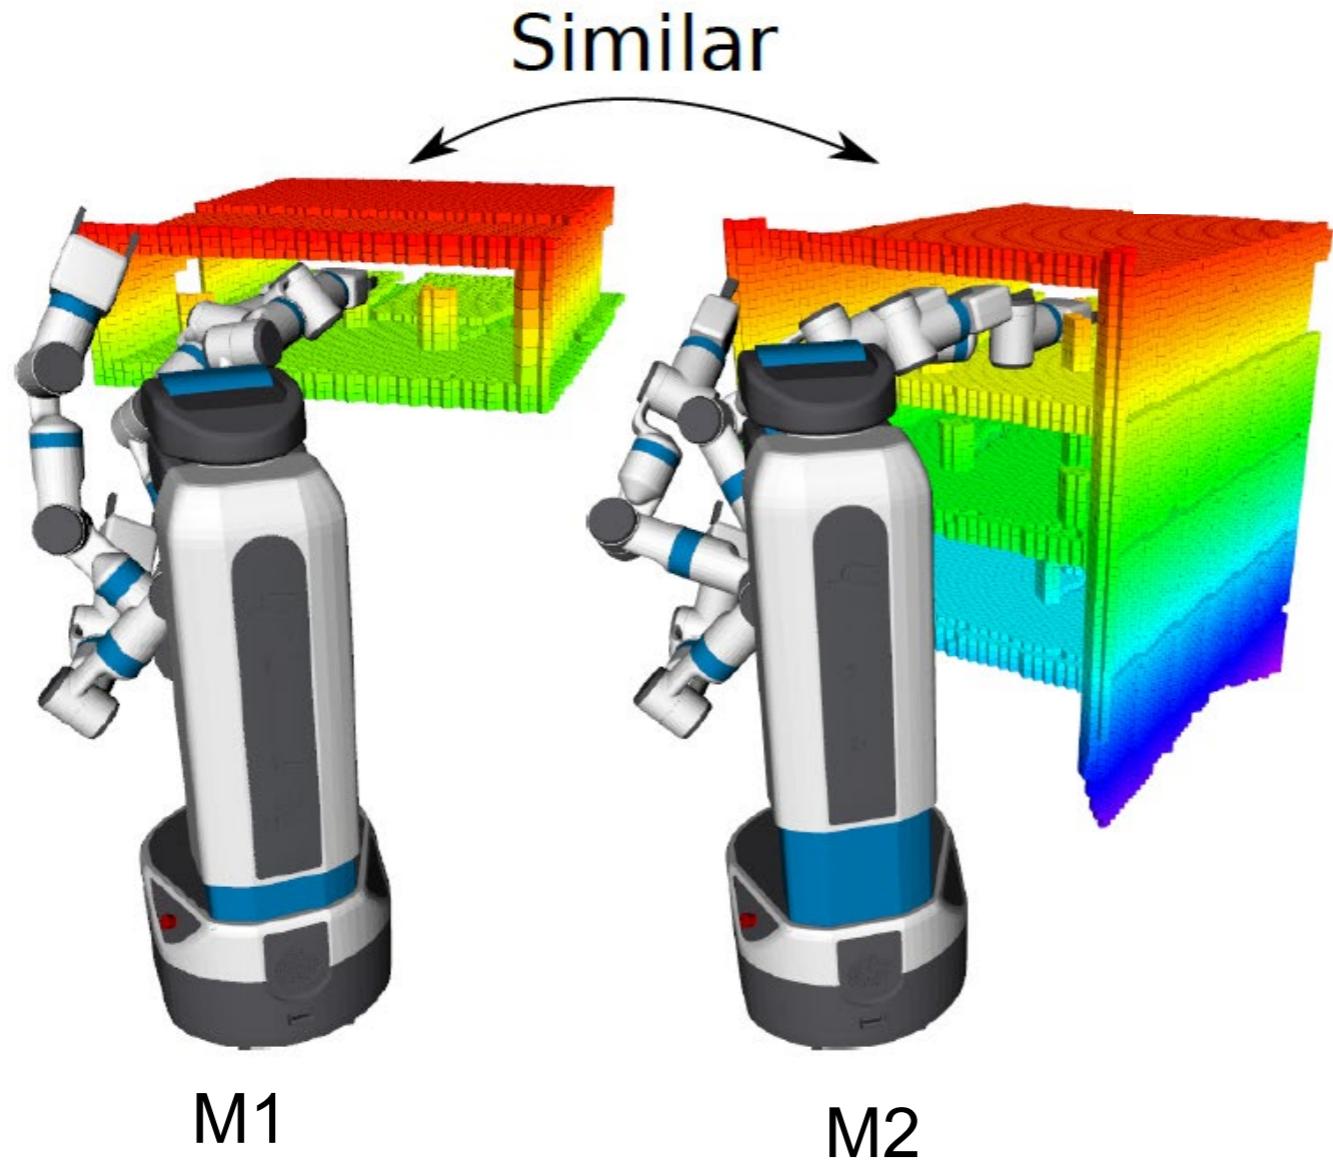
Today's Agenda

- Important Considerations
 - Locality of narrow passages
 - Discontinuity
- Biased Samplers (Cont.)
 - 2D Workspaces
 - **3D Workspaces**
 - Workplace, Start and Goal

Learn a sampling distribution in 3D workspaces



The Learning Considerations in 3D



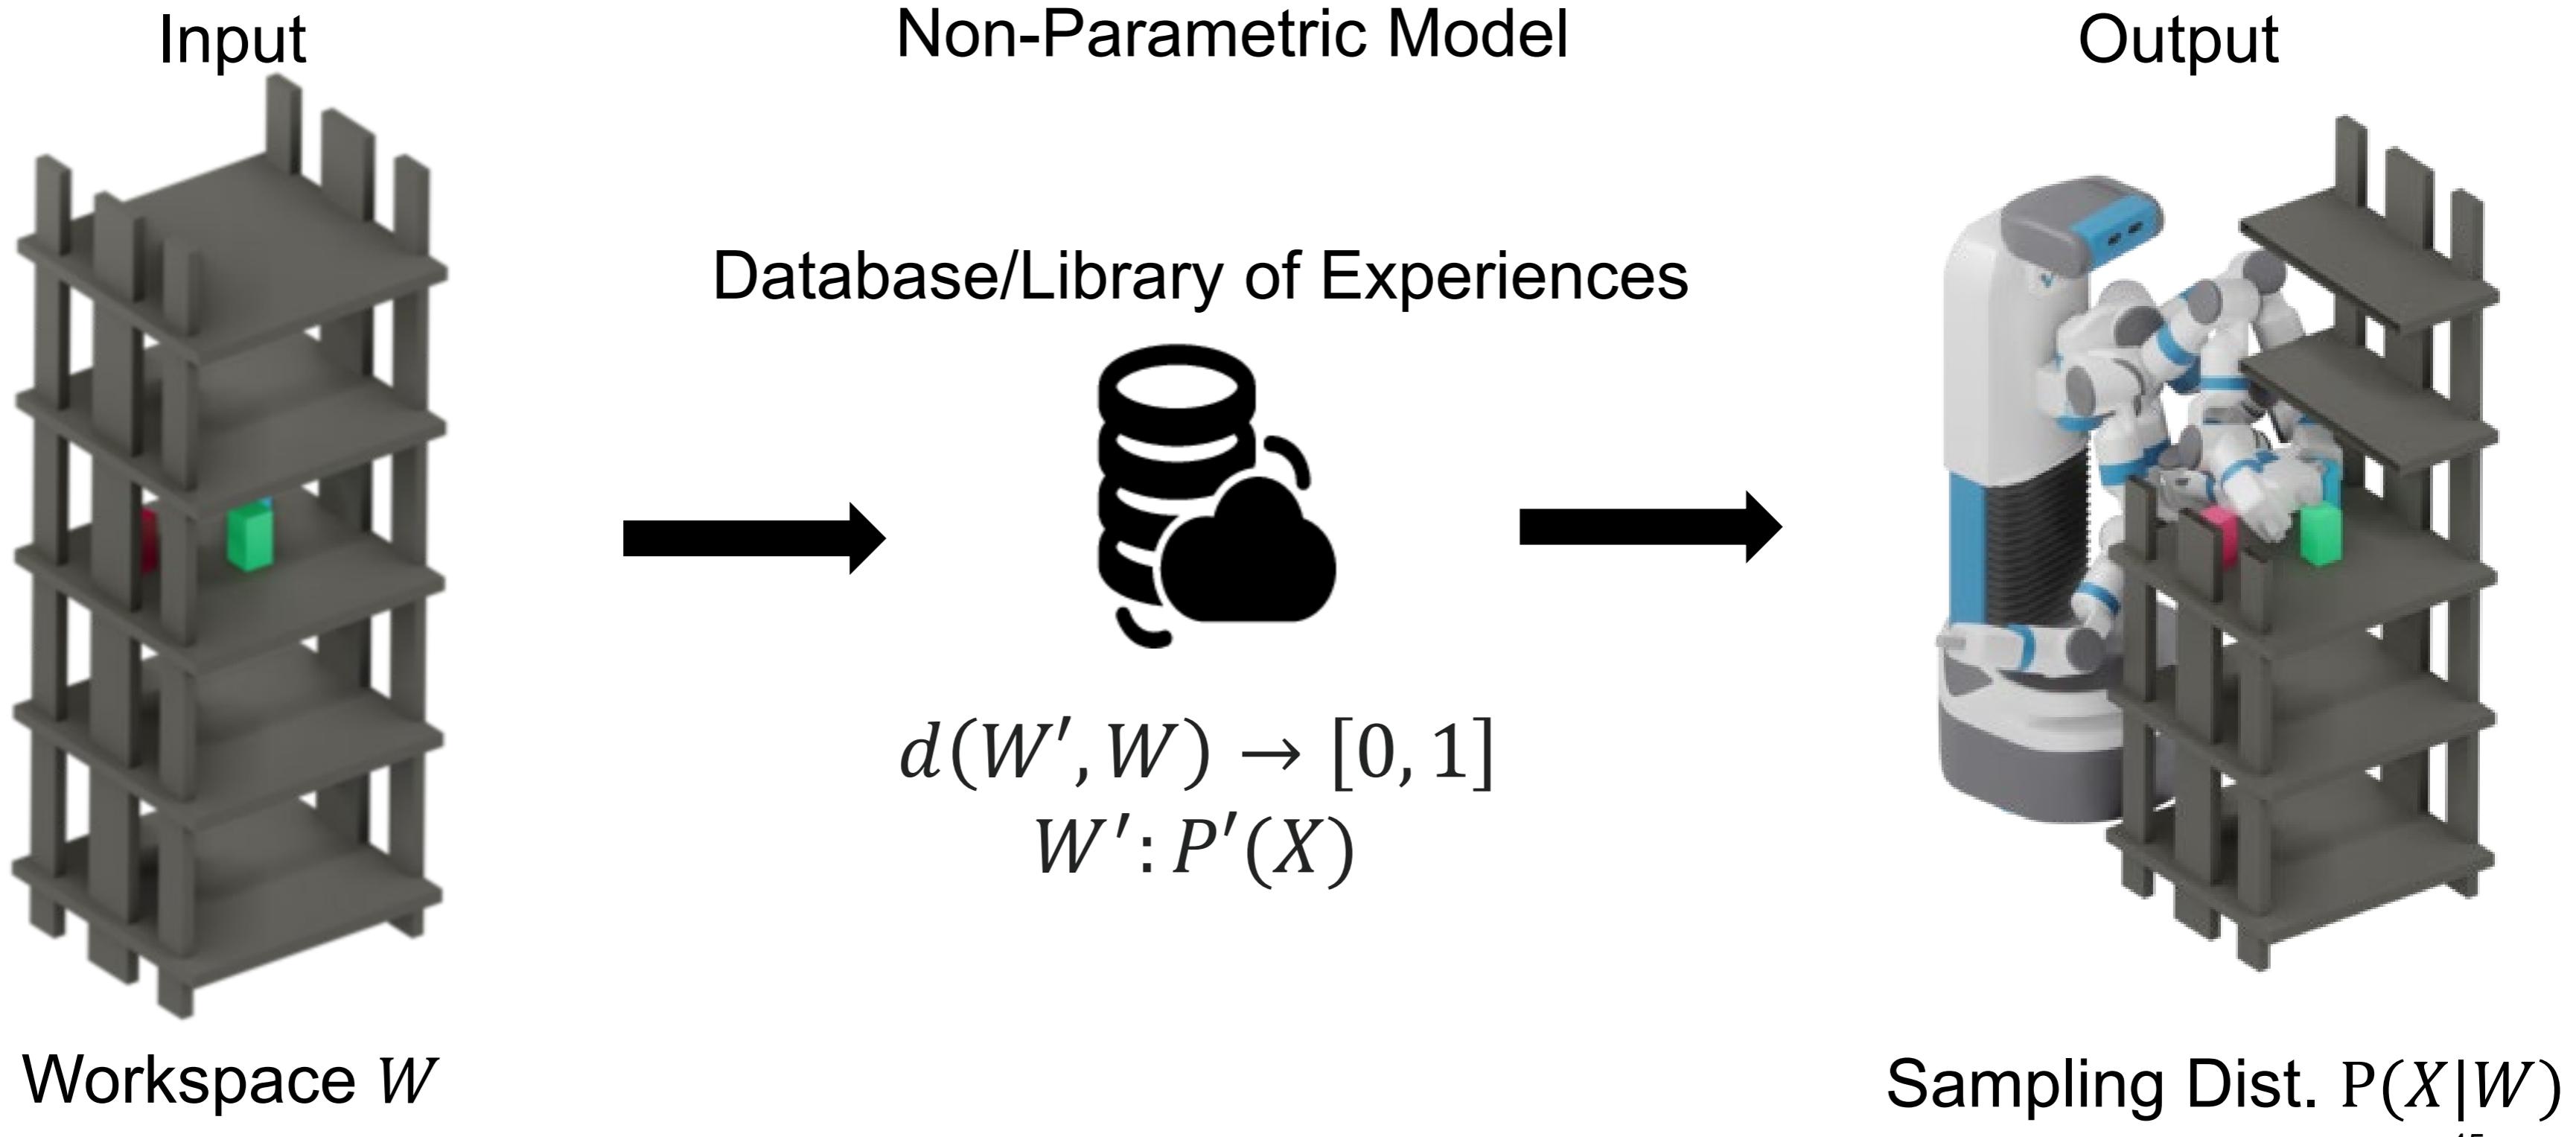
Locality (M1&M2):

- **similar** path solutions (output)
- workspaces **locally** similar (input)

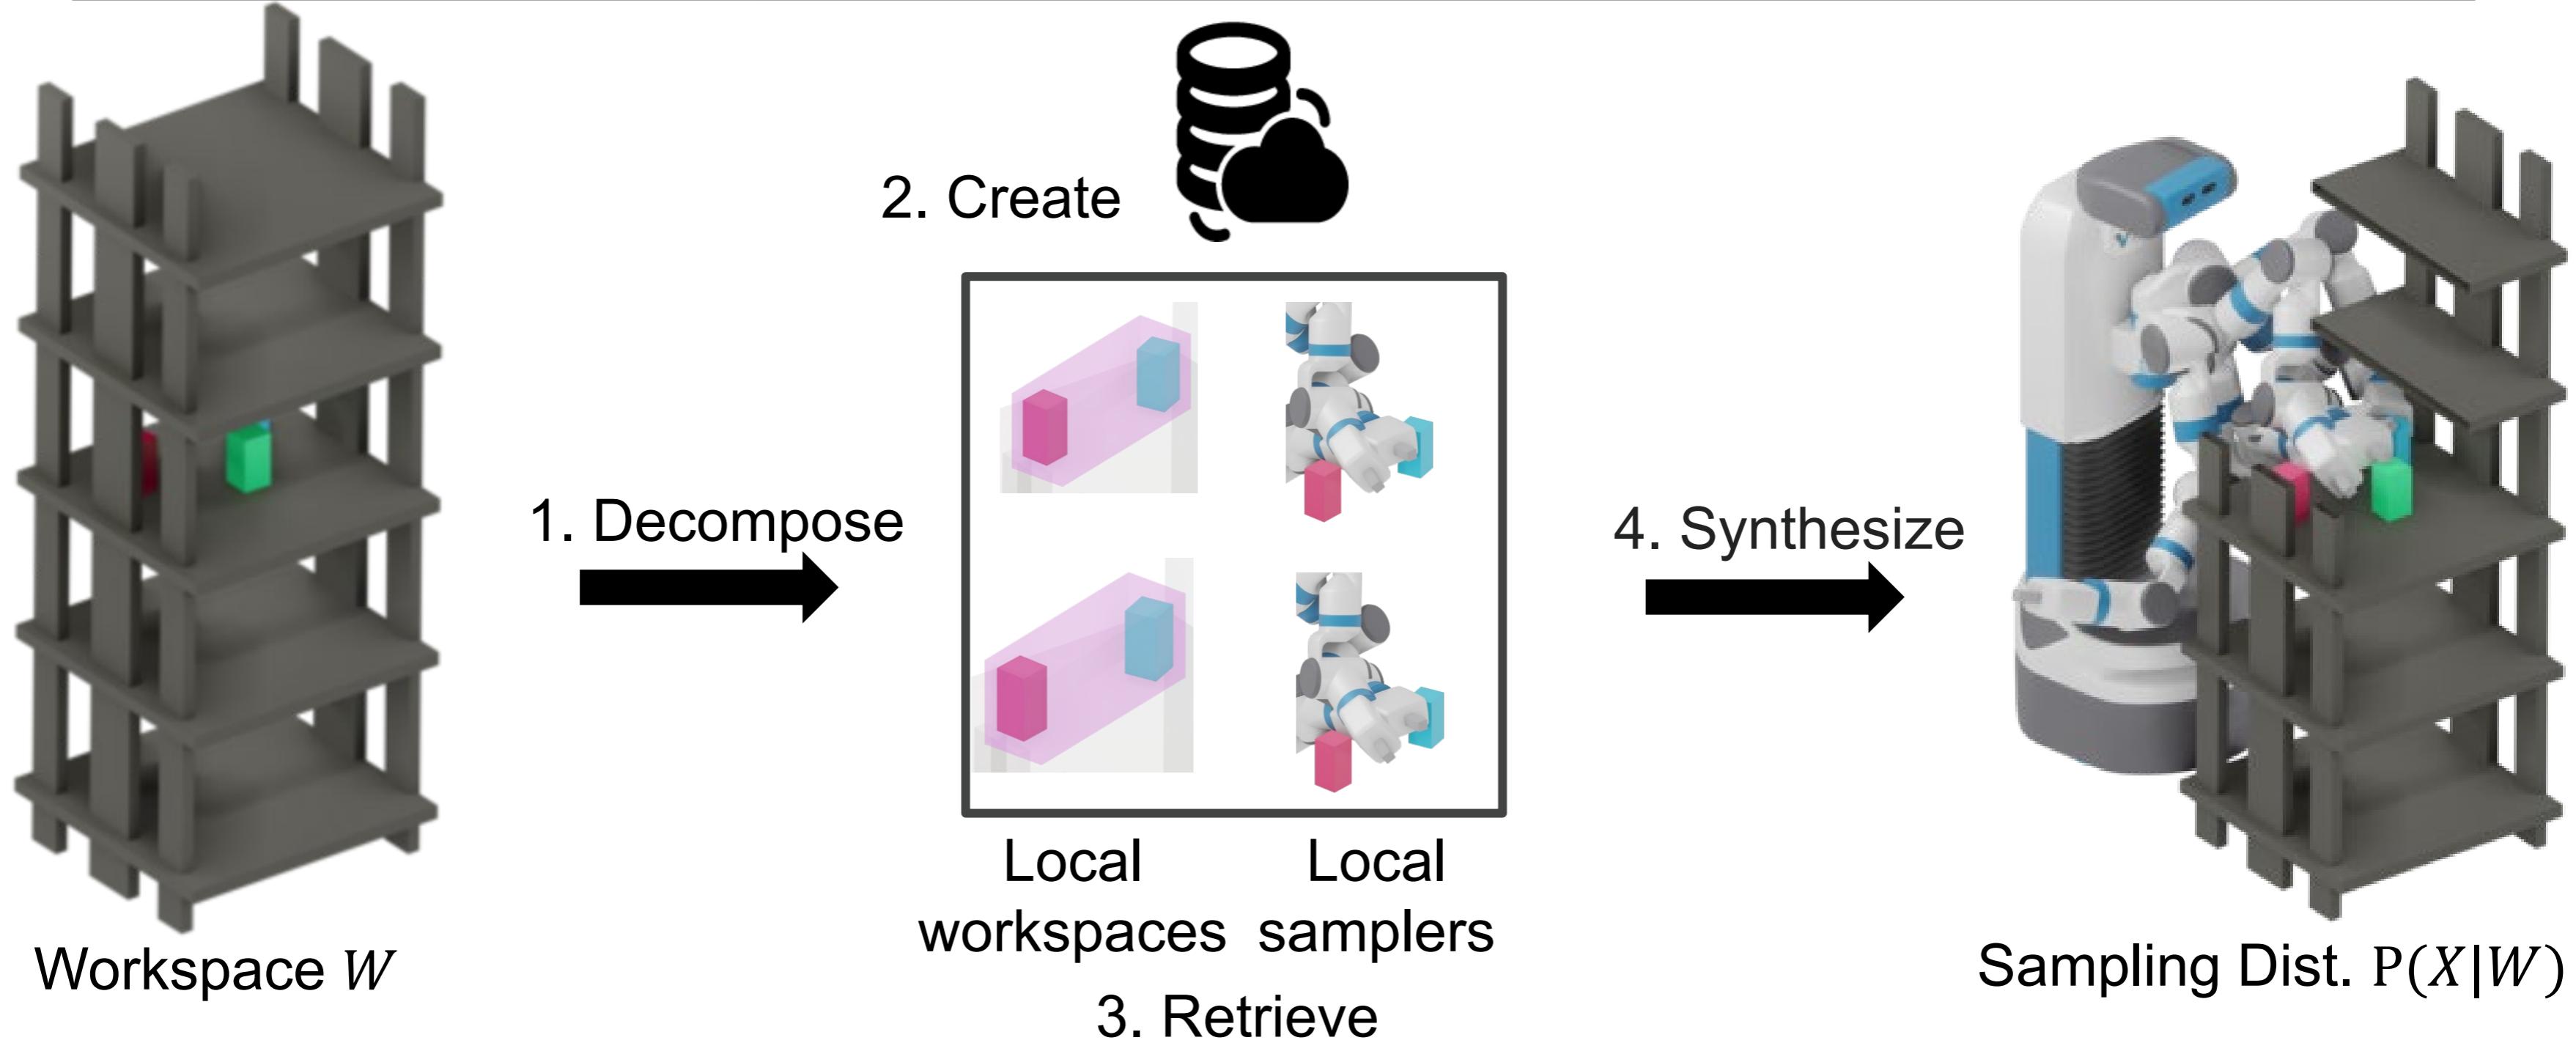
Discontinuous (M2& M3):

- **dissimilar** solution paths(output)
- very similar (input), due to **small** changes

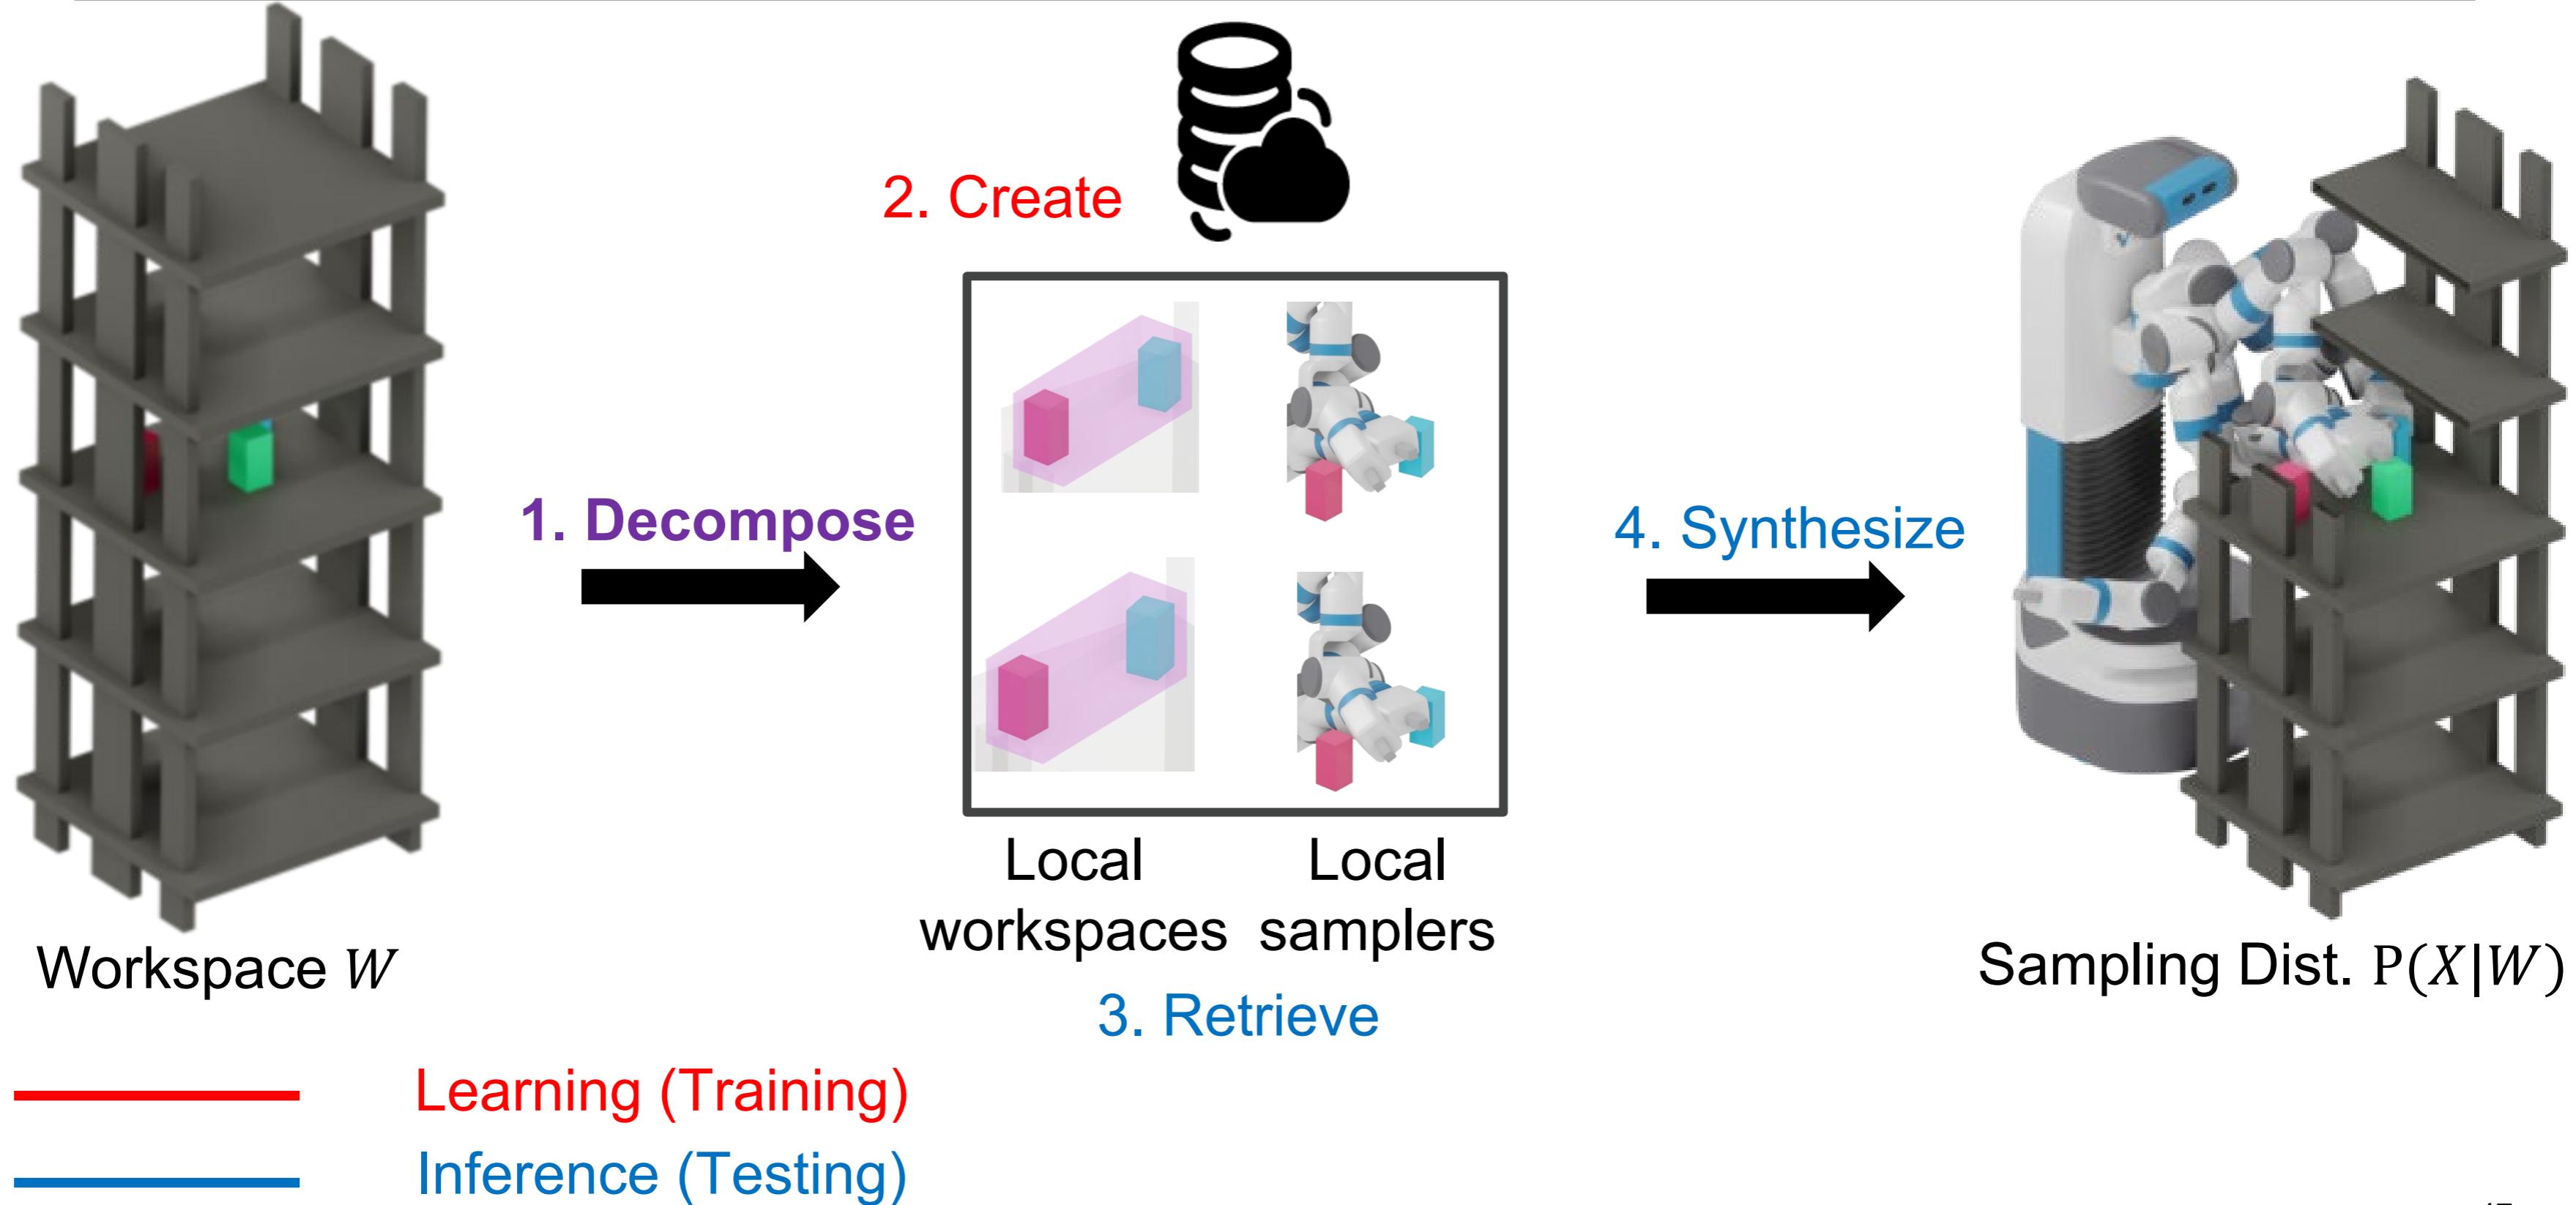
Use an experience DB to store/retrieve samplers



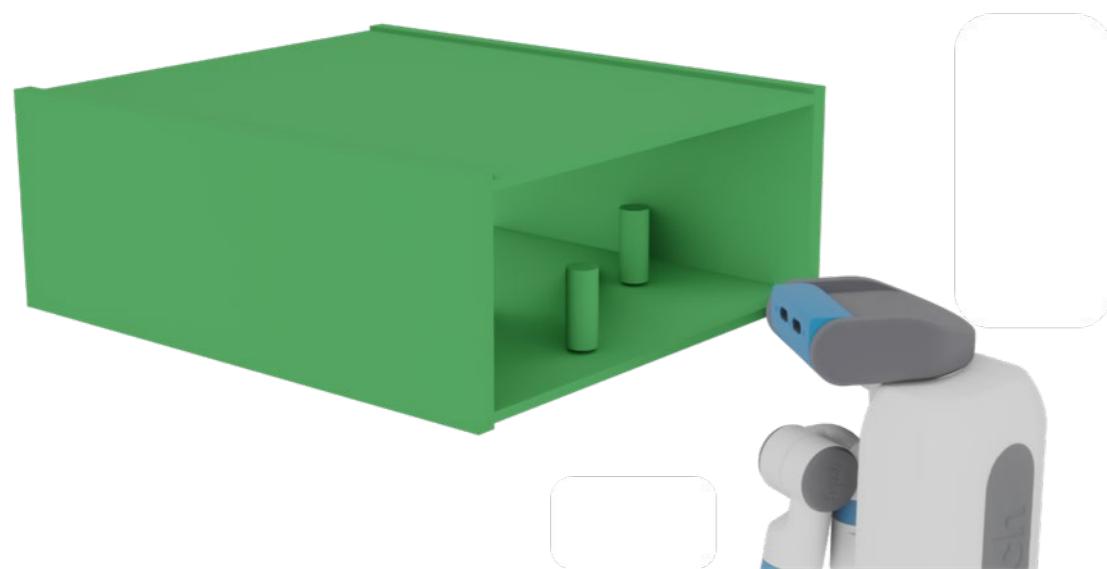
Proposed framework pipeline



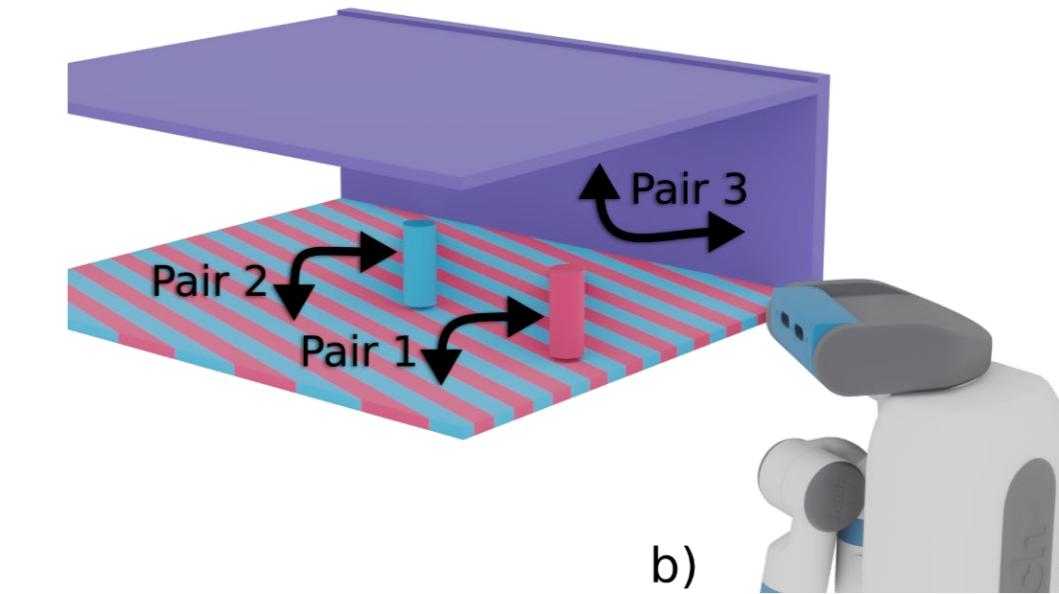
Proposed framework pipeline



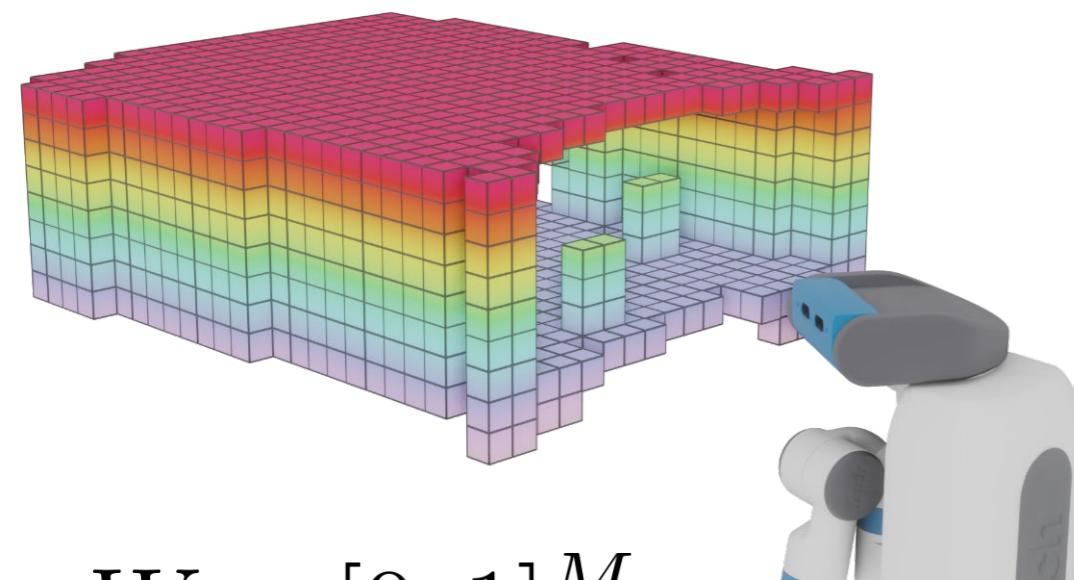
1) Decompose to Local Workspaces (Learning & Inference)



Geometric
Decomposition



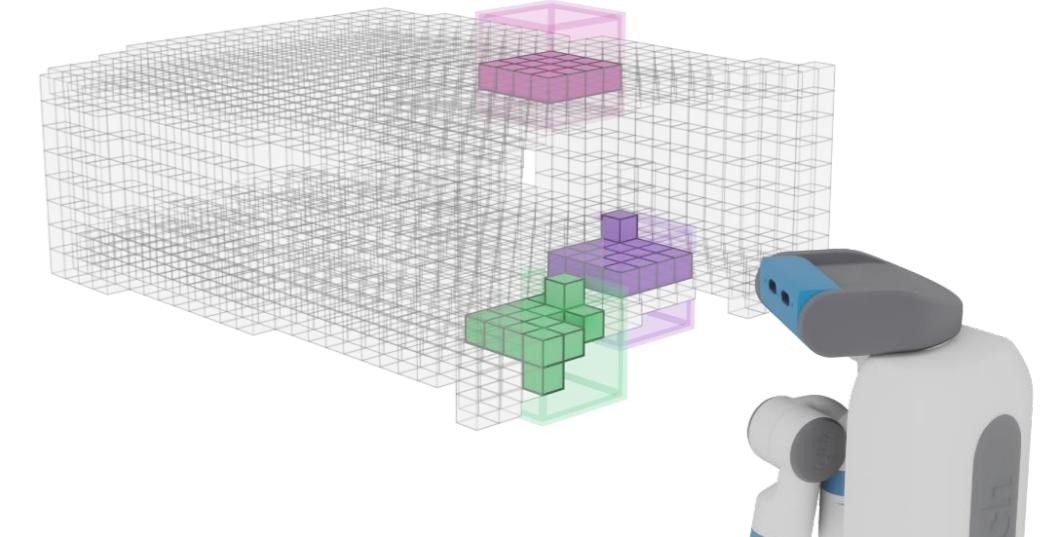
$$W = \{o : o = (\text{pose}, \text{geom})\}$$



Sensed
Decomposition

$$W \in [0, 1]^M$$

$$lw = \{x_a, q_a, s_a, x_b, q_b, s_b\}$$



$$lw = \{x_i, b_i\}, b_i \in [0, 1]^{64}$$

2) Create local workspaces/samplers (Learning)



A workspace and a
corresponding solution path

2) Create local workspaces/samplers (Learning)



Critical samples in
geometric workspace

- a) Associate critical samples
with local workspaces

2) Create local workspaces/samplers (Learning)



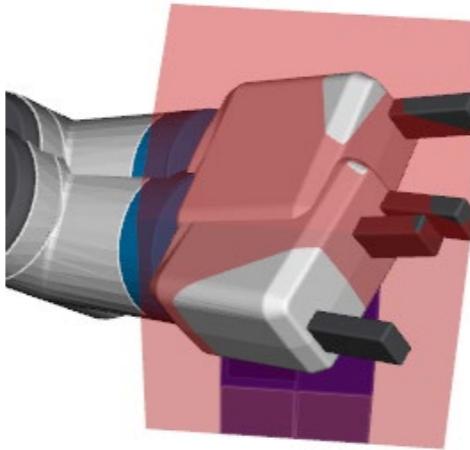
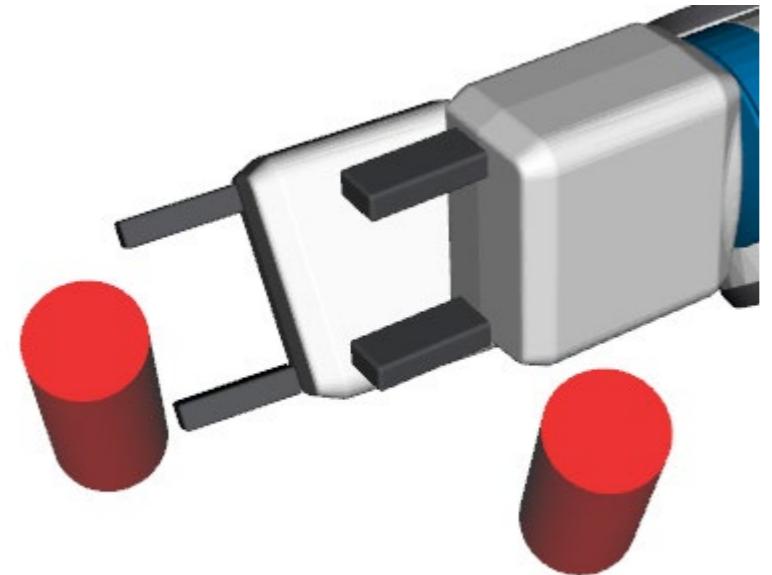
Critical samples in
geometric workspace



Critical samples in
sensed workspace

a) Associate critical samples
with local workspaces

2) Create local workspaces/samplers (Learning)



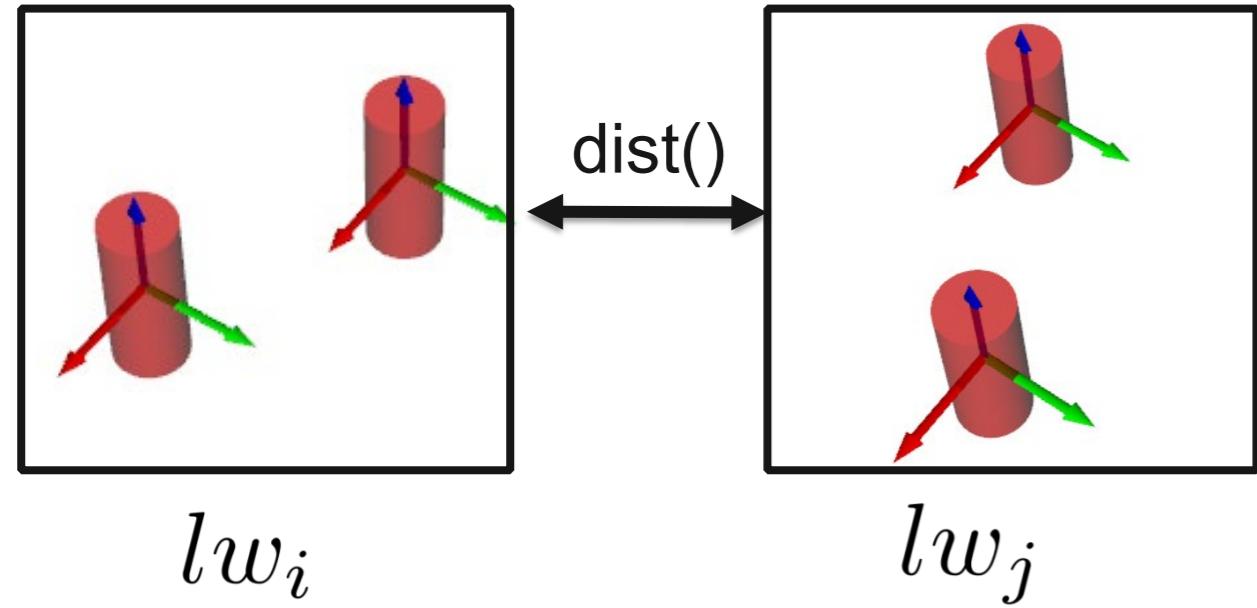
Local workspaces (lw) are shown in red with their associated critical samples (x)

a) Associate critical samples with local workspaces

$$p(x|lw) = GMM(\mu, \Sigma) = \frac{1}{M} \sum_{j=1}^M N(q_j, \Sigma_j)$$

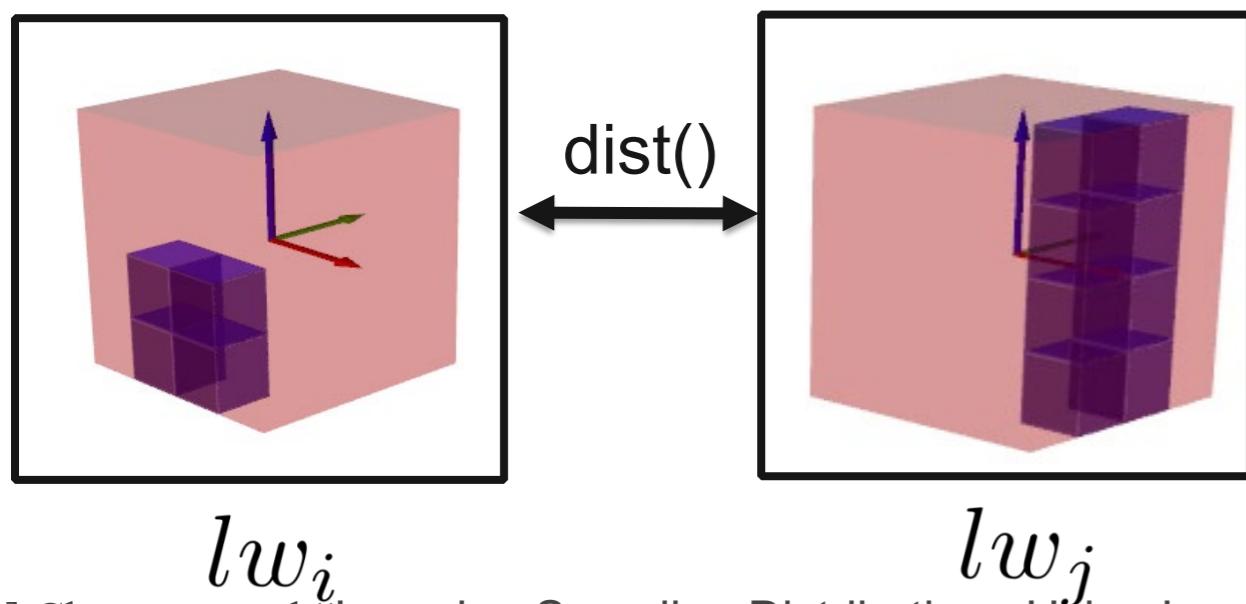
b) Fit critical samples with a GMM for each local workspace

3) Retrieve local samplers (Inference)



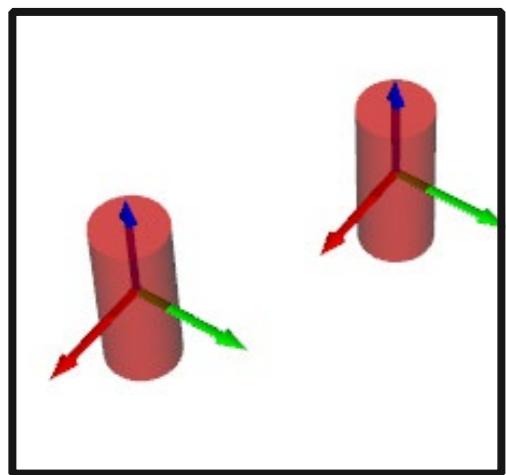
$$d_o(o_i, o_j) = w_x \|x_i - x_j\| + w_r(1 - \langle q_i, q_j \rangle) + w_s \|s_i - s_j\|$$

$$d(lw_i, lw_j) = \min \begin{cases} d_o(o_i^a, o_j^a) + d_o(o_i^b, o_j^b) \\ d_o(o_i^a, o_j^b) + d_o(o_i^b, o_j^a) \end{cases}$$



$$d(lw_i, lw_j) = \begin{cases} 0 & x_i = x_j, b_i = b_j, \\ \infty & \text{otherwise} \end{cases}$$

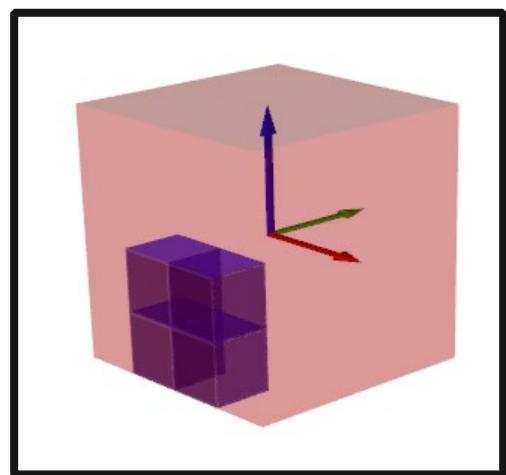
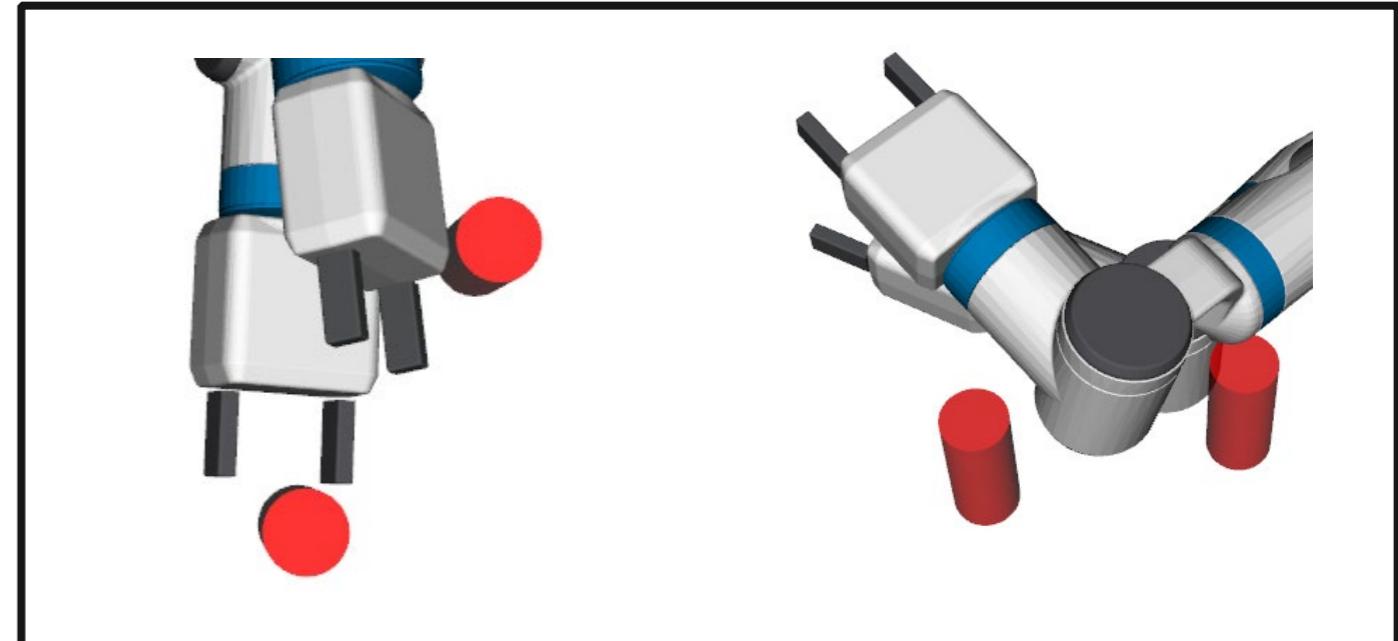
3) Retrieve local samplers (Inference)



Fixed-radius
NN search



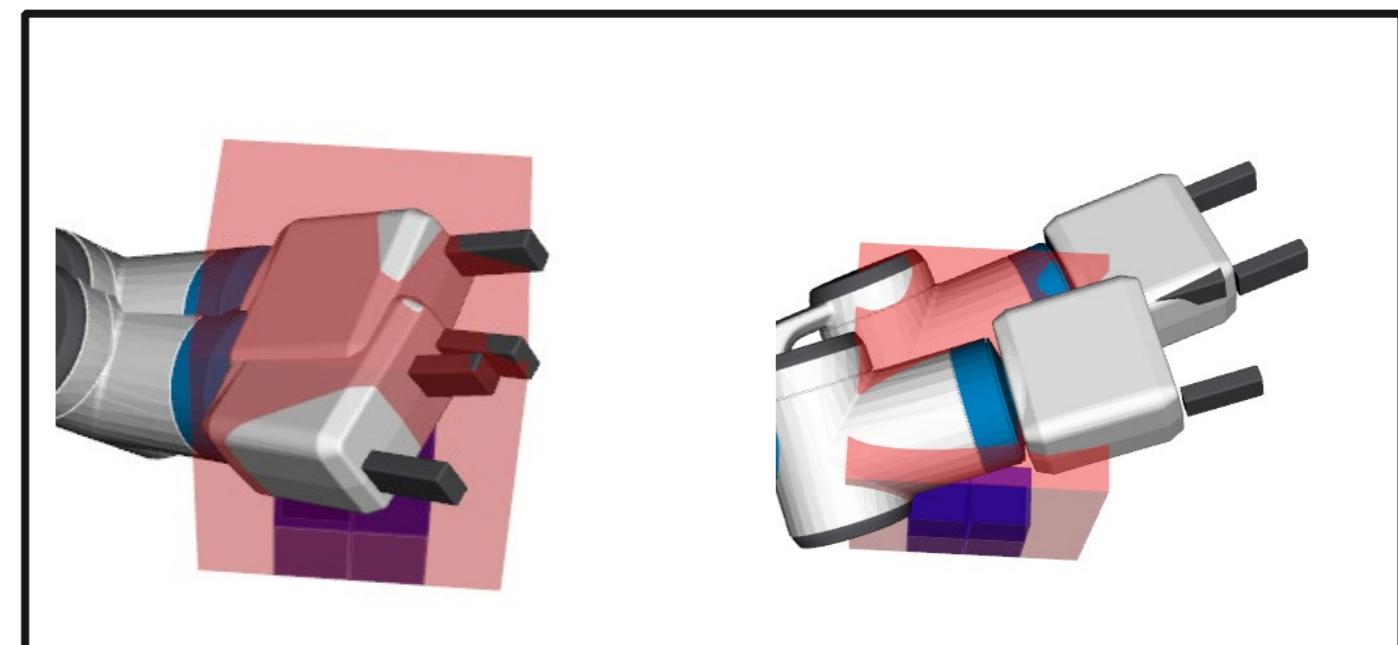
Retrieve



Fixed-radius
NN search

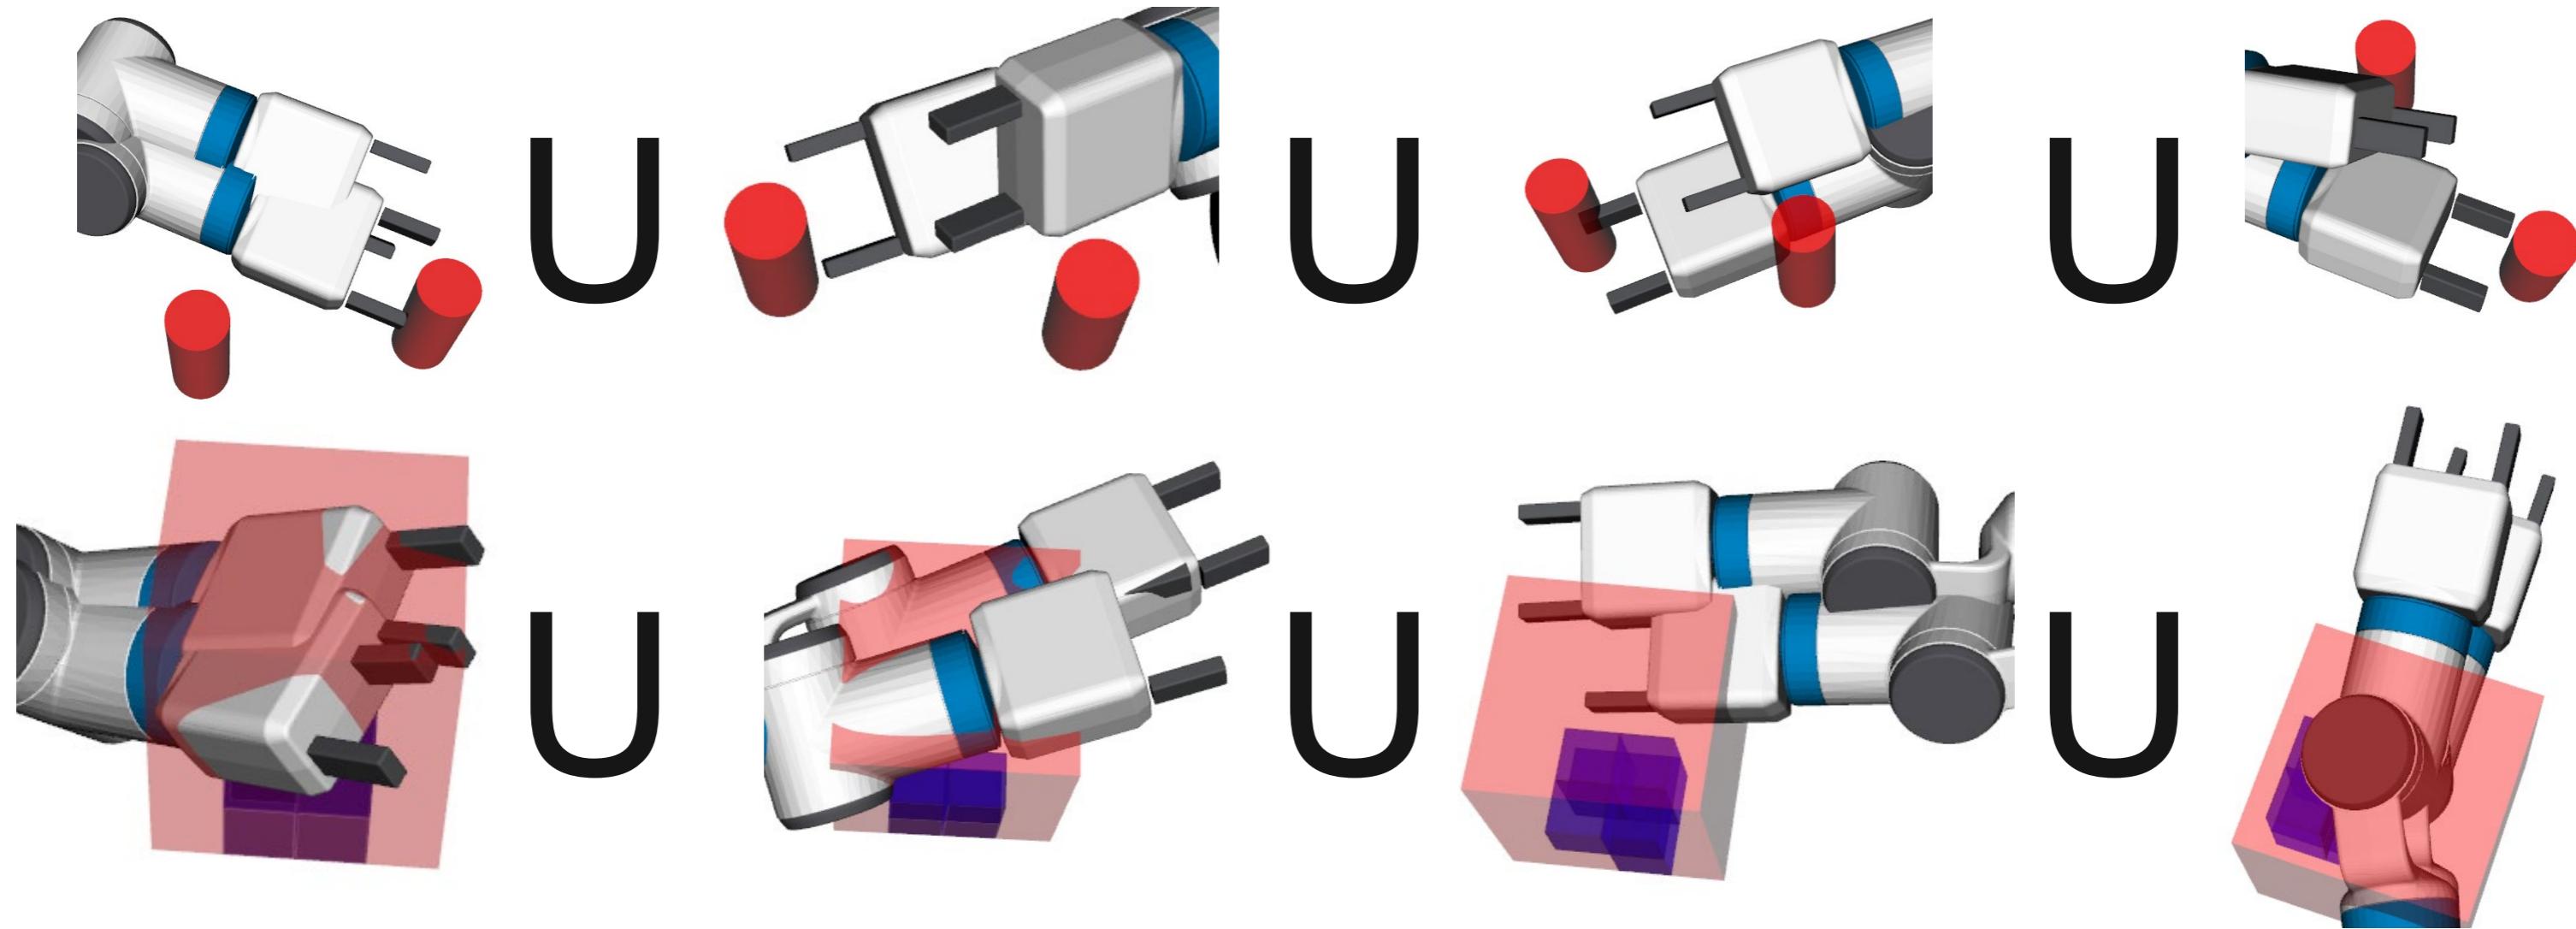


Retrieve



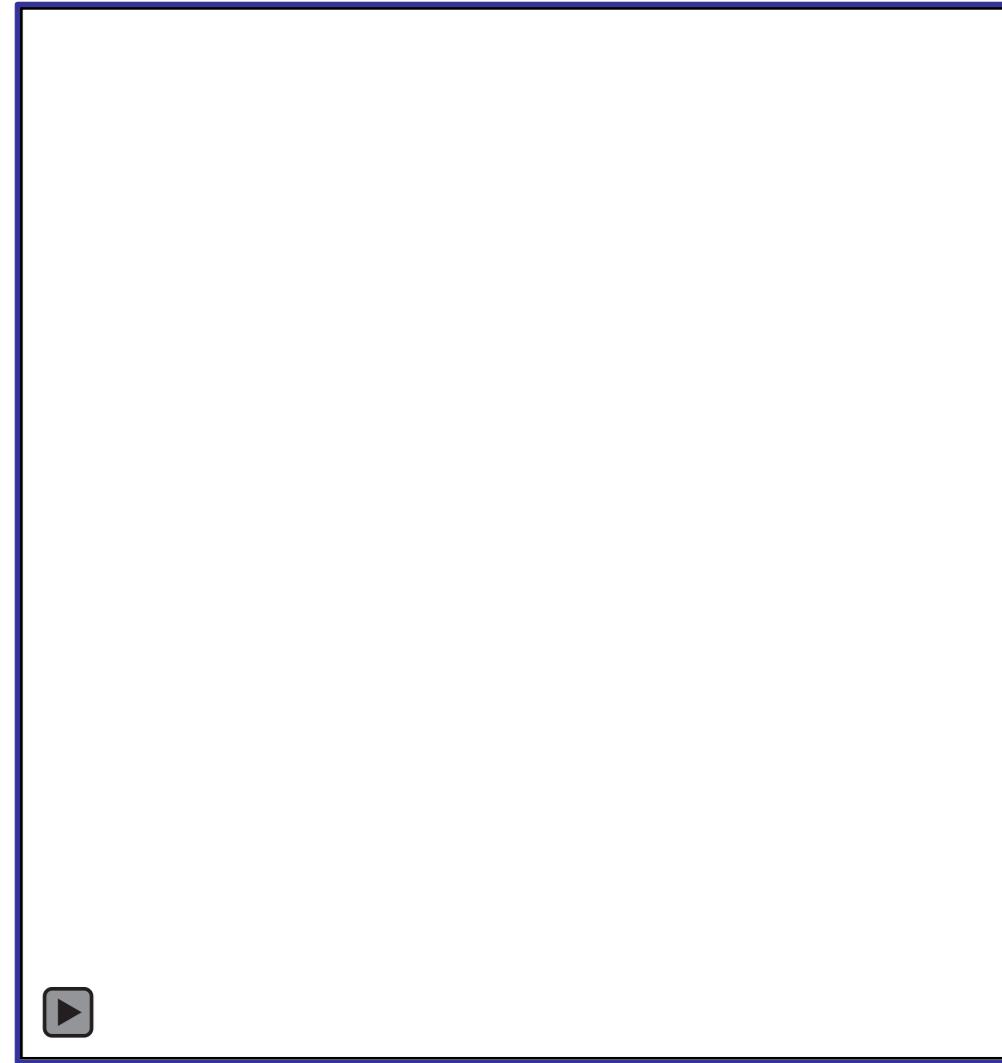
We used the GNAT structure which
has logarithmic retrieval time

4) Synthesize local samplers (Testing)

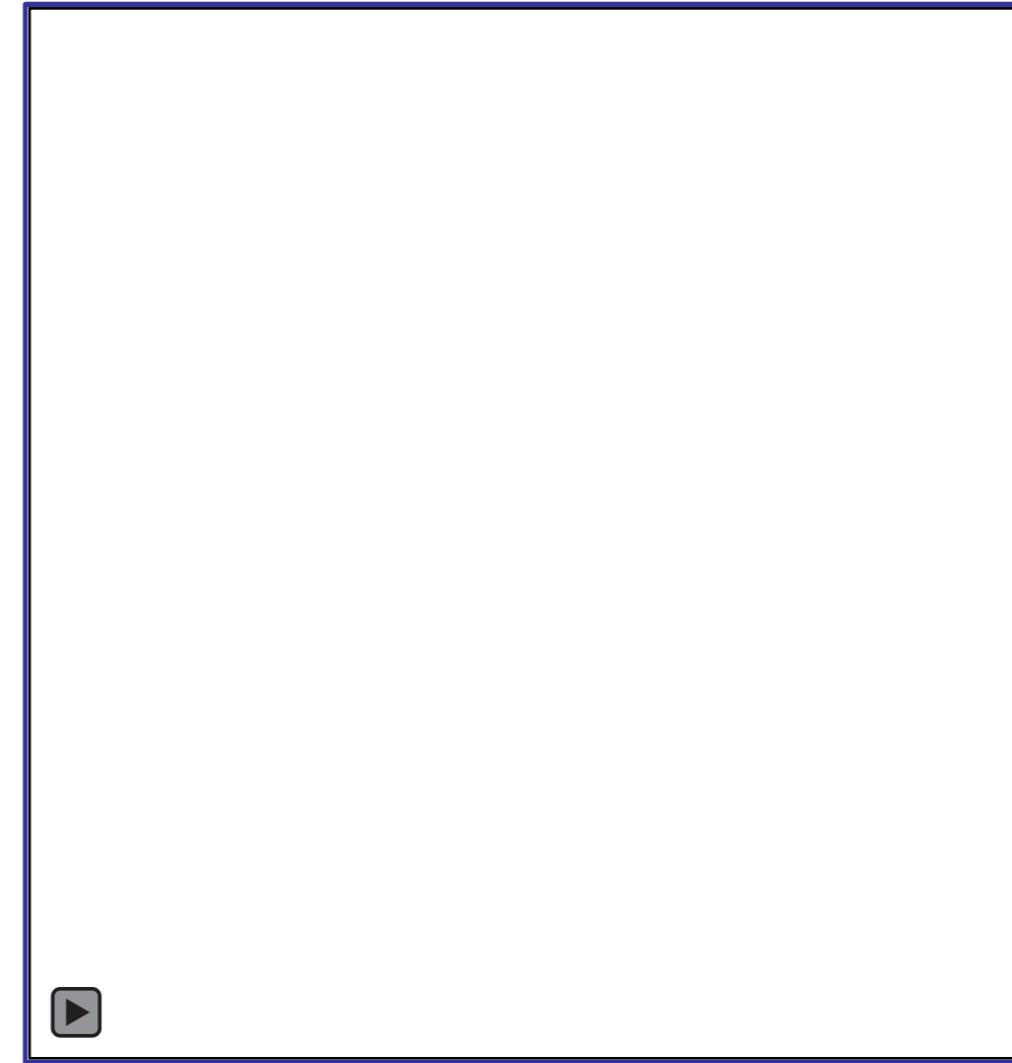


$$p(x|W) = \frac{1}{K} \sum_{i=1}^K \frac{1}{M_i} \sum_{j=1}^{M_i} \mathcal{N}(q_{ij}, \Sigma_{ij})$$

Inference with proposed methods SPARK & FLAME



SPARK works with geometrically
known environments



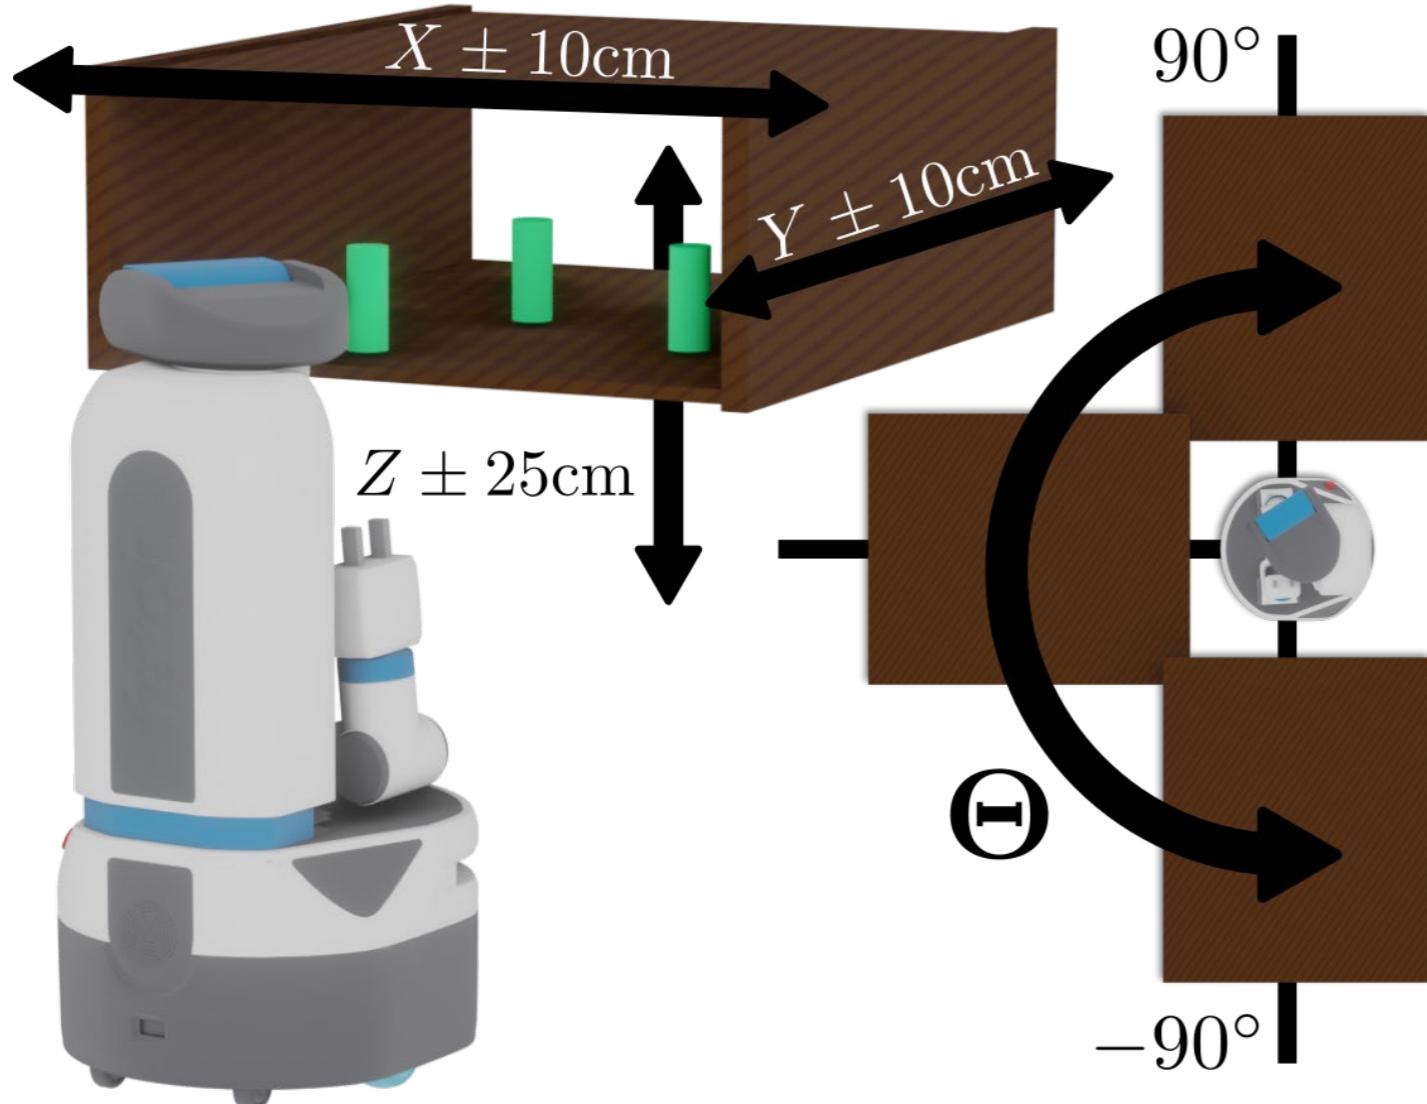
FLAME works with sensed
environments

Retrieved local
workspaces are
shown in **red**

Experiments

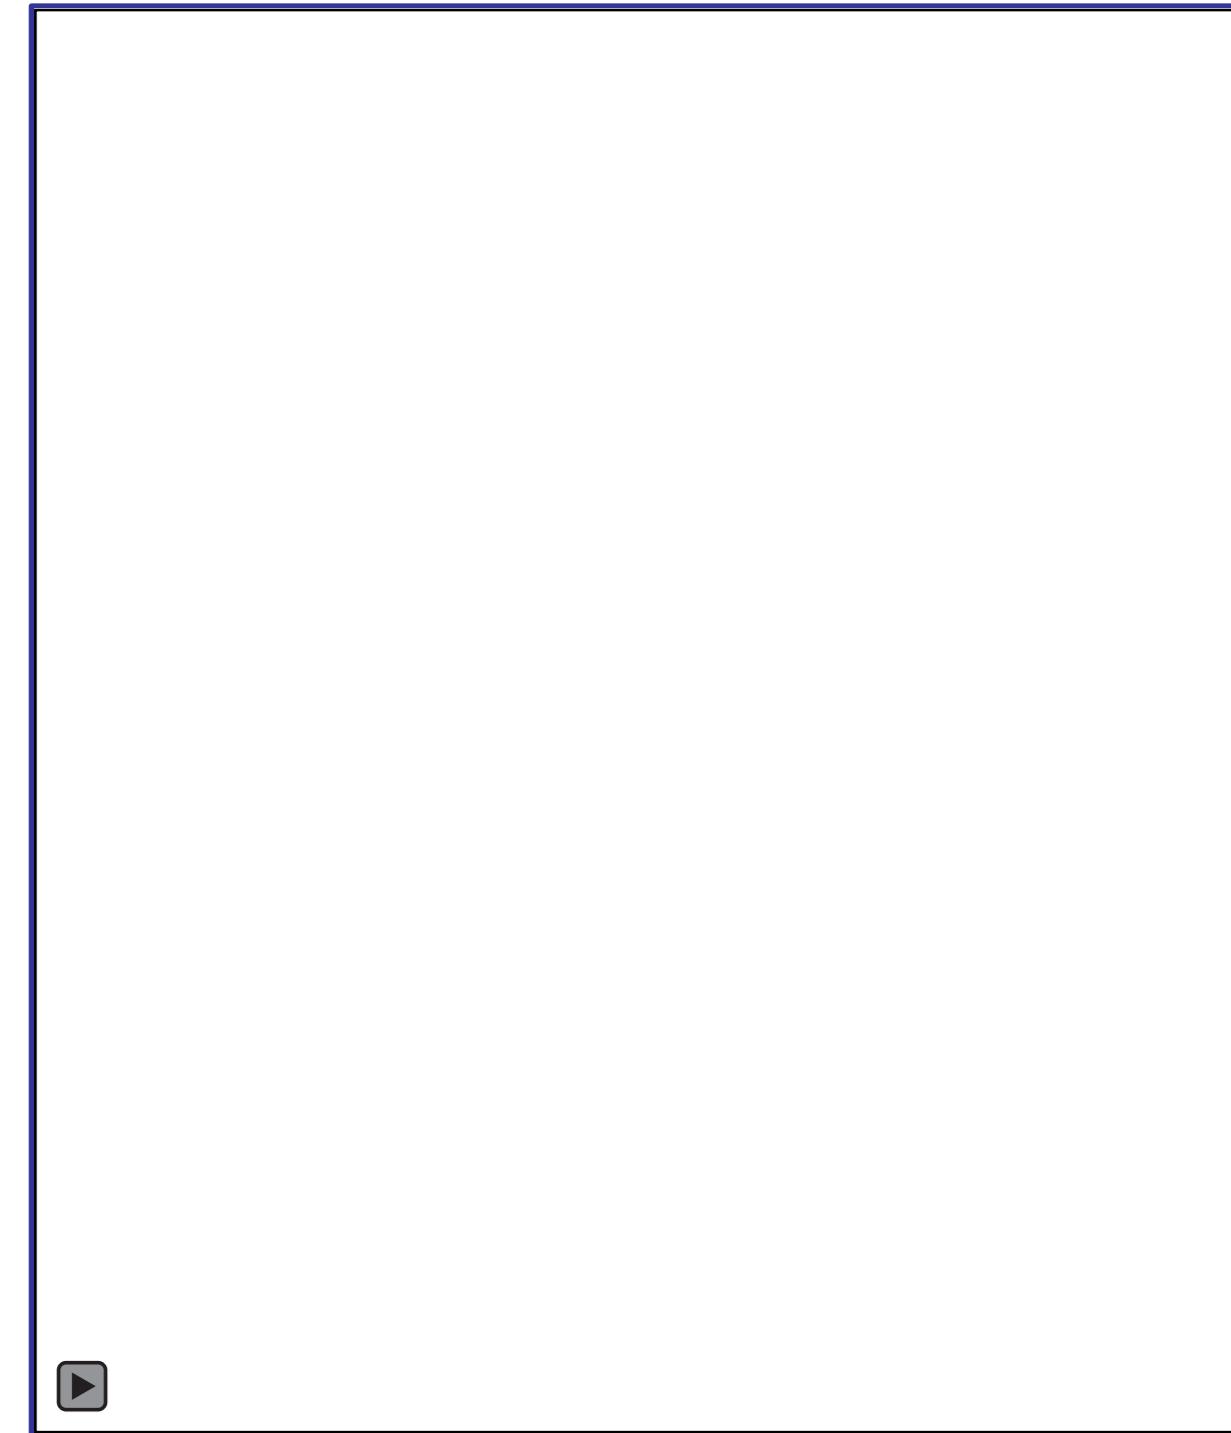
- 1. Robustness to environment variation**
2. Generalization to different tasks
3. Real robot application

SPARK/FLAME are robust to workspace variations

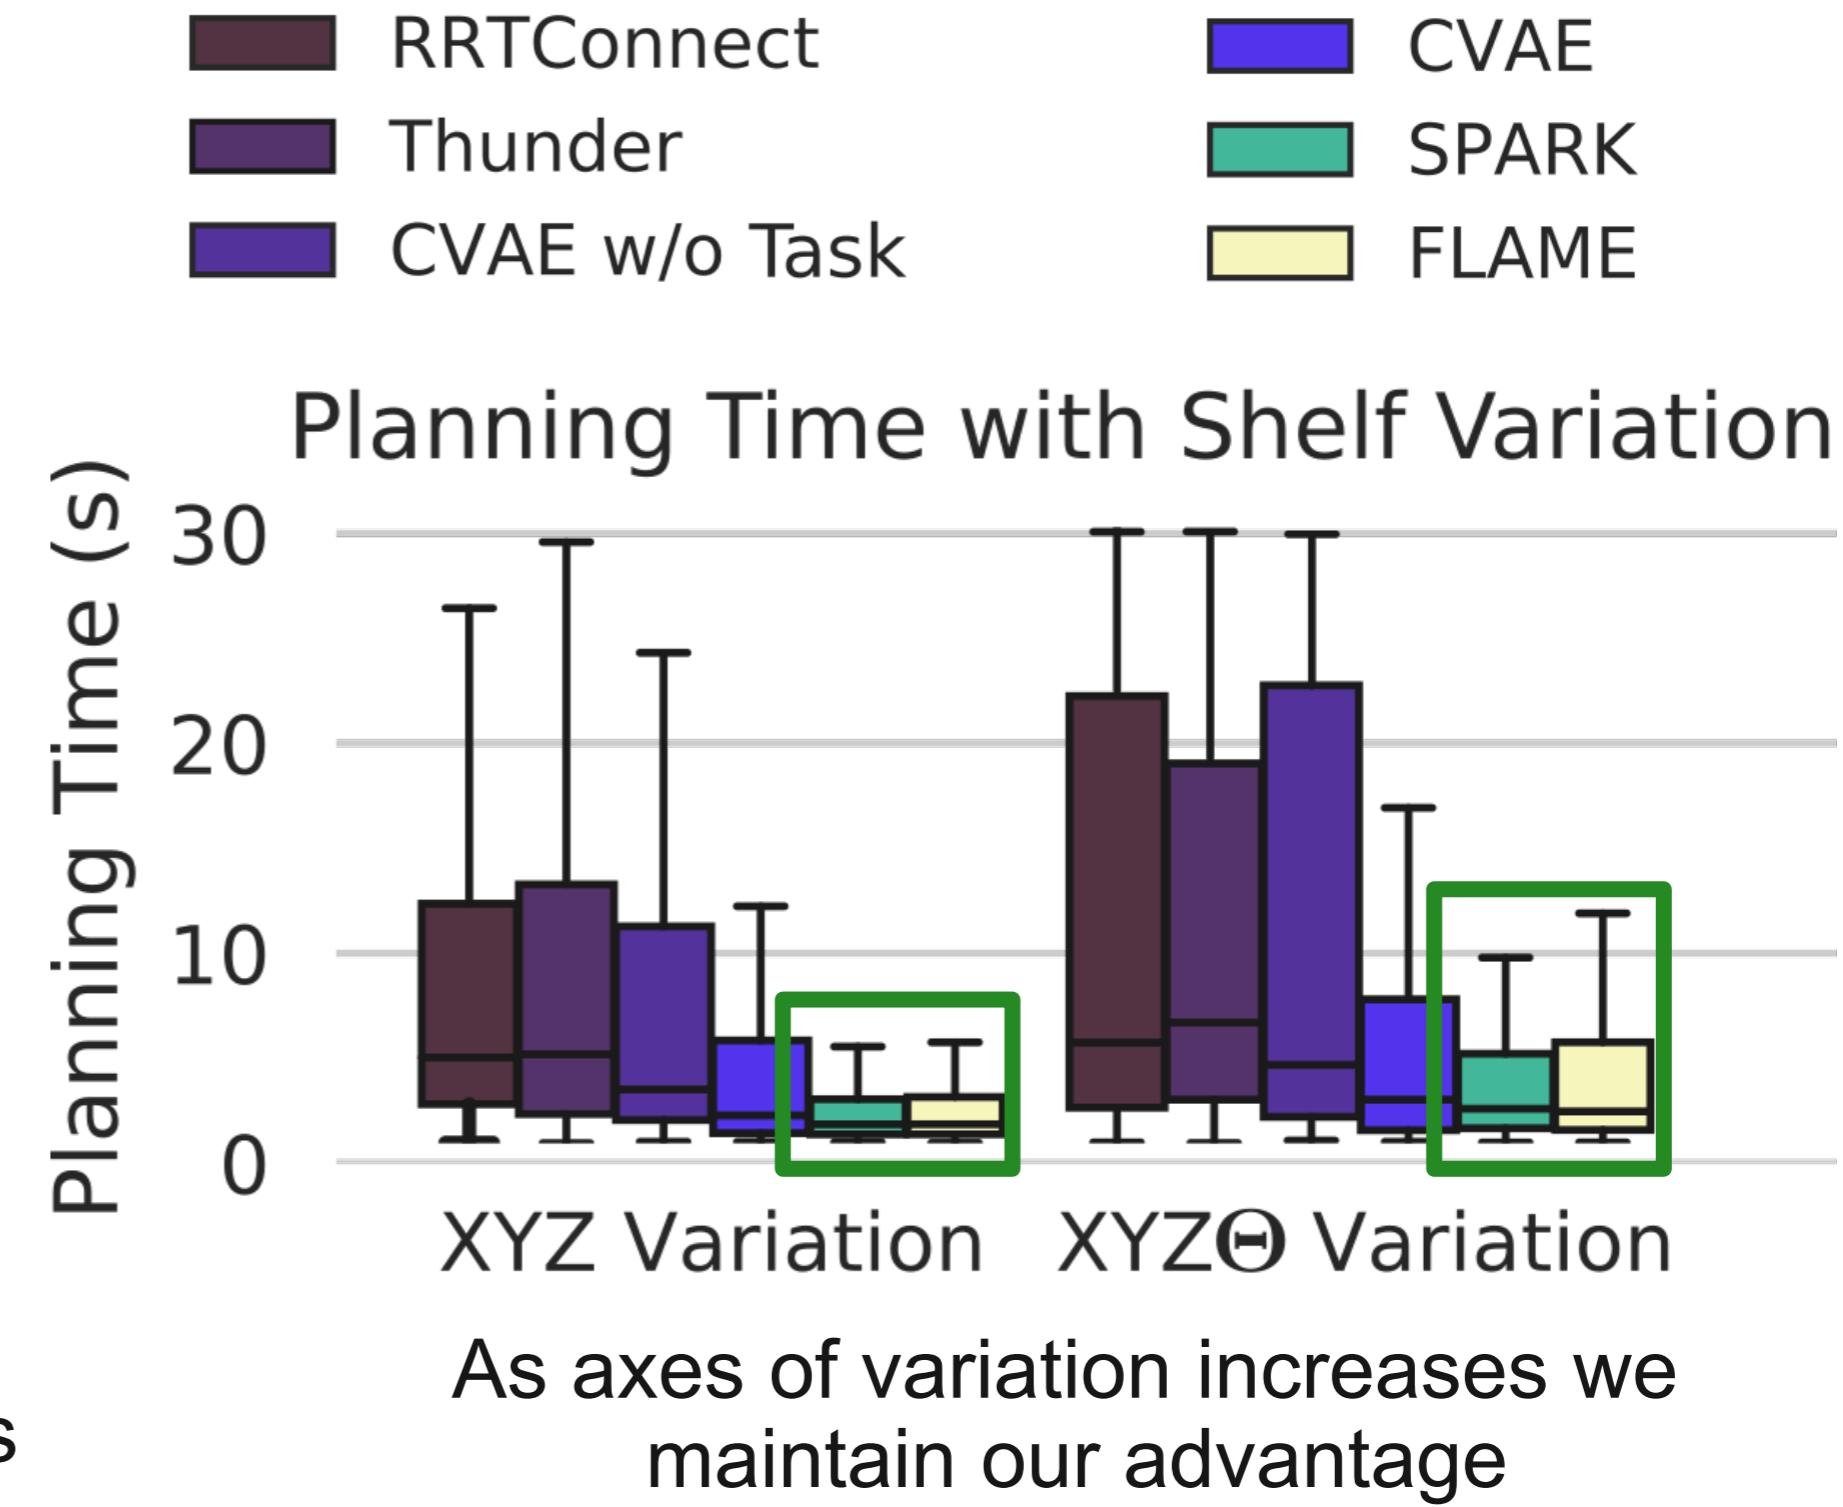


Realistic variation of shelf workspaces

SPARK/FLAME are robust to workspace variations



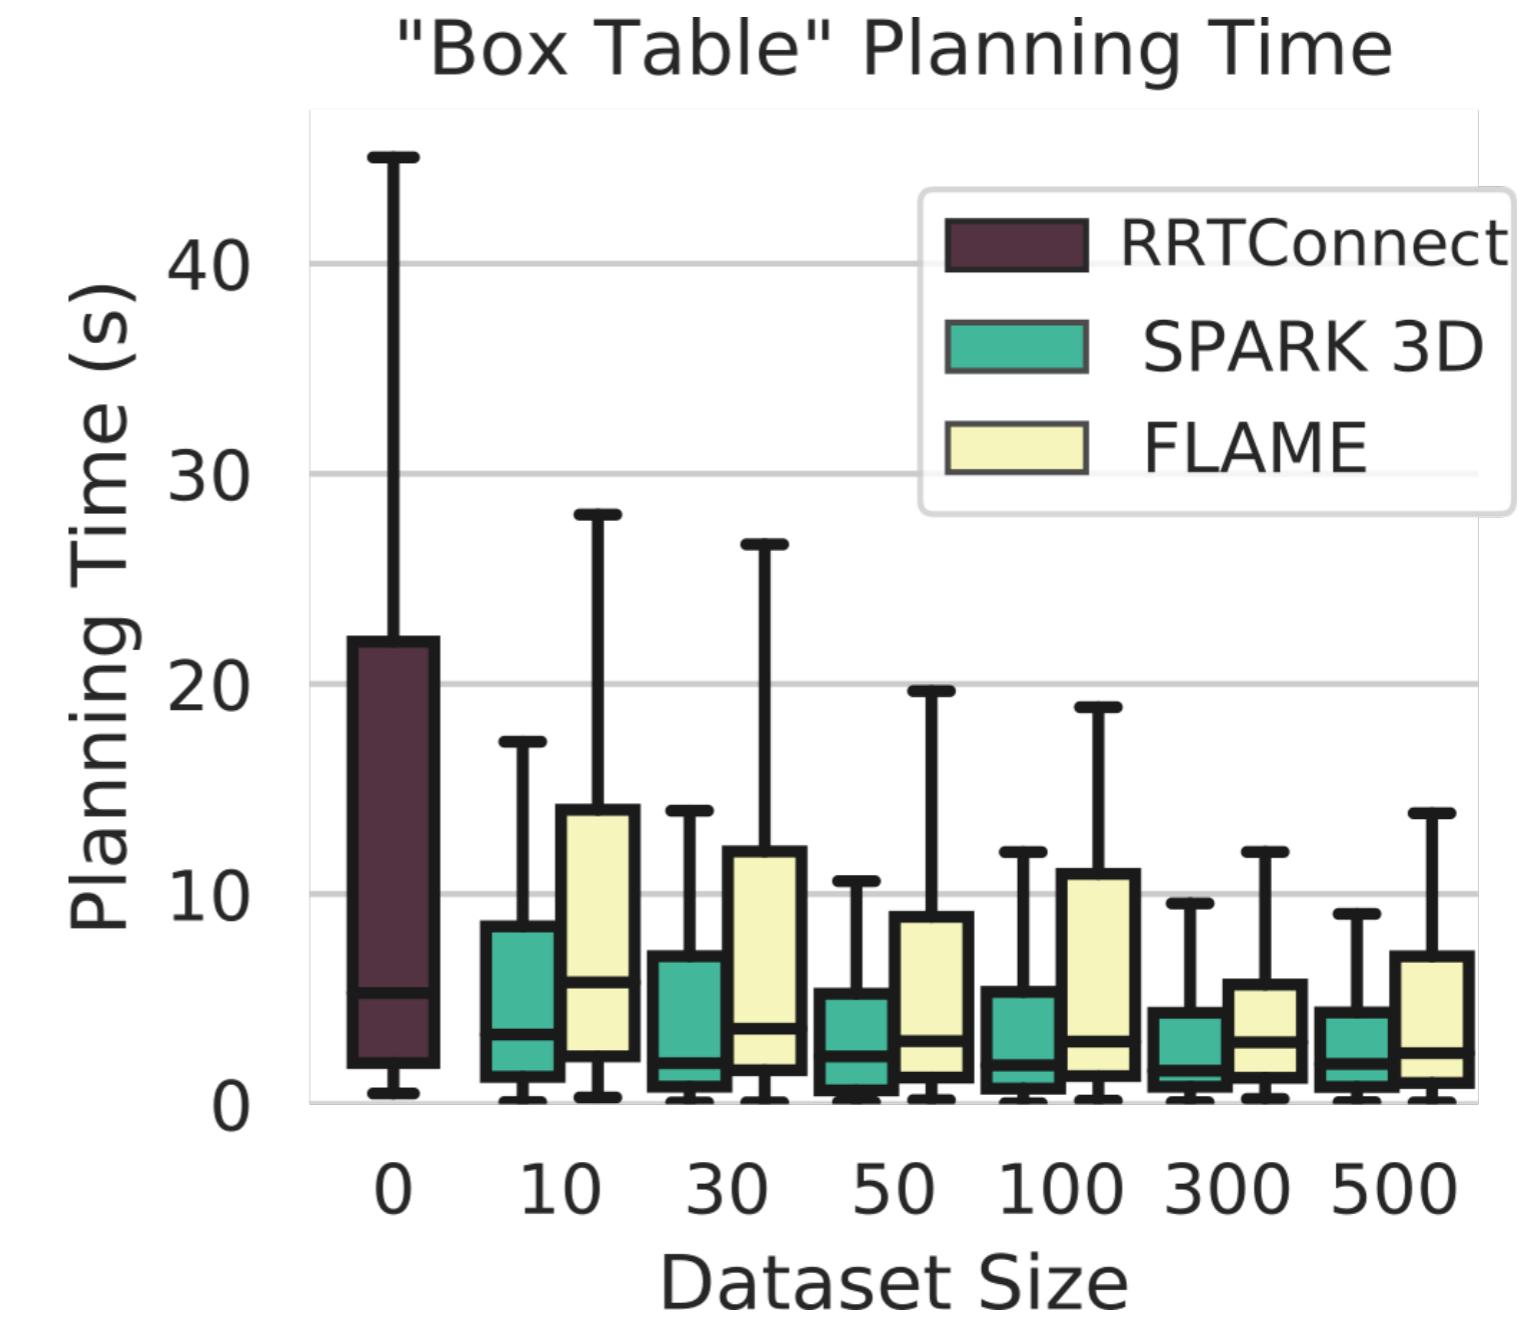
Realistic variation of shelf workspaces



SPARK/FLAME are robust to workspace variations



Realistic variation of
table workspaces



As number of experiences increase the
performance improves

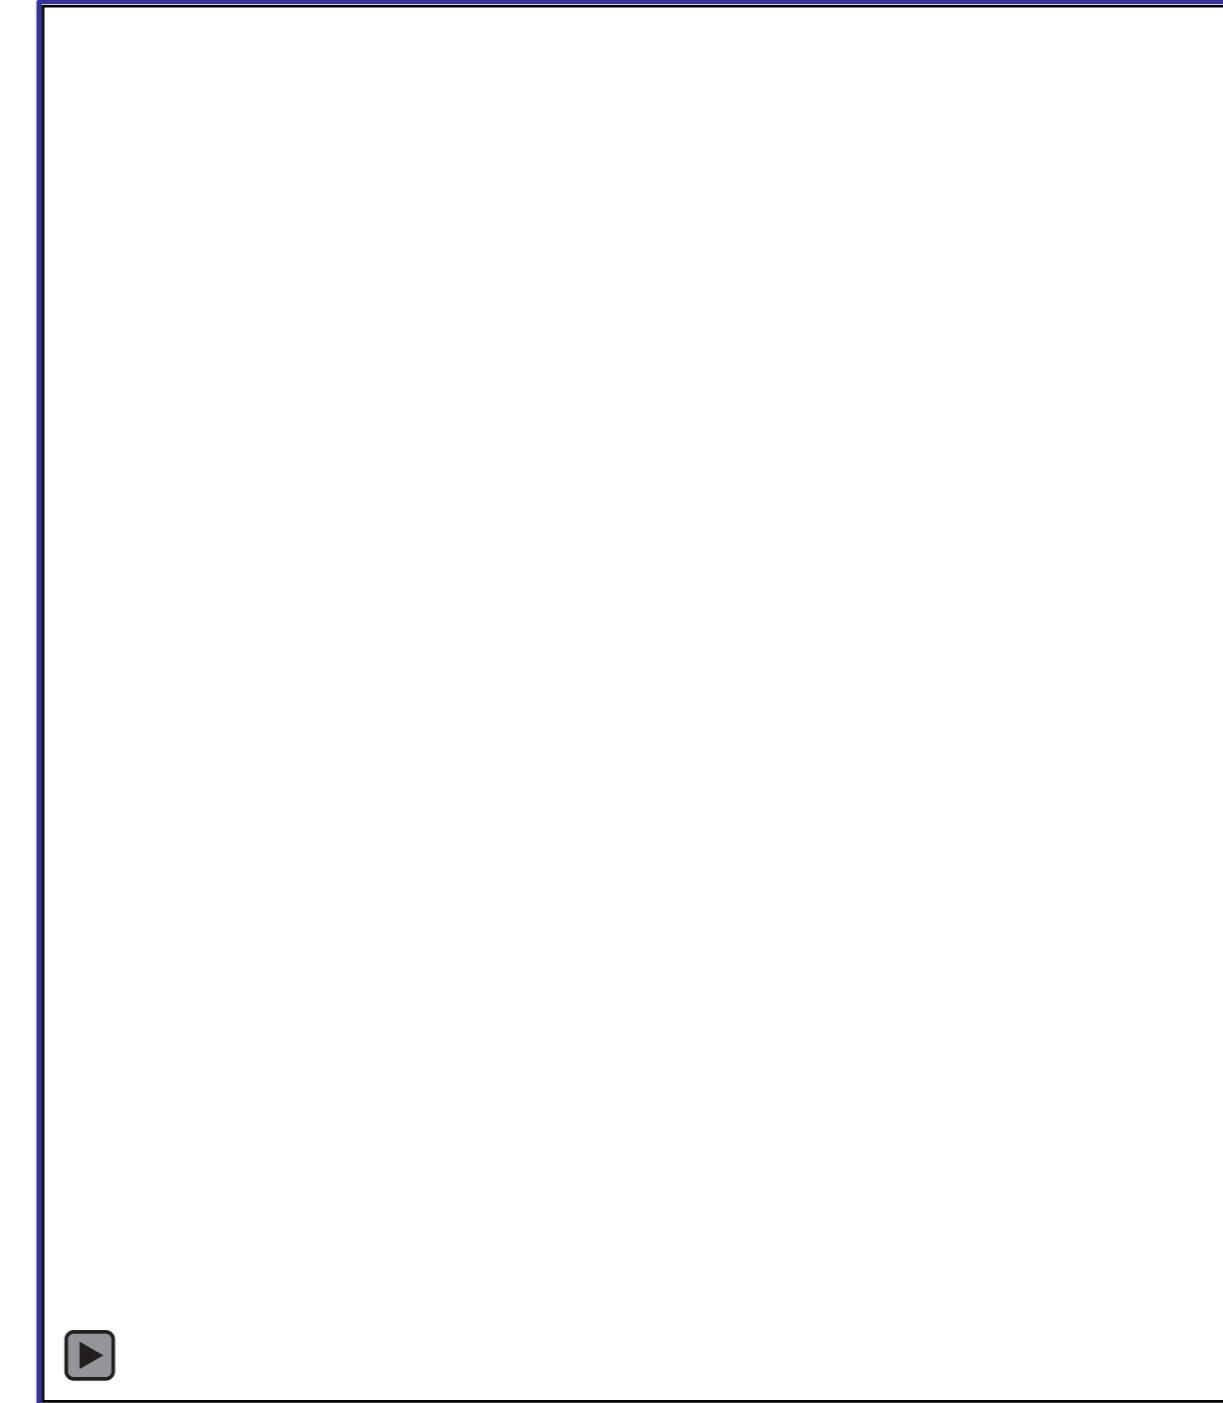
Experiments

1. Robustness to environment variation
2. Generalization to different tasks
3. Real robot application

SPARK/FLAME: Generalization to a different task



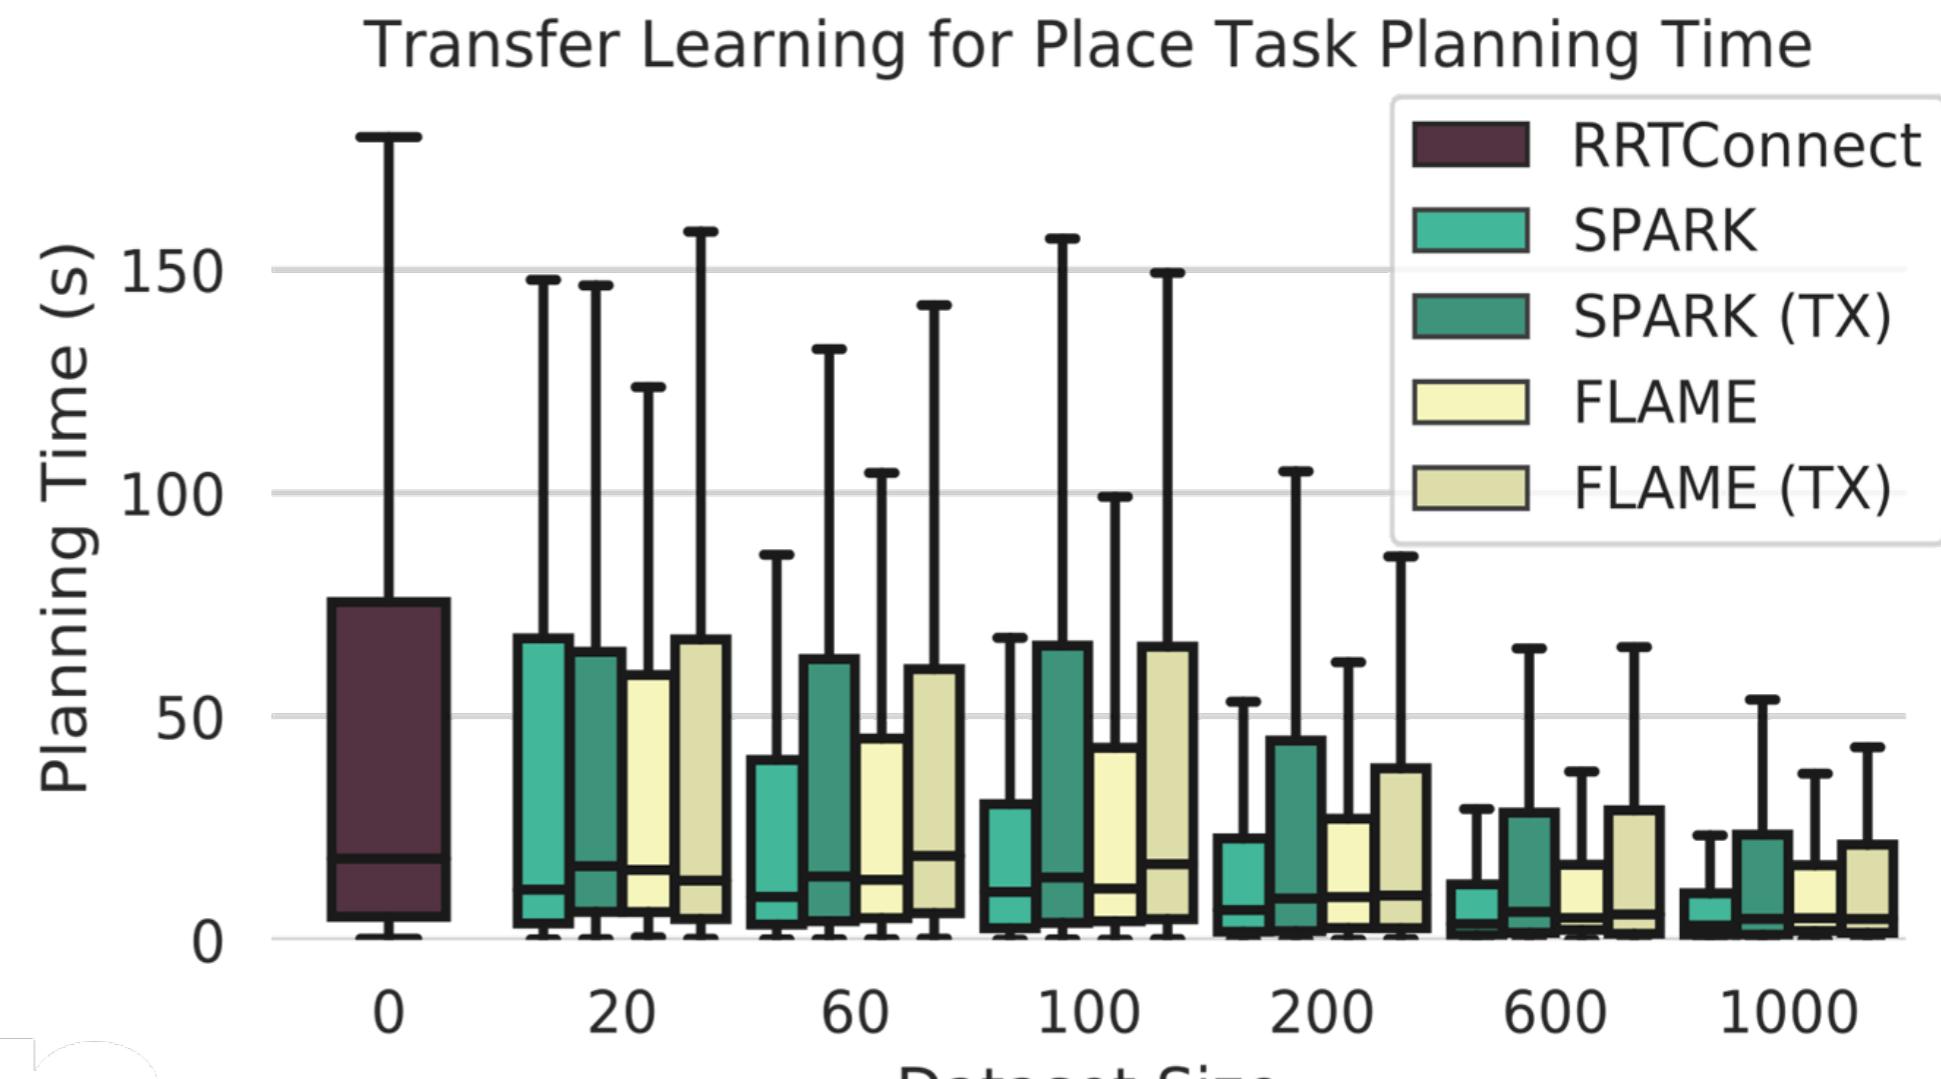
Pick Task



Place Task

SPARK/FLAME: Generalization from pick to place

Method	Train Task	Test Task
SPARK	Place	Place
FLAME	Place	Place
SPARK (TX)	Pick	Place
FLAME (TX)	Pick	Place

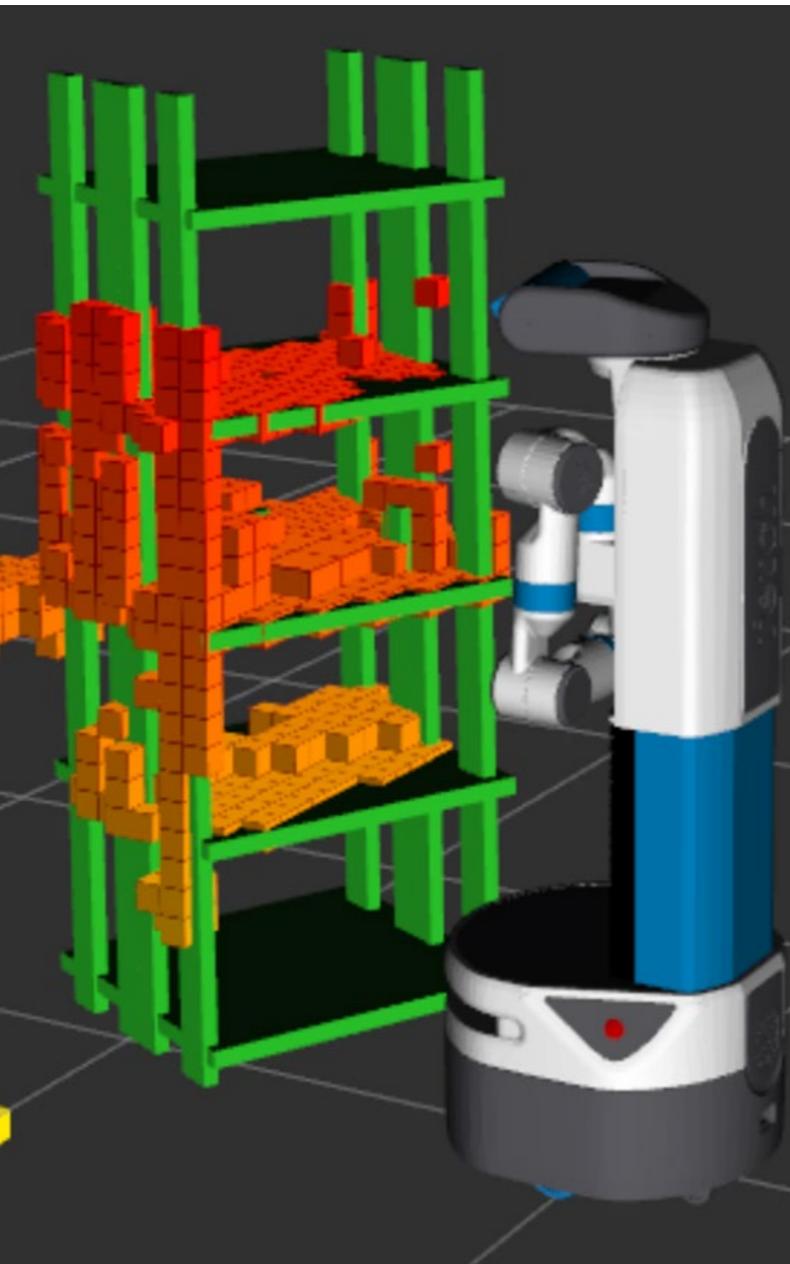


While trained only on pick tasks we generalize
to the place task as well

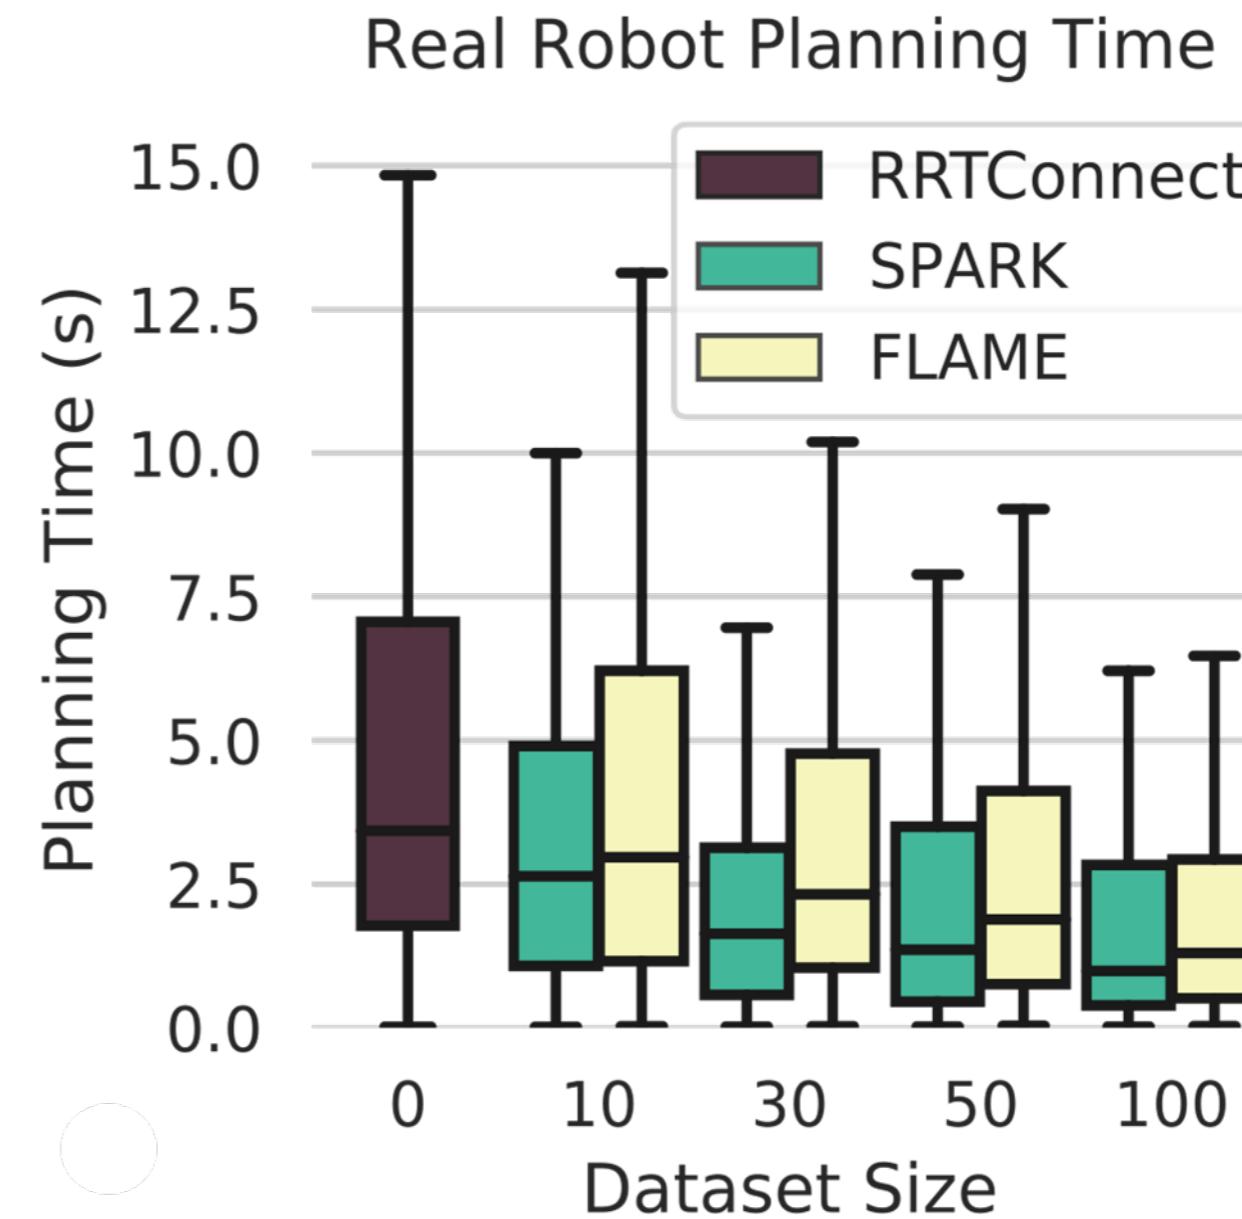
Experiments

- 1. Robustness to environment variation**
- 2. Generalization to different Tasks**
- 3. Real robot application**

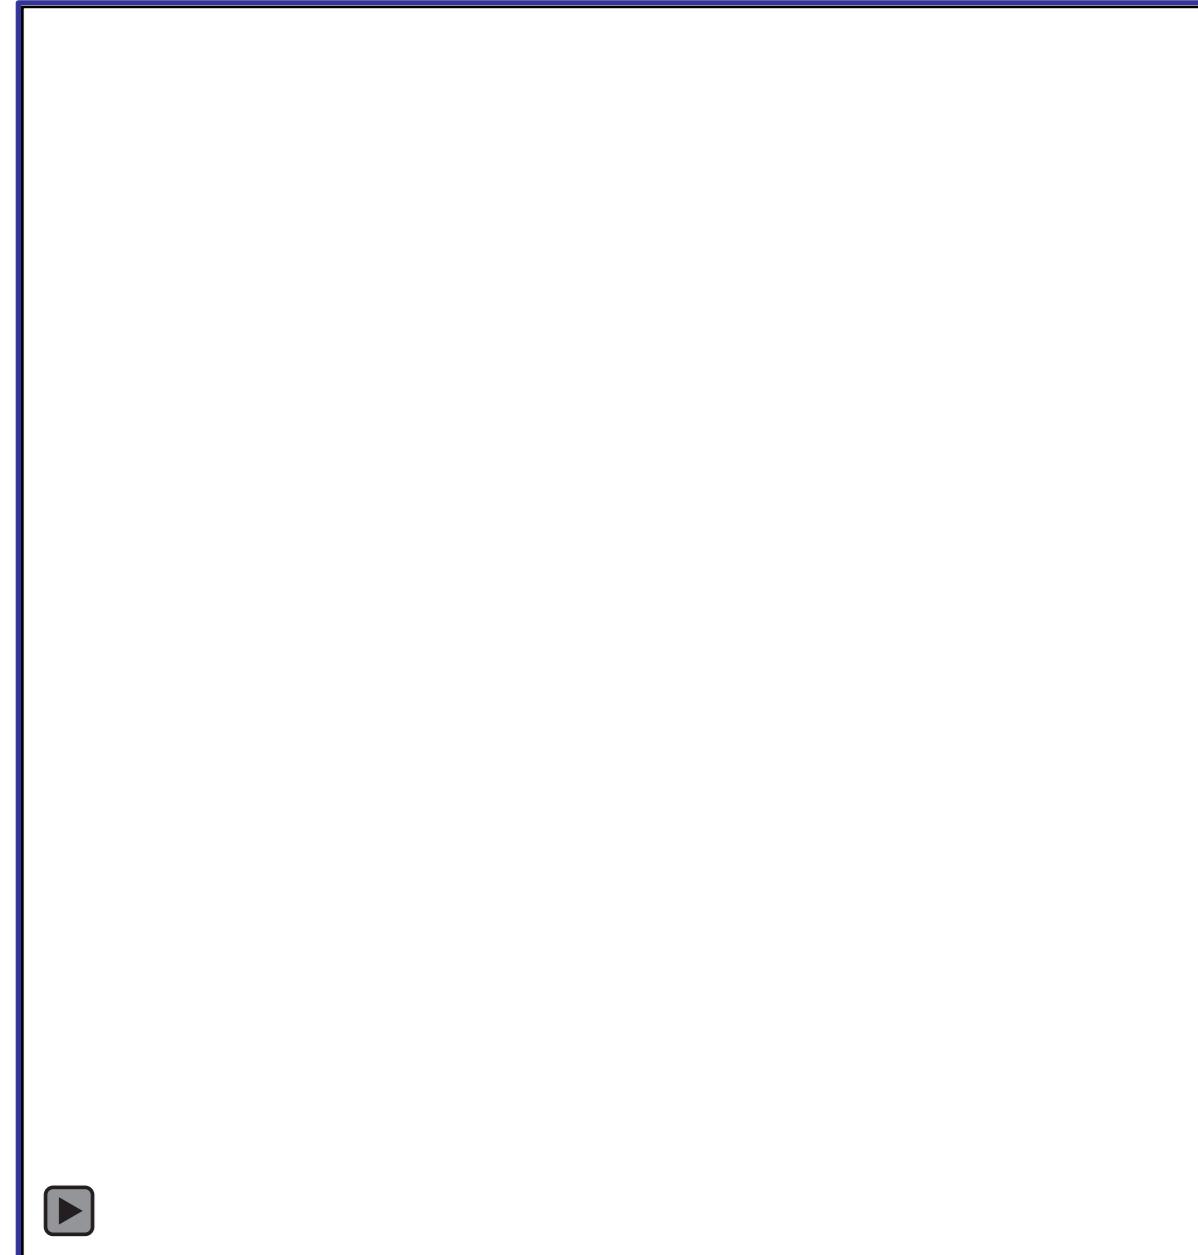
SPARK/FLAME: Real Fetch experiment



Application in a
real setting



SPARK/FLAME: Real Fetch experiment



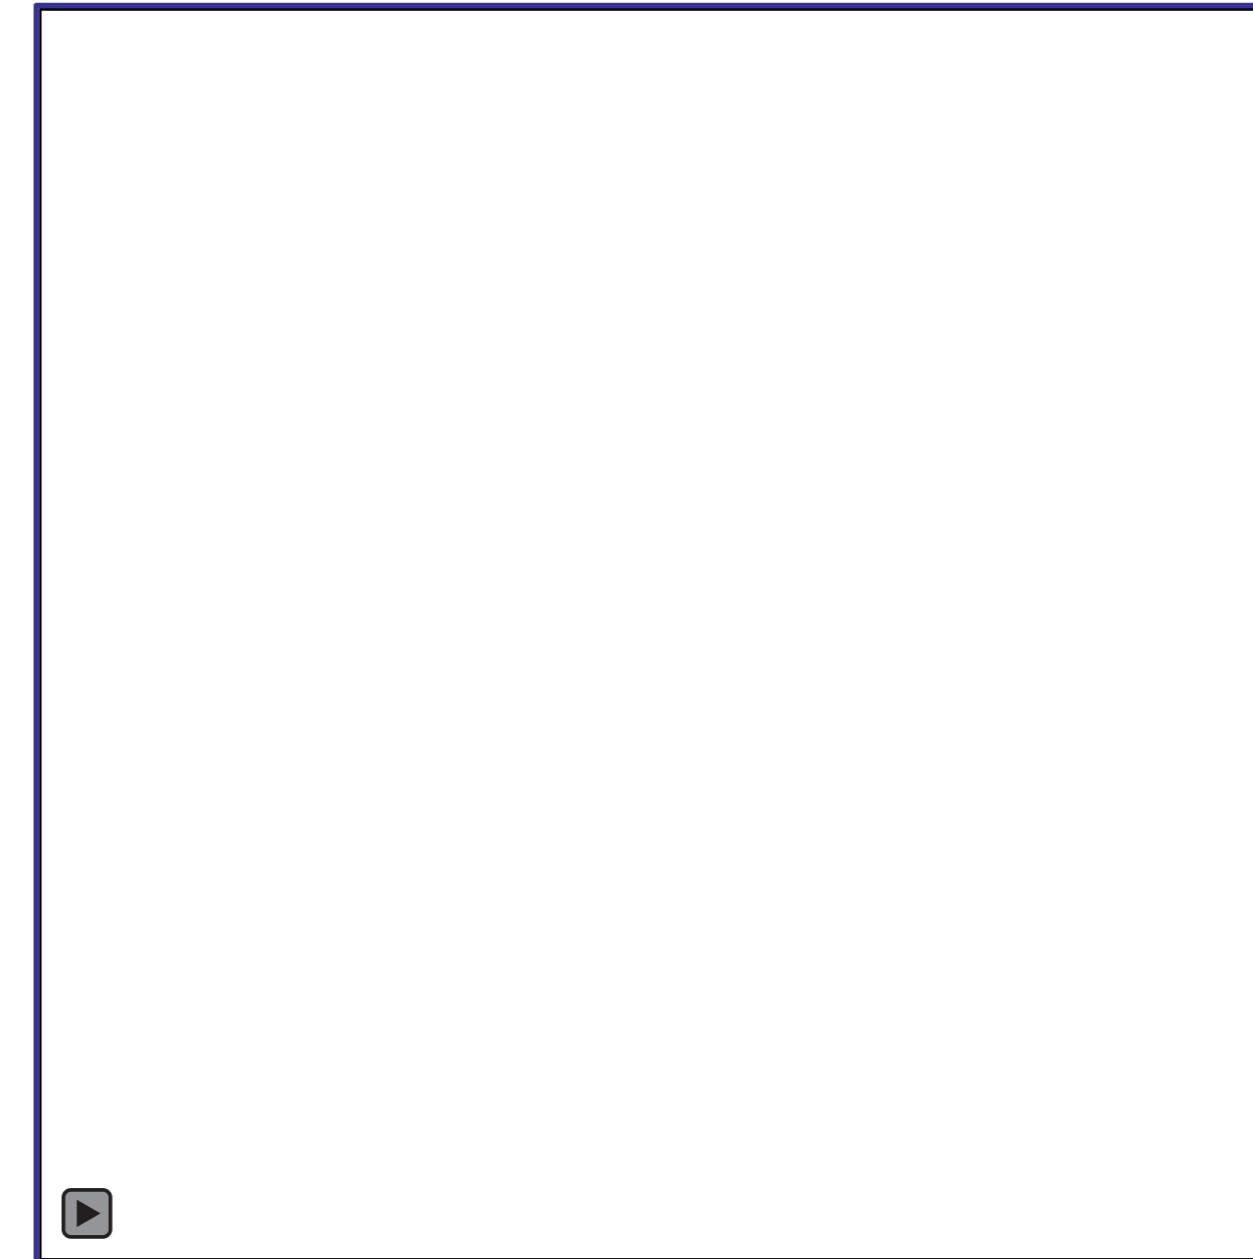
Method	Planning time
SPARK	~3.4 faster
FLAME	~2.6 faster

Example solution path
of pick and tuck

SPARK/FLAME: Real Fetch experiment

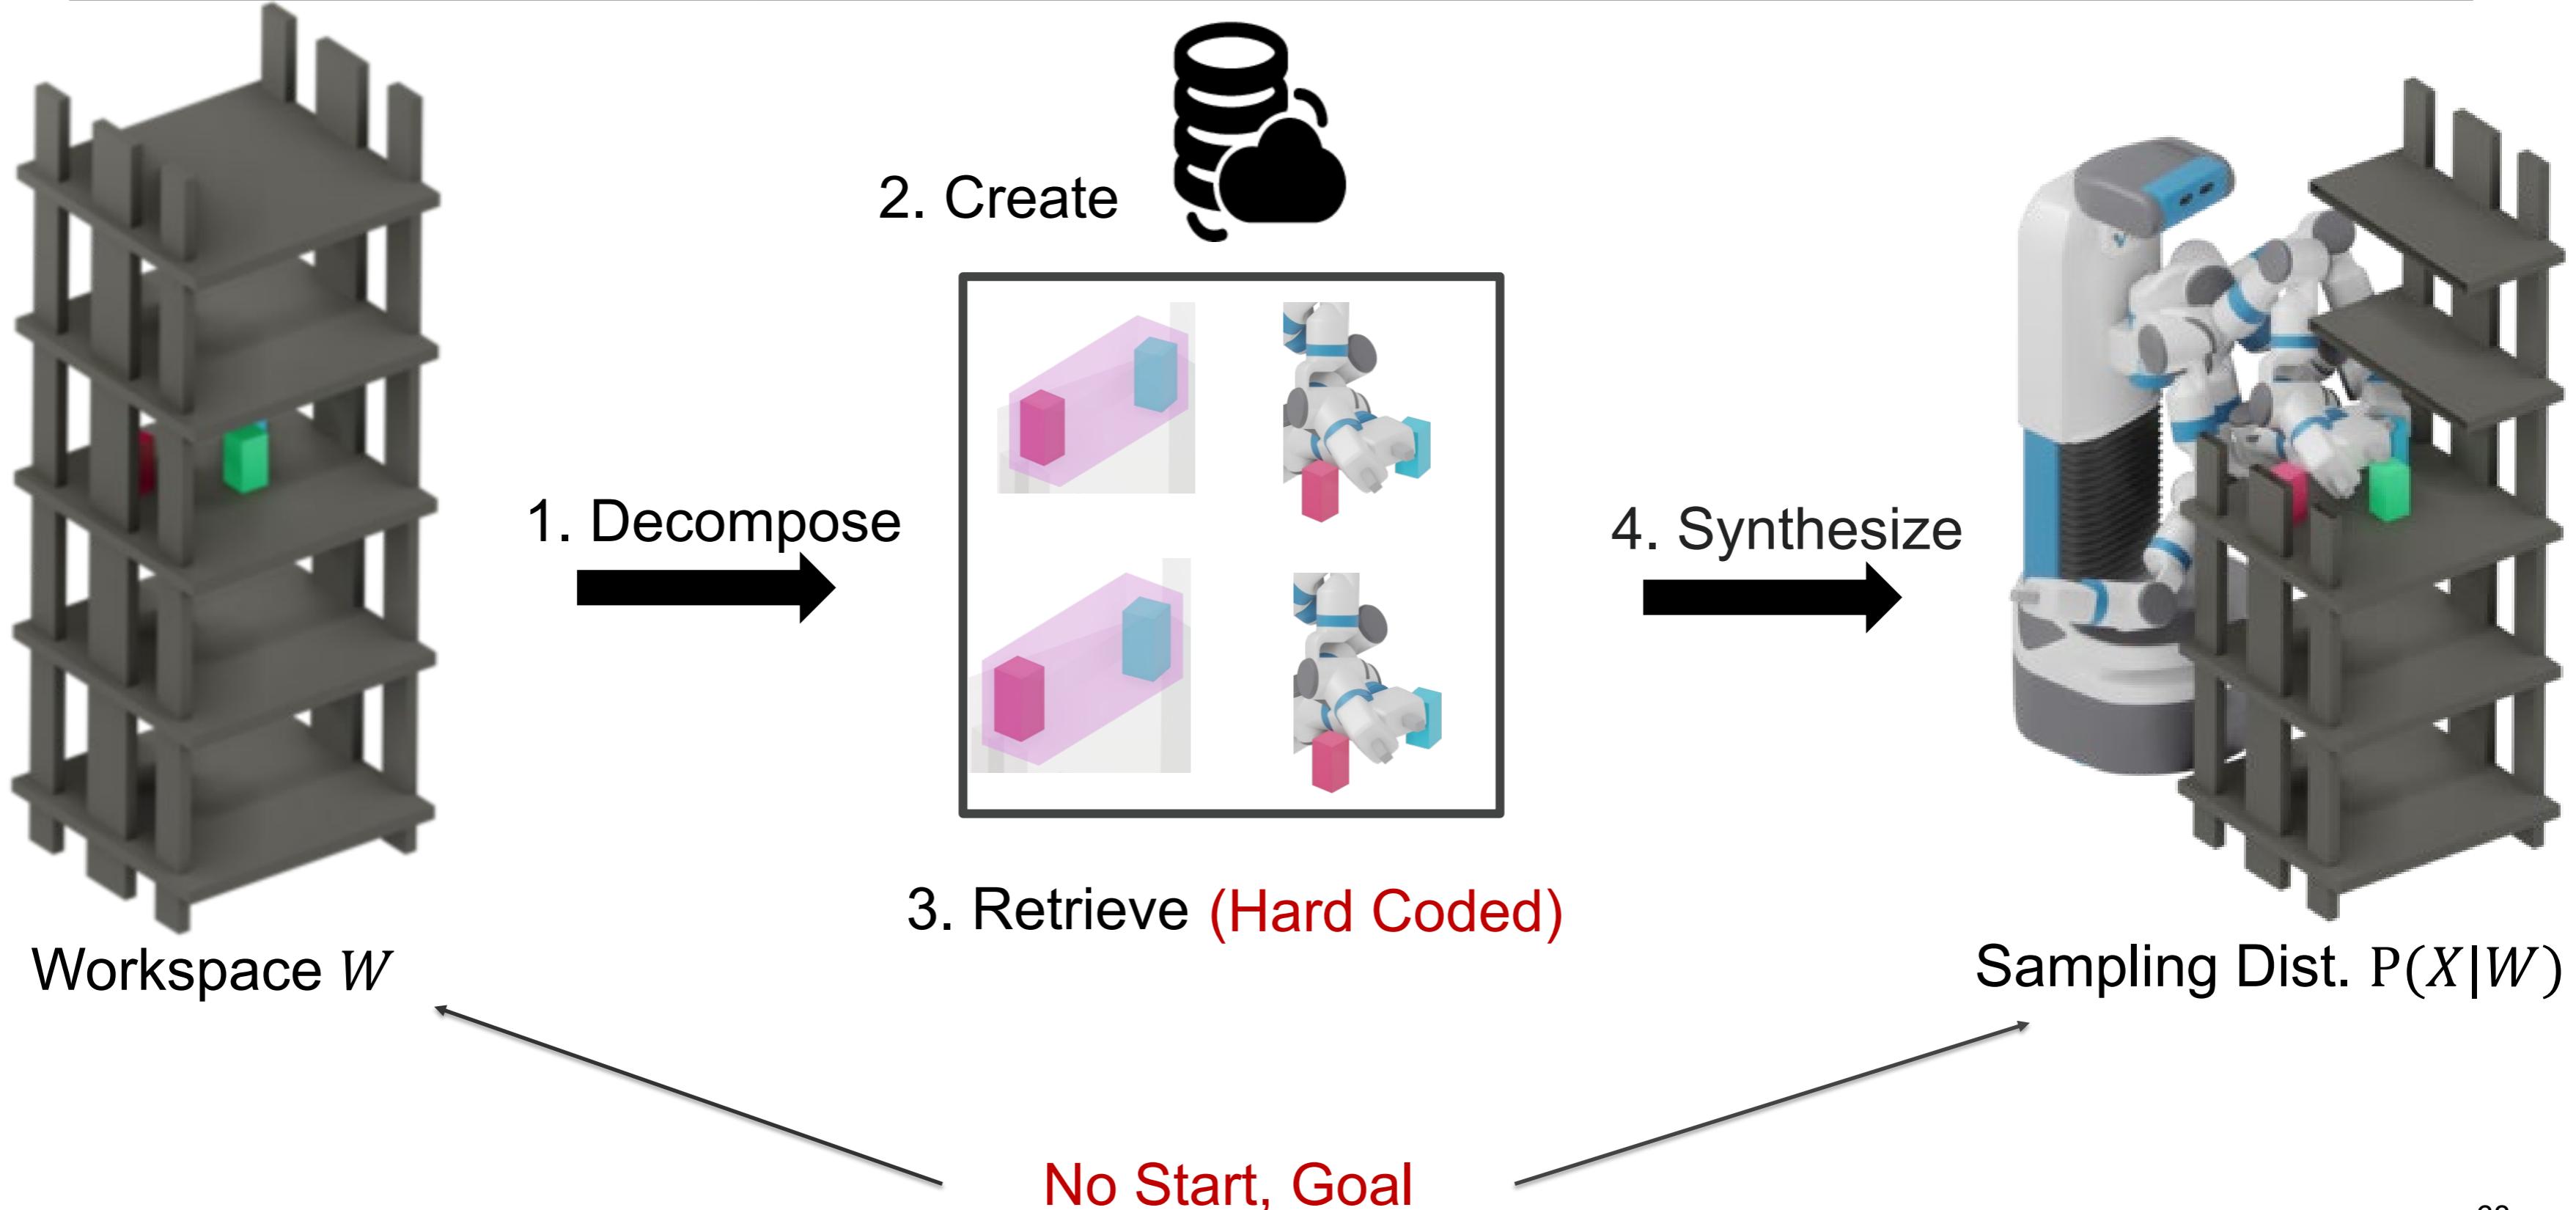


RRT-Connect returns only
approximate solution in 2s



Solution path

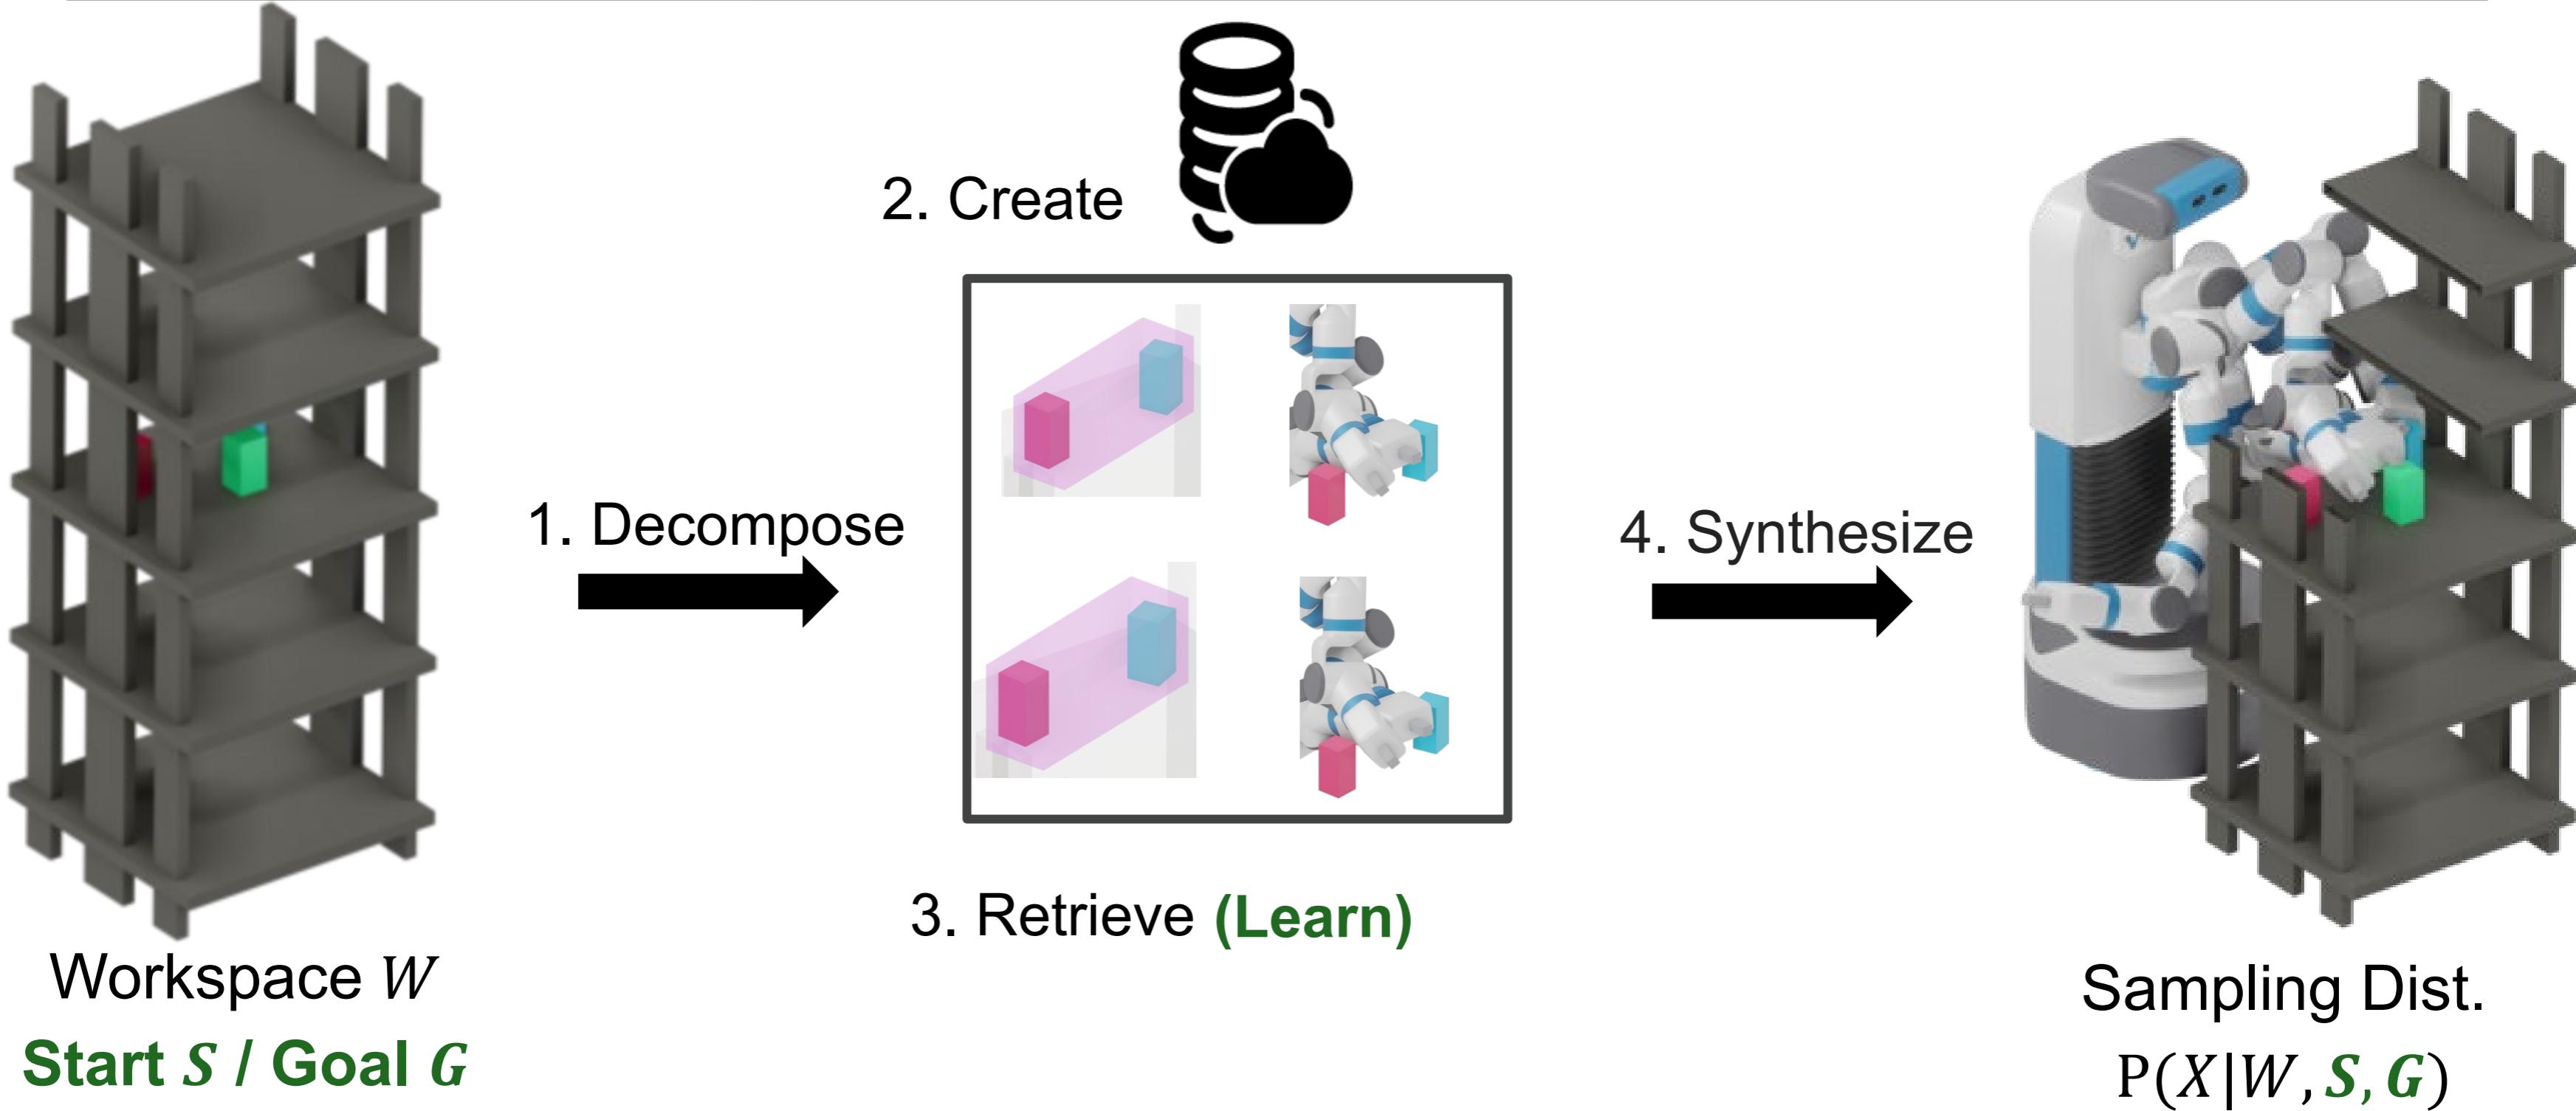
Limitations of framework



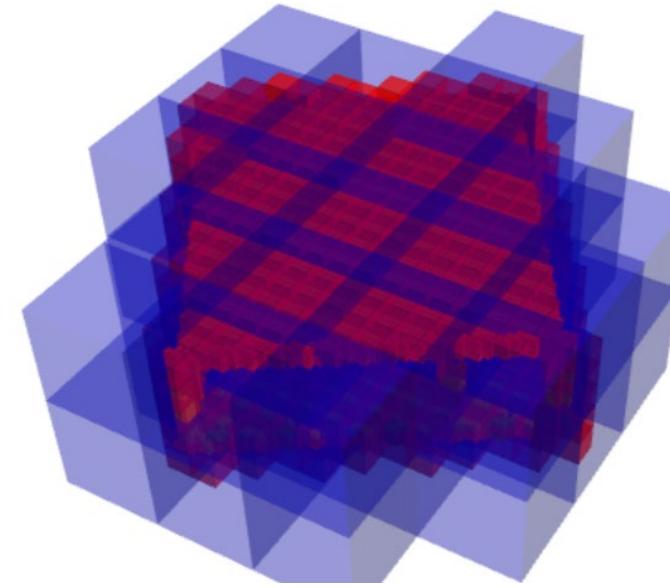
Today's Agenda

- Important Considerations
 - Locality of narrow passages
 - Discontinuity
- Biased Samplers (Cont.)
 - 2D Workspaces
 - 3D Workspaces
 - **Workplace, Start and Goal**

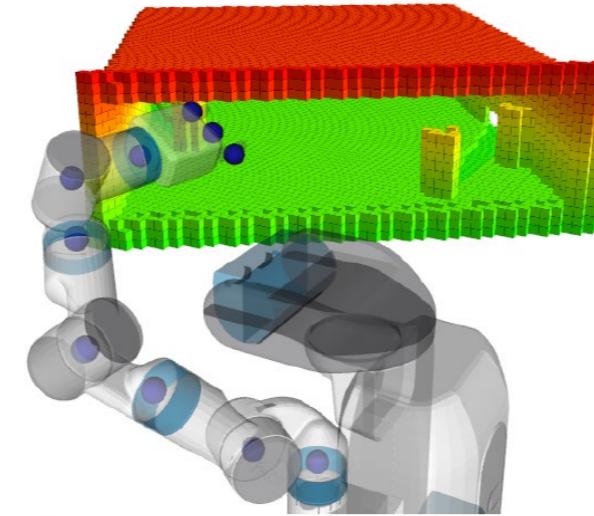
Learning the similarity function



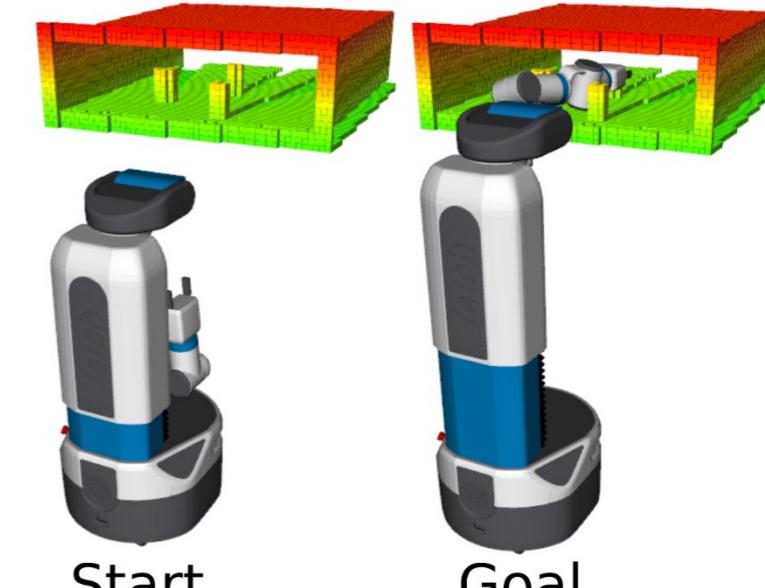
Local representations and similarity datasets



Local Features



Intermediate Features



Start

Goal

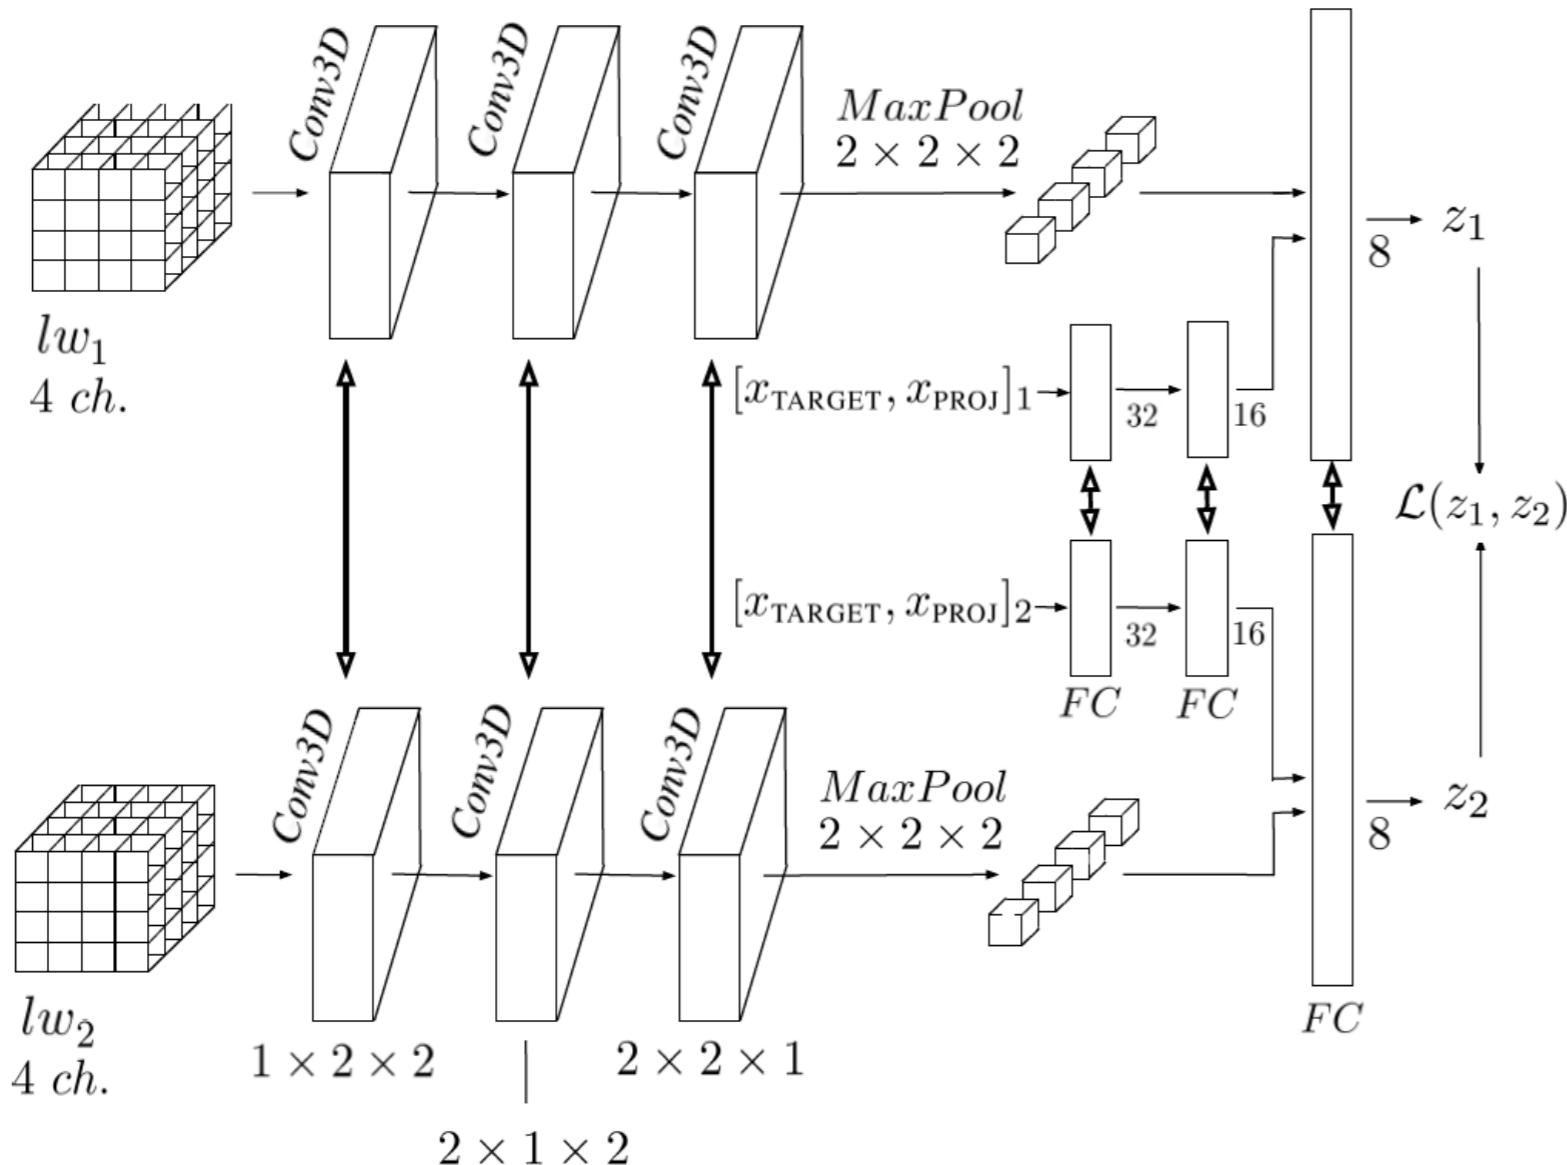
Global Features

Generating a dataset of similar/dissimilar pairs of local workspaces

Similar pairs: Locally compare past solution paths

Dissimilar pairs: Randomly sample local workspaces

Learn the similarity with a Siamese Network



Siamese Network architecture

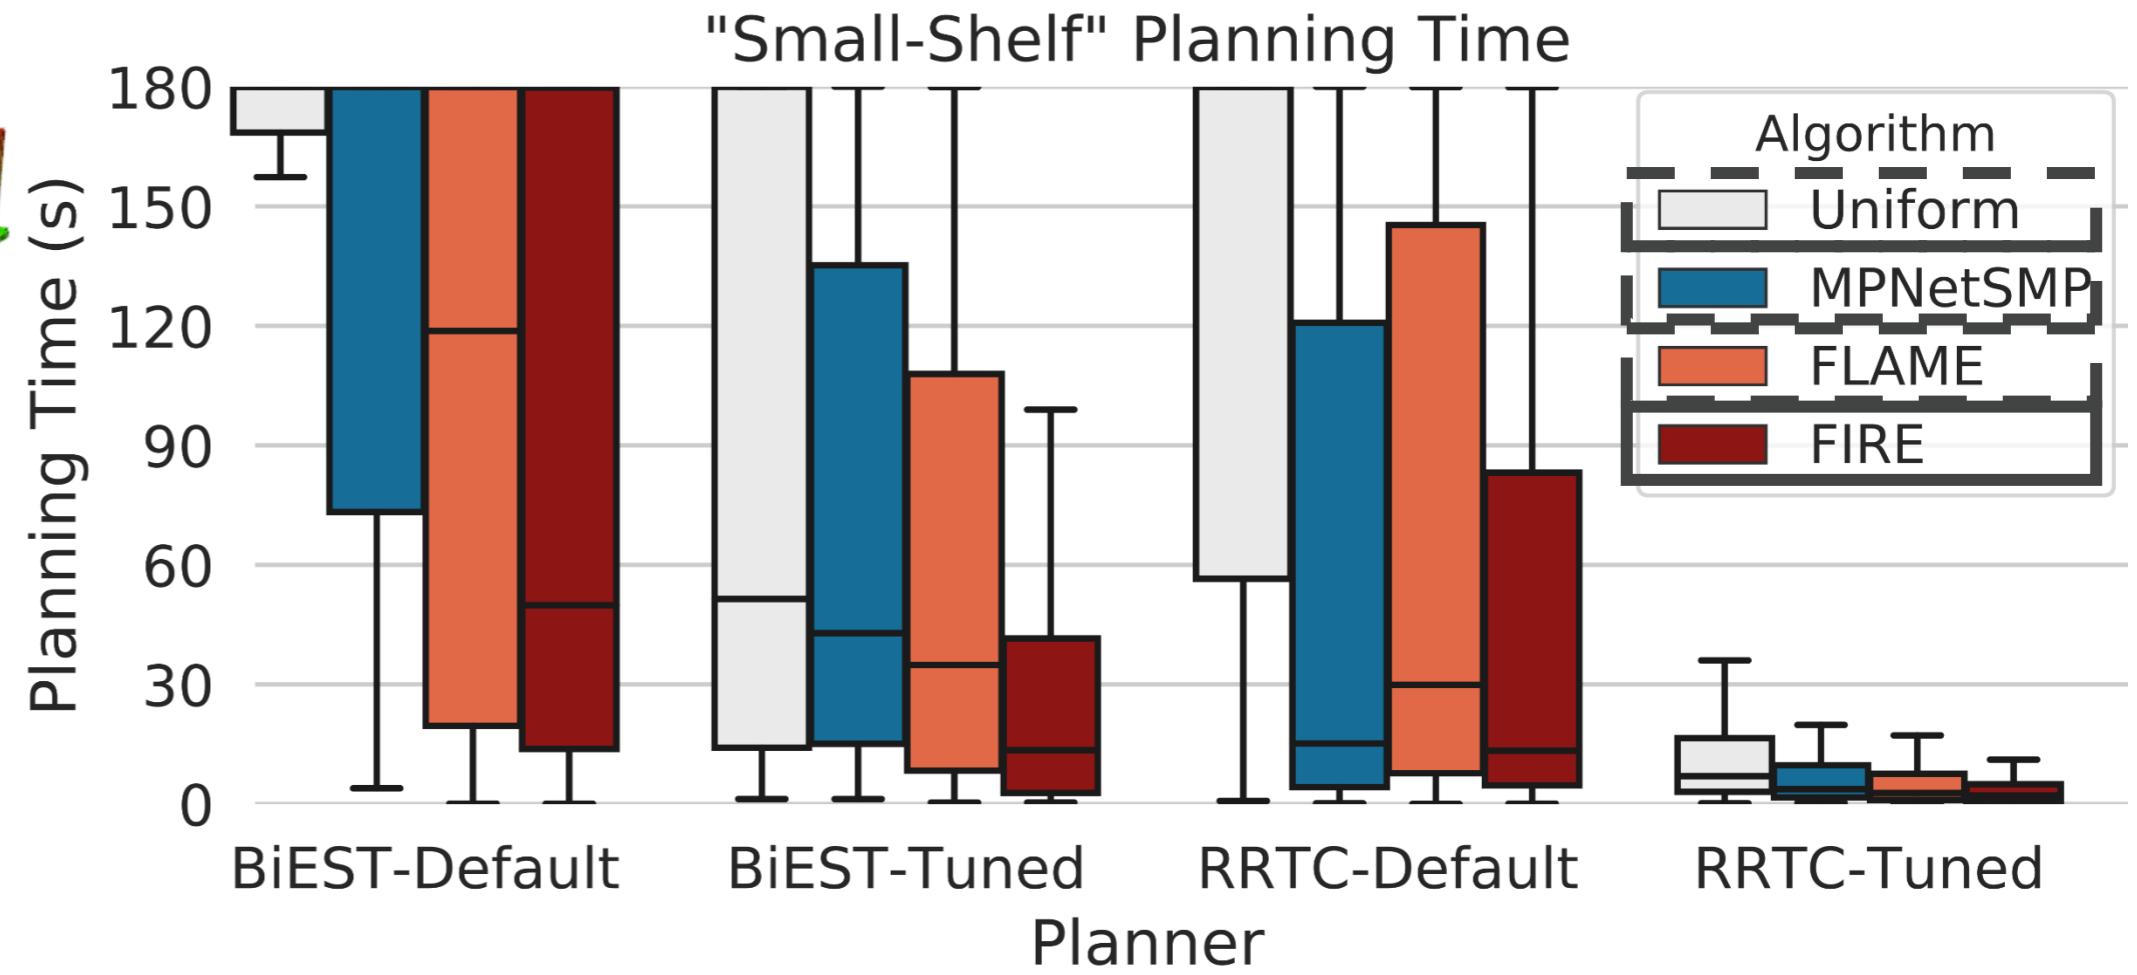
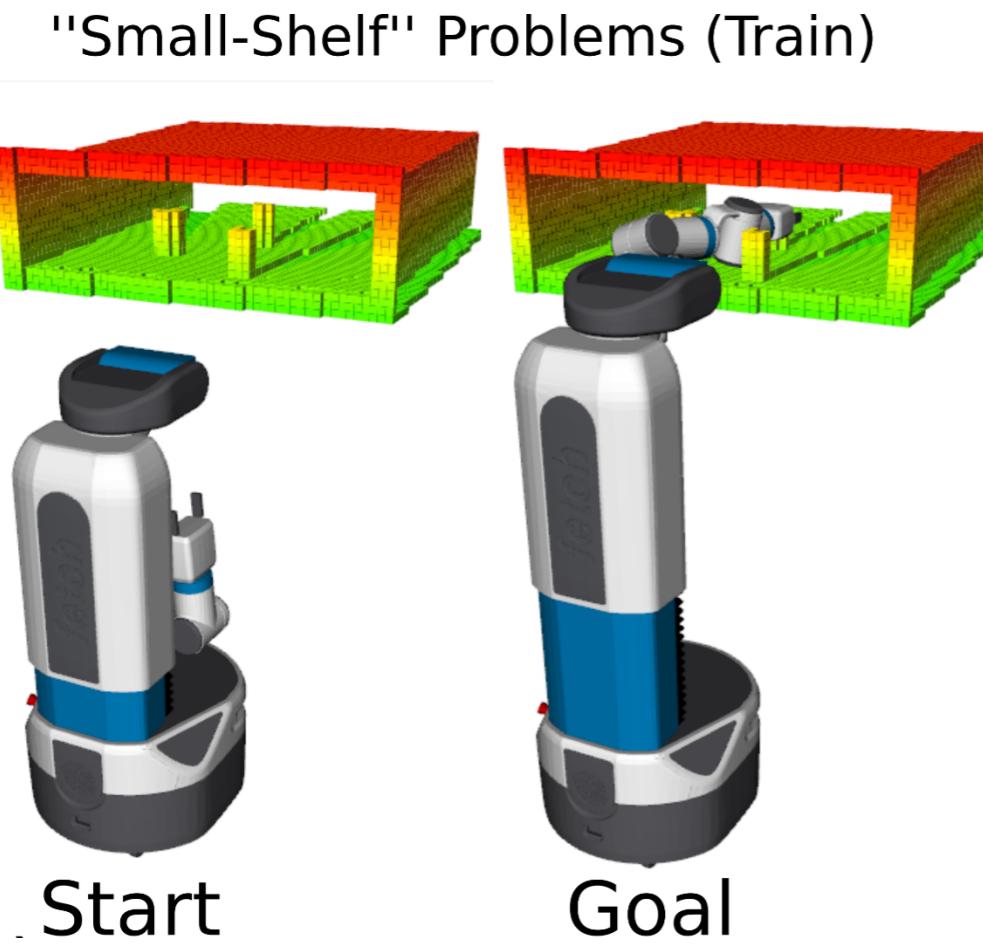
$$\mathcal{L}(\ell_i, \ell_j) = \begin{cases} \max(0, d_m - \|z_i - z_j\|^2) & \text{if } \langle \ell_j, \ell_i \rangle \in \mathcal{DS} \\ \|z_i - z_j\|^2 & \text{if } \langle \ell_j, \ell_i \rangle \in \mathcal{S} \end{cases}$$

Contrastive loss function

$$SIM(\ell_i, \ell_j) = \begin{cases} 1 & \text{if } \|z_i - z_j\|^2 < R \\ 0 & \text{otherwise} \end{cases}$$

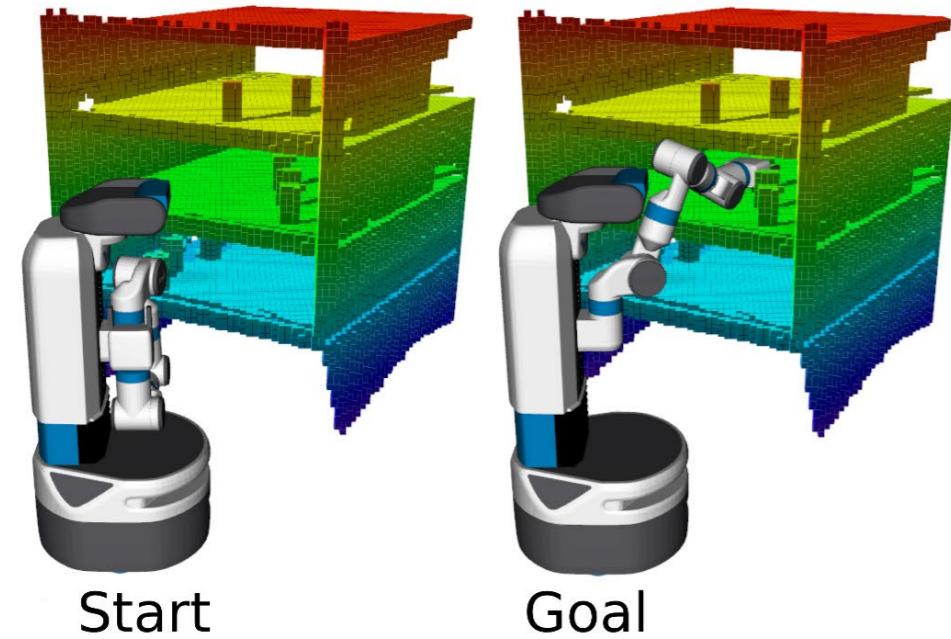
Realized similarity function

Generalizing within the training class of problems

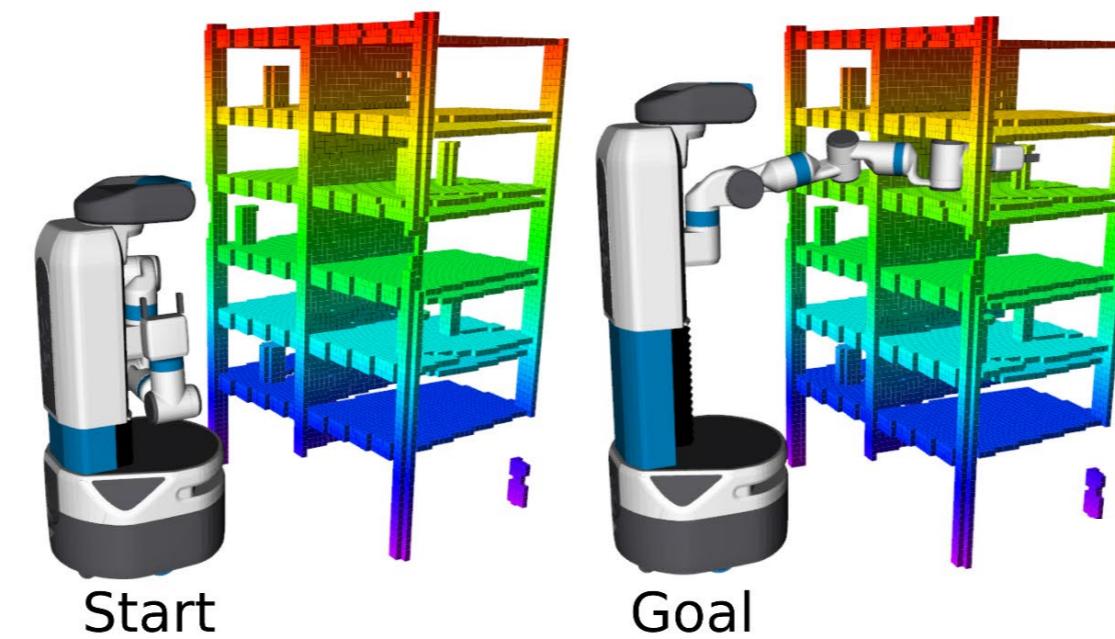


Generalizing outside the training class of problems

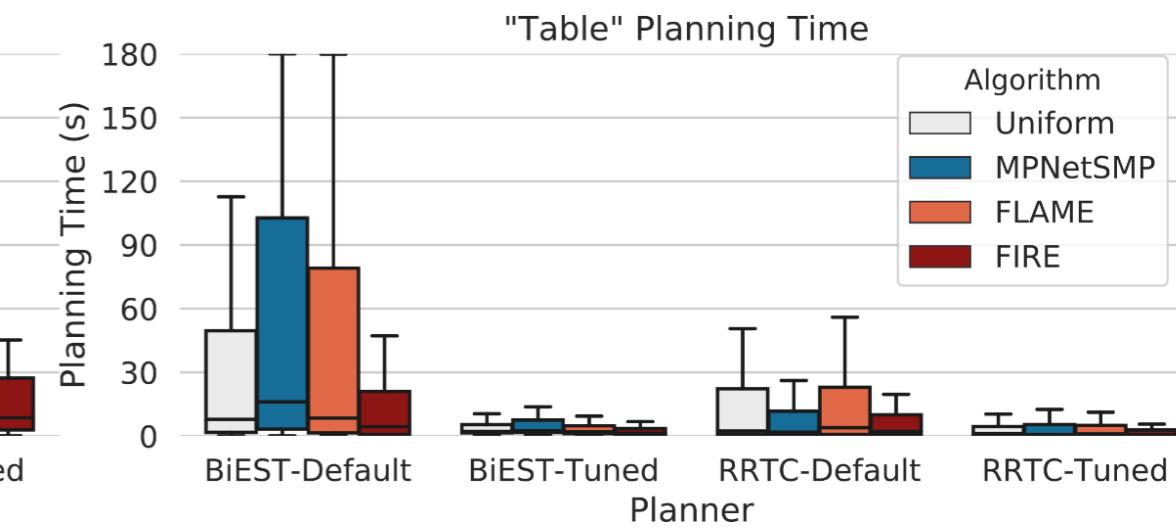
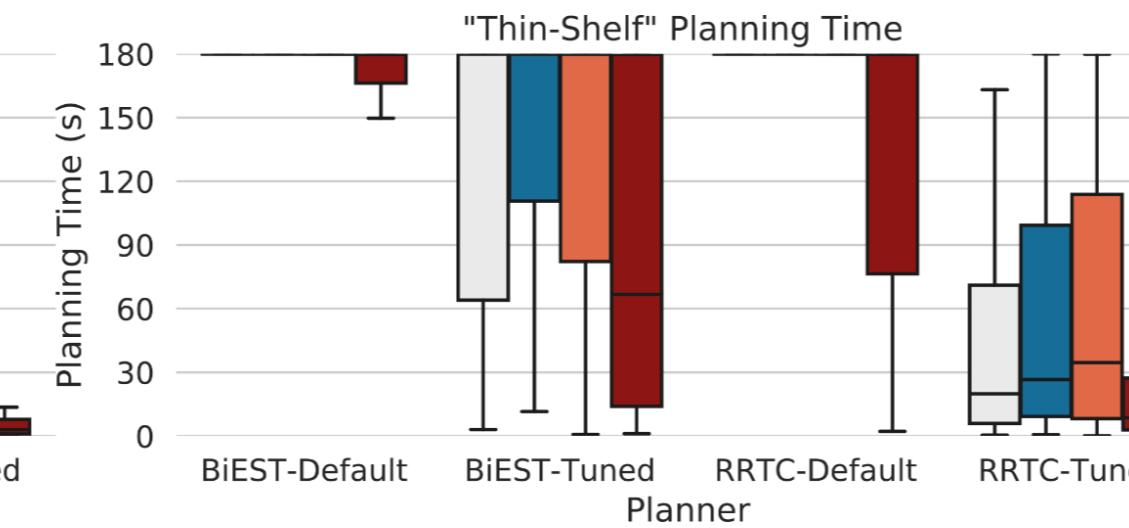
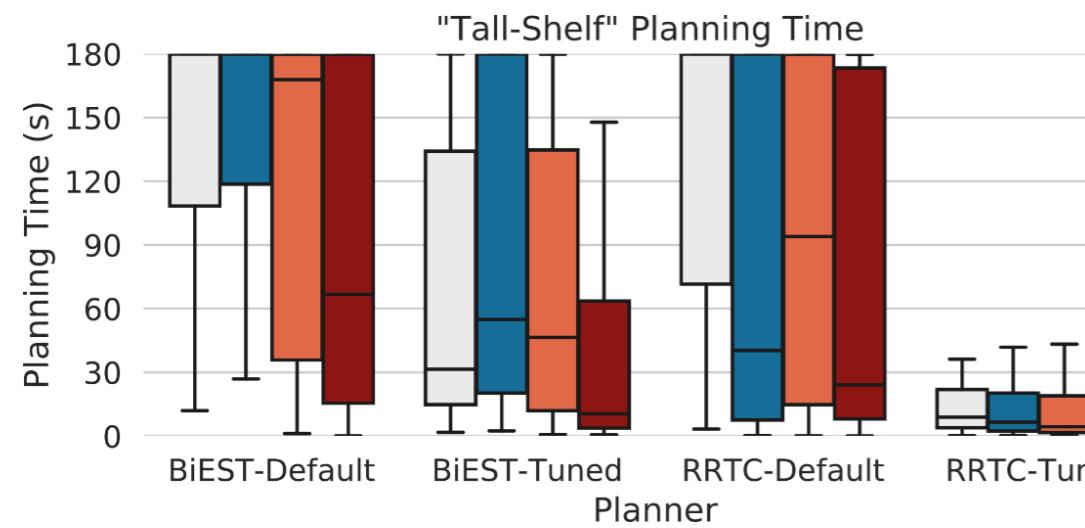
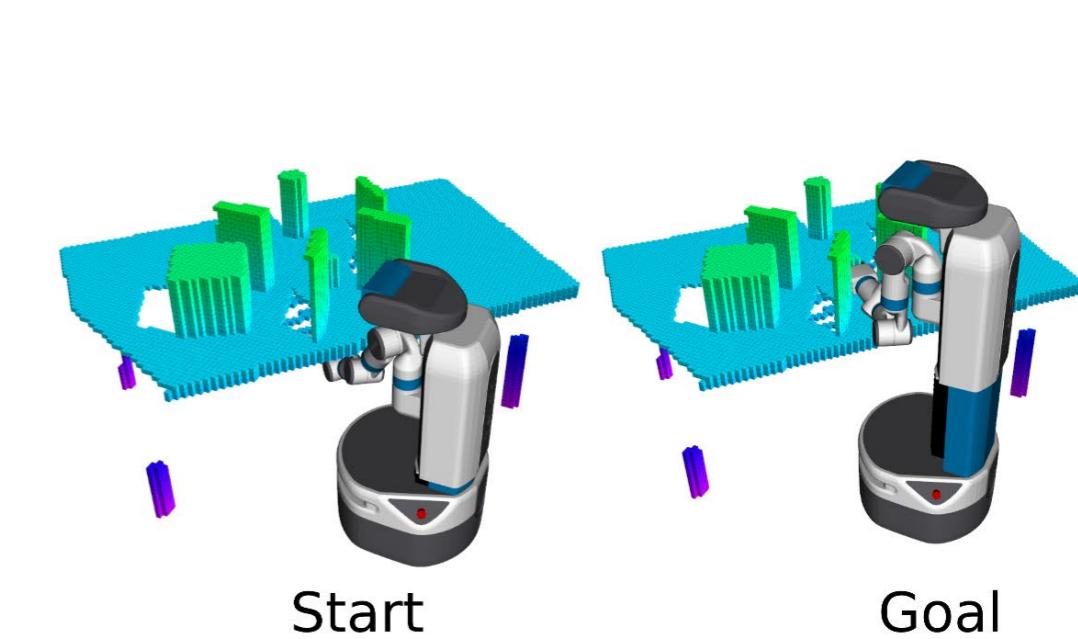
"Tall-Shelf" Dataset (Test)



"Thin-Shelf" Dataset (Test)

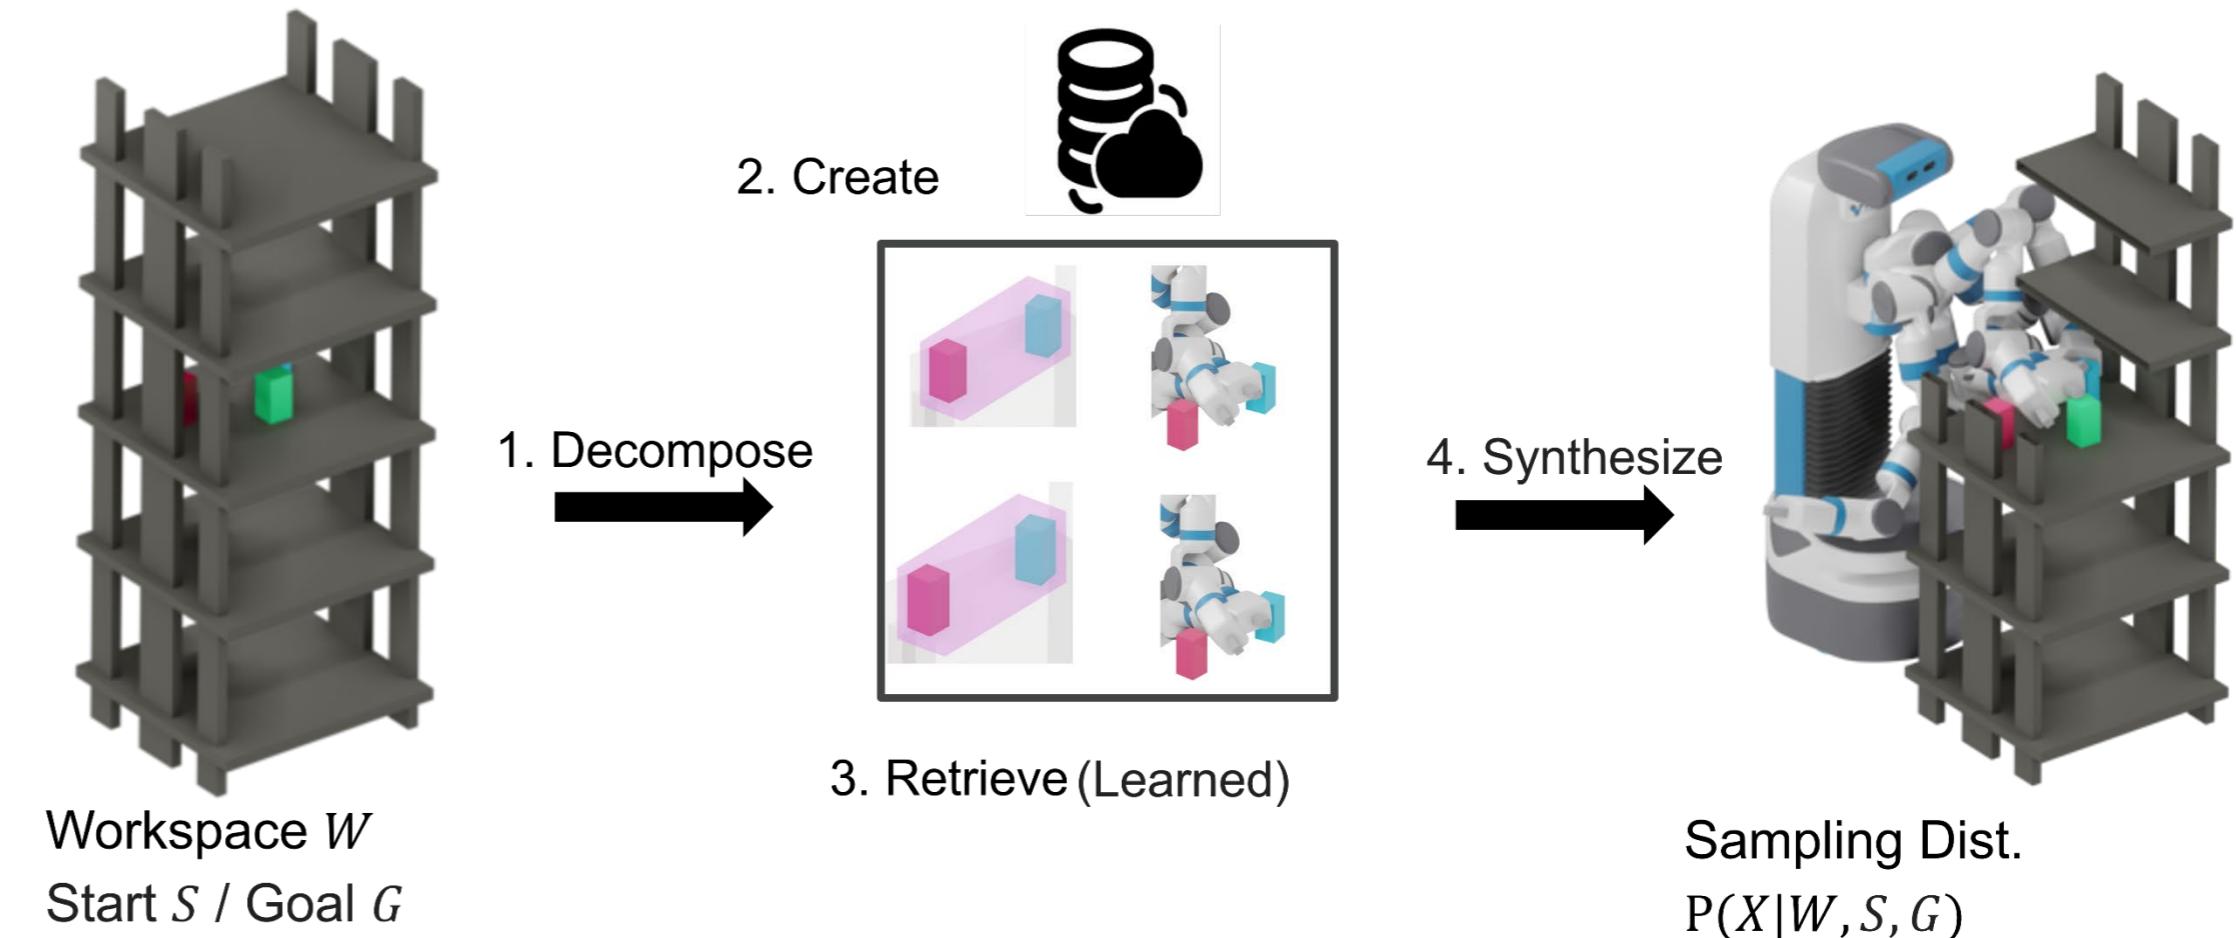


"Table" Dataset (Test)



Summary of SPARK+FLAME+FIRE = Pyre

- Improve planning time for high-dof robots
- Learns incrementally
- Apply to realistic robots and diverse workspaces
- Generalized to unseen tasks and environments



Open-source code: <https://github.com/KavrakiLab/pyre>

Today's Agenda

- Important Considerations
 - Locality of narrow passages
 - Discontinuity
- Biased Samplers (Cont.)
 - 2D Workspaces (SPARK2D)
 - 3D Workspaces (FLAME)
 - Workplace, Start and Goal (FIRE, CVAE)

References (Retrieve-and-Repair)

- [1] Berenson, Dmitry, Pieter Abbeel, and Ken Goldberg. "A robot path planning framework that learns from experience." *International Conference on Robotics and Automation*. IEEE, 2012.
- [2] Coleman, David, et al. "Experience-based planning with sparse roadmap spanners." *International Conference on Robotics and Automation (ICRA)*. IEEE, 2015.
- [3] Pairet, Èric, et al. "Path Planning for Manipulation using Experience-driven Random Trees." *IEEE Robotics and Automation Letters* 6.2 (2021): 3295-3302.
- [4] Lien, Jyh-Ming, and Yanyan Lu. "Planning motion in environments with similar obstacles." *Robotics: Science and systems*. 2009.
- [5] Jetchev, Nikolay, and Marc Toussaint. "Fast motion planning from experience: trajectory prediction for speeding up movement generation." *Autonomous Robots* 34.1 (2013): 111-127.

References (Biased Samplers 1)

- [6] Lehner, Peter, and Alin Albu-Schäffer. "Repetition sampling for efficiently planning similar constrained manipulation tasks." *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. IEEE, 2017.
- [7] Lehner, Peter, and Alin Albu-Schäffer. "The repetition roadmap for repetitive constrained motion planning." *IEEE Robotics and Automation Letters* 3.4 (2018): 3884-3891.
- [8] Zucker, Matt, James Kuffner, and J. Andrew Bagnell. "Adaptive workspace biasing for sampling-based planners." *International Conference on Robotics and Automation*. IEEE, 2008.
- [9] Chamzas, Constantinos, Anshumali Shrivastava, and Lydia E. Kavraki. "Using local experiences for global motion planning." *International Conference on Robotics and Automation (ICRA)*. IEEE, 2019.

References (Biased Samplers 2)

- [11] Ichter, Brian, James Harrison, and Marco Pavone. "Learning sampling distributions for robot motion planning." *International Conference on Robotics and Automation (ICRA)*. IEEE, 2018.
- [10] Chamzas, Constantinos, et al. "Learning sampling distributions using local 3D workspace decompositions for motion planning in high dimensions." *International Conference on Robotics and Automation (ICRA)*. IEEE, 2021
- [12] Chamzas, Constantinos, et al. "Learning to retrieve relevant experiences for motion planning." *International Conference on Robotics and Automation (ICRA)*. IEEE, 2022.