

RBE550

Motion Planning

Overview

Constantinos Chamzas
<https://www.cchamzas.com>

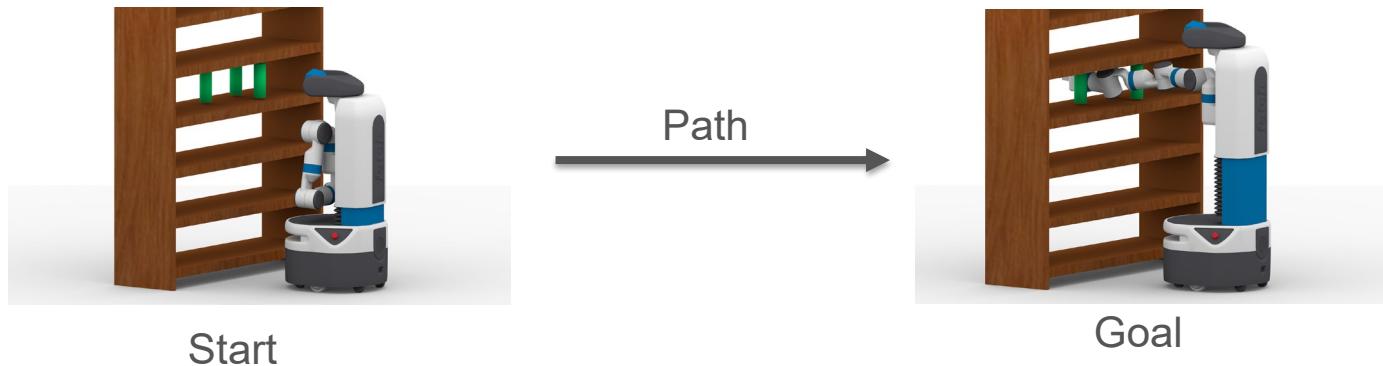
Overview

- Motion Planning Definition
- Bug Algorithms
- Discrete Search/Roadmaps
- Configuration Space
- Sampling-Based Planners
- Asymptotically Optimal Planning
- Kinodynamic Planning
- Learning And Planning
- Task and Motion Planning
- Planning Under Uncertainty



What is Motion Planning?

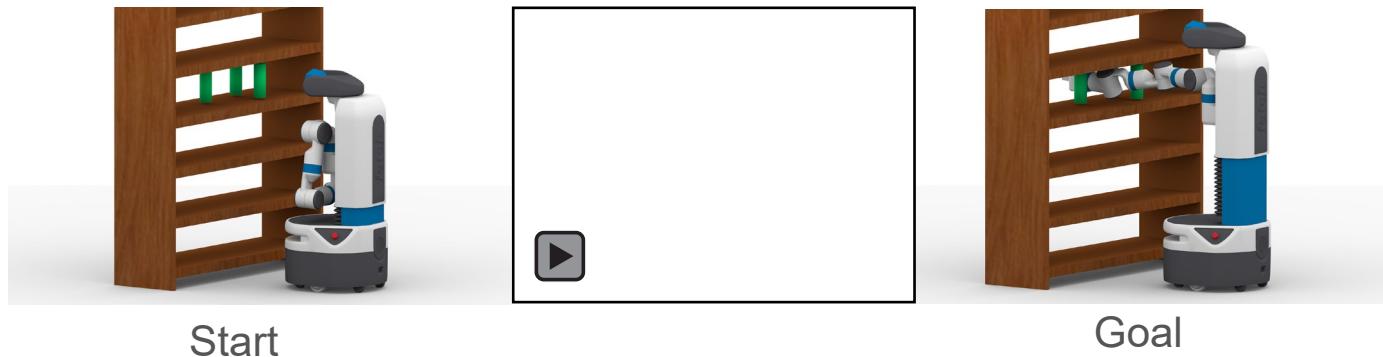
Research field that designs the trajectories for a robot to follow



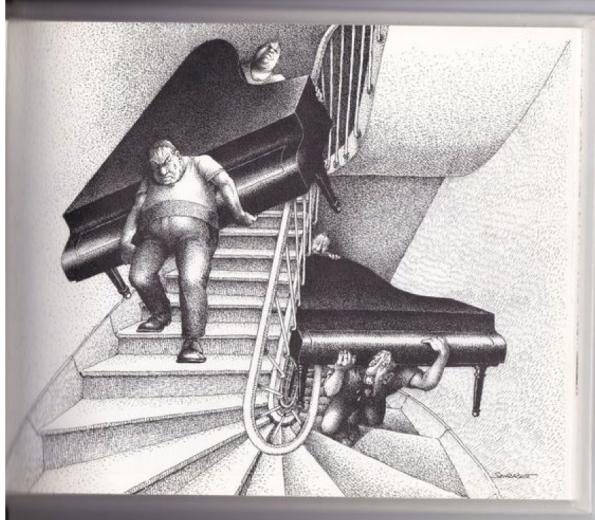


What is Motion Planning?

Research field that designs the trajectories for a robot to follow



Motion Planning- The Piano's Movers Problem



Picture credit: <https://www.pinterest.com/pin/178736678931496070/>



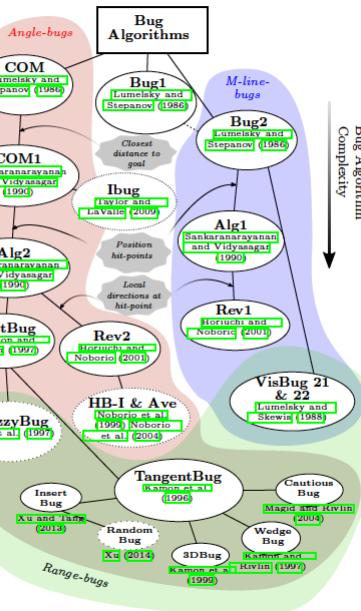
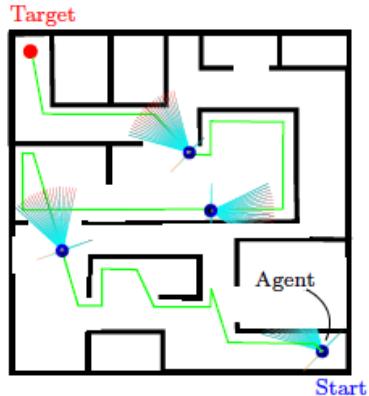
Motion Planning is Hard

Problem	Complexity
Sofa Mover (3 DOF)	$O(n^{2+\epsilon})$ not implemented
Piano Mover (6 DOF)	Polynomial – no practical algorithm known
n Disks in the Plane	NP-hard
n Link Planar Chain	PSPACE-Complete
Generalized Mover	PSPACE-Complete
Shortest Path for a Point in 3D	NP-hard
Curvature Constrained Point in 2D	NP-hard
Simplified Coulomb Friction	Undecidable

Overview

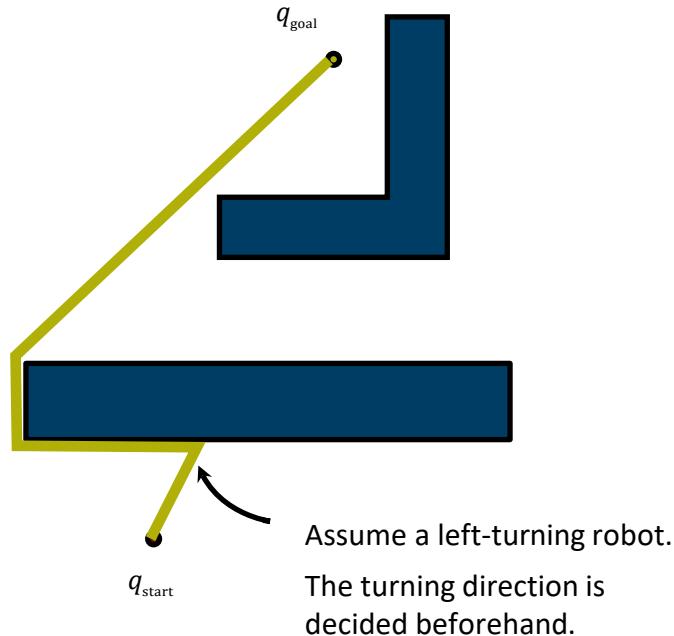
- Motion Planning Definition
- Bug Algorithms
- Discrete Search/Roadmaps
- Configuration Space
- Sampling-Based Planners
- Asymptotically Optimal Planning
- Kinodynamic Planning
- Learning And Planning
- Task and Motion Planning
- Planning Under Uncertainty

Bug Algorithms



From McGuire et al, 2018

Bug 0

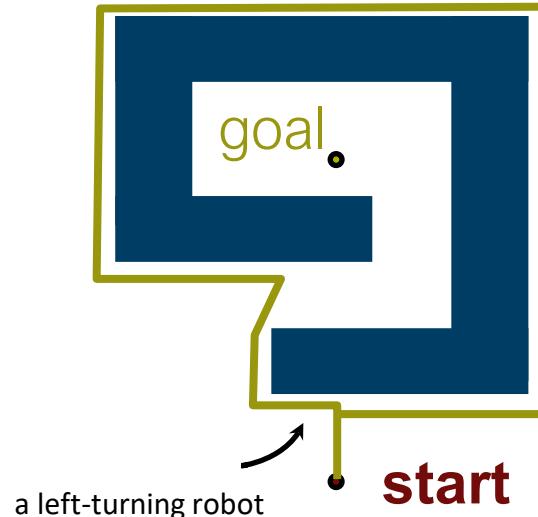


“Bug 0” algorithm

1. head toward goal
2. follow obstacles until you can head toward the goal again
3. continue

done?

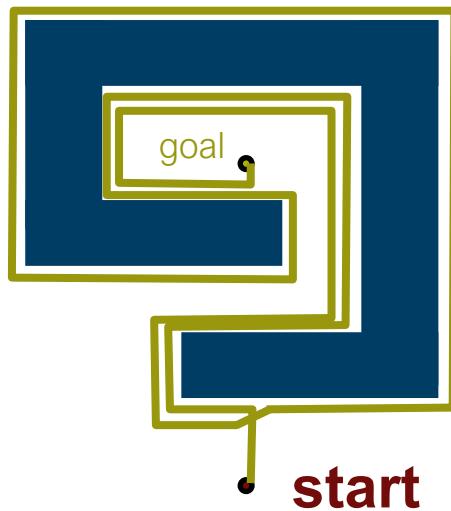
What map will foil Bug 0



“Bug 0” algorithm

1. head toward goal
2. follow obstacles until you can head toward the goal again
3. continue

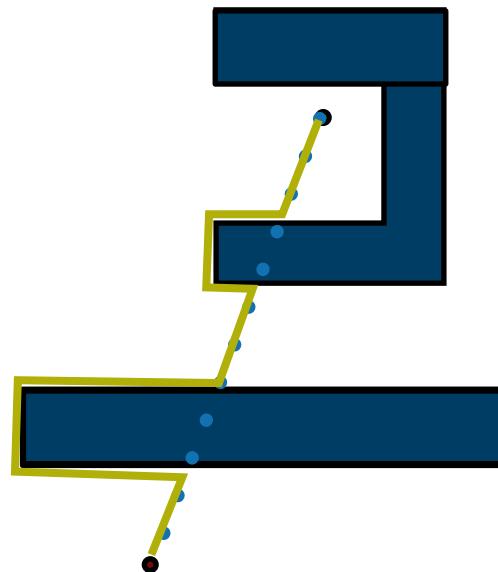
Bug 1 Algorithm



- **Assumptions:**
 - known direction to goal
 - otherwise local sensing walls/obstacles & encoders
 - Has memory

- **Strategy:**
 1. head toward goal
 2. if an obstacle is encountered, circumnavigate it and remember how close you get to the goal
 3. return to that closest point (by wall-following) and continue

Bug 2 Algorithm



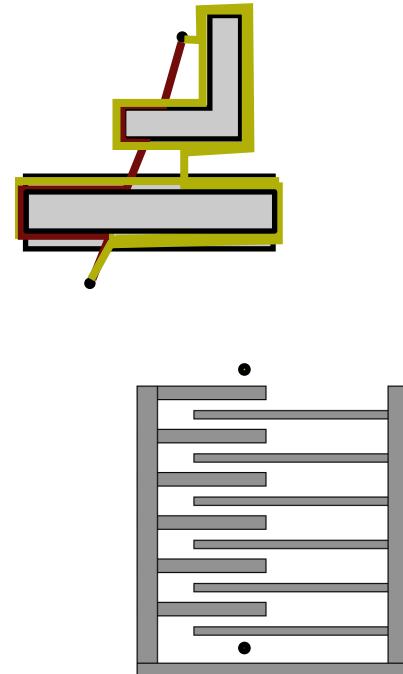
“Bug 2” algorithm

1. head toward goal on the m-line
2. if an obstacle is in the way, follow it until you encounter the m-line again [closer to the goal](#)
3. leave the obstacle and continue toward the goal

Better or worse than Bug1?

Bug 1 Vs Bug 2

- BUG 1 is an *exhaustive* search algorithm
 - it looks at **all** choices before committing
- BUG 2 is a *greedy* algorithm
 - it takes the **first** thing that looks better
- In many cases, BUG 2 will outperform BUG 1, but
- BUG 1 has a more predictable performance overall



Limitations

- Relies on knowing the boundary,
- Limited to 2D robots
- Highly non optimal and unsafe motions

Overview

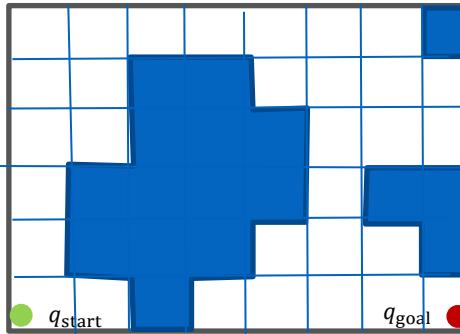
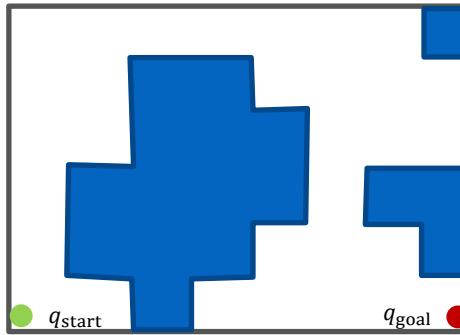
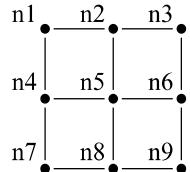
- Motion Planning Definition
- Bug Algorithms
- Discrete Search/Roadmaps
- Configuration Space
- Sampling-Based Planners
- Asymptotically Optimal Planning
- Kinodynamic Planning
- Learning And Planning
- Task and Motion Planning
- Planning Under Uncertainty

Discretization of Workspace Produces Graphs

n1	n2	n3
n4	n5	n6
n7	n8	n9

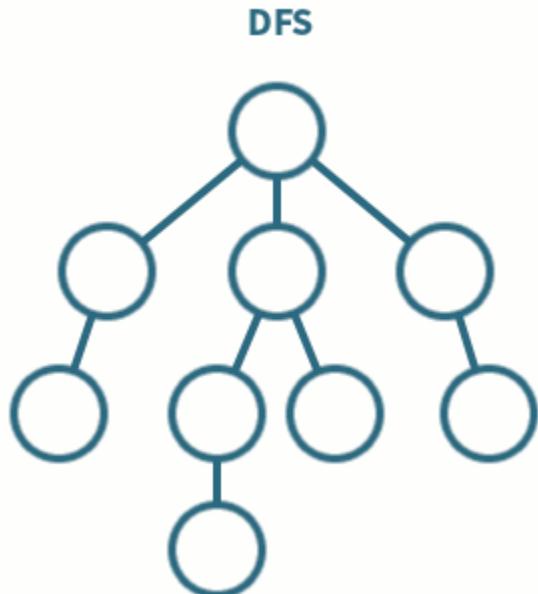


L_1 -norm
neighbors

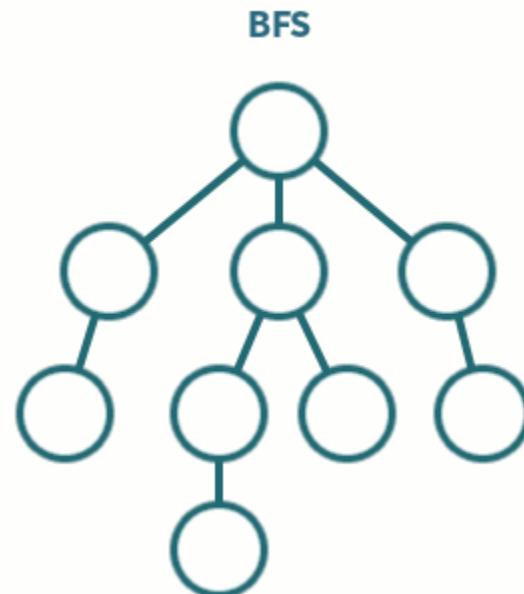


Graph Search Algorithms

- Depth-first search



- Breadth-first search



A* Algorithm

- A* algorithm searches a graph efficiently w.r.t. a given heuristic

- Heuristic: estimates the cost to the goal node
- Good heuristic → efficient search
- Bad heuristic → inefficient search

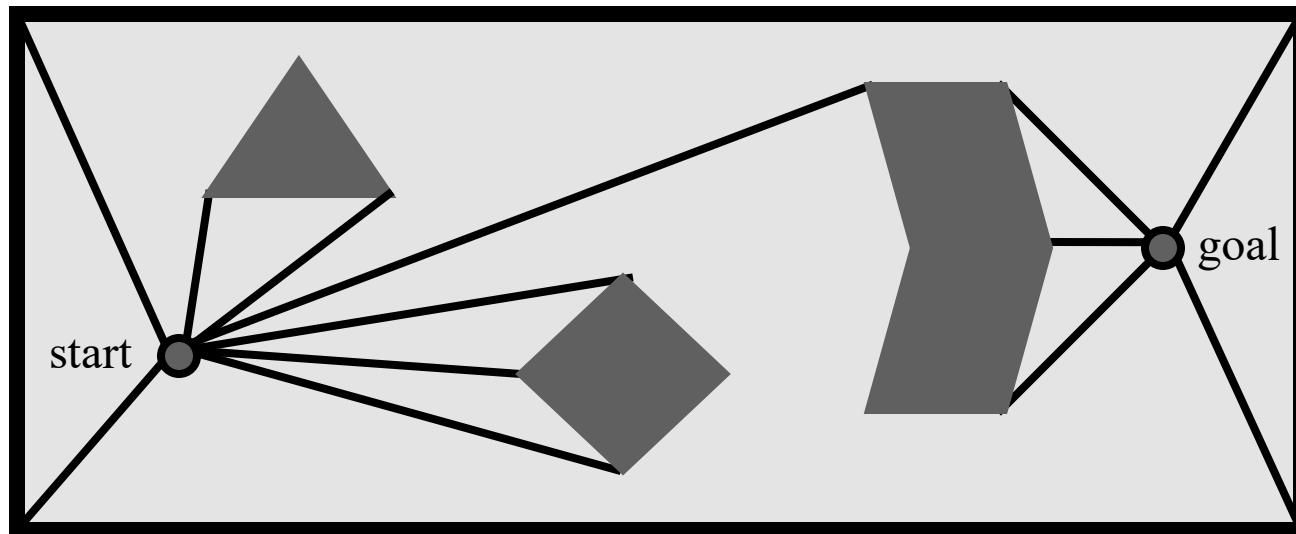


Limitations

- We need to create a grid
- Does not scale due to the course of dimensionality
- If resolution is not “fine” enough we will not be able to navigate through cluttered obstacles.

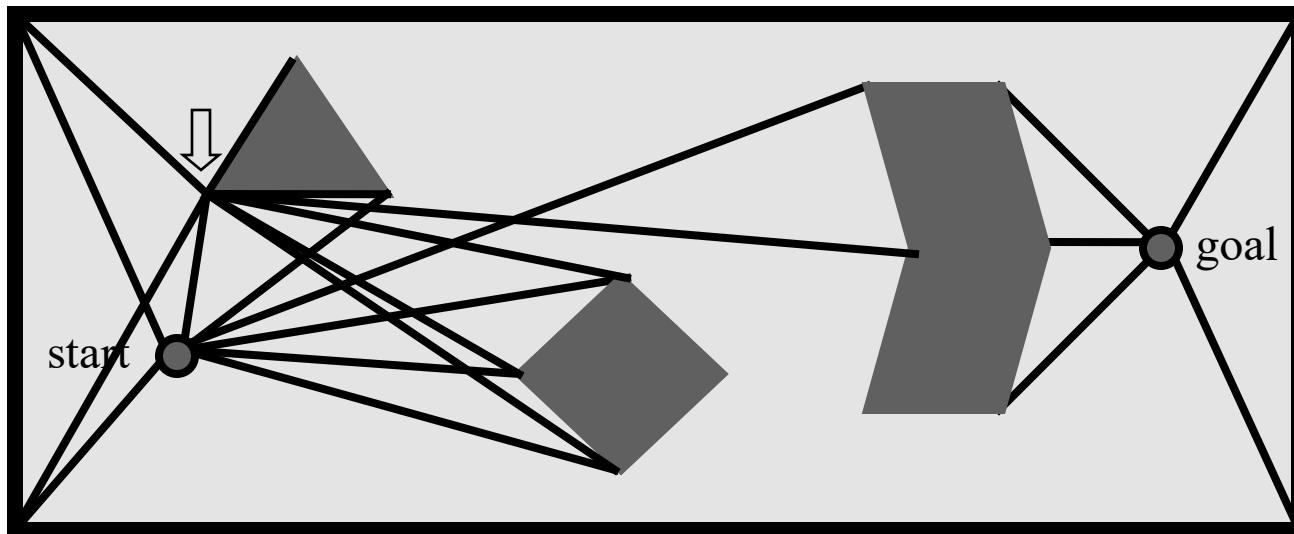
The Visibility Graph in Action (Part 1)

- First, draw lines of sight from the start and goal to all “visible” vertices and corners of the world.



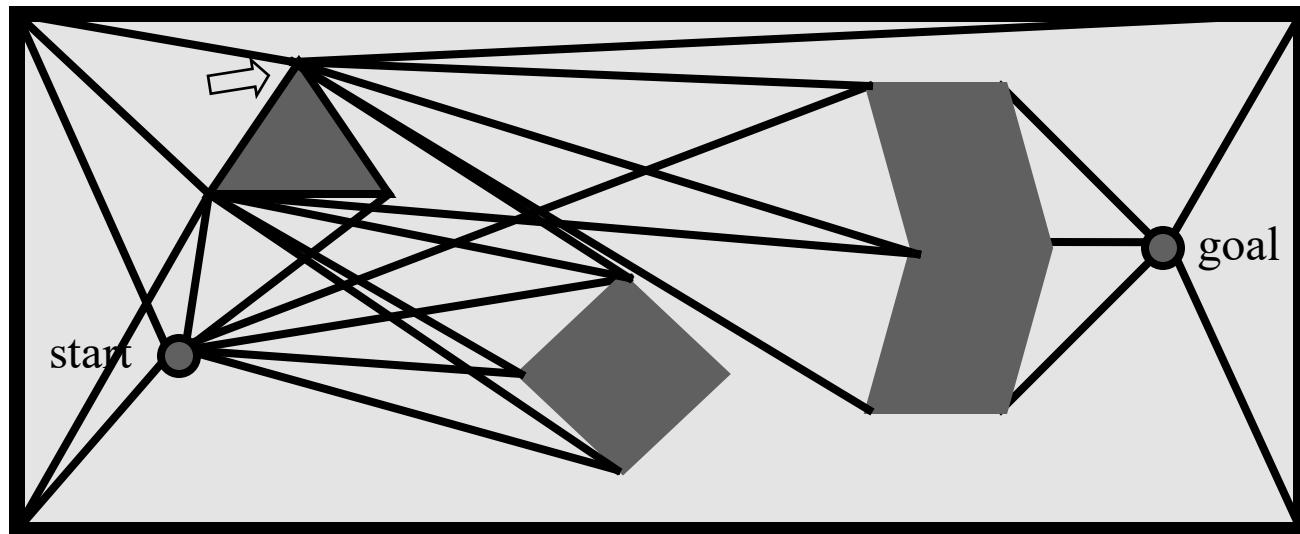
The Visibility Graph in Action (Part 2)

- Second, draw lines of sight from every vertex of every obstacle like before. Remember lines along edges are also lines of sight.



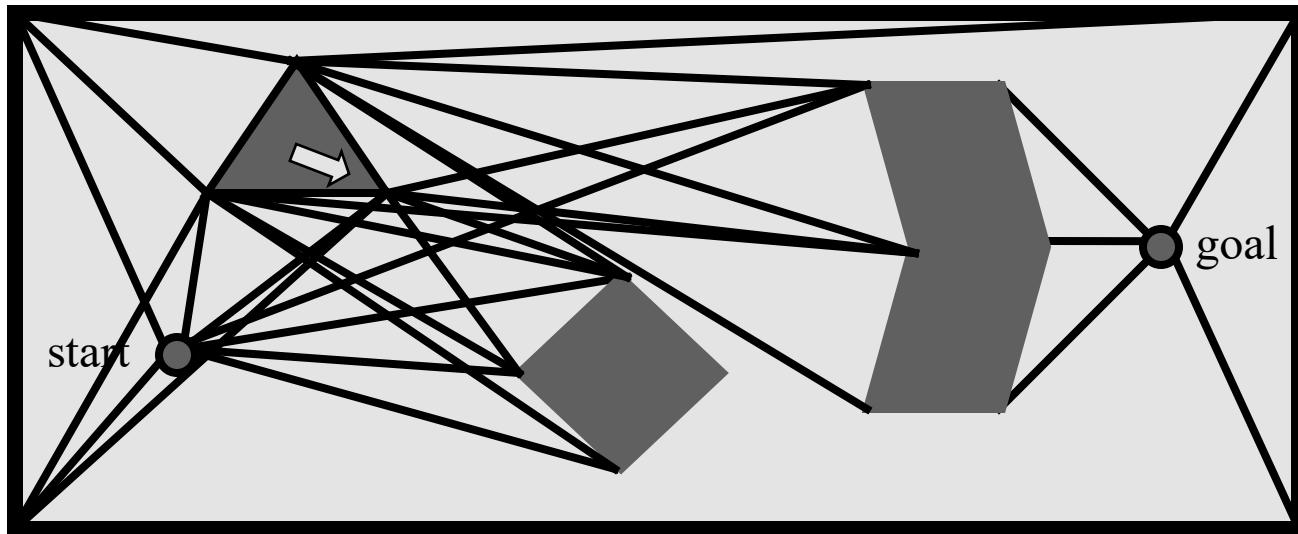
The Visibility Graph in Action (Part 3)

- Second, draw lines of sight from every vertex of every obstacle like before. Remember lines along edges are also lines of sight.



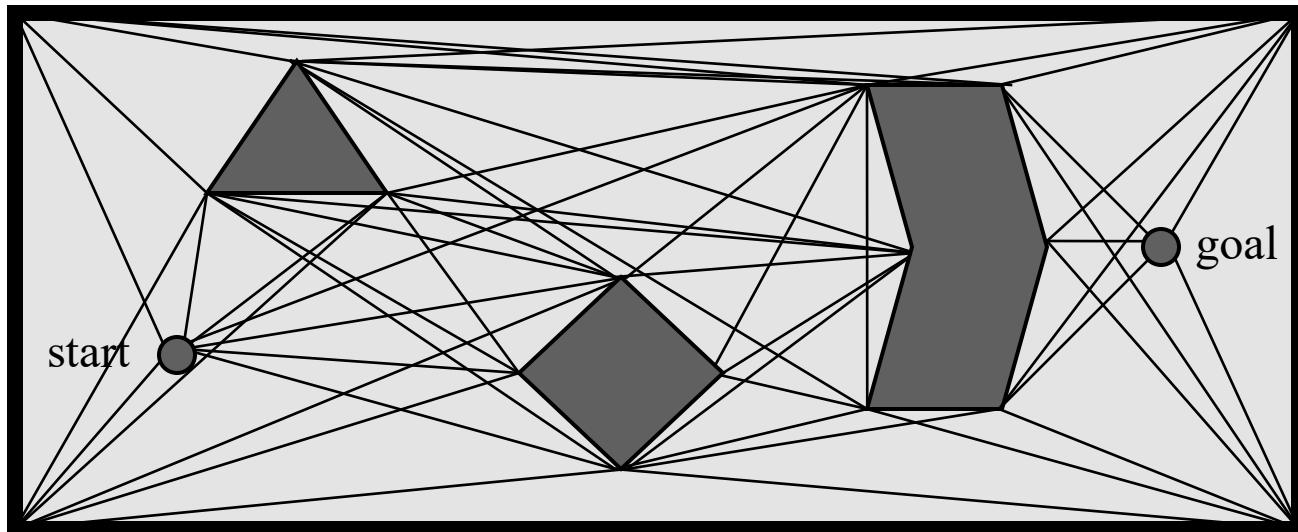
The Visibility Graph in Action (Part 4)

- Second, draw lines of sight from every vertex of every obstacle like before. Remember lines along edges are also lines of sight.

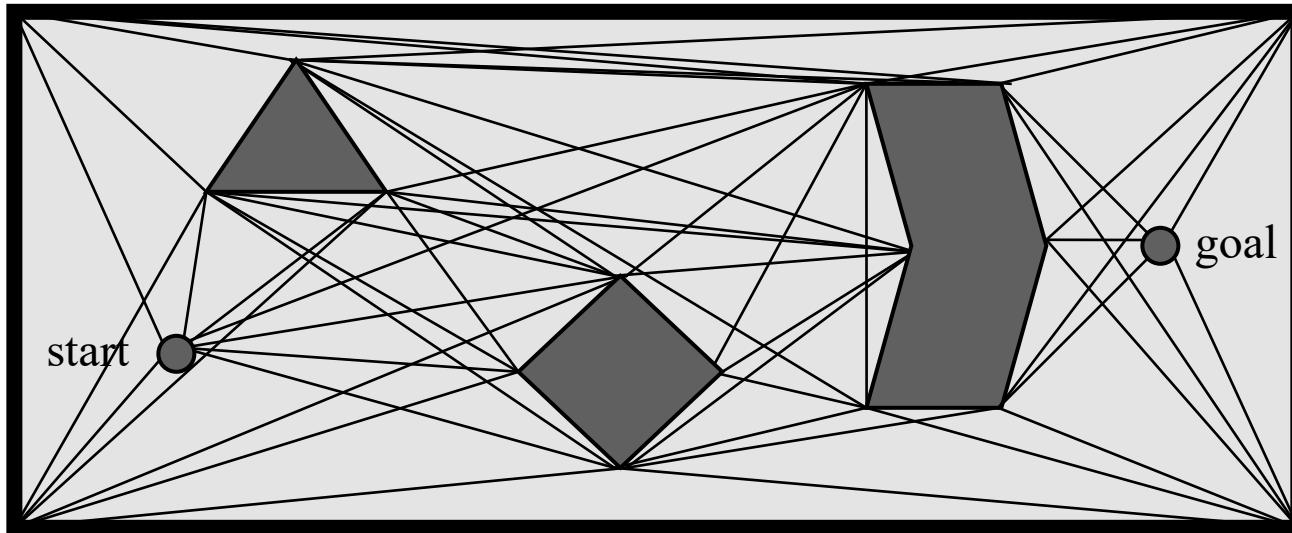


The Visibility Graph (Done)

- Repeat until you're done.



We can just run A* on this graph



This graph can be searched using A* and solves the problem continuously

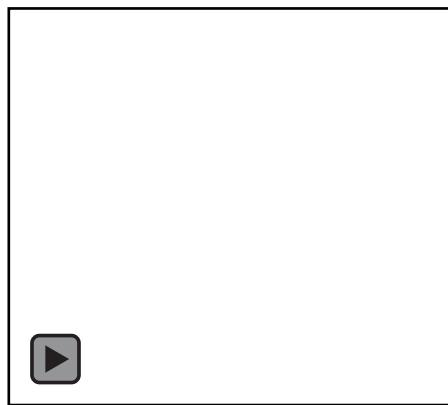
Limitations

- We achieved optimal paths in continuous spaces but only point robots so far
- ...

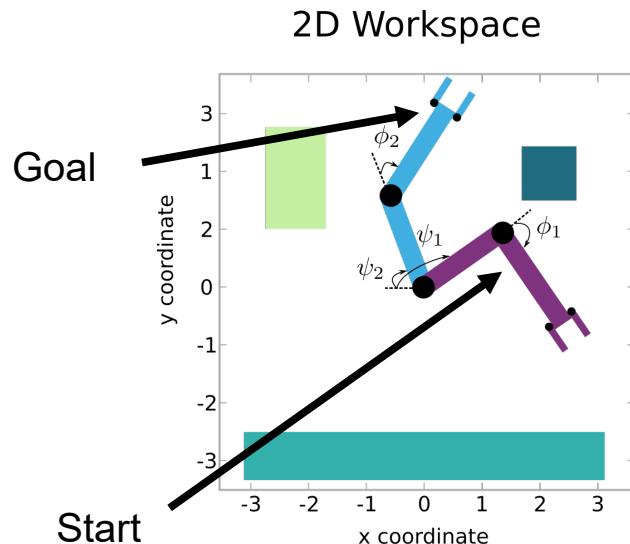
Overview

- Motion Planning Definition
- Bug Algorithms
- Discrete Search/Roadmaps
- Configuration Space
- Sampling-Based Planners
- Asymptotically Optimal Planning
- Kinodynamic Planning
- Learning And Planning
- Task and Motion Planning
- Planning Under Uncertainty

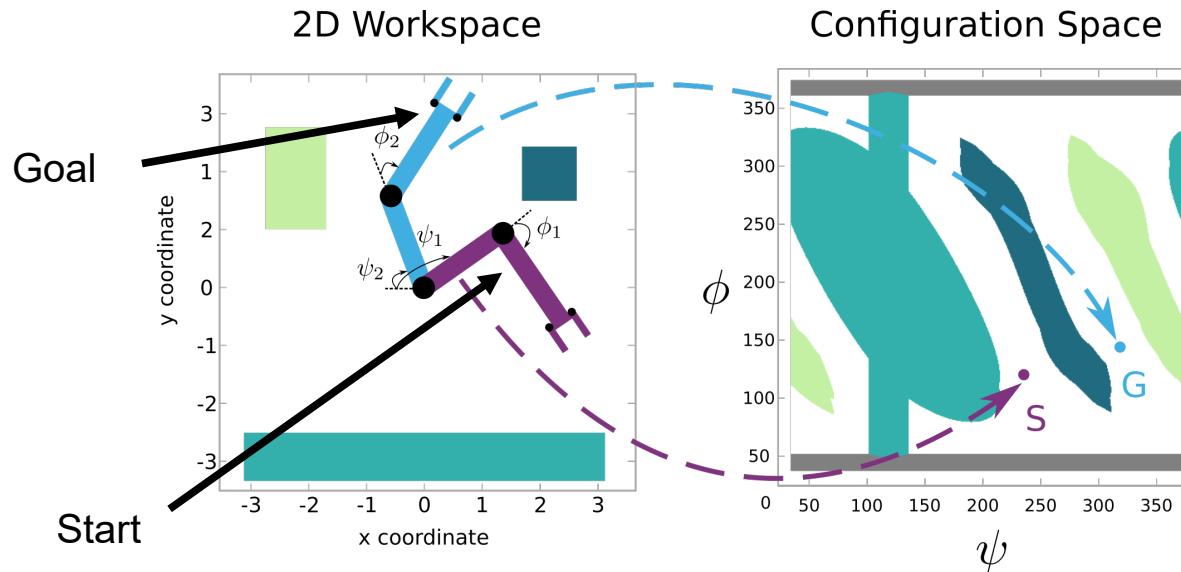
A Simple 2-link Articulated Robot



Big Idea: Configuration Space

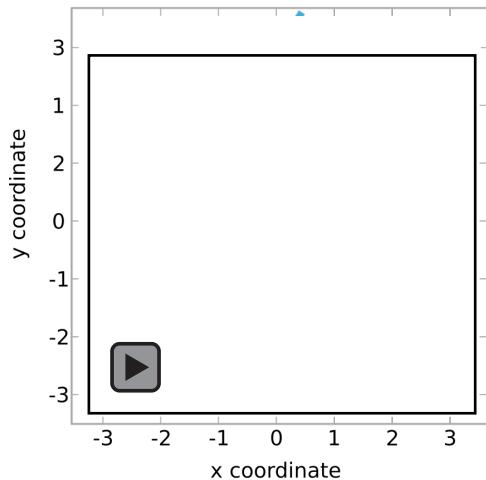


Big Idea: Configuration Space

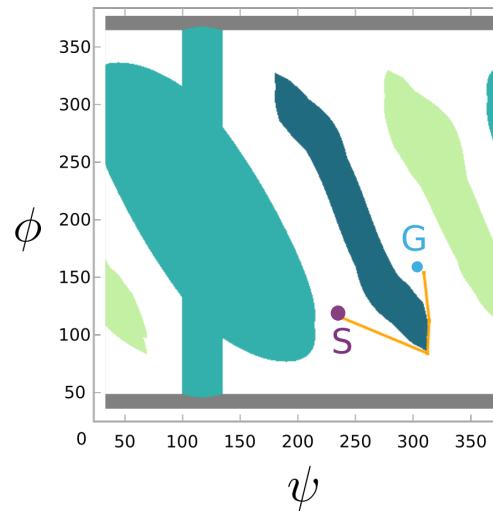


Path in Configuration space

2D Workspace

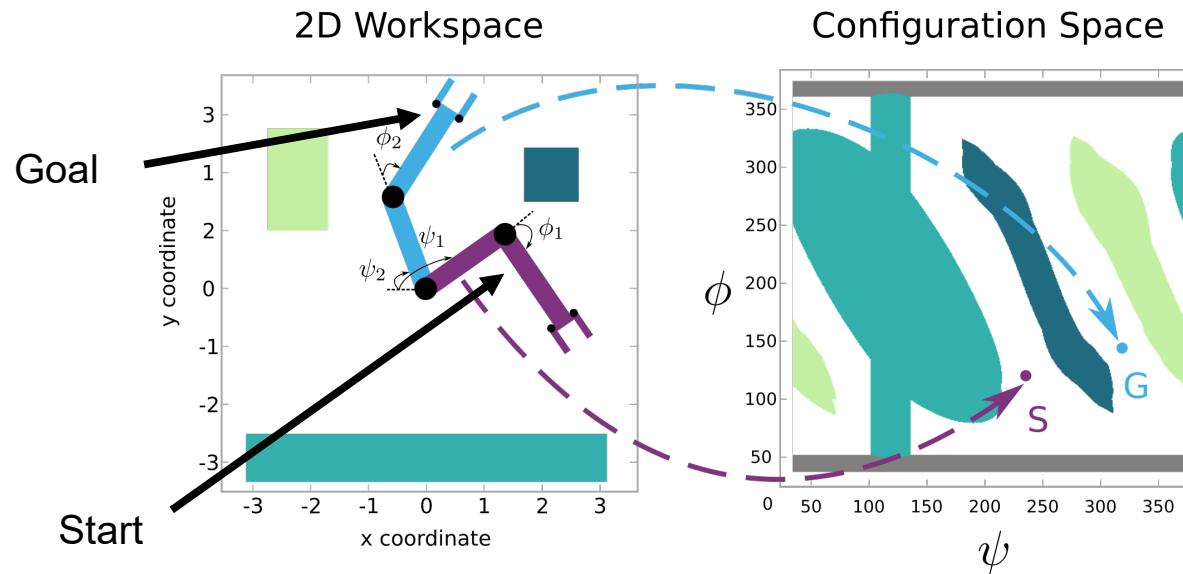


Configuration Space



Now we can use point Planning Algorithms to plan in the C-Space!

Workspace vs Configuration Space

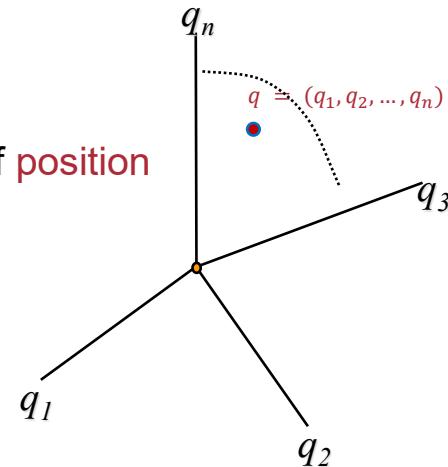


Configuration Space

- The **configuration** of a moving object is a specification of the position of every point on the object.

- Usually a configuration is expressed as a **vector of position & orientation** parameters:

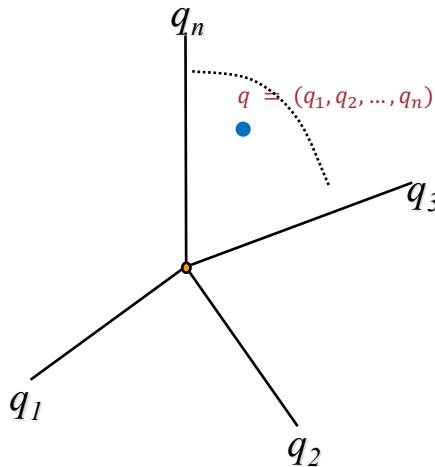
$$q = (q_1, q_2, \dots, q_n)$$



- The **configuration space C** is the set of all possible configurations.
- A configuration is a point in C

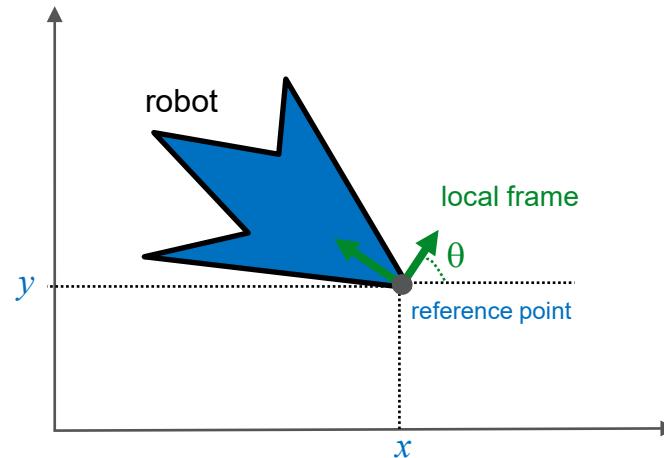
Configuration Space

- The **dimension** of a configuration space:
 - The **minimum** number of parameters needed to specify the configuration of the object completely.
 - also called the number of **degrees of freedom (DoF)** of a moving object.



Configuration Space

- Example: rigid body in 2-D workspace



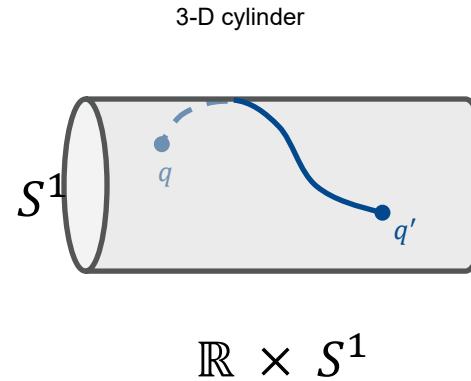
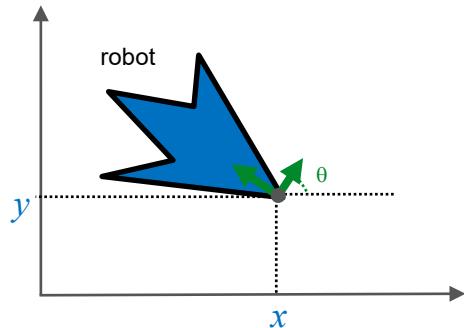
- 3-parameter specification:

$$q = (x, y, \theta) \text{ with } \theta \in [0, 2\pi)$$

- 3-D configuration space (3 DoF)

Topology of Configuration Space

- The topology of C is usually **not** that of a Cartesian space \mathbb{R}^n
- Here the robot **can only move in the x direction** and rotate $C = \mathbb{R} \times S^1$



(note this is not $\mathbb{R}^2 \times S^1$)

Configuration Space

- Example: rigid body in 3-D
 - How many parameters in C-space?

$$q = (x, y, z, \mathbf{q}')$$

- Parametrization of orientations by matrix:

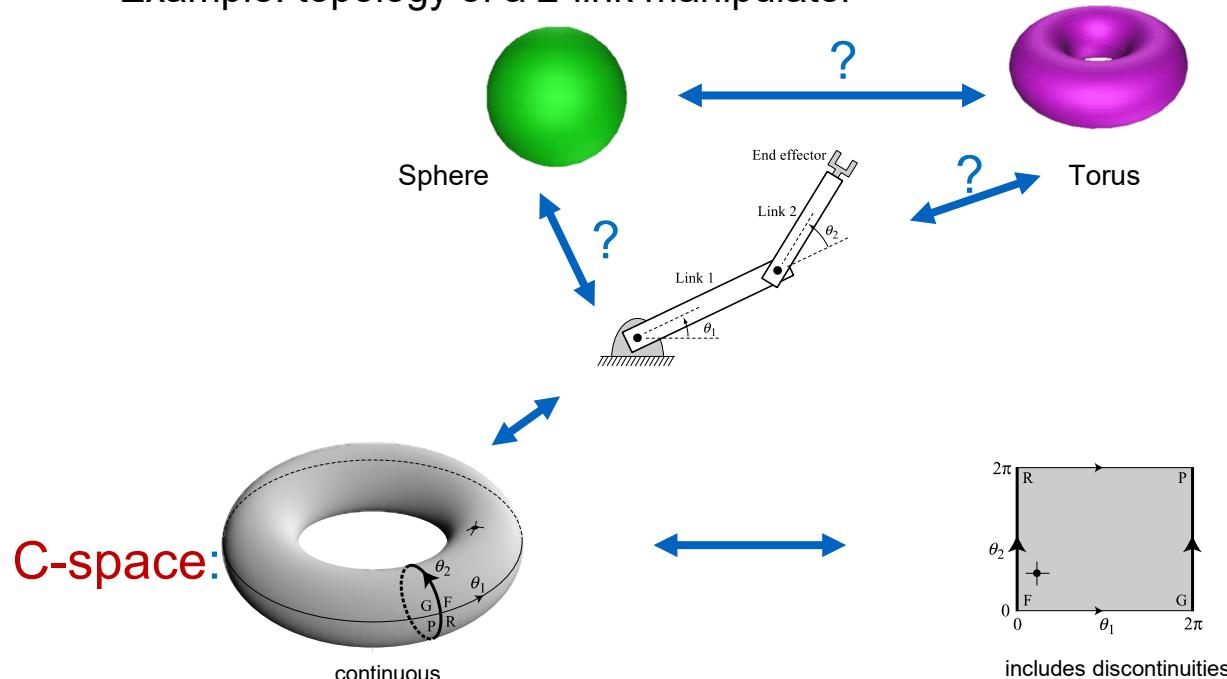
$$\mathbf{q}' = (r_{11}, r_{12}, \dots, r_{32}, r_{33})$$

where $r_{11}, r_{12}, \dots, r_{33}$ are the elements of rotation matrix

$$R = \begin{pmatrix} r_{11} & r_{12} & r_{13} \\ r_{21} & r_{22} & r_{23} \\ r_{31} & r_{32} & r_{33} \end{pmatrix} \quad \text{with} \quad \begin{array}{l} \bullet \quad r_{1i}^2 + r_{2i}^2 + r_{3i}^2 = 1 \text{ for all } i \\ \bullet \quad r_{1i}r_{1j} + r_{2i}r_{2j} + r_{3i}r_{3j} = 0 \text{ for all } i \neq j \\ \bullet \quad \det(R) = +1 \end{array}$$

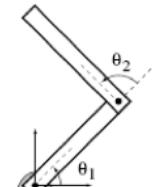
2-link manipulator robot Topology

- Example: topology of a 2-link manipulator

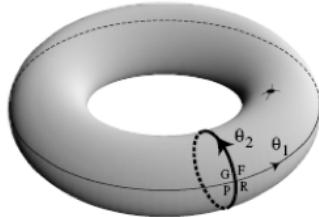


- Topology is the **intrinsic** character of the space

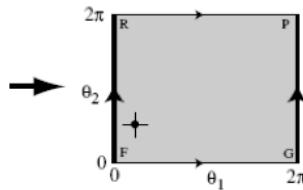
Articulated Robot



(a)



(b)



(c)

$$(\theta_1, \theta_2) \in \mathbb{R}^2,$$

problems at $\theta_i = \{0, 2\pi\}$.

Topology of Some C-Space

- Examples of some common robots

Type of robot	Representation of C-Space
Mobile robot translating in the plane	\mathbb{R}^2
Mobile robot translating and rotating in the plane	$SE(2)$ or $\mathbb{R}^2 \times S^1$
Rigid body translating in 3-D	\mathbb{R}^3
A spacecraft	$SE(3)$ or $\mathbb{R}^3 \times SO(3)$
An n-joint revolute arm	T^n
A planar mobile robot with an attached n -joint arm	$SE(2) \times T^n$

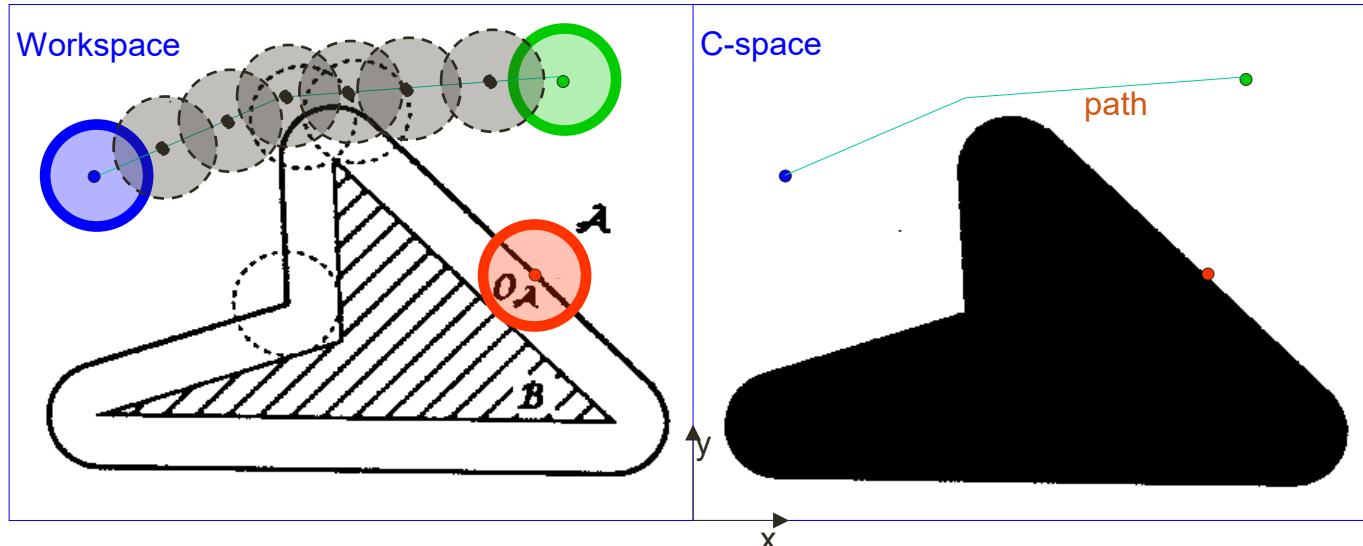
- Note that:
- $S^1 \times S^1 \times \dots \times S^1 = T^n$, n-dimensional torus
- $S^1 \times S^1 \times \dots \times S^1 \neq S^n$, n-dimensional sphere
- $S^1 \times S^1 \times S^1 \neq SO(3)$
- $SE(2) \neq \mathbb{R}^3$
- $SE(3) \neq \mathbb{R}^6$

Limitations

- Now we can use point algorithms to move any robot if we know the topology of its configuration space
- But what about avoiding obstacles?

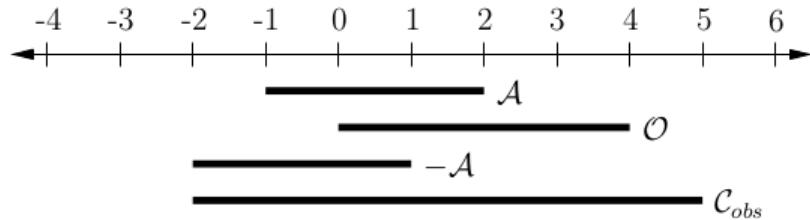
Obstacles in the Configuration Space

- 2-D disc robot (translation only)



- Configuration: $q = (x, y)$ coordinates of robot's center
- configuration space $C = \mathbb{R}^2$
- free space C_{free} = the set of collision-free configurations

C-Space Obstacle in a 1D Case



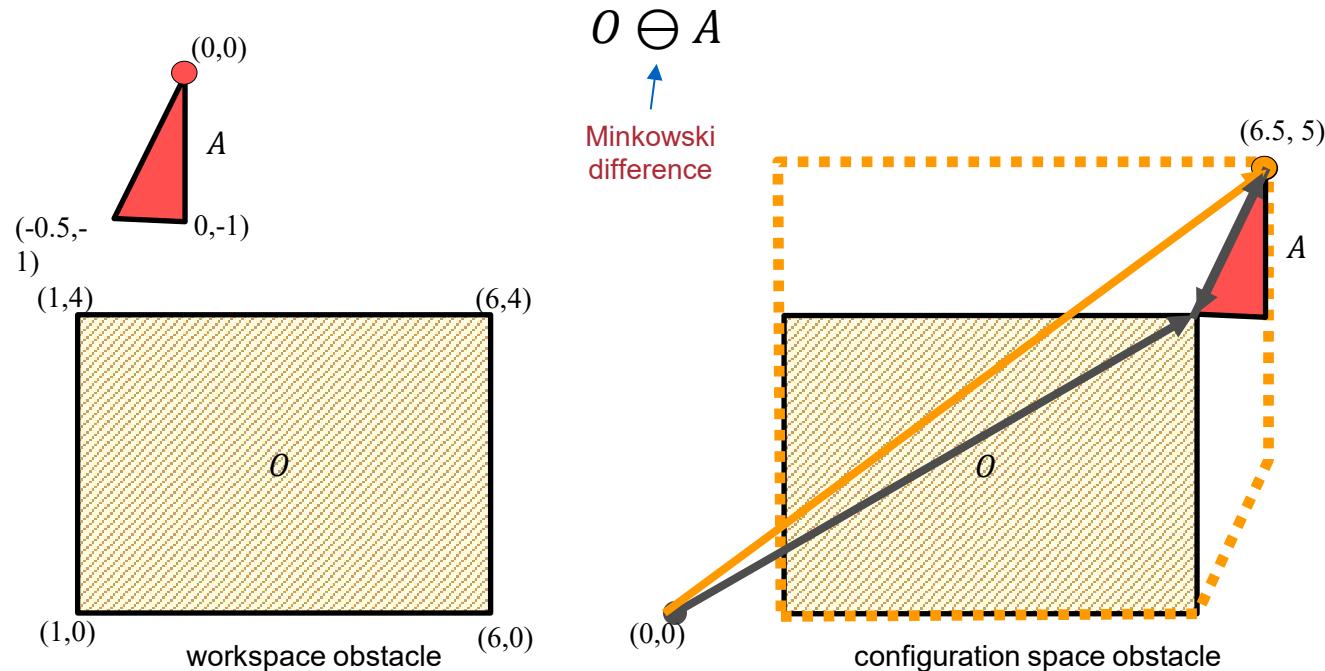
$$A = [-1, 2] \quad O = [0, 4]$$

$$-A = [-2, 1]$$

$$\mathcal{C} = O \ominus A = O + (-A) = [-2, 5]$$

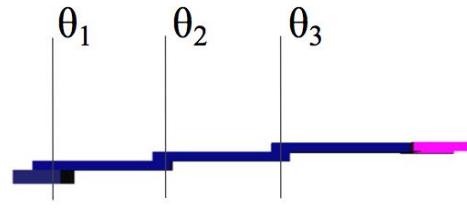
C-Space obstacles

- If O is an obstacle in the workspace and A is a moving object, then the C-space obstacle corresponding to O is



Articulated Robot in 2D workspace

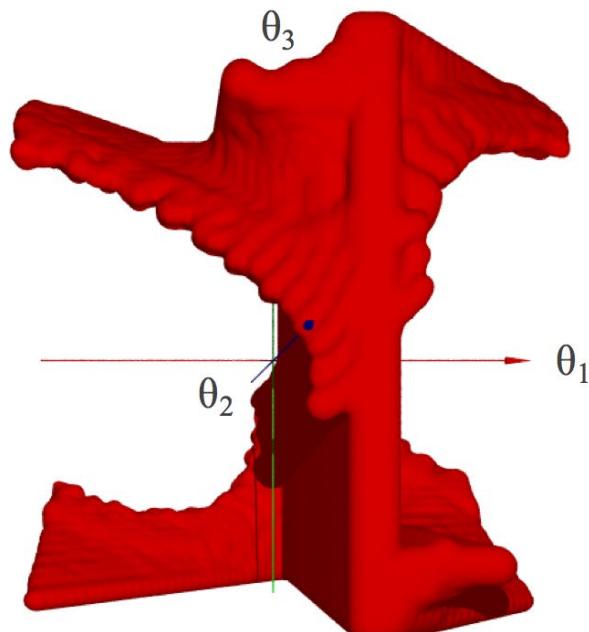
- 3-link arm



side view



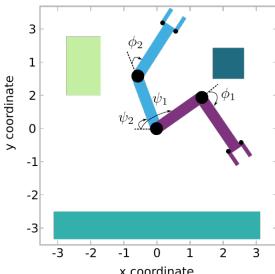
top view



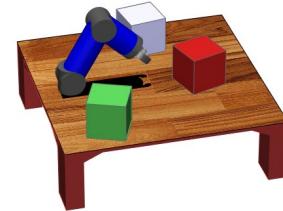
High-Dof Robots configuration space

Workspace

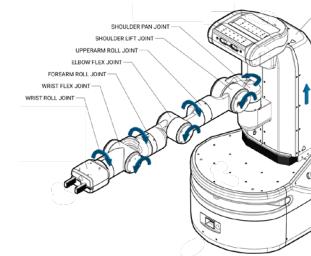
2-Dof Robot



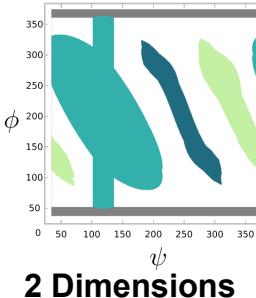
3-Dof Robot



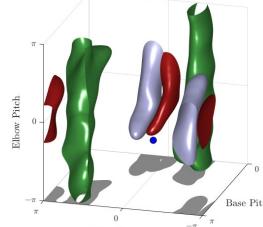
8-Dof Robot



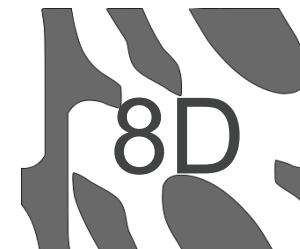
Configuration Space



2 Dimensions



3 Dimensions



8 Dimensions

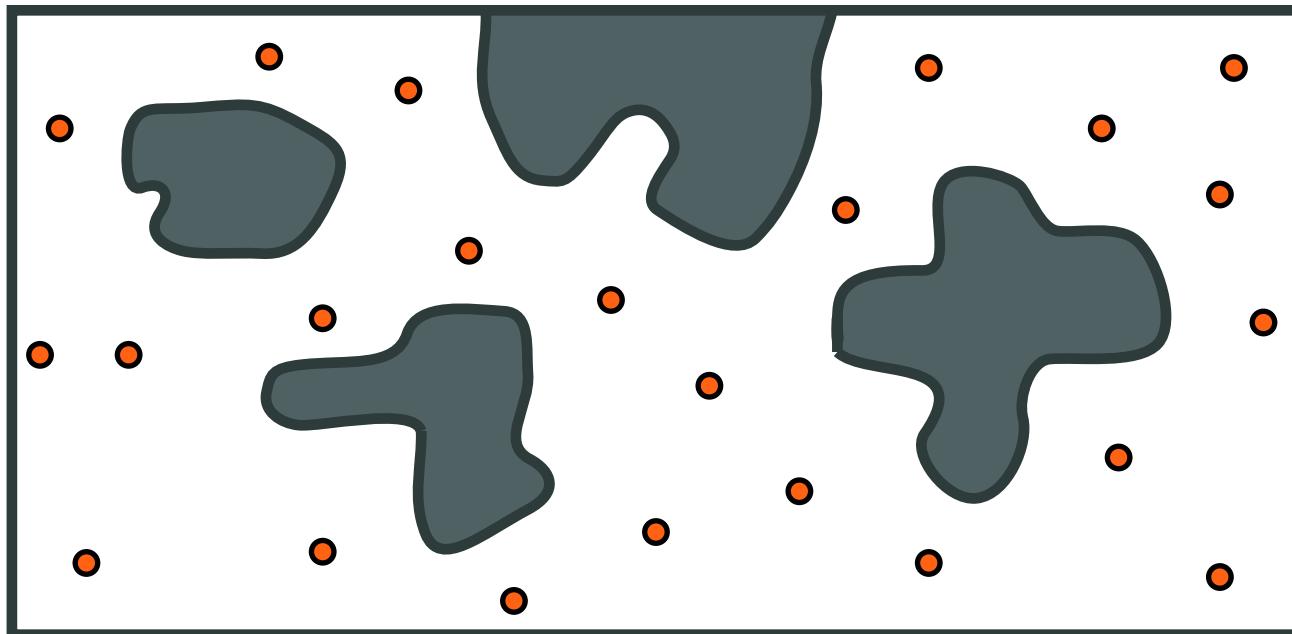
Limitations

- How can we create this high-dimensional configuration obstacles?

Overview

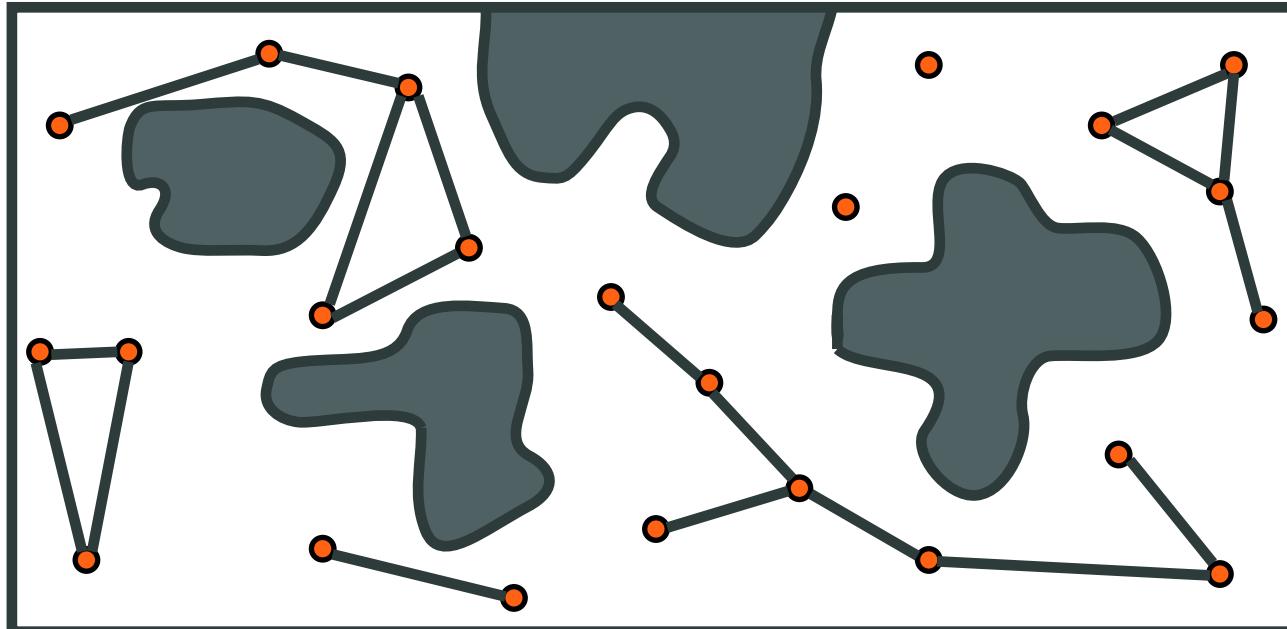
- Motion Planning Definition
- Bug Algorithms
- Discrete Search/Roadmaps
- Configuration Space
- Sampling-Based Planners
- Asymptotically Optimal Planning
- Kinodynamic Planning
- Learning And Planning
- Task and Motion Planning
- Planning Under Uncertainty

Random Sampling!



● : nodes, random states

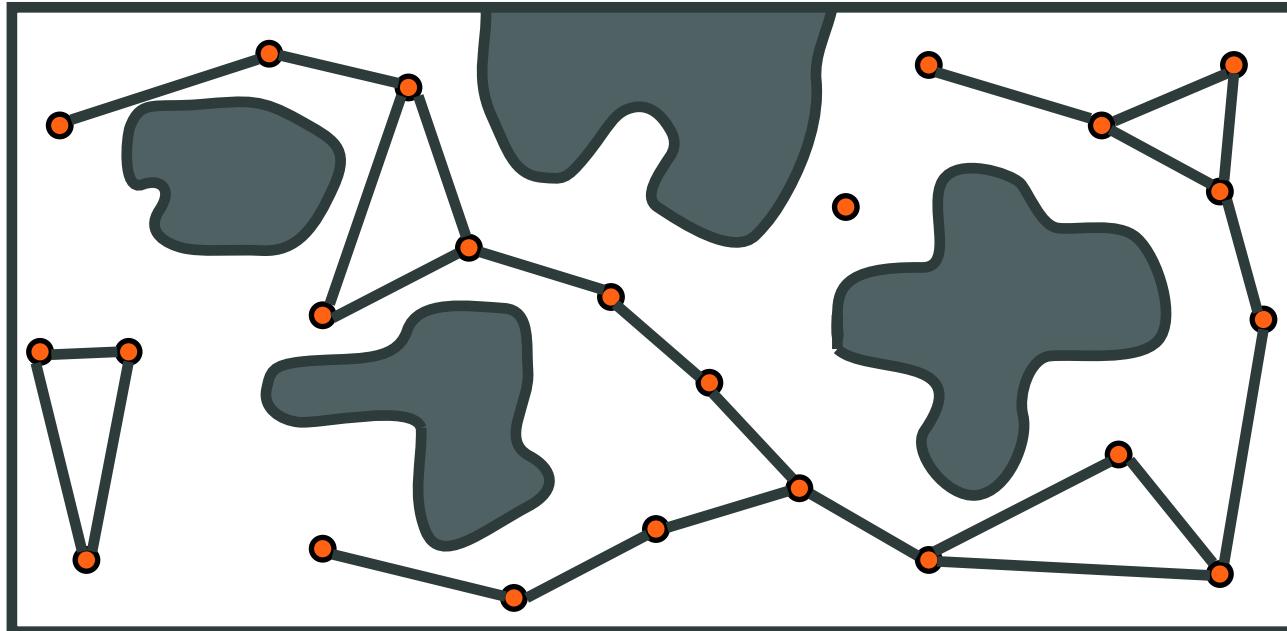
Operation of PRM



● : nodes, random states

— : edges, paths computed by local planner

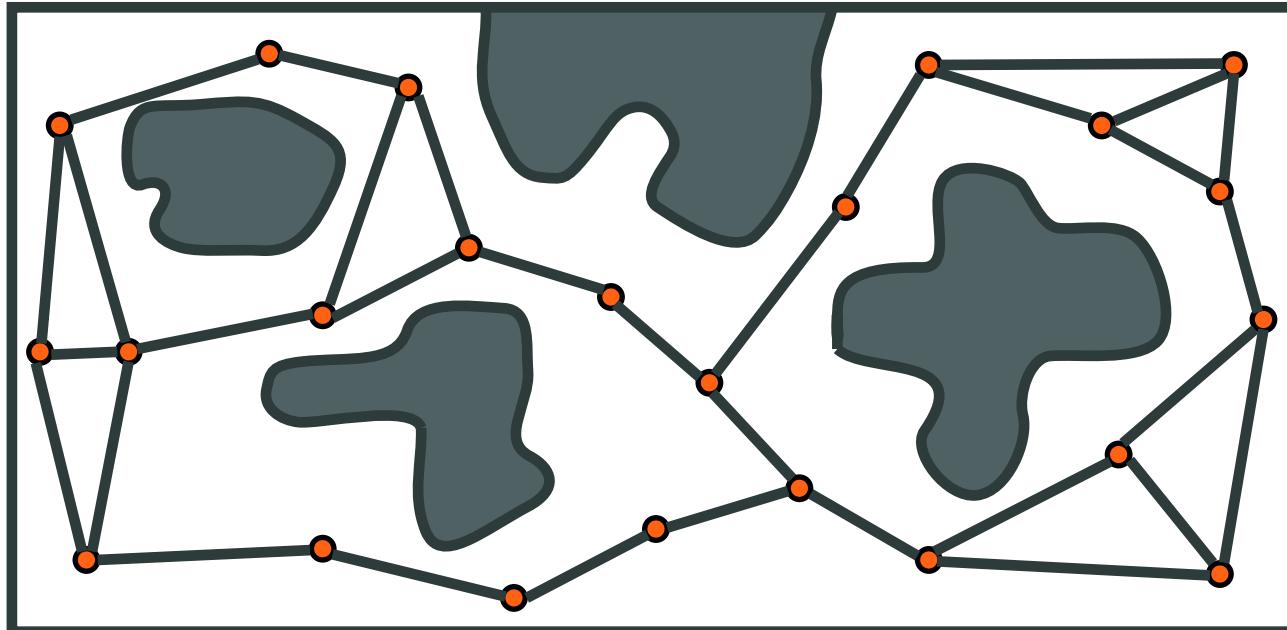
Operation of PRM



● : nodes, random states

— : edges, paths computed by local planner

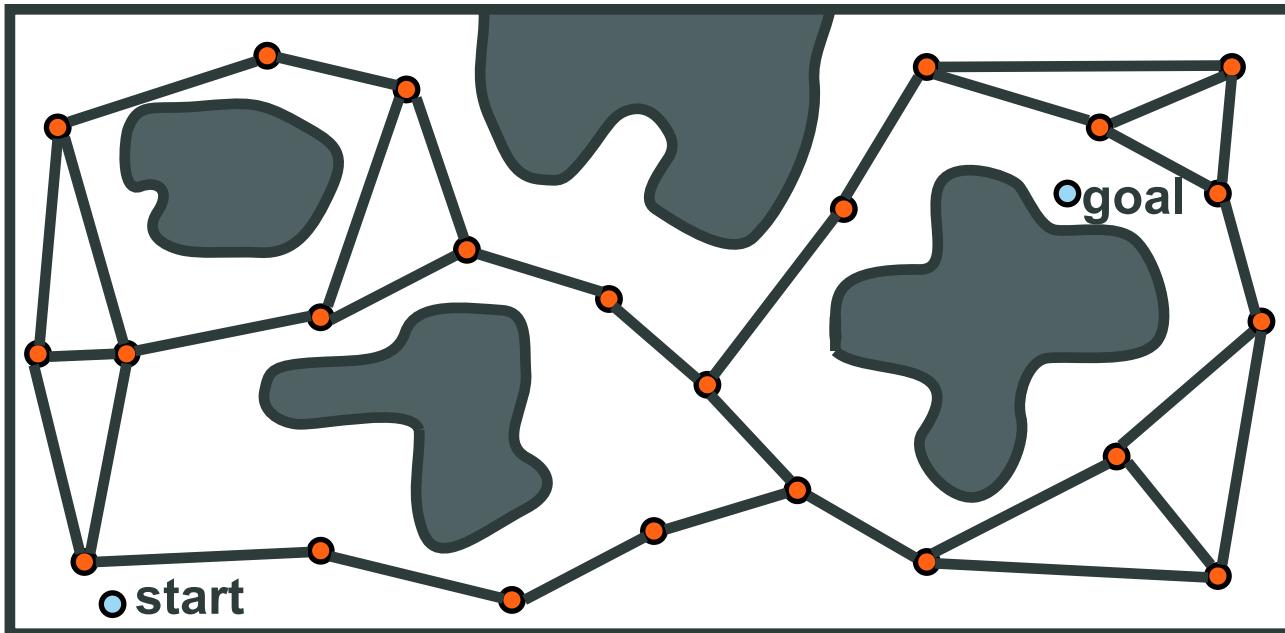
Operation of PRM



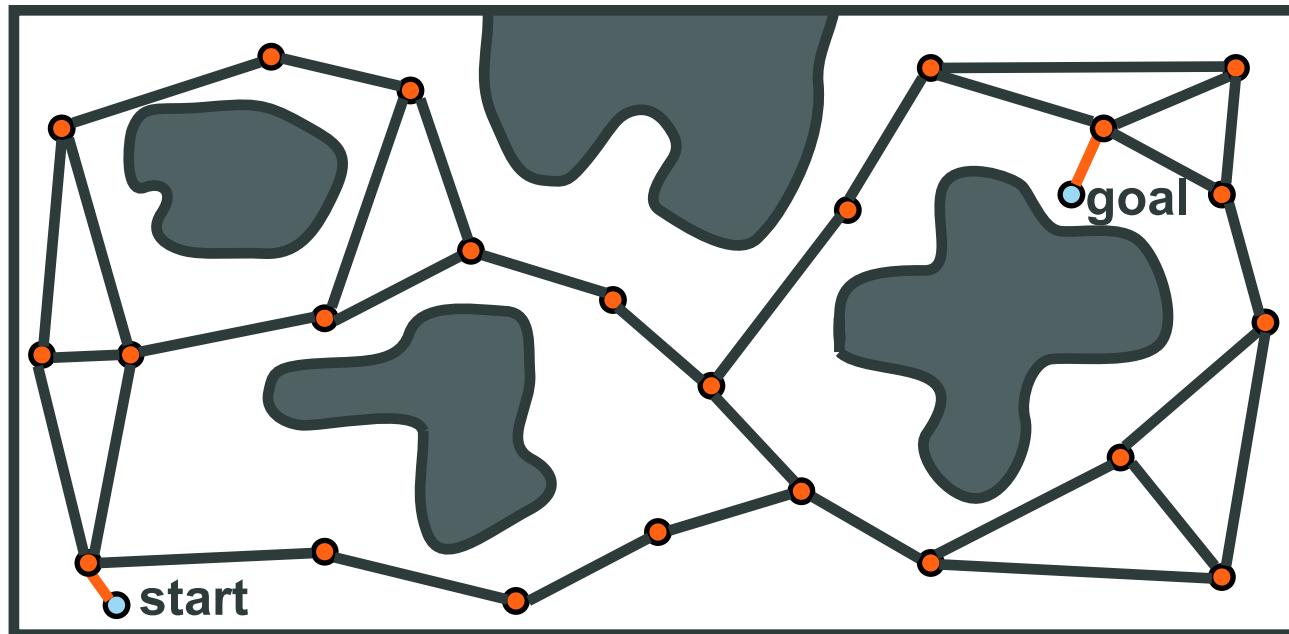
● : nodes, random states

— : edges, paths computed by local planner

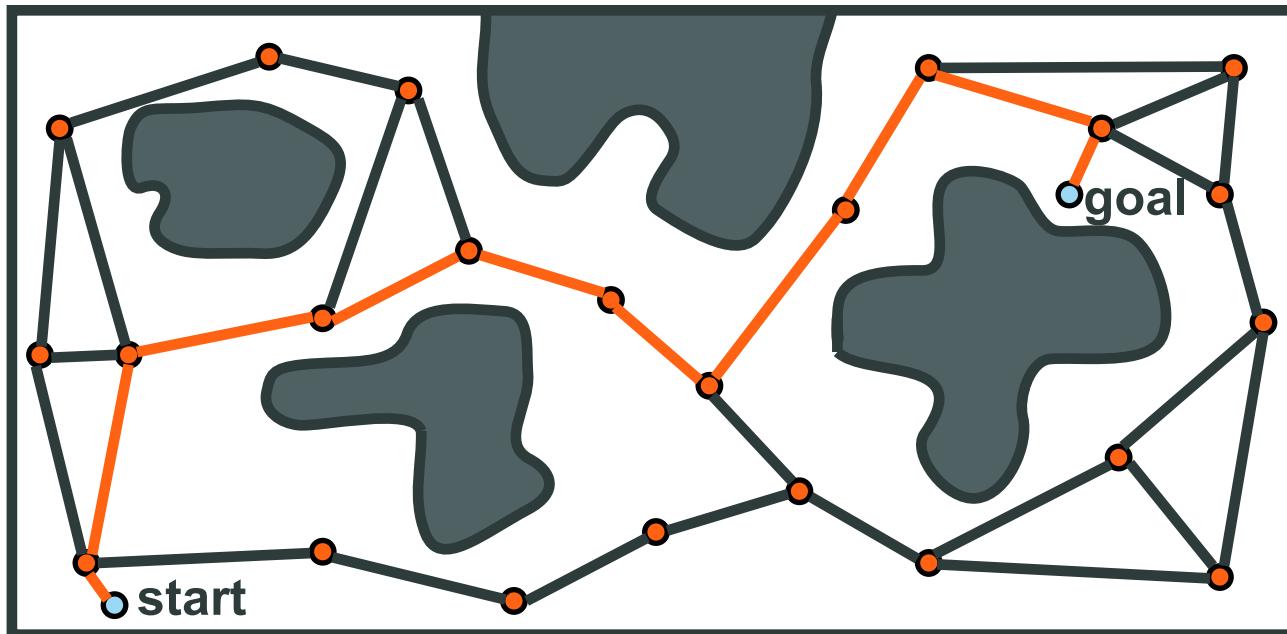
Answering a Query with PRM



Answering a Query with PRM



Answering a Query with PRM



Probabilistic Completeness

- Tradeoff between efficiency and completeness
- **Probabilistic completeness definition:**
 - Suppose $a, b \in Q_{free}$ can be connected by a path in Q_{free} . A planner is probabilistically complete if

$$\lim_{N \rightarrow \infty} \Pr((a, b) \text{ Failure}) = 0$$



Probability that PRM fails to answer query (a, b) with N samples

What if we only have a single-Query?

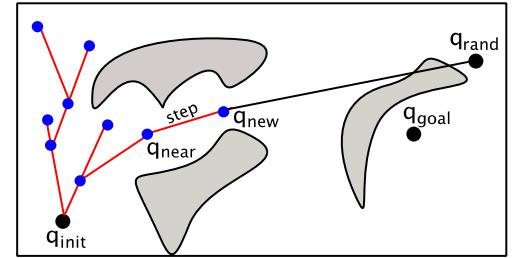
Rapidly-exploring Random Tree (RRT)

- **Idea:** Pull the tree toward random samples in the configuration space

RRT (q_{init}, q_{goal})

//initialize tree

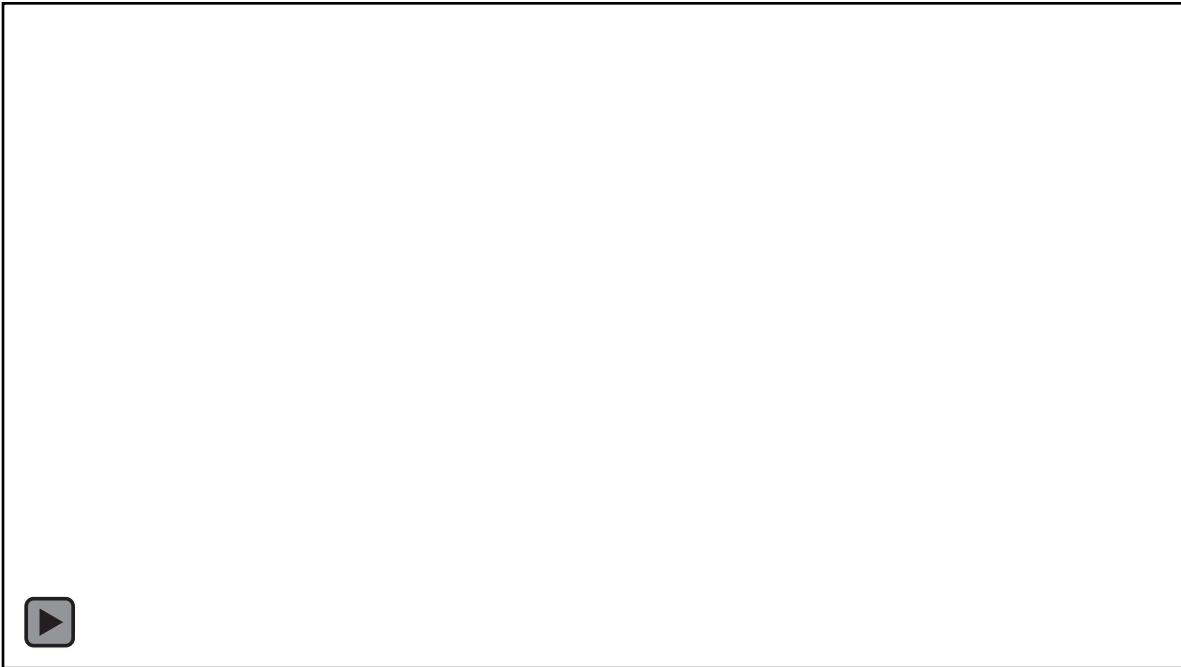
```
1:  $T \leftarrow$  create tree rooted at  $q_{init}$ 
2: while solution not found do
   // select configuration from tree
3:    $q_{rand} \leftarrow$  generate a random sample
4:    $q_{near} \leftarrow$  nearest configuration in  $T$  to  $q_{rand}$  according to distance  $\rho$ 
   // add new branch to tree from selected configuration
5:    $path \leftarrow$  generate path (not necessarily collision free) from  $q_{near}$  to  $q_{rand}$ 
6:   if IsSubpathCollisionFree( $path, 0, step$ ) then
7:      $q_{new} \leftarrow path(step)$ 
8:     add configuration  $q_{new}$  and edge  $(q_{near}, q_{new})$  to  $T$ 
   //check if a solution is found
9:   if  $\rho(q_{new}, q_{goal}) \approx 0$  then
10:    return solution path from root to  $q_{new}$ 
```



- RRT relies on nearest neighbors and distance metric $\rho : Q \times Q \leftarrow \mathbb{R}^{\geq 0}$

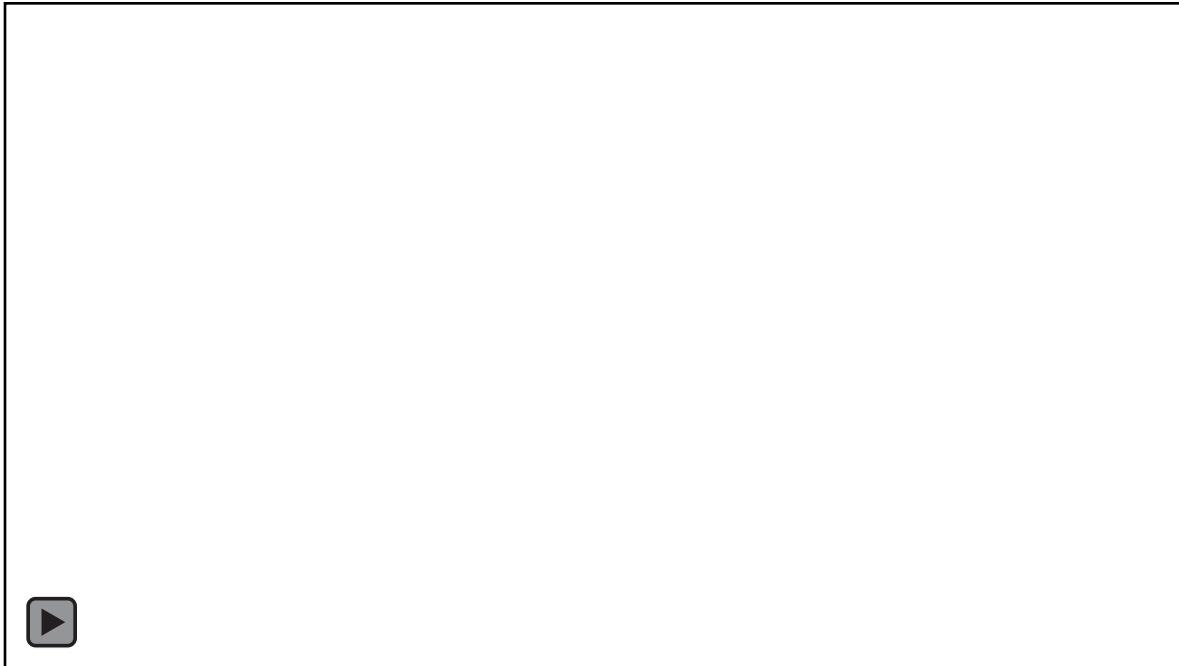
- RRT adds Voronoi bias to tree growth space

RRT-Expansion Step



Video Credit: Peter Murray

RRT- Selecting node to extend



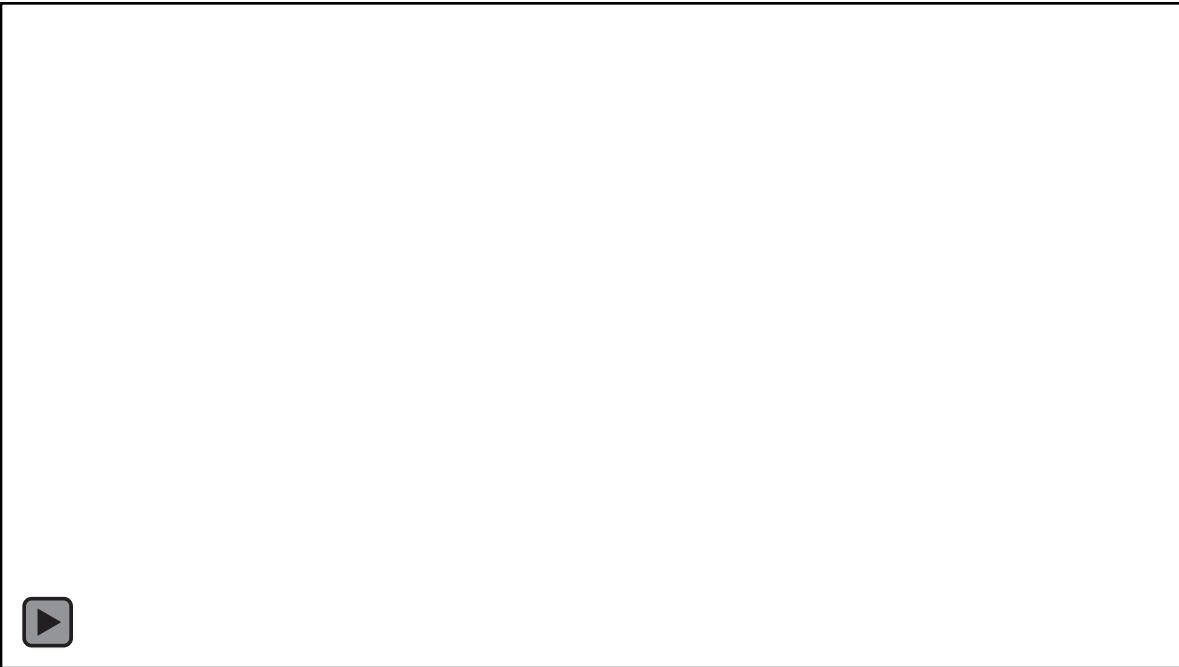
Video Credit: Peter Murray

RRT- Growth Illustration



Video Credit: Peter Murray

RRT- Voronoi -Bias



Video Credit: Peter Murray

Goal-Bias RRT-Improvements

Aspects for Improvement

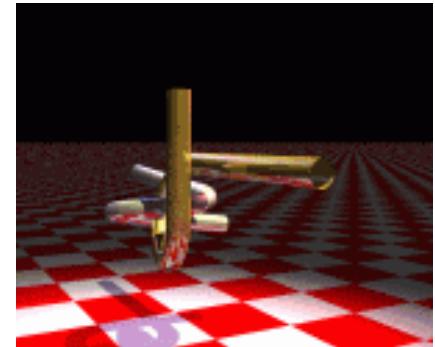
- BasicRRT does not take advantage of q_{goal}
- Tree is pulled towards random directions based on the uniform sampling of Q
- In particular, tree growth is not directed towards q_{goal}

Suggested Improvements in the Literature

- Introduce goal-bias to tree growth (known as GoalBiasRRT)
- q_{rand} is selected as q_{goal} with probability p
- q_{rand} is selected based on uniform sampling of Q with probability $1 - p$
- Probability p is commonly set to ≈ 0.05

RRT Examples: The Alpha Puzzle

- VERY hard 6DOF motion planning problem (long, winding narrow passage)
- *“In 2001, it was solved by using a balanced bidirectional RRT, developed by James Kuffner and Steve LaValle. There are no special heuristics or parameters that were tuned specifically for this problem. On a current PC (circa 2003), it consistently takes a few minutes to solve” –RRT website*
- RRT became famous in large part because it was able to solve this puzzle



RRT-Connect

```
BUILD_RRT( $q_{init}$ )
1    $\mathcal{T}$ .init( $q_{init}$ );
2   for  $k = 1$  to  $K$  do
3        $q_{rand} \leftarrow$  RANDOM_CONFIG();
4       EXTEND( $\mathcal{T}$ ,  $q_{rand}$ );
5   Return  $\mathcal{T}$ 
```

```
EXTEND( $\mathcal{T}, q$ )
1    $q_{near} \leftarrow$  NEAREST_NEIGHBOR( $q, \mathcal{T}$ );
2   if NEW_CONFIG( $q, q_{near}, q_{new}$ ) then
3        $\mathcal{T}$ .add_vertex( $q_{new}$ );
4        $\mathcal{T}$ .add_edge( $q_{near}, q_{new}$ );
5       if  $q_{new} = q$  then
6           Return Reached;
7       else
8           Return Advanced;
9   Return Trapped;
```

Figure 2: The basic RRT construction algorithm.

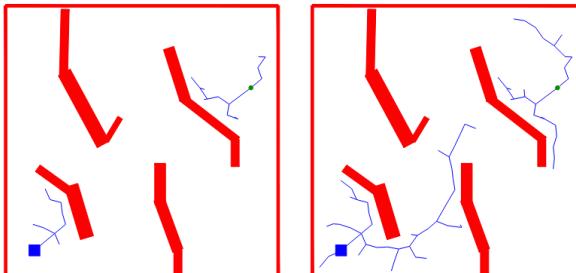
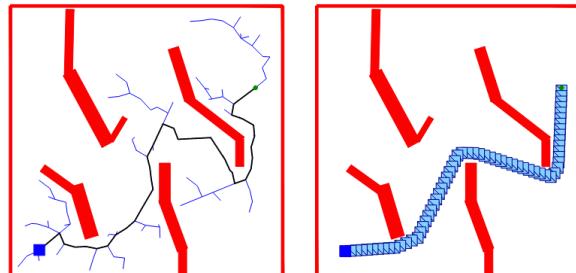


Figure 6: Growing two trees towards each other.

```
CONNECT( $\mathcal{T}, q$ )
1   repeat
2        $S \leftarrow$  EXTEND( $\mathcal{T}, q$ );
3   until not ( $S =$  Advanced)
4   Return  $S$ ;
```

```
RRT_CONNECT_PLANNER( $q_{init}, q_{goal}$ )
1    $\mathcal{T}_a$ .init( $q_{init}$ );  $\mathcal{T}_b$ .init( $q_{goal}$ );
2   for  $k = 1$  to  $K$  do
3        $q_{rand} \leftarrow$  RANDOM_CONFIG();
4       if not (EXTEND( $\mathcal{T}_a, q_{rand}$ ) = Trapped) then
5           if (CONNECT( $\mathcal{T}_b, q_{new}$ ) = Reached) then
6               Return PATH( $\mathcal{T}_a, \mathcal{T}_b$ );
7       SWAP( $\mathcal{T}_a, \mathcal{T}_b$ );
8   Return Failure
```

Figure 5: The RRT-Connect algorithm.



Expansive-Space Tree (EST)

- **Idea:** Push the tree frontier in the free configuration space
- EST relies on a probability distribution to guide tree growth
- EST associates a weight $w(q)$ with each tree configuration q
- $w(q)$ is a running estimate on importance of selecting q as the tree configuration from which to add a new tree branch
 - $w(q) = \frac{1}{1+\deg(q)}$
 - $w(q) = 1/(1 + \text{number of neighbors near } q)$
 - combination of different strategies

What if the problem is challenging?

- Biased Sampling,
 - Bridge Sampling
 - Gaussian Sampling
 - Importance Sampling

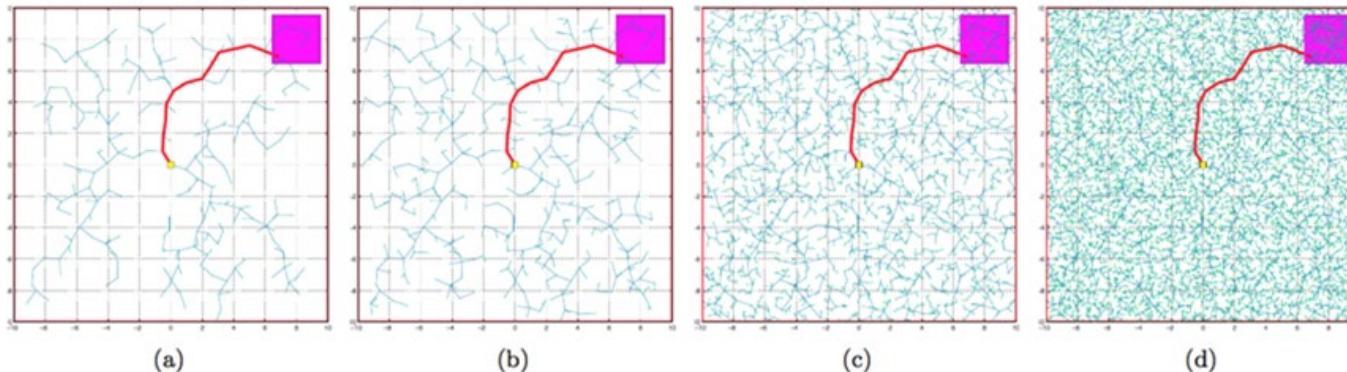
Limitations

- But what about the optimal solution?

Overview

- Motion Planning Definition
- Bug Algorithms
- Discrete Search/Roadmaps
- Configuration Space
- Sampling-Based Planners
- Asymptotically Optimal Planning
- Kinodynamic Planning
- Learning And Planning
- Task and Motion Planning
- Planning Under Uncertainty

More samples do not improve the solution



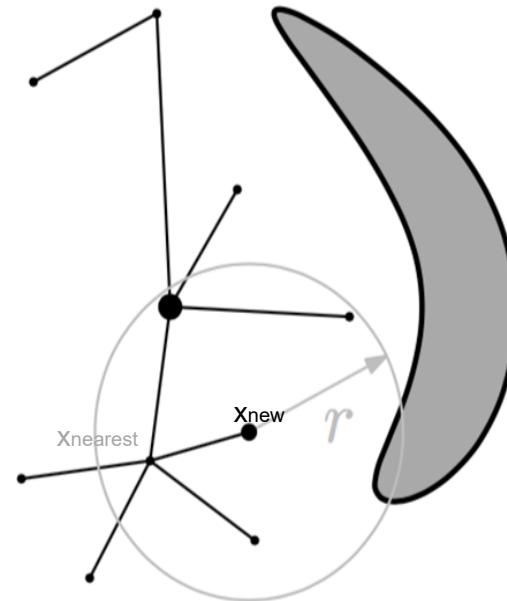
RRT* Extend

ExtendRRT*(G,x):

```
V' ← V; E' ← E;  
xnearest ← Nearest(G, x);  
xnew ← Steer(xnearest, x);  
if ObstacleFree(xnearest, xnew) then  
    V' ← V' ∪ {xnew};  
    xmin ← xnearest;  
    Xnear ← Near(G, xnew, |V|);  
    for all xnear ∈ Xnear do  
        if ObstacleFree(xnear, xnew) then  
            c' ← Cost(xnear) + c(Line(xnear, xnew));  
            if c' < Cost(xnew) then  
                xmin ← xnear;  
E' ← E' ∪ {(xmin, xnew)};  
for all xnear ∈ Xnear \ {xmin} do  
    if ObstacleFree(xnew, xnear) and  
    Cost(xnear) >  
    Cost(xnew) + c(Line(xnew, xnear)) then  
        xparent ← Parent(xnear);  
        E' ← E' \ {(xparent, xnear)};  
        E' ← E' ∪ {(xnew, xnear)};  
  
return G' = (V', E')
```

Karaman, Frazzoli. 2010.

Graphic from Gammell, 2016. Thesis.





RRT* Extend

define $\text{Near}(G, x, n)$ to be the set of all vertices within the closed ball of radius r_n centered at x , where

$$r_n = \min \left\{ \left(\frac{\gamma}{\zeta_d} \frac{\log n}{n} \right)^{1/d}, \eta \right\},$$

ExtendRRT*(G,x):

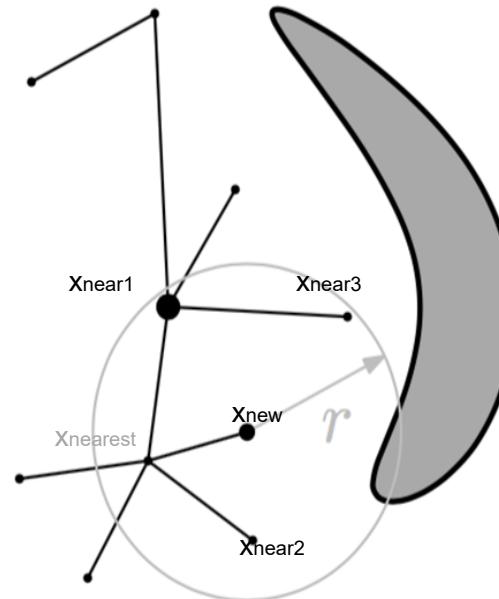
```

 $V' \leftarrow V; E' \leftarrow E;$ 
 $x_{\text{nearest}} \leftarrow \text{Nearest}(G, x);$ 
 $x_{\text{new}} \leftarrow \text{Steer}(x_{\text{nearest}}, x);$ 
if ObstacleFree( $x_{\text{nearest}}, x_{\text{new}}$ ) then
     $V' \leftarrow V' \cup \{x_{\text{new}}\};$ 
     $x_{\text{min}} \leftarrow x_{\text{nearest}};$ 
     $X_{\text{near}} \leftarrow \text{Near}(G, x_{\text{new}}, |V|);$ 
for all  $x_{\text{near}} \in X_{\text{near}}$  do
    if ObstacleFree( $x_{\text{near}}, x_{\text{new}}$ ) then
         $c' \leftarrow \text{Cost}(x_{\text{near}}) + c(\text{Line}(x_{\text{near}}, x_{\text{new}}));$ 
        if  $c' < \text{Cost}(x_{\text{new}})$  then
             $x_{\text{min}} \leftarrow x_{\text{near}};$ 
     $E' \leftarrow E' \cup \{(x_{\text{min}}, x_{\text{new}})\};$ 
for all  $x_{\text{near}} \in X_{\text{near}} \setminus \{x_{\text{min}}\}$  do
    if ObstacleFree( $x_{\text{new}}, x_{\text{near}}$ ) and
         $\text{Cost}(x_{\text{near}}) > \text{Cost}(x_{\text{new}}) + c(\text{Line}(x_{\text{new}}, x_{\text{near}}))$  then
             $x_{\text{parent}} \leftarrow \text{Parent}(x_{\text{near}});$ 
             $E' \leftarrow E' \setminus \{(x_{\text{parent}}, x_{\text{near}})\};$ 
             $E' \leftarrow E' \cup \{(x_{\text{new}}, x_{\text{near}})\};$ 
return  $G' = (V', E')$ 

```

Karaman, Frazzoli. 2010.

Graphic from Gammell, 2016. Thesis.



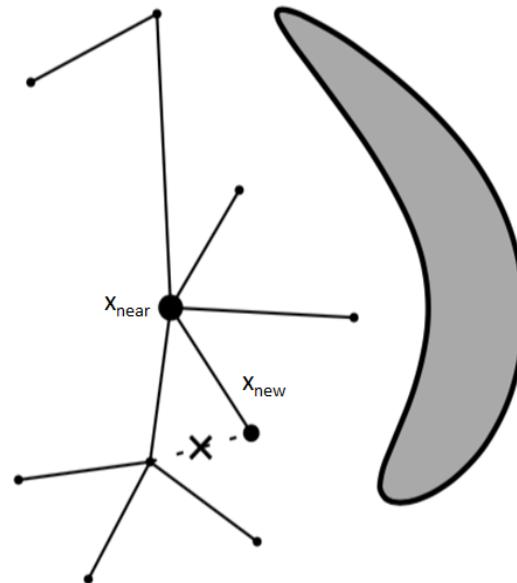
RRT* Extend

ExtendRRT*(G,x):

```
V' ← V; E' ← E;  
xnearest ← Nearest(G, x);  
xnew ← Steer(xnearest, x);  
if ObstacleFree(xnearest, xnew) then  
    V' ← V' ∪ {xnew};  
    xmin ← xnearest;  
    Xnear ← Near(G, xnew, |V|);  
for all xnear ∈ Xnear do  
    if ObstacleFree(xnear, xnew) then  
        c' ← Cost(xnear) + c(Line(xnear, xnew));  
        if c' < Cost(xnew) then  
            xmin ← xnear;  
E' ← E' ∪ {(xmin, xnew)};  
for all xnear ∈ Xnear \ {xmin} do  
    if ObstacleFree(xnew, xnear) and  
    Cost(xnear) >  
    Cost(xnew) + c(Line(xnew, xnear)) then  
        xparent ← Parent(xnear);  
        E' ← E' \ {(xparent, xnear)};  
        E' ← E' ∪ {(xnew, xnear)};  
return G' = (V', E')
```

Karaman, Frazzoli. 2010.

Graphic from Gammell, 2016. Thesis.

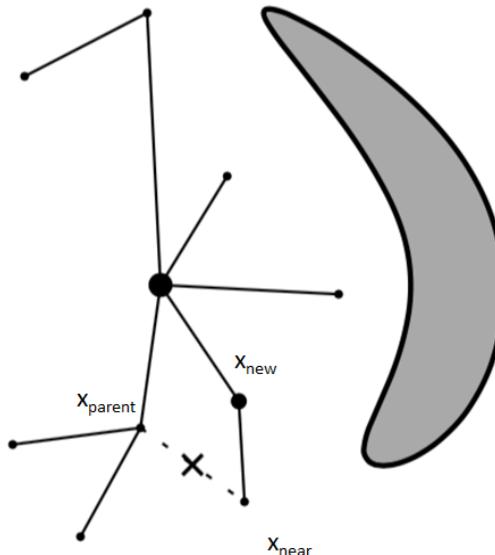


RRT* Extend

```
V' ← V; E' ← E;  
xnearest ← Nearest( $G, x$ );  
xnew ← Steer( $x_{nearest}, x$ );  
if ObstacleFree( $x_{nearest}, x_{new}$ ) then  
    | V' ← V' ∪ { $x_{new}$ };  
    | xmin ←  $x_{nearest}$ ;  
    | Xnear ← Near( $G, x_{new}, |V|$ );  
    | for all  $x_{near} \in X_{near}$  do  
    |   | if ObstacleFree( $x_{near}, x_{new}$ ) then  
    |   |   | c' ← Cost( $x_{near}$ ) + c(Line( $x_{near}, x_{new}$ ));  
    |   |   |   | if c' < Cost( $x_{new}$ ) then  
    |   |   |   |   | xmin ←  $x_{near}$ ;  
    |   |   | | E' ← E' ∪ {( $x_{parent}, x_{near}$ )}.  
    |   | for all  $x_{near} \in X_{near} \setminus \{x_{min}\}$  do  
    |   |   | if ObstacleFree( $x_{new}, x_{near}$ ) and  
    |   |   | Cost( $x_{near}$ ) >  
    |   |   | Cost( $x_{new}$ ) + c(Line( $x_{new}, x_{near}$ )) then  
    |   |   |   | xparent ← Parent( $x_{near}$ );  
    |   |   |   | E' ← E' \ {( $x_{parent}, x_{near}$ )};  
    |   |   |   | E' ← E' ∪ {( $x_{new}, x_{near}$ )};  
  
return  $G' = (V', E')$ 
```

Karaman, Frazzoli. 2010.

Graphic from Gammell, 2016. Thesis.



RRT* Extend vs RRT Extend

```

 $V' \leftarrow V; E' \leftarrow E;$ 
 $x_{\text{nearest}} \leftarrow \text{Nearest}(G, x);$ 
 $x_{\text{new}} \leftarrow \text{Steer}(x_{\text{nearest}}, x);$ 
 $\text{if ObstacleFree}(x_{\text{nearest}}, x_{\text{new}}) \text{ then}$ 
     $V' \leftarrow V' \cup \{x_{\text{new}}\};$ 
     $x_{\text{min}} \leftarrow x_{\text{nearest}};$ 
     $X_{\text{near}} \leftarrow \text{Near}(G, x_{\text{new}}, |V|);$ 
     $\text{for all } x_{\text{near}} \in X_{\text{near}} \text{ do}$ 
         $\quad \text{if ObstacleFree}(x_{\text{near}}, x_{\text{new}}) \text{ then}$ 
             $\quad \quad c' \leftarrow \text{Cost}(x_{\text{near}}) + c(\text{Line}(x_{\text{near}}, x_{\text{new}}));$ 
             $\quad \quad \text{if } c' < \text{Cost}(x_{\text{new}}) \text{ then}$ 
                 $\quad \quad \quad x_{\text{min}} \leftarrow x_{\text{near}};$ 
     $E' \leftarrow E' \cup \{(x_{\text{min}}, x_{\text{new}})\};$ 
     $\text{for all } x_{\text{near}} \in X_{\text{near}} \setminus \{x_{\text{min}}\} \text{ do}$ 
         $\quad \text{if ObstacleFree}(x_{\text{new}}, x_{\text{near}}) \text{ and}$ 
         $\quad \quad \text{Cost}(x_{\text{near}}) >$ 
         $\quad \quad \text{Cost}(x_{\text{new}}) + c(\text{Line}(x_{\text{new}}, x_{\text{near}})) \text{ then}$ 
             $\quad \quad \quad x_{\text{parent}} \leftarrow \text{Parent}(x_{\text{near}});$ 
             $\quad \quad \quad E' \leftarrow E' \setminus \{(x_{\text{parent}}, x_{\text{near}})\};$ 
             $\quad \quad \quad E' \leftarrow E' \cup \{(x_{\text{new}}, x_{\text{near}})\};$ 
 $\text{return } G' = (V', E')$ 

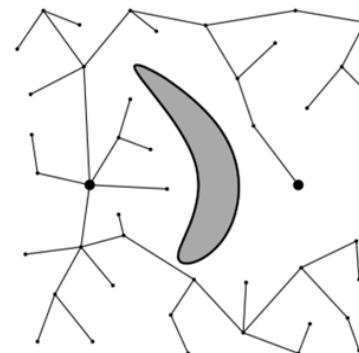
```

Karaman, Frazzoli. 2010.

```

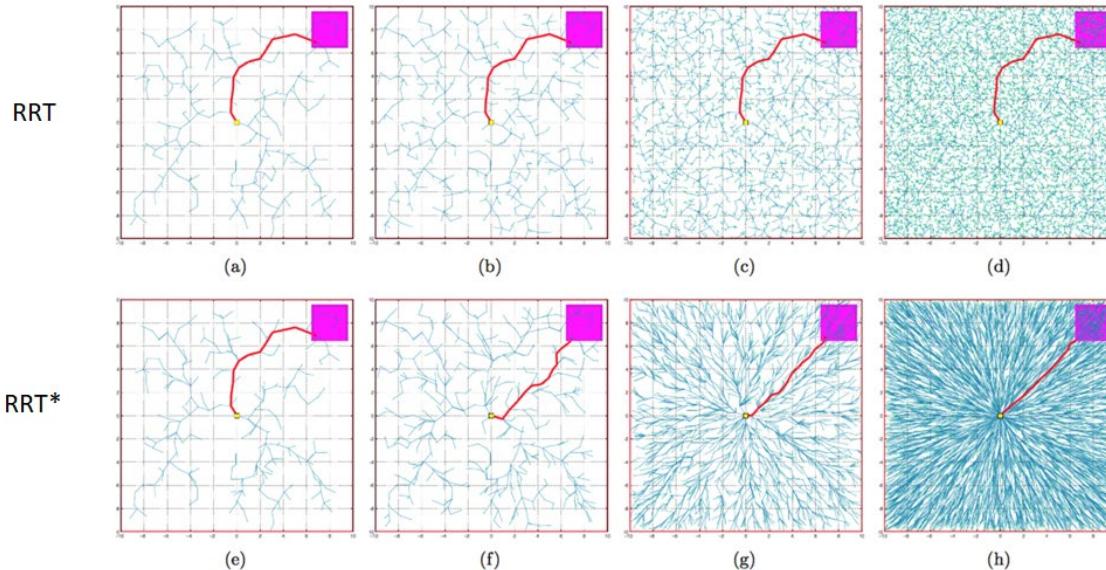
 $V' \leftarrow V; E' \leftarrow E;$ 
 $x_{\text{nearest}} \leftarrow \text{Nearest}(G, x);$ 
 $x_{\text{new}} \leftarrow \text{Steer}(x_{\text{nearest}}, x);$ 
 $\text{if ObstacleFree}(x_{\text{nearest}}, x_{\text{new}}) \text{ then}$ 
     $V' \leftarrow V' \cup \{x_{\text{new}}\};$ 
     $E' \leftarrow E' \cup \{(x_{\text{nearest}}, x_{\text{new}})\};$ 
 $\text{return } G' = (V', E')$ 

```





RRT*



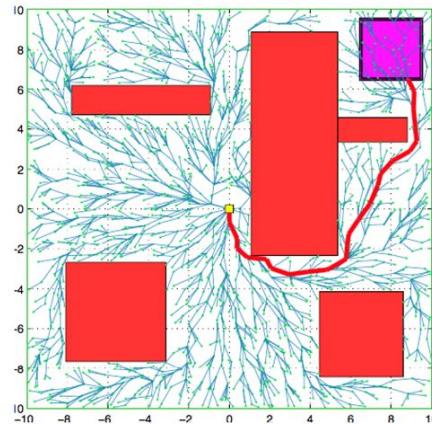
Karaman, Frazzoli. 2010.



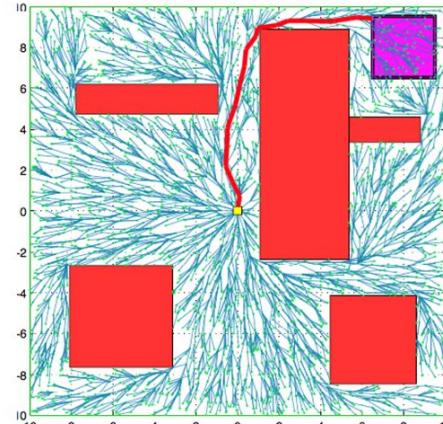
RRT*

RRT*

2500 Iterations



5000 Iterations



Karaman, Frazzoli. 2010.



PRM*

- “Simple” PRM

```
 $V \leftarrow \{x_{\text{init}}\} \cup \{\text{SampleFree}_i\}_{i=1,\dots,n}; E \leftarrow \emptyset;$ 
foreach  $v \in V$  do
   $U \leftarrow \text{Near}(G = (V, E), v, r) \setminus \{v\};$ 
  foreach  $u \in U$  do
    if CollisionFree( $v, u$ ) then  $E \leftarrow E \cup \{(v, u), (u, v)\}$ 
return  $G = (V, E);$ 
```

- PRM*

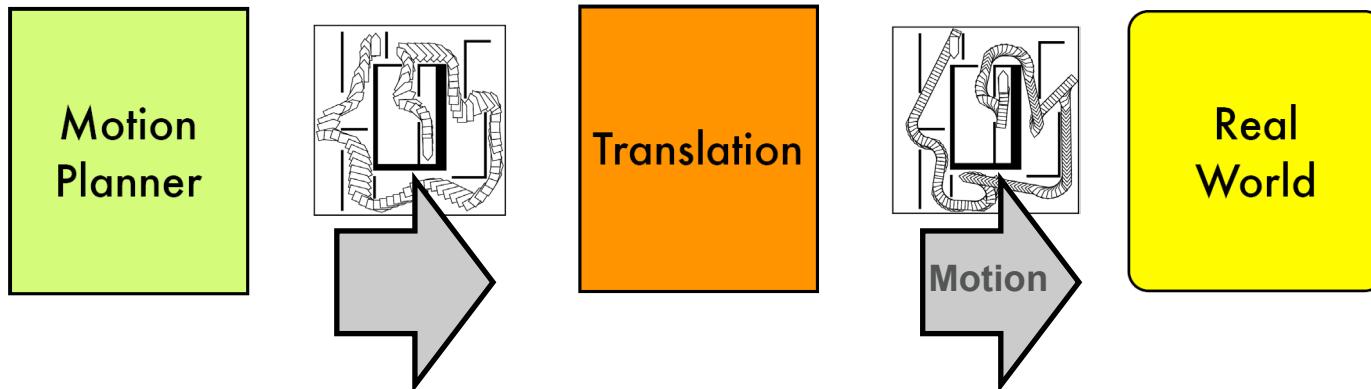
- Radius shrinks with more vertices
- One change, and its' optimal!

```
 $V \leftarrow \{x_{\text{init}}\} \cup \{\text{SampleFree}_i\}_{i=1,\dots,n}; E \leftarrow \emptyset;$ 
foreach  $v \in V$  do
   $U \leftarrow \text{Near}(G = (V, E), v, \gamma_{\text{PRM}}(\log(n)/n)^{1/d}) \setminus \{v\};$ 
  foreach  $u \in U$  do
    if CollisionFree( $v, u$ ) then  $E \leftarrow E \cup \{(v, u), (u, v)\}$ 
return  $G = (V, E);$ 
```

Overview

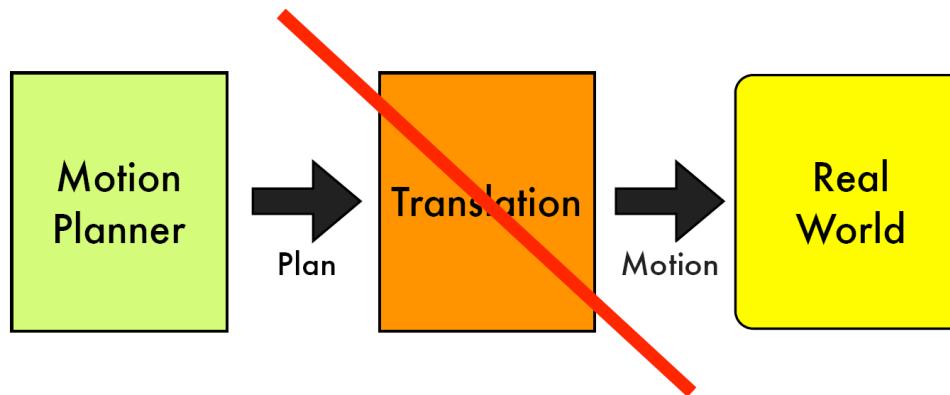
- Motion Planning Definition
- Bug Algorithms
- Discrete Search/Roadmaps
- Configuration Space
- Sampling-Based Planners
- Asymptotically Optimal Planning
- Kinodynamic Planning
- Learning And Planning
- Task and Motion Planning
- Planning Under Uncertainty

Decoupled Motion Planning Paradigm



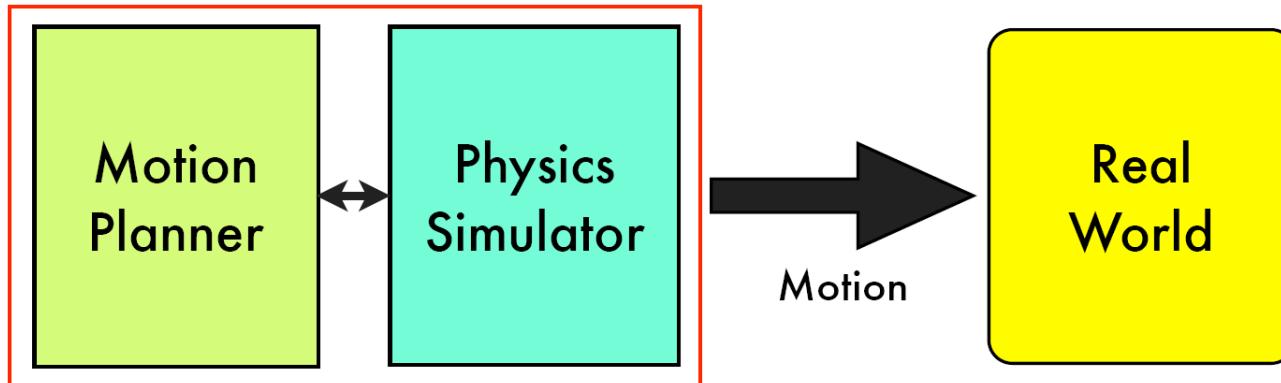
- Finds a geometrically feasible plan with a planner that is approximately correct
- Satisfy any remaining physical constraints using a translation step
- Execute in the real world

Shortcomings of Decoupled Motion Planning



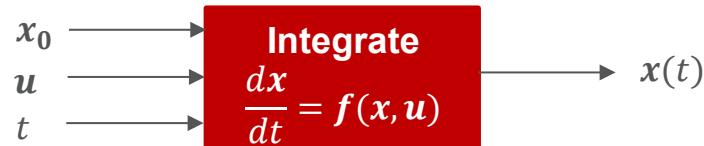
- Motions are often low quality or inadmissible
- Translating “collision-free” paths to physical motion is hard

Another approach (Native Approach)



- Reason directly about physical constraints in a simulation

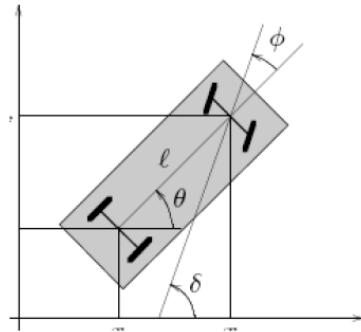
Motion Planning with Physics-based Simulations



- Physics engines can be used to simulate robot motions
- Physics engines provide greater simulation accuracy
- Physics engines can take into account friction, gravity, and interactions of the robot with objects in the environment
- Existing tools



Simple Car VS Real World Car



simple Car:

$$\dot{x} = \begin{pmatrix} \dot{x} \\ \dot{y} \\ \dot{\theta} \end{pmatrix}$$

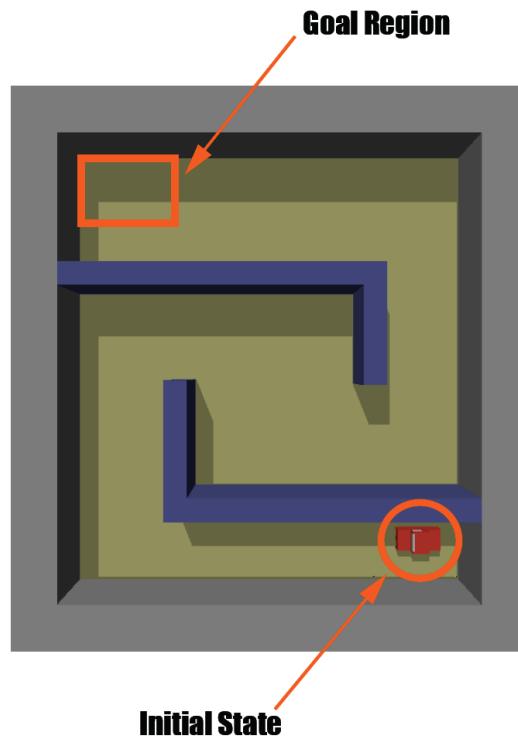
$$= \begin{pmatrix} u_v \cos \theta \\ u_v \sin \theta \\ u_v \tan u_\phi \end{pmatrix} \frac{L}{L}$$



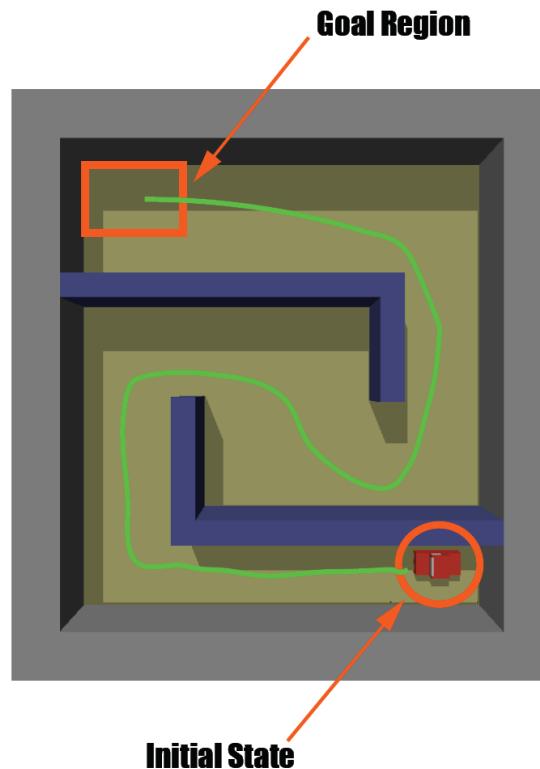
Real Car:

- A real car is even more complex
- Turning is controlled by steering
- Car motion is determined by tire contact
- Skidding, slip and other behavior

Example Problem

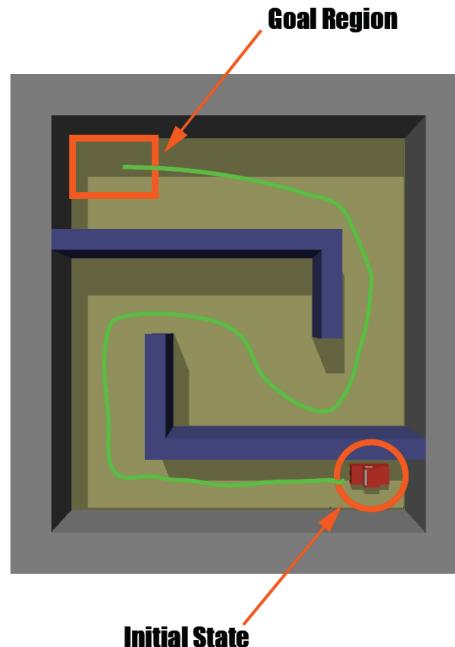


Example Problem



This is a Search Problem

- State space is enormous, continuous and hard to discretize
- Solution paths are usually “deep”
- There are many solutions



Moving between two states (without obstacles)

- Two-Point Boundary Value Problem (BVP):
 - Find a control sequence to take system from state X_I to state X_G while obeying kinematic constraints.
 - How to solve this problem?





Shooting Method

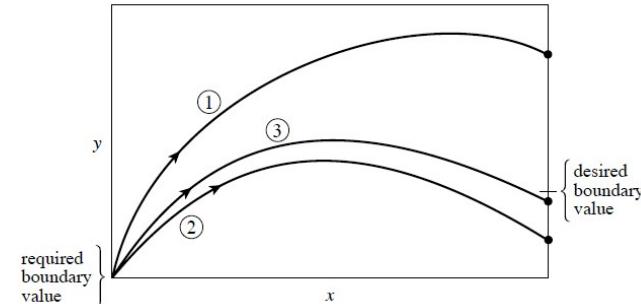
- “Shoot” out trajectories in different directions until a trajectory of the desired boundary value is found.

- System

$$\frac{dy}{dx} + f(x, y) = 0$$

- Boundary condition

$$y(0) = 0, y(1) = 1$$



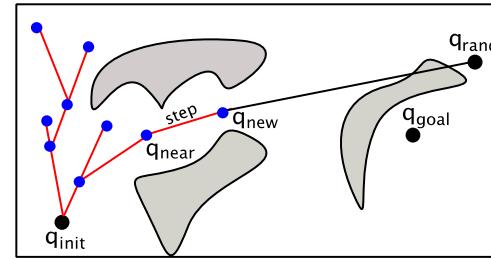
- Convert to an initial boundary problem: Pick initial guess, iteratively get closer to goal

Kinodynamic Motion Planning with Tree-Based Methods

Tree-based Approaches

RRT

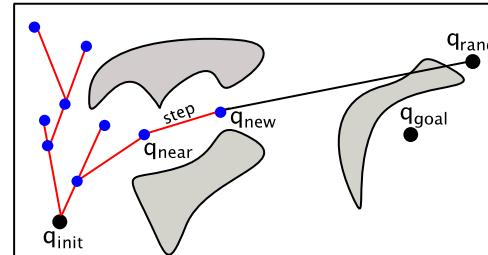
- 1: $T \leftarrow$ create tree rooted at x_0
- 2: **While** solution not found **do**
 - \|\ select state from tree
 - 3: $x_{rand} \leftarrow$ StateSample()
 - 4: $x_{near} \leftarrow$ nearest state in T to x_{rand} according to distance ρ
 - \|\ add new branch to tree from selected state
 - 5: $\lambda \leftarrow$ GenerateLocalTrajectory(x_{near}, x_{rand})
 - 6: **if** IsSubTrajectoryValid($\lambda, 0, step$) **then**
 - 7: $x_{new} \leftarrow \lambda(step)$
 - 8: add configuration x_{new} and edge (x_{near}, x_{new}) to T
 - \|\ check if a solution is found
 - 9: **if** $\rho(x_{new}, x_{goal}) \approx 0$ **then**
 - 10: **return** solution trajectory from root to x_{new}



Kinodynamic Motion Planning with Tree-Based Methods

Implementation of Tree-based Approaches

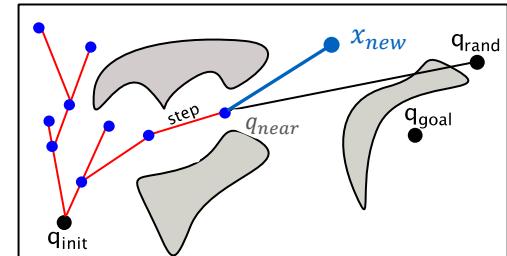
- `SampleState()`
 - random values for state components
- $\rho: X \times X \rightarrow \mathbb{R}^{\geq 0}$
 - Distance metric between states
- $\lambda \leftarrow \text{GenerateLocalTrajectory}(x_{near}, x_{rand})$
 - does it not create the same two-boundary value problems as in PRM?
 - is it necessary to connect to x_{rand} ?
 - does it suffice to just come close to x_{rand} ?
- `IsSubTrajectoryValid(λ , 0, steps)`
 - Incremental approach



Kinodynamic Motion Planning with Tree-Based Methods

Idea for avoiding the 2-point Boundary

Problem: Rather than computing a trajectory from x_{near} to x_{rand} compute a trajectory that starts at x_{near} and extends toward x_{rand}

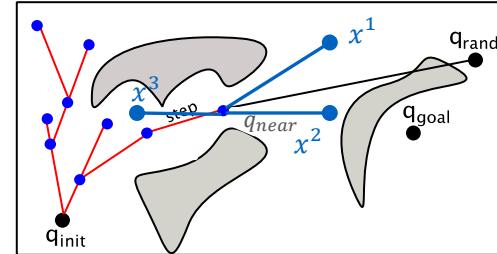


- Approach 1 – extend according to random control
 - Sample random control u in U
 - Integrate equations of motions when applying u to x_{near} for Δt units of time, i.e.,

$$\lambda \leftarrow x(t) = x_{near} + \int_0^{\Delta t} f(x(\tau), u) d\tau$$

Kinodynamic Motion Planning with Tree-Based Methods

Idea for avoiding the 2-point Boundary Problem: Rather than computing a trajectory from x_{near} to x_{rand} compute a trajectory that starts at x_{near} and extends toward x_{rand}

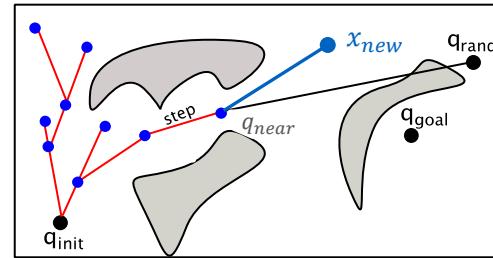


- Approach 2 – find the best-out-of-many random controls

```
1: for  $i = 1, \dots, m$  do
2:    $u \leftarrow$  sample random control in  $U$ 
3:    $\lambda \leftarrow x(t) = x_{near} + \int_0^{\Delta t} f(x(\tau), u) d\tau$ 
4:    $d_i \leftarrow \rho(x_{rand}, \lambda_i(\Delta t))$ 
5: return  $\lambda_i$  with minimum  $d_i$ 
```

Kinodynamic Motion Planning with Tree-Based Methods

Idea for avoiding the 2-point Boundary Problem: Rather than computing a trajectory from x_{near} to x_{rand} compute a trajectory that starts at x_{near} and extends toward x_{rand}



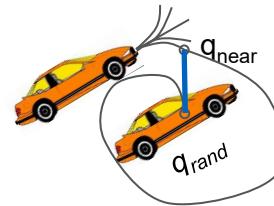
- Approach 2 – find the best-out-of-many random controls

```
1: for  $i = 1, \dots, m$  do
2:    $u \leftarrow$  sample random control in  $U$ 
3:    $\lambda \leftarrow x(t) = x_{near} + \int_0^{\Delta t} f(x(\tau), u) d\tau$ 
4:    $d_i \leftarrow \rho(x_{rand}, \lambda_i(\Delta t))$ 
5: return  $\lambda_i$  with minimum  $d_i$ 
```

RRTs and Distance Metrics

- Hard to define d , the distance metric
 - Mixing velocity, position, rotation ,etc.

How do you pick a good q_{near} ?



Configurations are close according to
Euclidian metric, but actual distance is large

Overview

- Motion Planning Definition
- Bug Algorithms
- Discrete Search/Roadmaps
- Configuration Space
- Sampling-Based Planners
- Asymptotically Optimal Planning
- Kinodynamic Planning
- Learning And Planning
- Task and Motion Planning
- Planning Under Uncertainty

Machine Learning and Motion Planning

Training/Learning Phase



Testing/Inference Phase



Biasing the Sampling Distribution

- Theoretical Insights supporting biased sampling

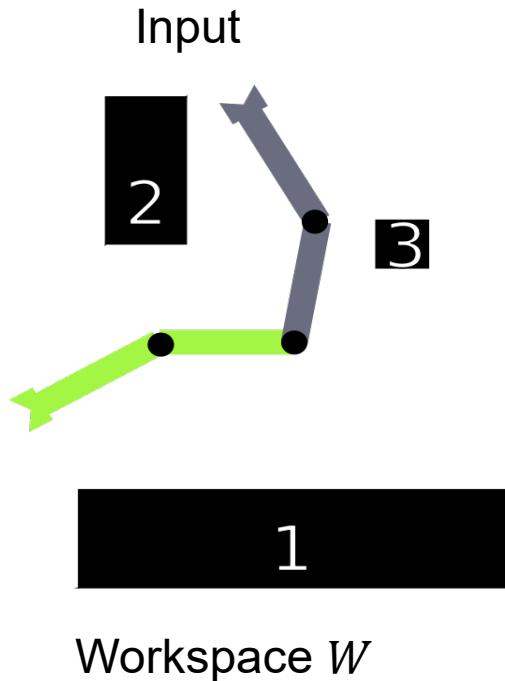
Algorithm 1: General Sampling Based Planner

Input : Number of iterations N

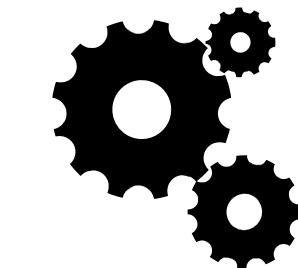
Output : Graph structure G

```
1 while  $i \leq N$  or solutionFound() do
2   | x ~ Uniform()
3   | update(G, x)
4 end
5 return G
```

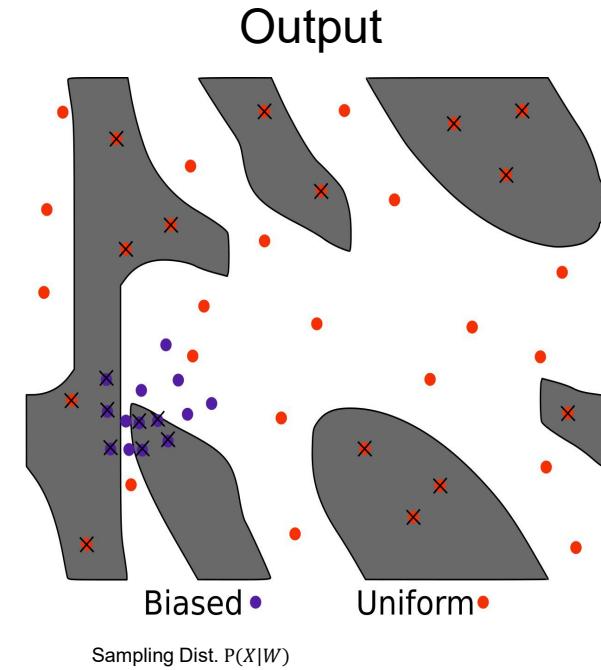
Learning a biased sampling distribution



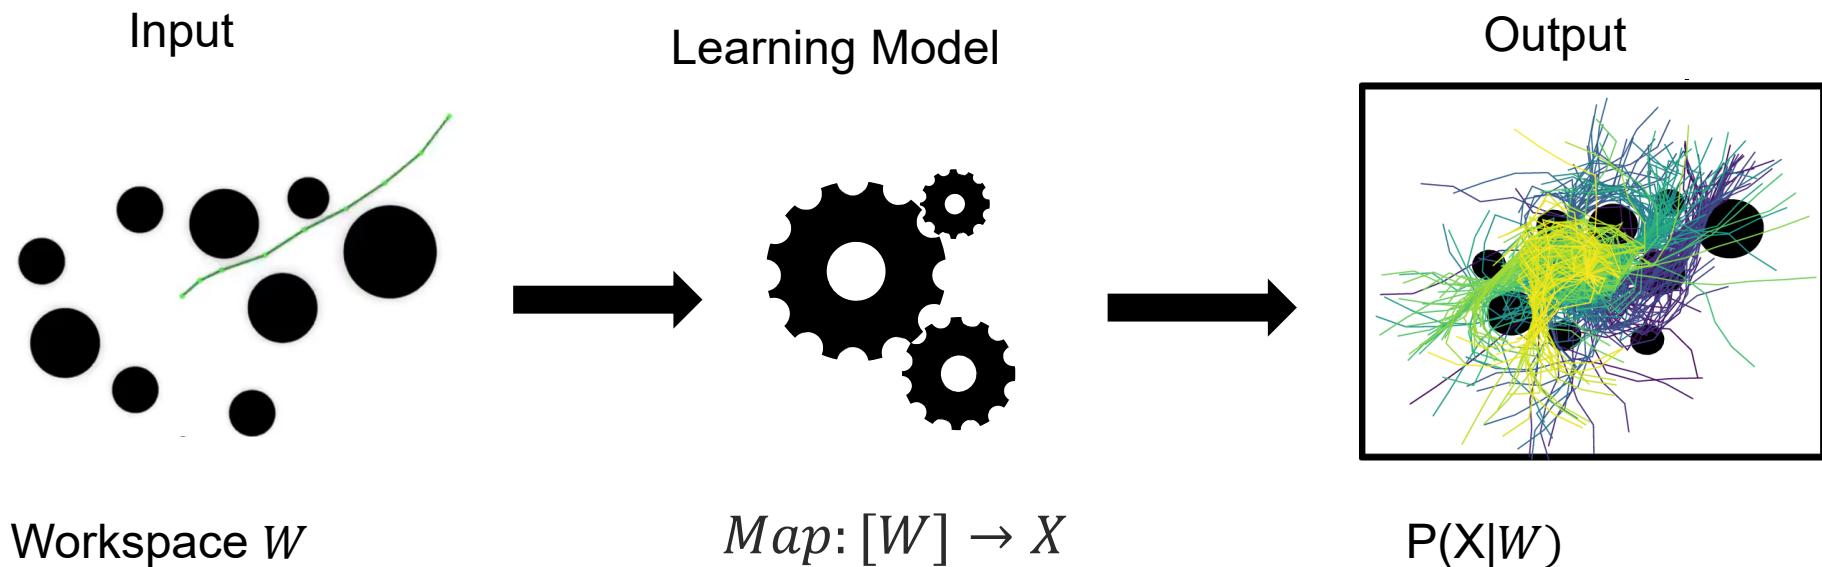
Learning Model



$Map: [W] \rightarrow X$

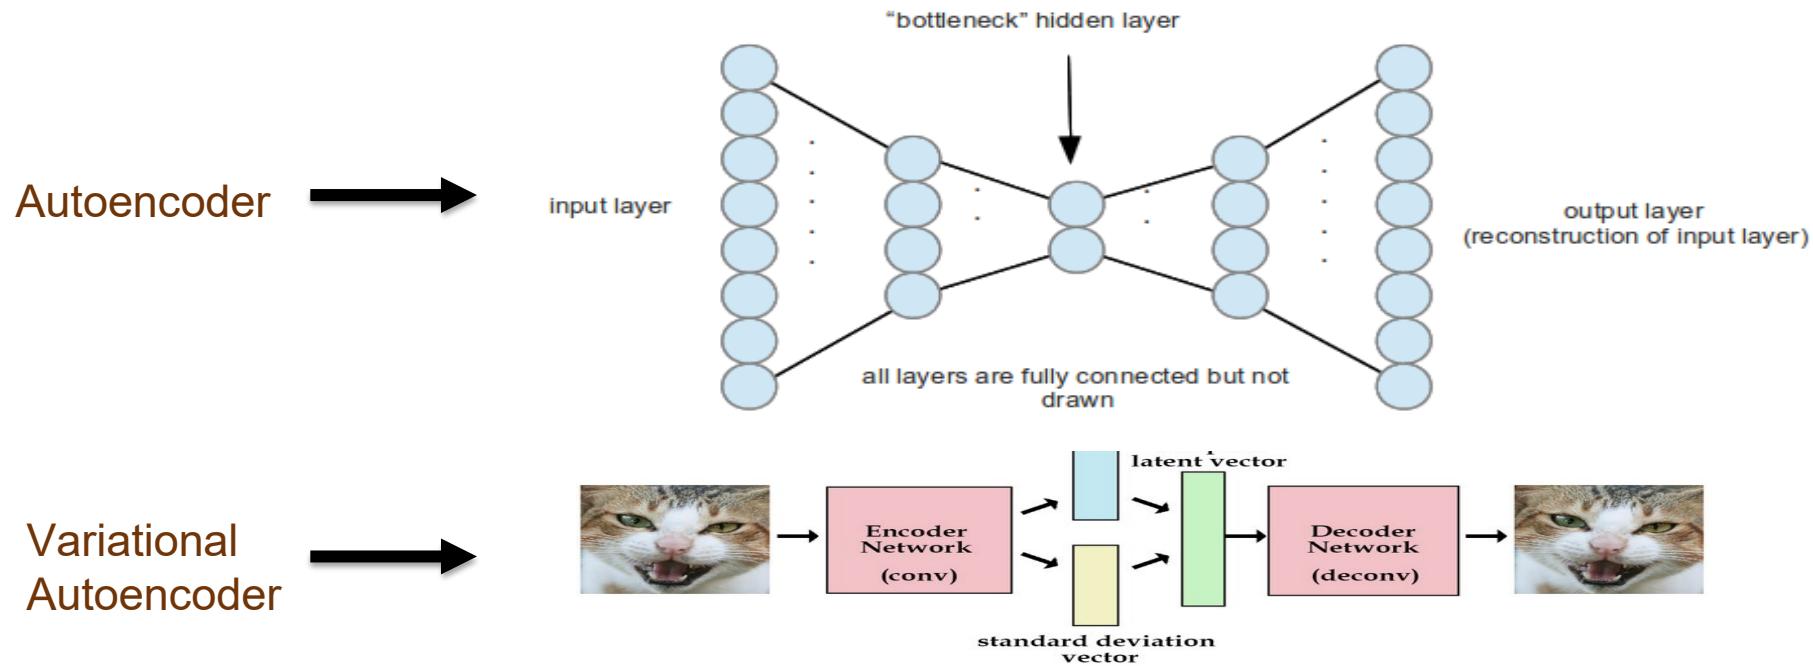


Retrieval of Database Samples

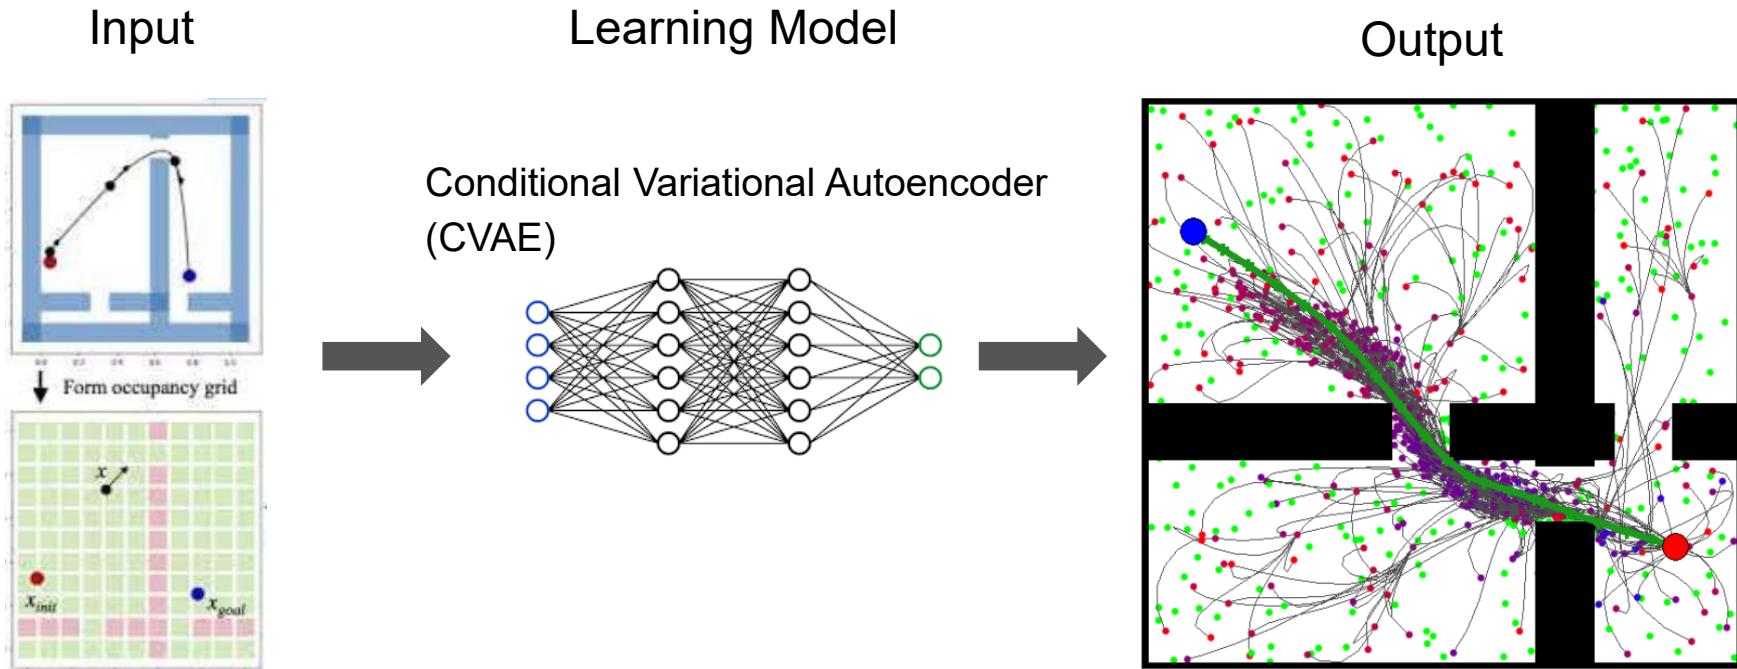


Variational Autoencoders

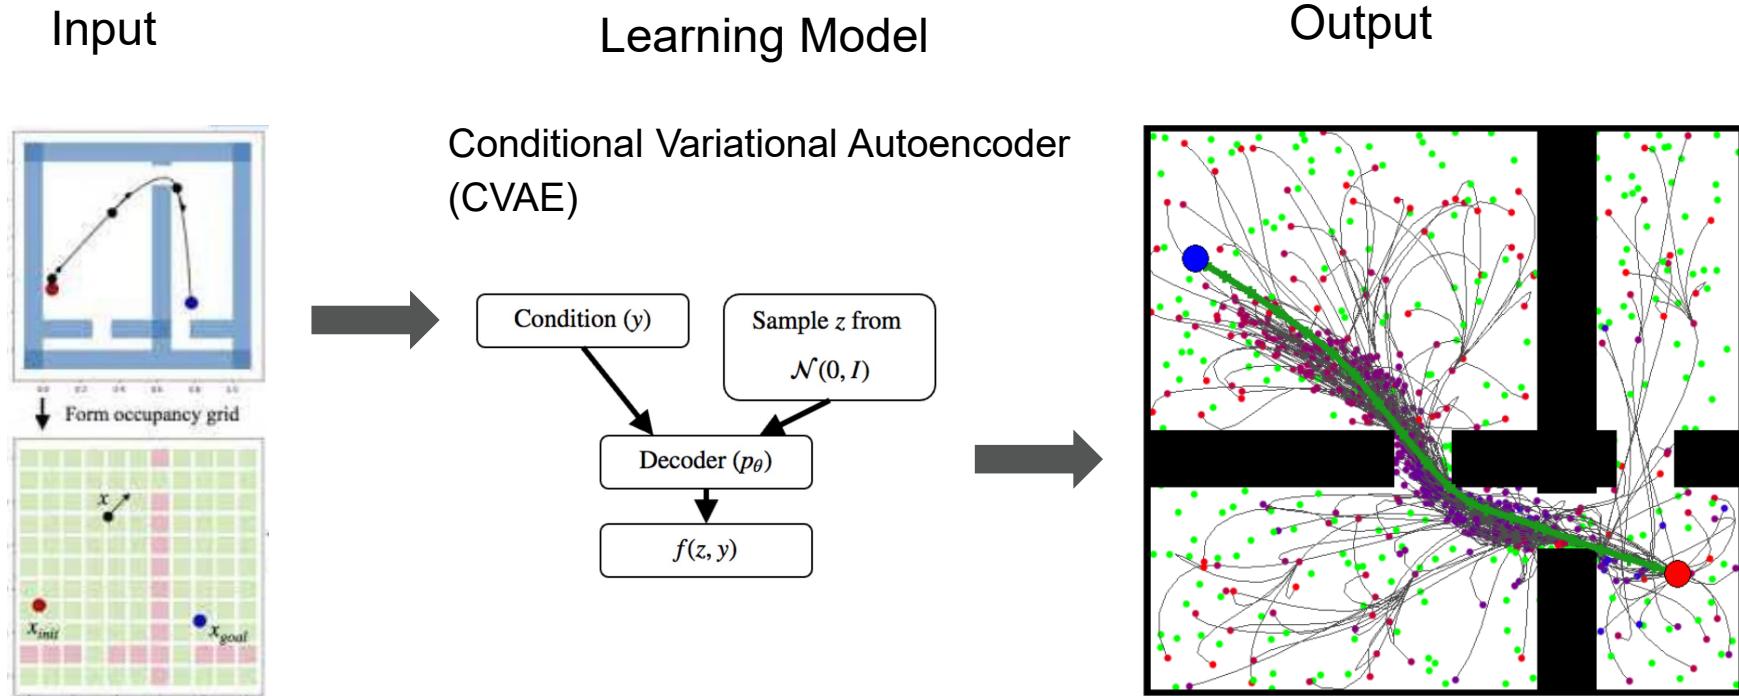
How to learn a sampling distribution with a neural network?



A neural network baseline

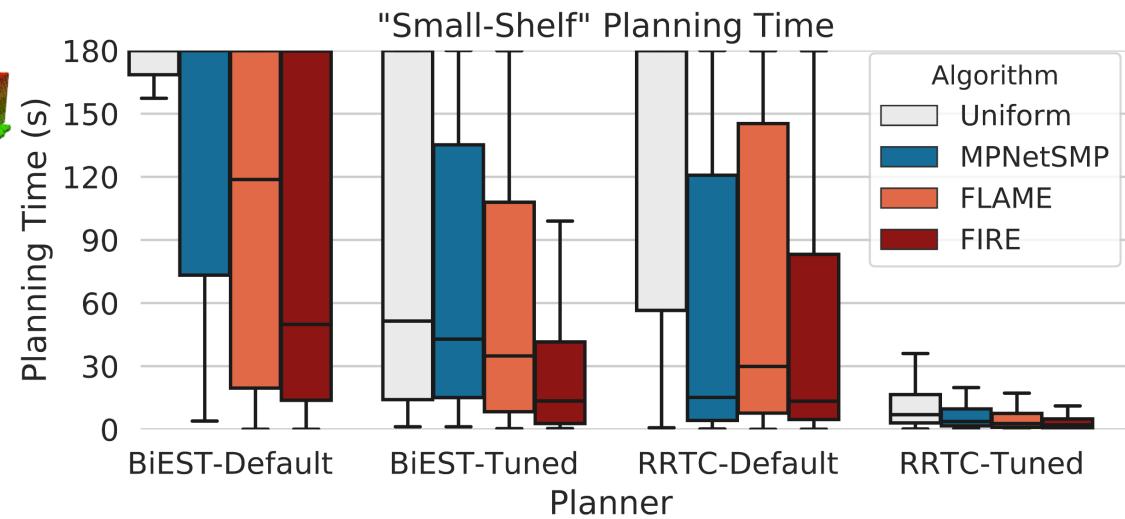
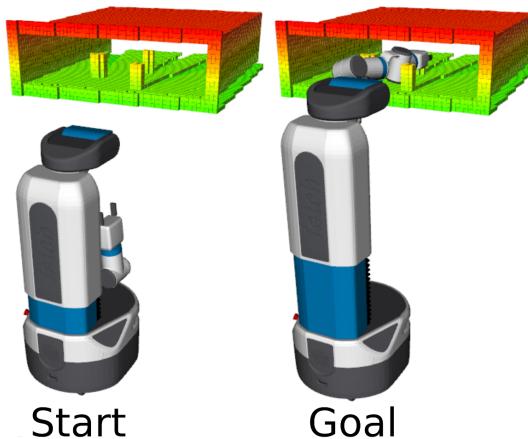


A neural network baseline (CVAE)



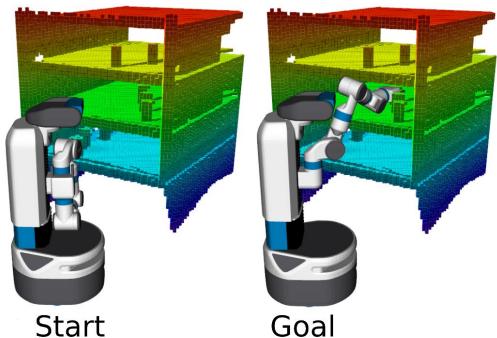
Generalizing within the training class of problems

"Small-Shelf" Problems (Train)

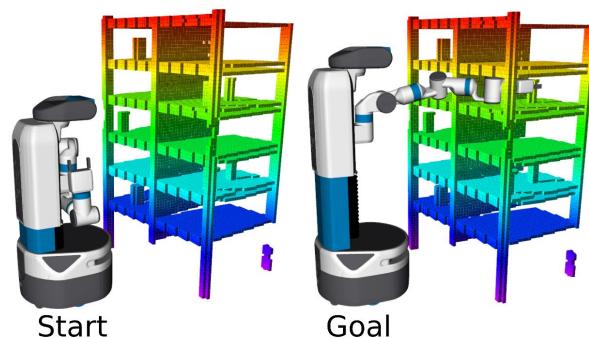


Generalizing outside the training class of problems

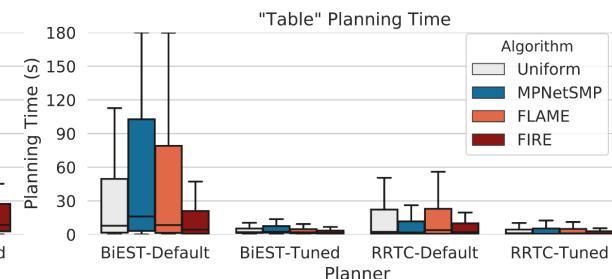
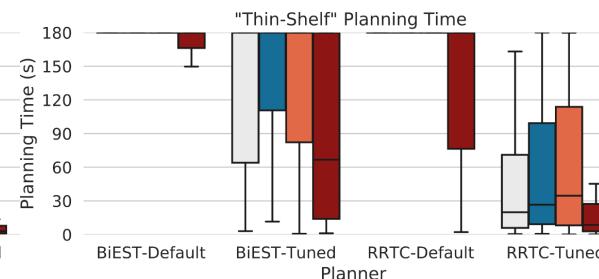
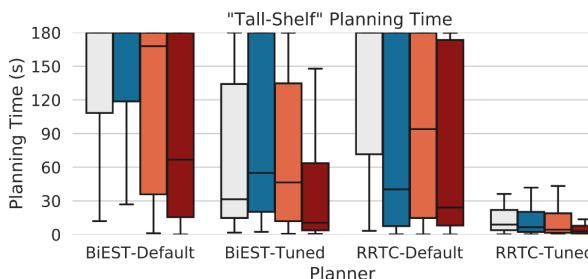
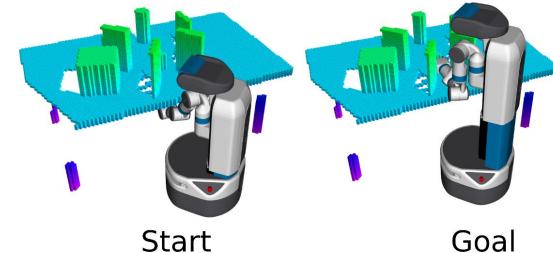
"Tall-Shelf" Dataset (Test)



"Thin-Shelf" Dataset (Test)



"Table" Dataset (Test)



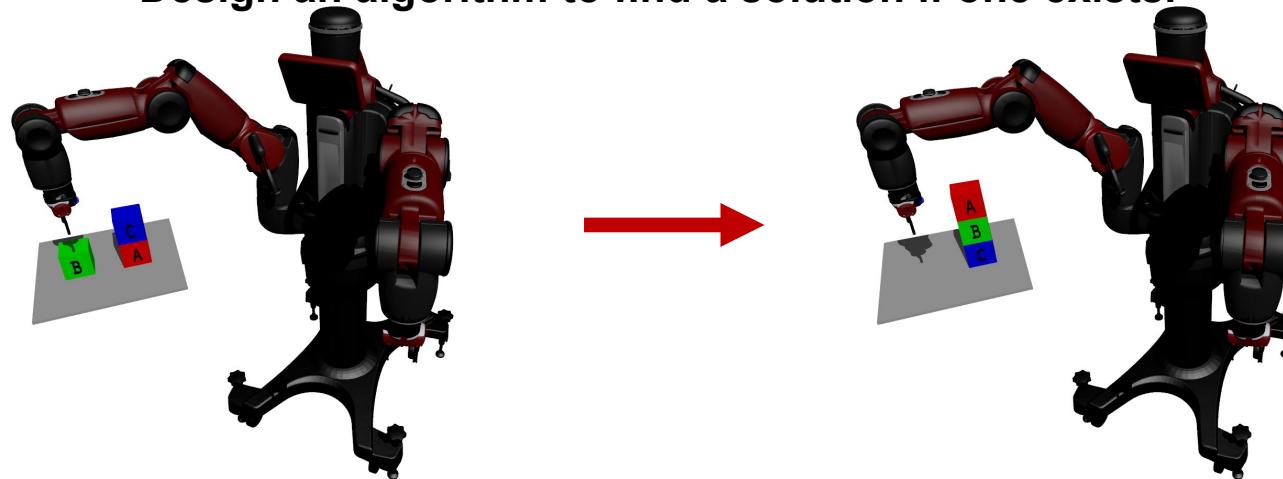
Overview

- Motion Planning Definition
- Bug Algorithms
- Discrete Search/Roadmaps
- Configuration Space
- Sampling-Based Planners
- Asymptotically Optimal Planning
- Kinodynamic Planning
- Learning And Planning
- **Task and Motion Planning**
- Planning Under Uncertainty

Task and Motion Planning

Informal	General Case
Tell robot where everything is	Initial state
Tell it where everything needs to end up	Goal state(s)
Let the robot figure it out	Transitions between states

Design an algorithm to find a solution if one exists.



Planning Domain Definition Language

- Planning Domain Definition Language (PDDL) is a way to express a task planning problem
- Use prefix notation:
 - `(+ 1 2) returns 3; (+ 1 2 (* 3 4)) returns 15`
- Define:
 - **Predicates:** `on (?x, ?y)`, `clear (?x)`, etc.
 - **Actions:** `stack (?x, ?y)`, etc.
 - **Initial state:** `(ontable a)`, etc.
 - **Goal state:** `(on a b)`, `(on b c)`, `(ontable a)`

Planning Domain Definition Language

Operators

```
(define (domain blocks)
  (:predicates (on ?x ?y) (ontable ?x)
               (clear ?x) (handempty) (holding ?x))
  (:action pick-up :parameters (?x)
    :precondition (and (clear ?x) (ontable ?x)
                       (handempty))
    :effect (and (not (ontable ?x)) (not (clear ?x))
                  (not (handempty)) (holding ?x)))
  (:action put-down :parameters (?x)
    :precondition (holding ?x)
    :effect (and (not (holding ?x)) (clear ?x)
                  (handempty) (ontable ?x)))
  (:action stack :parameters (?x ?y)
    :precondition (and (holding ?x) (clear ?y))
    :effect (and (not (holding ?x)) (not (clear ?y))
                  (clear ?x) (handempty) (on ?x ?y)))
  (:action unstack :parameters (?x ?y)
    :precondition (and (on ?x ?y) (clear ?x)
                       (handempty))
    :effect (and (holding ?x) (clear ?y)
                  (not (clear ?x))
                  (not (handempty))
                  (not (on ?x ?y))))
```

Facts

```
(define
  (problem sussman-anomaly)
  (:domain blocks)
  (:objects a b c)
  (:init (on c a)
         (ontable a)
         (ontable b)
         (clear c)
         (clear b)
         (handempty))
  (:goal (and (on b c)
              (on a b))))
```

Grounding on First-Order Logic

- Rely on the fact that there are finite:
 - Constants
 - Objects
 - locations
 - steps
- Construct a list of things that could be true or false **for each time step**
 - At (Obj?, Loc?) → PLATE_2_AT_LOC_3
 - At (Obj?, Loc?) → PLATE_1_AT_LOC_5
 - Etc.

Graph Search Approach

- Search state space to find a sequence of actions from start to goal
- Find successor states by looking at all actions that can be applied from a given state:
 - this means preconditions are met
- Search until we reach the goal:
 - this search space is typically very large
 - general heuristic search for task planning is an open research problem
 - use heuristic search
 - e.g., A*, Dijkstra's, FF, ...

Task Planning Approaches

- **Graph approach (employed so far)**
 - From PDDL description build a graph
 - Apply heuristic search algorithms
 - Search in graph implied by logical description
 - Heuristics are essential for good performance
- **Another approach that can be very efficient**
 - Reduce planning to a **SAT problem** instance (Boolean logic)
 - Solve with off-the-shelf solver
 - Avoid relying on task planning heuristics

SAT-Plan Encoding

- We want to feed a problem into the SAT solver
 - If the problem is **SAT**,
 - We get an assignment we can convert to a **task plan**
 - If the problem is **UNSAT**,
 - it means there is **no task plan within** that **time horizon**
- We need to **eliminate** quantifiers \forall and \exists
 - Quantifiers are used with variables
 - so we will **ground variables** and **time steps**

SAT-Plan Encoding

1. **(Initial State)** Need to be in the start state at step 0:

- Start State: start-step-0

2. **(Goal State)** Need to get to the goal by step k (a no-op can be used if we get there too early):

- Goal State: $\text{goal-step-}k$

3. **(Operator Encodings)** Applying an operator at step i implies preconditions and, on the next step, effects

- $\text{op-step-}i \rightarrow \text{pre(op)-step-}i \wedge \text{eff(op)-step-}(i+1)$

SAT-Plan Encoding

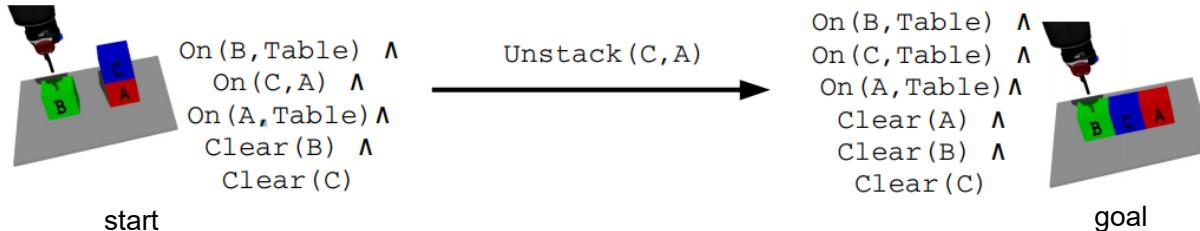
- 4. (Frame Axioms) Nothing changes unless we change it:
 - Frame Axioms: $(P\text{-step-}i == P\text{-step-}(i+1)) \vee (op\text{-}P\text{-}0 \vee \dots \vee op\text{-}P\text{-}n)$
- 5. (Complete Exclusion Axion) Can take only one action per step:
 - Operator Exclusion: $op\text{-}0 \rightarrow ((\neg op\text{-}1) \wedge \dots \wedge (\neg op\text{-}n))$
 - We ground all possible operators and predicates:
 - Relations: $rel_0\text{-arg_}0, \dots, rel_0\text{-arg_}K, \dots, rel_M\text{-arg_}0, \dots, rel_M\text{-arg_}K$
 - Operators: $op_0\text{-arg_}0, \dots, op_0\text{-arg_}K, \dots, op_M\text{-arg_}0, \dots, op_M\text{-arg_}K$

SAT-Plan Encoding

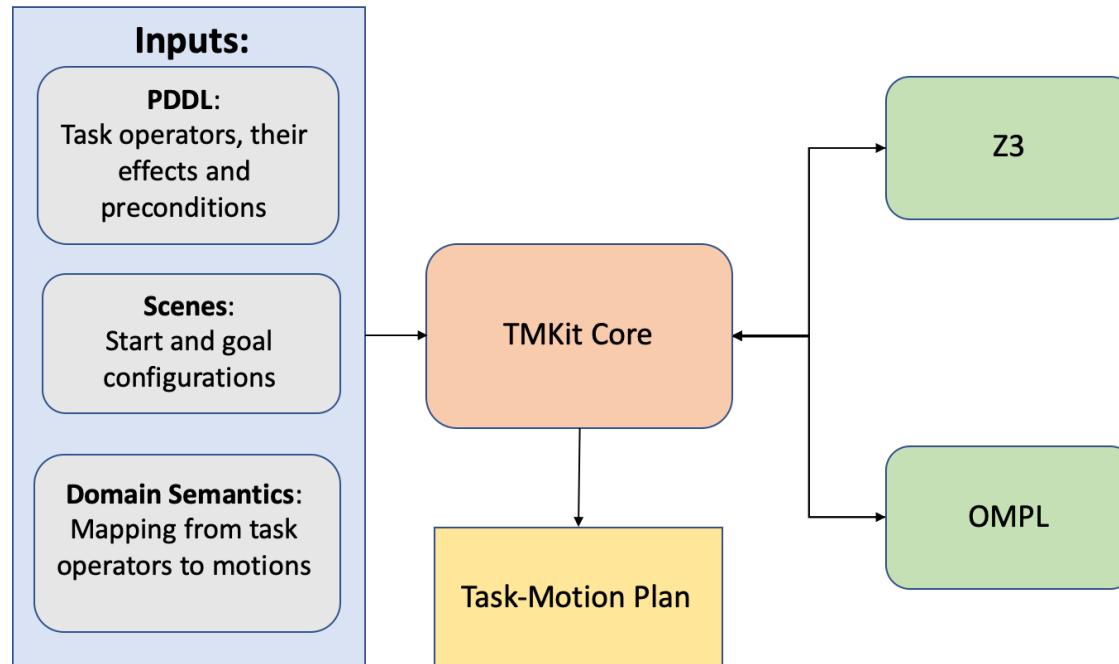
- Assume a value k for the number of steps in the solution
- Construct a huge formula with is the conjunction of 1-5 above
- A solution plan exists if and only if the resulting Boolean formula is satisfiable
- Give formula to a SAT solver
- Get the values that need to be assigned truth values and these will indicate the steps than need to be taken (steps is indicated in the encoding)

SAT-Plan Encoding

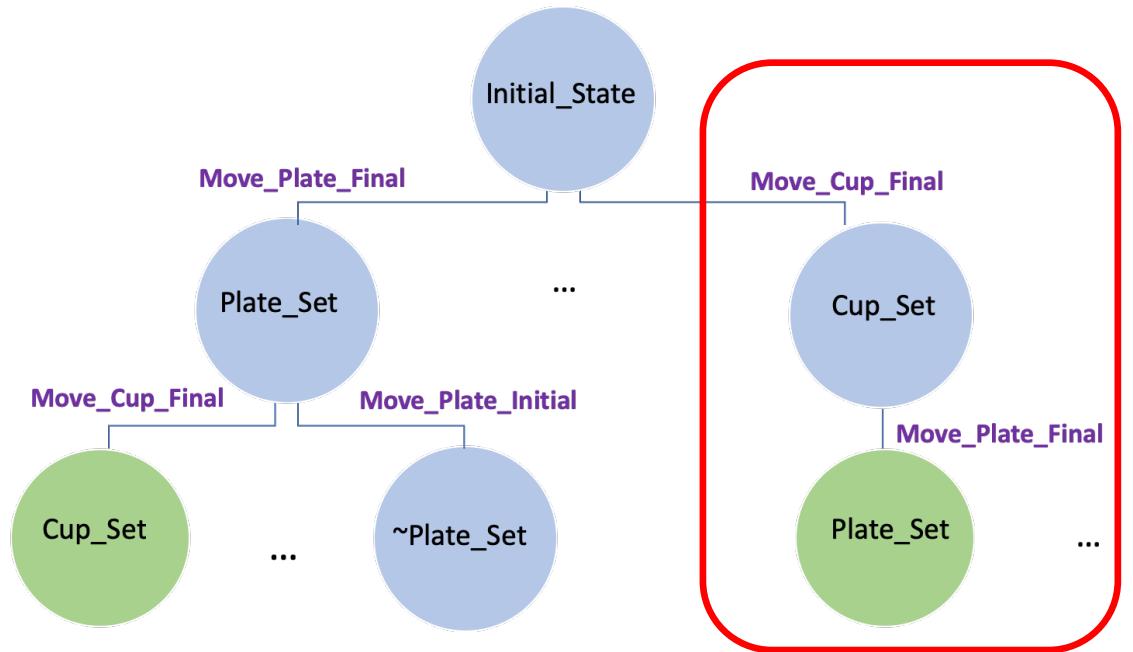
- Start State: start-step-0
- Goal State: goal-step-k
- Operator:
 - $\text{Unstack_C_A_step_0} \rightarrow \text{pre_unstack_C_A_step_0} \wedge \text{eff_unstack_C_A_step_1}$
- Frame axiom:
 - $(\text{on_C_A_step_0} == \text{on_C_A_step_1}) \vee (\text{unstack_C_A_step_0}) \vee \dots$
- Operator exclusion:
 - $\text{Unstack_C_A_step_0} \rightarrow (\neg \text{stack_C_B_step_0}) \wedge (\neg \text{stack_B_C_step_0}) \wedge \dots$



Case Study – TMKit architecture



Task and Motion Planning

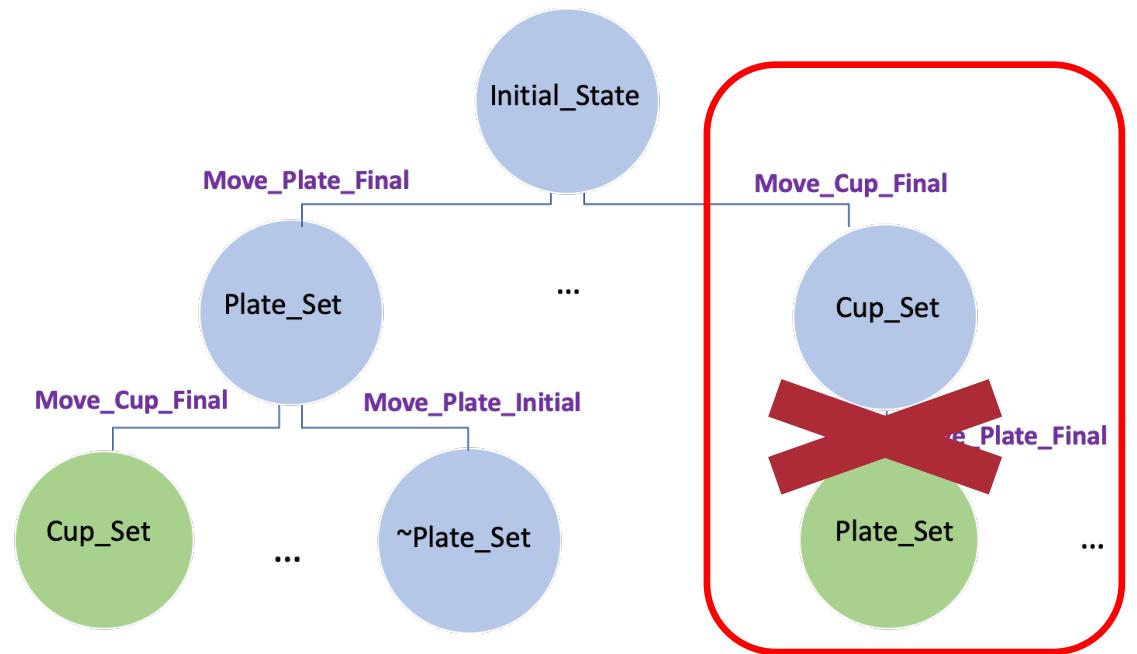


Initial_State:

- ¬ Plate_Set
- ¬ Cup_Set
- Plate_Initial
- Cup_Initial
- Hand_Free

...

Task and Motion Planning



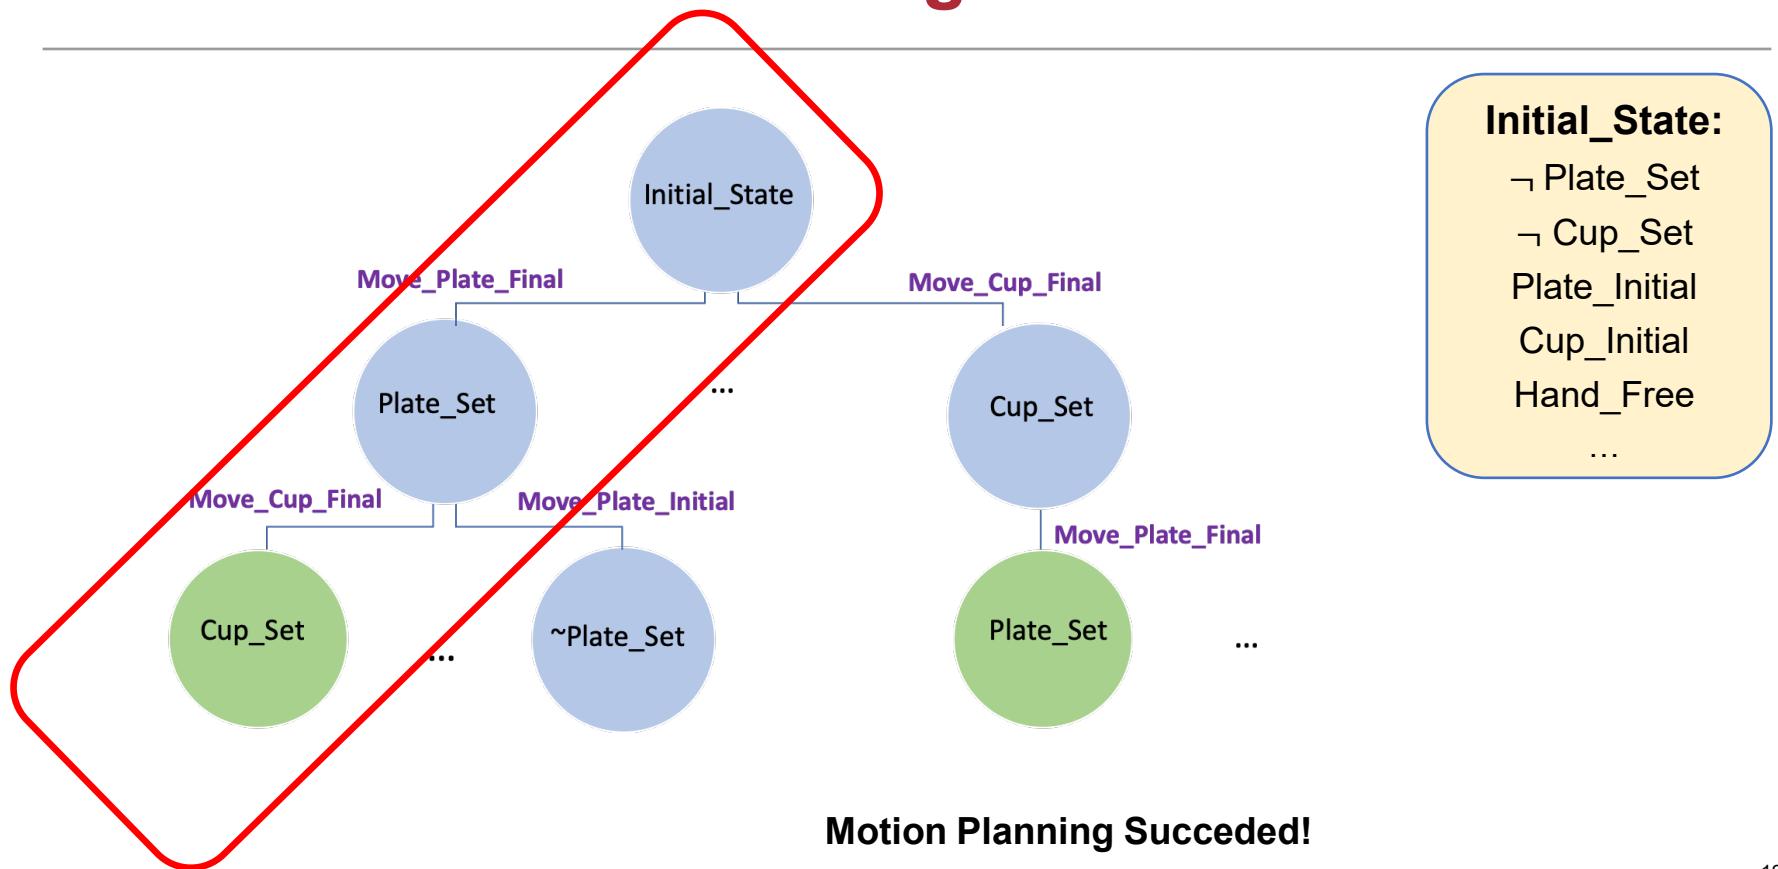
Initial_State:

- ¬ Plate_Set
- ¬ Cup_Set
- Plate_Initial
- Cup_Initial
- Hand_Free

...

Motion Planning Failed!

Task and Motion Planning



Overview

- Motion Planning Definition
- Bug Algorithms
- Discrete Search/Roadmaps
- Configuration Space
- Sampling-Based Planners
- Asymptotically Optimal Planning
- Kinodynamic Planning
- Learning And Planning
- Task and Motion Planning
- **Planning Under Uncertainty**

Source of Uncertainty

- What are the sources of uncertainty

- Sensing (State)

- Where is the robot?
- Where are the obstacles?
- What is the map of the environment?

- Motion (Actions)

- How will objects in the environment move?

- Given a command action, where will the robot end up?



Boston Dynamics

Markov Decision Processes

An MDP is defined by:

A set of states $s \in S$

A set of actions $a \in A$

- A transition function $T(s, a, s')$

- Probability that a from s leads to s' , i.e., $P(s'| s, a)$

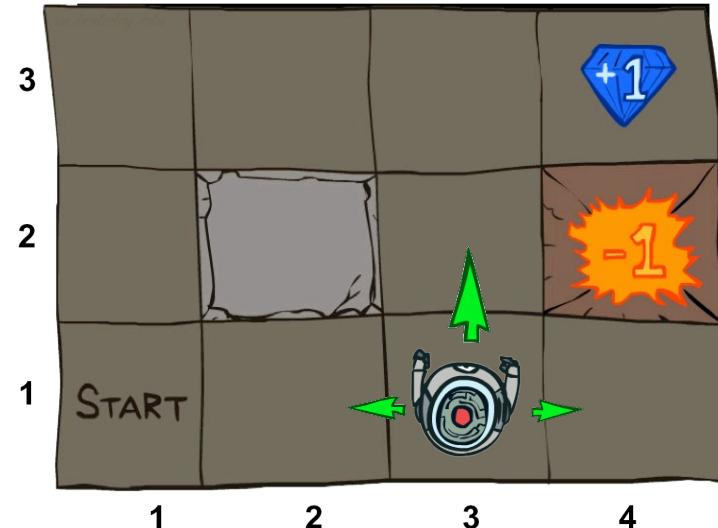
- Also called the model or the dynamics

- A reward function $R(s, a, s')$

- Sometimes just $R(s)$ or $R(s')$

- A start state

- Maybe a terminal state





Policies

In deterministic single-agent search problem
of actions, from start to a goal

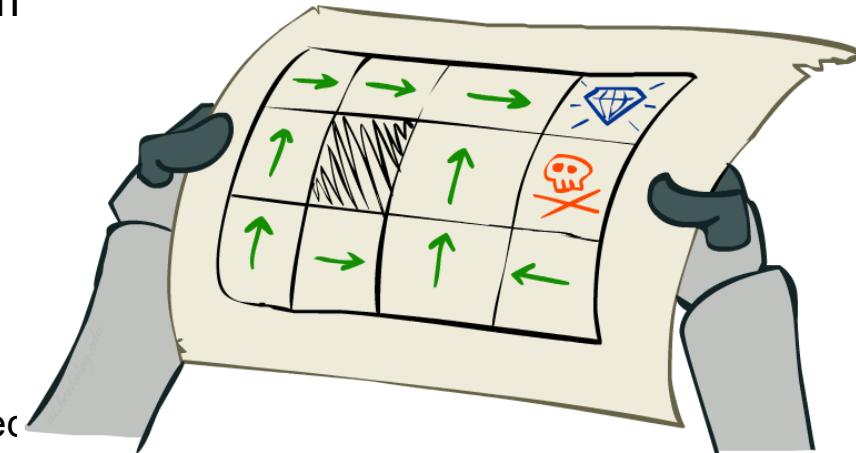
For MDPs, we want a policy $\pi: S \rightarrow A$

A policy π gives an action for each state

An optimal policy π^* is one that maximizes expected

An explicit policy defines a reflex agent

It computed the action for a single state only



Optimal policy when $R(s, a, s') = -0.03$
for all non-terminals s

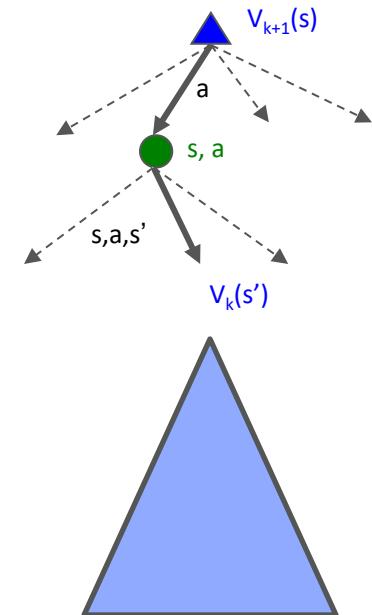


Value Iteration

- Start with $V_0(s) = 0$: no time steps left means an expected reward sum of zero
- Given vector of $V_k(s)$ values, do one step from each state:

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

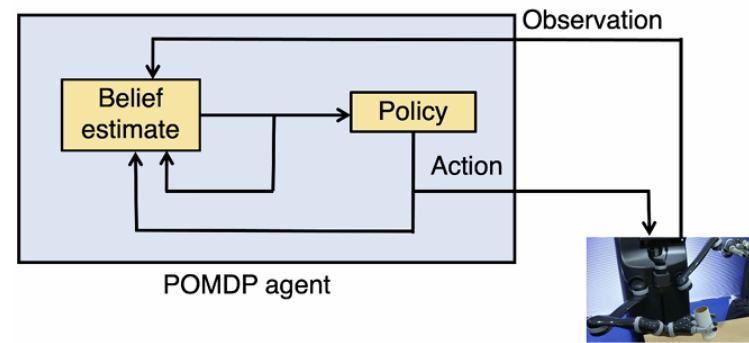
- Repeat until convergence
- Complexity of each iteration: $O(S^2A)$
- Theorem: will converge to unique optimal values
 - Basic idea: approximations get refined towards optimal values
 - Policy may converge long before values do





POMDP Definition

- 6-tuple model:
 - States: S
 - Actions: A
 - Transition function: $T(s, a, s')$ or $P(s'|s, a)$
 - Represents the non-deterministic effects of a
 - Reward: $R(s, a, s')$
 - Observation space: O
 - Observation function: $Z(s', a, o)$ or $P(o|s', a)$
 - Represent errors and noise in measurement and perception



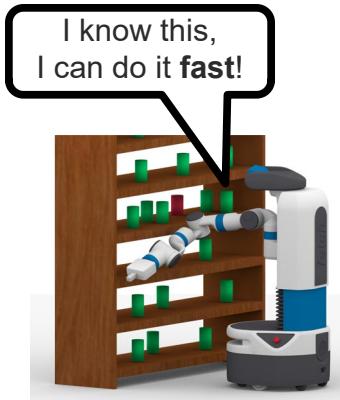
Overview

- Motion Planning Definition
- Bug Algorithms
- Discrete Search/Roadmaps
- Configuration Space
- Sampling-Based Planners
- Asymptotically Optimal Planning
- Kinodynamic Planning
- Learning And Planning
- Task and Motion Planning
- OMPL and Other Tools

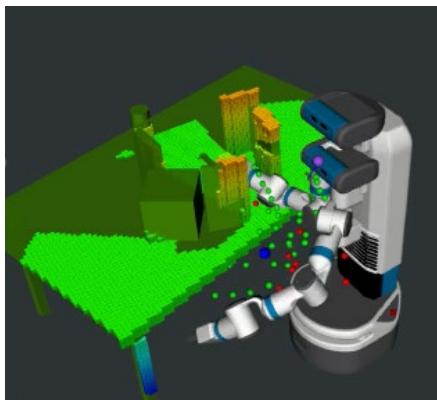
ELPIS Lab

Efficient Learning and Planning for Intelligent Systems

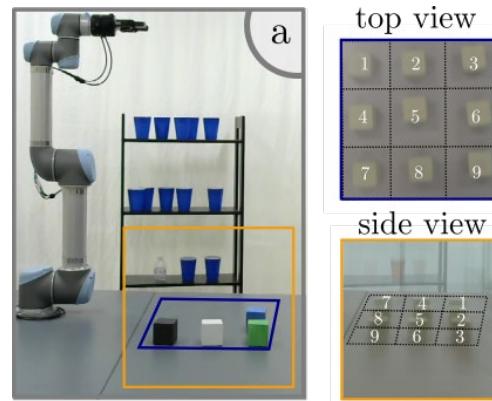
Research Directions/Projects



Learning For Planning Efficiency



Planning Robustly with Uncertainty



Vision-Based Planning



Available Hardware



Universal Robotics UR-10
6-Dof Industrial Manipulator



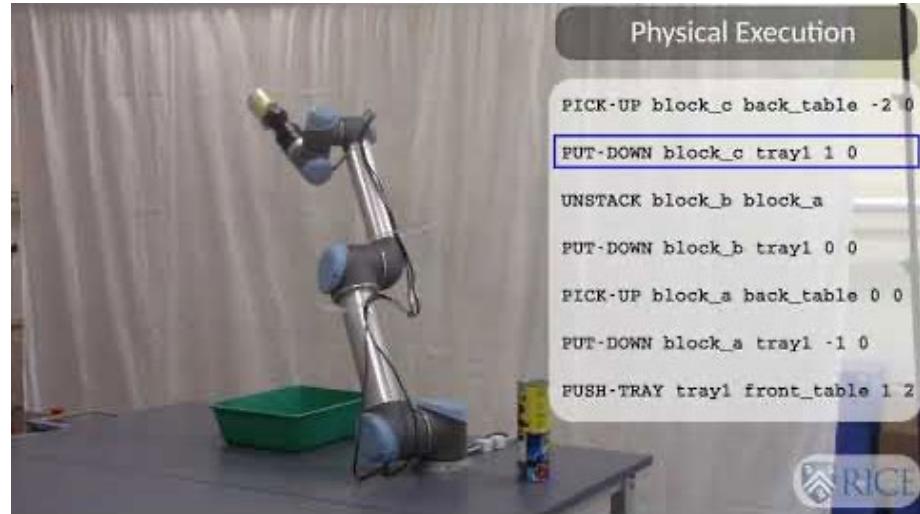
Deep Learning Workstation:
4xA6000 Nvidia GPU, 28-core Intel Xeon

... and more coming soon !

Planned Projects



Self-Supervised Learning Setup



Task and Motion Planning Setup

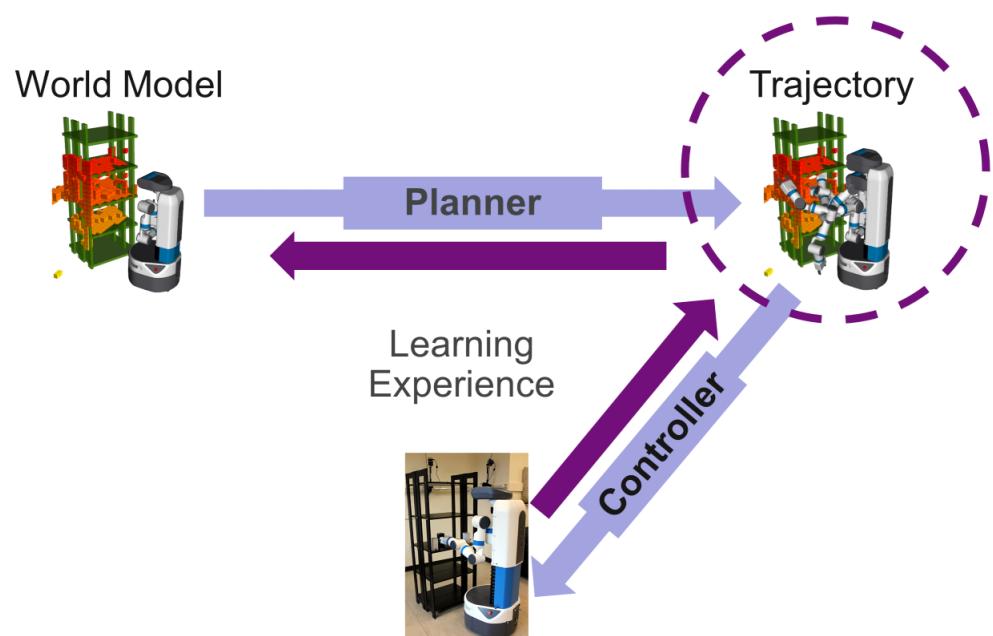
Future Work on the Intersection of Planning and Control

Kinodynamic manipulation

- Non prehensile manipulation
- Kinodynamic TAMP

Interleaving plan/sense actions for manipulation

- POMDP planning
- Ensure safety



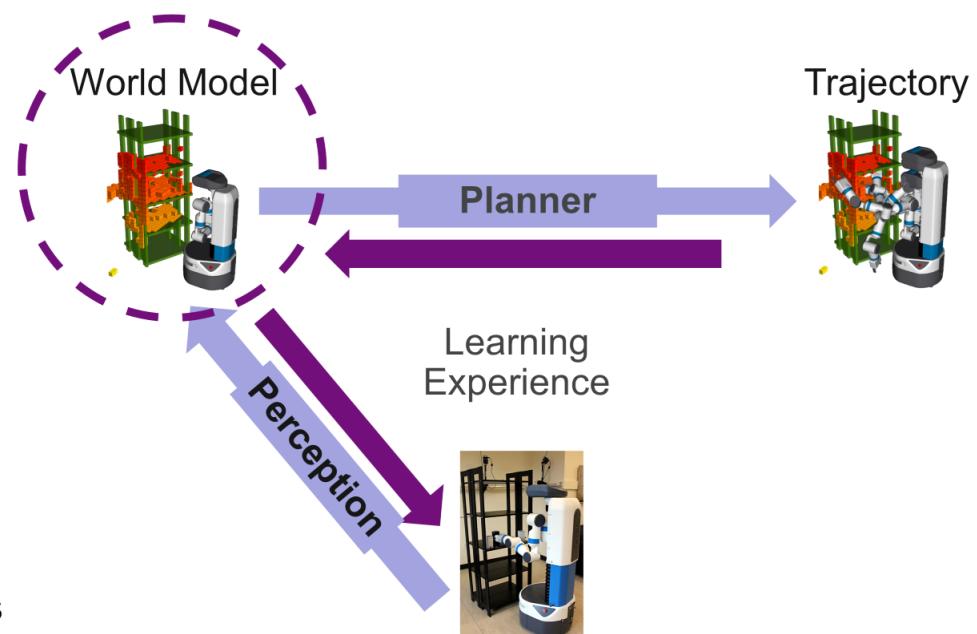
Future Work on the Intersection of Planning and Perception

Learn symbolic logic abstractions

- Image to symbols
- Learn through self-supervision

Learn forward dynamic models of environment objects

- Predict future environment states
- Model deformable/articulated objects



Next weeks Presentation Schedule!

- Groups 9 (Inclusive) to 19 will present on Monday (Dec 10)
- Groups 1-8 will present on Thursday (Dec 13)